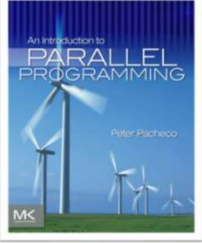



An Introduction to Parallel Programming  
Peter Pacheco



## Chapter 3

### Distributed Memory Programming with MPI



Copyright © 2010, Elsevier Inc. All rights Reserved


1

1

## Roadmap

- Writing your first MPI program.
- Using the common MPI functions.
- The Trapezoidal Rule in MPI.
- Collective communication.
- MPI derived datatypes.
- Performance evaluation of MPI programs.
- Parallel sorting.
- Safety in MPI programs.

# Chapter Subtitle

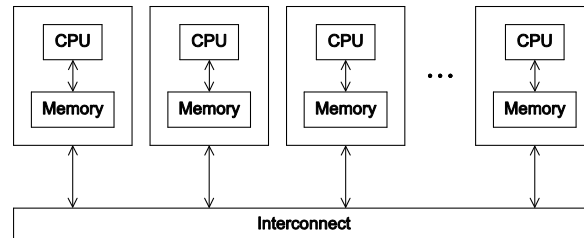


Copyright © 2010, Elsevier Inc. All rights Reserved

2

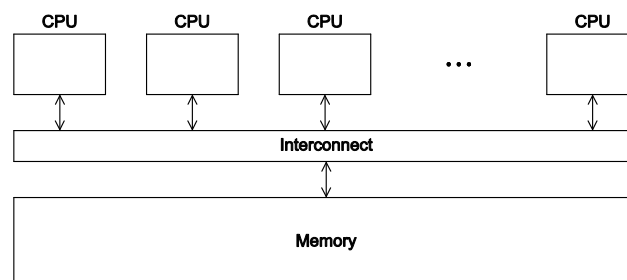
2

## A distributed memory system



3

## A shared memory system



4

## Hello World!

```
#include <stdio.h>

int main(void) {
    printf("hello, world\n");

    return 0;
}
```



**(a classic)**



Copyright © 2010, Elsevier Inc. All rights Reserved

5

5

## Identifying MPI processes

- Common practice to identify processes by nonnegative integer ranks.
- $p$  processes are numbered  $0, 1, 2, \dots, p-1$



Copyright © 2010, Elsevier Inc. All rights Reserved

6

6

## Our first MPI program

```

1 #include <stdio.h>
2 #include <string.h> /* For strlen */
3 #include <mpi.h> /* For MPI functions, etc */
4
5 const int MAX_STRING = 100;
6
7 int main(void) {
8     char    greeting[MAX_STRING];
9     int      comm_sz; /* Number of processes */
10    int      my_rank; /* My process rank */
11
12    MPI_Init(NULL, NULL);
13    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
14    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
15
16    if (my_rank != 0) {
17        sprintf(greeting, "Greetings from process %d of %d!",
18                my_rank, comm_sz);
19        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0,
20                MPI_COMM_WORLD);
21    } else {
22        printf("Greetings from process %d of %d!\n", my_rank, comm_sz);
23        for (int q = 1; q < comm_sz; q++) {
24            MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q,
25                    0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
26            printf("%s\n", greeting);
27        }
28    }
29
30    MPI_Finalize();
31    return 0;
32 } /* main */

```



Copyright © 2010, Elsevier Inc. All rights Reserved

7

7

## Compilation

*wrapper script to compile*

*source file*

**mpicc -g -Wall -o mpi\_hello mpi\_hello.c**

*produce debugging information*

*create this executable file name (as opposed to default a.out)*

*turns on all warnings*



Copyright © 2010, Elsevier Inc. All rights Reserved

8

8

## Execution

```
mpiexec -n <number of processes> <executable>
```

---

```
mpiexec -n 1 ./mpi_hello
```

*run with 1 process*

```
mpiexec -n 4 ./mpi_hello
```

*run with 4 processes*



Copyright © 2010, Elsevier Inc. All rights Reserved

9

9

## Execution

```
mpiexec -n 1 ./mpi_hello
```

Greetings from process 0 of 1 !

```
mpiexec -n 4 ./mpi_hello
```

Greetings from process 0 of 4 !

Greetings from process 1 of 4 !

Greetings from process 2 of 4 !

Greetings from process 3 of 4 !



Copyright © 2010, Elsevier Inc. All rights Reserved

10

10

## MPI Programs

- Written in C.
  - Has main.
  - Uses `stdio.h`, `string.h`, etc.
- Need to add `mpi.h` header file.
- Identifiers defined by MPI start with “MPI\_”.
- First letter following underscore is uppercase.
  - For function names and MPI-defined types.
  - Helps to avoid confusion.



Copyright © 2010, Elsevier Inc. All rights Reserved

11

11

## MPI Components

- `MPI_Init`
  - Tells MPI to do all the necessary setup.

```
int MPI_Init(
    int*      argc_p /* in/out */,
    char***   argv_p /* in/out */);
```

- `MPI_Finalize`
  - Tells MPI we're done, so clean up anything allocated for this program.

```
int MPI_Finalize(void);
```



Copyright © 2010, Elsevier Inc. All rights Reserved

12

12

## Basic Outline

```

1 2 3 4
#include <mpi.h>

int main(int argc, char* argv[]) {
    /* No MPI calls before this */
    MPI_Init(&argc, &argv);

    MPI_Finalize();
    /* No MPI calls after this */

    return 0;
}

```



Copyright © 2010, Elsevier Inc. All rights Reserved

13

13

## Communicators

- A collection of processes that can send messages to each other.
- MPI\_Init defines a communicator that consists of all the processes created when the program is started.
- Called **MPI\_COMM\_WORLD**.



Copyright © 2010, Elsevier Inc. All rights Reserved

14

14

## Communicators



```
int MPI_Comm_size(
    MPI_Comm comm      /* in */,
    int* comm_sz_p     /* out */);
```

*number of processes in the communicator*

```
int MPI_Comm_rank(
    MPI_Comm comm      /* in */,
    int* my_rank_p     /* out */);
```

*my rank  
(the process making this call)*



Copyright © 2010, Elsevier Inc. All rights Reserved

15

15

## SPMD

- Single-Program Multiple-Data
- We compile one program.
- Process 0 does something different.
  - Receives messages and prints them while the other processes do the work.
- The **if-else** construct makes our program SPMD.



Copyright © 2010, Elsevier Inc. All rights Reserved

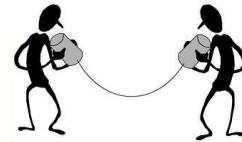
16

16



## Communication

```
int MPI_Send(
    void*      msg_buf_p    /* in */,
    int        msg_size     /* in */,
    MPI_Datatype msg_type    /* in */,
    int        dest         /* in */,
    int        tag          /* in */,
    MPI_Comm   communicator /* in */);
```



Copyright © 2010, Elsevier Inc. All rights Reserved

17

17

## Data types

MPI datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_LONG_LONG	signed long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	



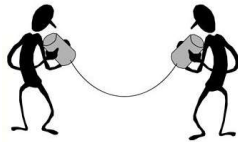
Copyright © 2010, Elsevier Inc. All rights Reserved

18

18

## Communication

```
int MPI_Recv(
    void*      msg_buf_p    /* out */,
    int        buf_size     /* in  */,
    MPI_Datatype buf_type    /* in  */,
    int        source       /* in  */,
    int        tag          /* in  */,
    MPI_Comm   communicator /* in  */,
    MPI_Status* status_p    /* out */);
```



Copyright © 2010, Elsevier Inc. All rights Reserved

19

19

## Message matching

```
MPI_Send(send_buf_p, send_buf_sz, send_type, dest, send_tag,
        send_comm);
```

*MPI\_Send*  
*src = q*



*MPI\_Recv*  
*dest = r*

```
MPI_Recv(recv_buf_p, recv_buf_sz, recv_type, src, recv_tag,
        recv_comm, &status);
```



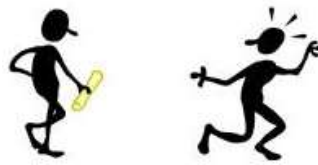
Copyright © 2010, Elsevier Inc. All rights Reserved

20

20

## Receiving messages

- A receiver can get a message without knowing:
  - the amount of data in the message,
  - the sender of the message,
  - or the tag of the message.



21

## status\_p argument

```
MPI_Recv(recv_buf_p, recv_buf_sz, recv_type, src, recv_tag,
recv_comm, &status);
```

**MPI\_Status\***

**MPI\_Status\* status;**

**status.MPI\_SOURCE**  
**status.MPI\_TAG**

*MPI\_SOURCE*  
*MPI\_TAG*  
*MPI\_ERROR*

22

## How much data am I receiving?

```
int MPI_Get_count(
    MPI_Status* status_p /* in */,
    MPI_Datatype type /* in */,
    int* count_p /* out */);
```



Copyright © 2010, Elsevier Inc. All rights Reserved

23

23

## Issues with send and receive

- Exact behavior is determined by the MPI implementation.
- MPI\_Send may behave differently with regard to buffer size, cutoffs and blocking.
- MPI\_Recv always blocks until a matching message is received.
- Know your implementation; don't make assumptions!



Copyright © 2010, Elsevier Inc. All rights Reserved

24

24