

# STATS 3001 / STATS 4104 / STATS 7054

## Statistical Modelling III

### Practical 1 - Linear model recap

### Solutions

#### Week 1

The following exercises are intended as revision of basic regression calculation and matrix manipulation functions in R.

For the purpose of this exercise, we will use built-in data set **Rubber** that is provided with the **MASS** library. The data set comprises three variables recorded on thirty samples of tyre rubber that were being tested for durability:

- **loss**: The abrasion loss in grams per hour;
- **hard**: The hardness in Shore units; and
- **tens**: The tensile strength in kg per square metre.

#### STEPS

- Load packages using the command

```
pacman::p_load(tidyverse, gglm)
```


- Load the rubber dataset from the MASS package

```
data(Rubber, package = "MASS")
```

- Obtain scatter-plots of loss against each of the other predictors.

```
Rubber %>%  
  pivot_longer(-loss) %>%  
  ggplot(aes(value, loss)) +  
  geom_point() +  
  facet_wrap(~name, scales = "free") +  
  geom_smooth(method = lm)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



- Use the `lm()` function to fit the following model.

$$E(\text{loss}_i) = \beta_0 + \beta_1 \times \text{hard}_i + \beta_2 \times \text{tens}_i.$$

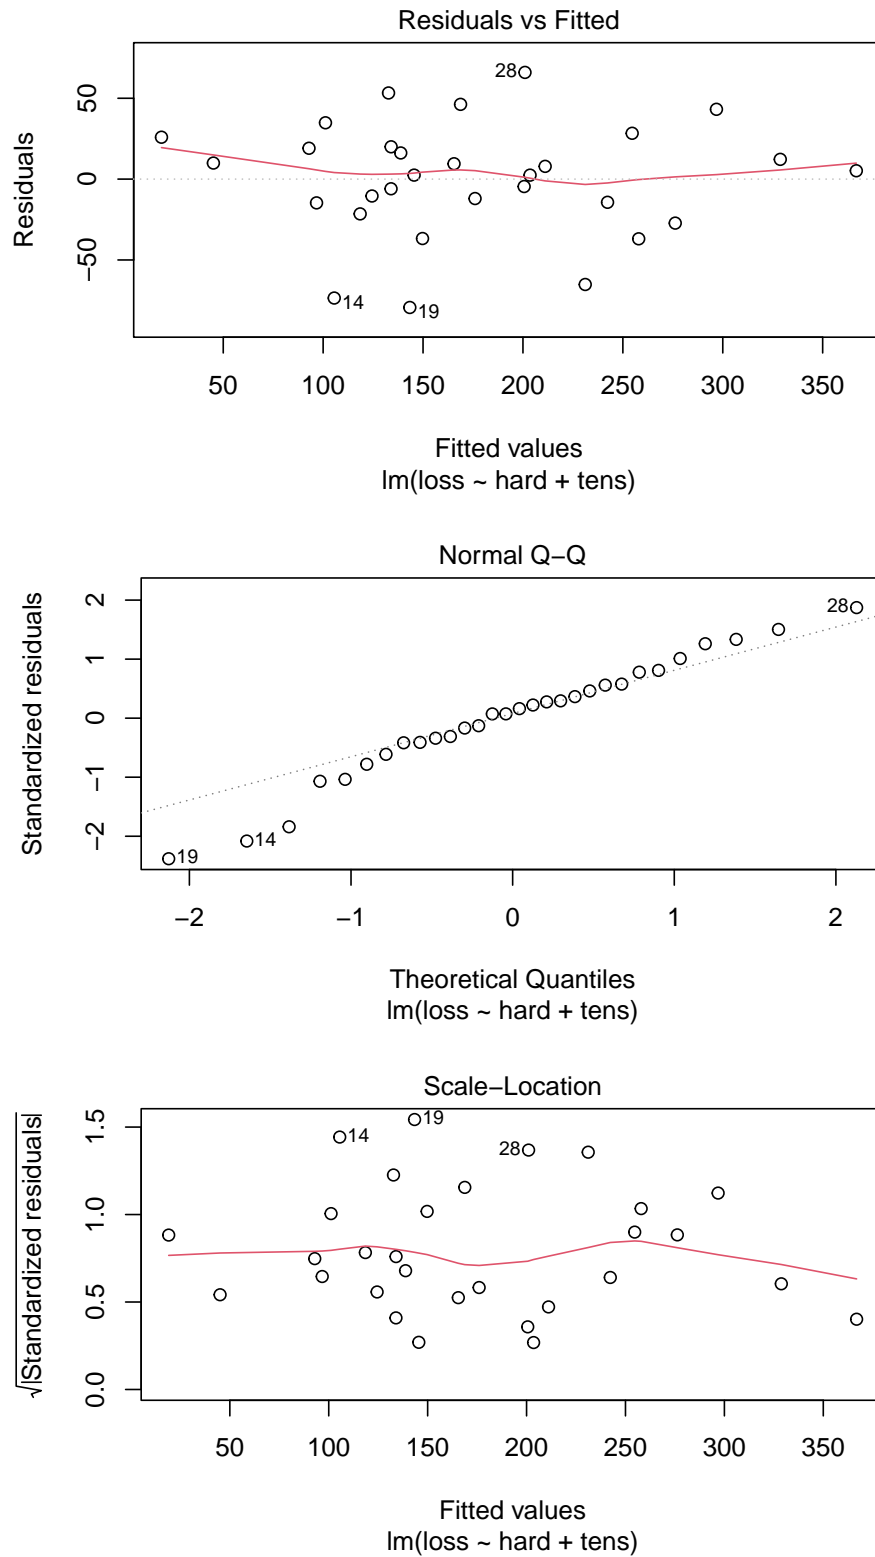
```
rubber_lm <- lm(loss ~ hard + tens, data = Rubber)
summary(rubber_lm)
```


```
##
## Call:
## lm(formula = loss ~ hard + tens, data = Rubber)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -79.385 -14.608   3.816  19.755  65.981
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  885.1611    61.7516  14.334 3.84e-14 ***
## hard         -6.5708     0.5832 -11.267 1.03e-11 ***
## tens         -1.3743     0.1943  -7.073 1.32e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 36.49 on 27 degrees of freedom
## Multiple R-squared:  0.8402, Adjusted R-squared:  0.8284
## F-statistic:    71 on 2 and 27 DF,  p-value: 1.767e-11
```

- Interpret the output
- The least squares estimate of  $\beta_1$  is  $\hat{\beta}_1 = -6.57$ . The interpretation of this estimate is that and increase of 1 unit of hardness, with tensile strength kept constant, would reduce the rate of abrasion by 6.57 grams per hour.
- The least squares estimate of  $\beta_2$  is  $\hat{\beta}_2 = -1.37$ . The interpretation of this estimate is that and increase of 1 unit of tensile strength, with hardness kept constant, would reduce the rate of abrasion by 1.37 grams per hour.

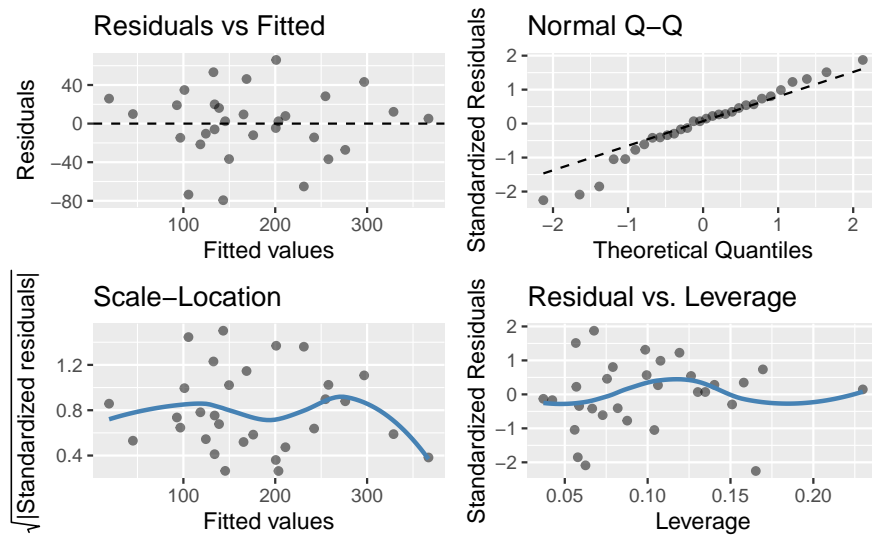
- What are the assumptions of the linear regression model. Produce appropriate plots to check them.

```
plot(rubber_lm)
```





```
gglm(rubber_lm)
```



### Linearity:

The residuals are roughly randomly scattered about the zero line in the residuals versus fitted values plot, apart from slight curvature near the endpoints. The residuals versus hardness plot shows a random scatter but the slight curvature is apparent in the residuals versus tensile strength plot. On balance, the linearity assumption is close to reasonable.

### Homoscedasticity

The spread about the zero line appears roughly constant in all three plots indicating that the assumption of constant variance is reasonable.

### Normality

There is some departure from normality in the lower tail of the distribution of residuals, with more large negative residuals than expected for a normal distribution. The bulk of the data is close to normally distributed however.

- Predict the loss for the following points:

hard	tens
50	200
65	190

- Calculate 95% CI and PI for the above points.

```
new_pts <- tibble(
  hard = c(50, 65),
  tens = c(200, 190)
)
new_pts
```

```
## # A tibble: 2 x 2
##   hard tens
##   <dbl> <dbl>
## 1    50  200
## 2    65  190
```

```
predict(rubber_lm)
```

```
##           1           2           3           4           5           6           7           8
## 366.83526 203.55082 165.50016 134.02032 101.16617  92.92030  45.07805  19.09546
##           9          10          11          12          13          14          15          16
## 257.92184 231.16639 176.02253 149.73921  96.70044 105.54777 242.33228 200.58874
##          17          18          19          20          21          22          23          24
## 133.97826 118.51804 143.38498 276.21795 211.11111 132.73328 138.83198 124.44535
##          25          26          27          28          29          30
## 328.78459 296.83263 254.65903 201.01880 168.76611 145.53215
```

```
predict(rubber_lm, newdata = new_pts)
```

```
##           1           2
## 281.7573 196.9379
```

```
predict(rubber_lm, newdata = new_pts, interval = "confidence")
```

```
##           fit          lwr          upr
## 1 281.7573 254.8762 308.6383
## 2 196.9379 181.8821 211.9938
```

```
predict(rubber_lm, newdata = new_pts, interval = "prediction")
```

```
##           fit          lwr          upr
## 1 281.7573 202.2079 361.3066
## 2 196.9379 120.5692 273.3067
```

- Get the design matrix of the model using the command `model.matrix()`.

```
X <- model.matrix(rubber_lm)
X
```

```
##      (Intercept) hard tens
## 1           1    45  162
## 2           1    55  233
## 3           1    61  232
## 4           1    66  231
## 5           1    71  231
## 6           1    71  237
## 7           1    81  224
## 8           1    86  219
## 9           1    53  203
## 10          1    60  189
## 11          1    64  210
## 12          1    68  210
## 13          1    79  196
## 14          1    81  180
## 15          1    56  200
## 16          1    68  173
## 17          1    75  188
## 18          1    83  161
## 19          1    88  119
## 20          1    59  161
## 21          1    71  151
## 22          1    80  165
## 23          1    82  151
## 24          1    89  128
## 25          1    51  161
## 26          1    59  146
## 27          1    65  148
## 28          1    74  144
## 29          1    81  134
## 30          1    86  127
## attr("assign")
## [1] 0 1 2
```

- Assign the response variable `loss` to a variable `Y`.

```
Y <- Rubber$loss
Y
```

```
## [1] 372 206 175 154 136 112  55  45 221 166 164 113  82  32 228 196 128  97  64
## [20] 249 219 186 155 114 341 340 283 267 215 148
```

- Using the R commands `%*%`, `solve()`, and `t()`, calculate

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

```
beta_hat <- solve(t(X) %*% X) %*% t(X) %*% Y
beta_hat
```

```
##           [,1]
## (Intercept) 885.161109
## hard        -6.570830
## tens        -1.374312
```

```
summary(rubber_lm)
```

```
##
## Call:
## lm(formula = loss ~ hard + tens, data = Rubber)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -79.385 -14.608   3.816  19.755  65.981
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 885.1611    61.7516  14.334 3.84e-14 ***
## hard        -6.5708     0.5832 -11.267 1.03e-11 ***
## tens        -1.3743     0.1943  -7.073 1.32e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 36.49 on 27 degrees of freedom
## Multiple R-squared:  0.8402, Adjusted R-squared:  0.8284
## F-statistic:    71 on 2 and 27 DF,  p-value: 1.767e-11
```

Is it the same as the answer from `lm()`?

- Calculate the fitted values

$$\hat{\eta} = X\hat{\beta}$$

```
eta <- X %*% beta_hat
head(eta)
```

```
##           [,1]
## 1 366.8353
## 2 203.5508
## 3 165.5002
## 4 134.0203
## 5 101.1662
## 6  92.9203
```

- Calculate the residual variance,  $s_e^2$  directly from the observed and fitted values. Compare the result to the residual standard error produced by `lm()`

```
n <- nrow(Rubber)
se <- sqrt(sum((Y - eta)^2) / (n - 3))
se
```

```
## [1] 36.48934
```

- Calculate the estimated variance matrix for  $\hat{\beta}$  using

$$(X^T X)^{-1} \times s_e^2$$

Compare this to the result of the built-in calculation `vcov()`.

```
vcov(rubber_lm)
```

```
##           (Intercept)           hard           tens
## (Intercept)  3813.25821 -30.01766345 -9.19635018
## hard         -30.01766   0.34010794  0.03390882
## tens         -9.19635   0.03390882  0.03775595
```

```
solve(t(X) %% X) * se^2
```

```
##           (Intercept)           hard           tens
## (Intercept)  3813.25821 -30.01766345 -9.19635018
## hard         -30.01766   0.34010794  0.03390882
## tens         -9.19635   0.03390882  0.03775595
```