

STATS 3001 / STATS 4104 / STATS 7054

Statistical Modelling III

Practical 3 - Assumption checking

Solutions

Week 5

GOAL

The purpose of this practical is to explore the application of influence diagnostics in multiple regression. The learning objectives of the practical are

- To demonstrate the practical application of influence diagnostics;
- To demonstrate the correspondence between leverage and the distribution of the x -values;
- To verify the built-in function for calculating leverage;
- To demonstrate the correspondence between Cook's Distance and changes in the parameter estimates.

DATA

The file `hills.csv` contains the record times in 1984 for 35 Scottish hills races.

The dataset contains the following variables:

- `dist`: The total distance in miles
- `climb`: The total climb in feet
- `time`: The record time in minutes

Interest is focused on modelling `time` using `dist` and `climb` as predictors.

STEPS

- Read in the data

```
pacman::p_load(tidyverse, ggrepel)
```

```
##
## The downloaded binary packages are in
## /var/folders/pr/yfj4d0gn5gv9gy0pnjtv1nt40000gp/T//RtmpEX0pgT/downloaded_packages
##
## ggrepel installed
```

```
hills <- read_csv(here::here("data", "hills.csv"))
```

```
## Rows: 35 Columns: 4
```

```
## -- Column specification -----
## Delimiter: ","
## chr (1): race
## dbl (3): dist, climb, time
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

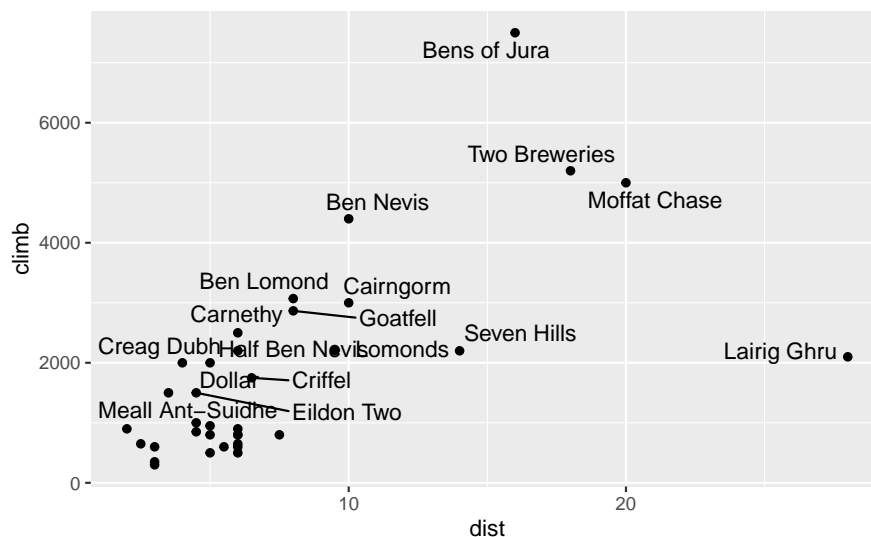
```
hills
```

```
## # A tibble: 35 x 4
##   race      dist climb  time
##   <chr>    <dbl> <dbl> <dbl>
## 1 Greenmantle  2.5   650  16.1
## 2 Carnethy     6  2500  48.4
## 3 Craig Dunain  6   900  33.6
## 4 Ben Rha     7.5   800  45.6
## 5 Ben Lomond   8  3070  62.3
## 6 Goatfell     8  2866  73.2
## 7 Bens of Jura 16  7500 205.
## 8 Cairnpapple  6   800  36.4
## 9 Scolty       5   800  29.8
## 10 Traprain    6   650  39.8
## # ... with 25 more rows
```

- Obtain a scatter plot of `dist` vs `climb`. Identify the points that you believe will have high leverage. The package `ggrepel` is very cool for this.

```
hills %>%
  ggplot(aes(dist, climb)) +
  geom_point() +
  geom_text_repel(aes(label = race))
```

```
## Warning: ggrepel: 18 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps
```



- Calculate the leverage values from the design matrix for the model $\{ \sim \text{climb} + \text{dist} \}$ using the matrix expression given in lectures. Note the command `diag(H)` will extract the diagonal values of a square matrix `H`.


Identify the points with leverage greater than $2p/n$.

Check whether the points you identified on the scatter plot do have high leverage.

```
X <- model.matrix(~ climb + dist, data = hills)
H <- X %>% solve( t(X) %>% X ) %>% t(X)
hills <-
  hills %>%
  add_column(
    leverage = diag(H)
  )
hills
```

```
## # A tibble: 35 x 5
##   race      dist climb  time leverage
##   <chr>    <dbl> <dbl> <dbl>    <dbl>
## 1 Greenmantle  2.5   650  16.1  0.0538
## 2 Carnethy     6  2500  48.4  0.0495
## 3 Craig Dunain  6   900  33.6  0.0384
## 4 Ben Rha     7.5   800  45.6  0.0485
## 5 Ben Lomond   8  3070  62.3  0.0553
## 6 Goatfell     8  2866  73.2  0.0468
## 7 Bens of Jura 16  7500 205.   0.420
## 8 Cairnpapple  6   800  36.4  0.0410
## 9 Scolty       5   800  29.8  0.0403
## 10 Traprain    6   650  39.8  0.0457
## # ... with 25 more rows
```

```
p <- ncol(X)
n <- nrow(X)
hills %>%
  ggplot(aes(leverage, fct_reorder(race, leverage))) +
  geom_col() +
  theme(axis.text.x = element_text(angle = -90, hjust=0)) +
  geom_vline(xintercept = 2 * p / n) +
  labs(y = "race")
```






- Calculate the leverage values using the built-in `hatvalues()` function in R and check that they agree with those calculated from the formula.
(Note: R also provides functions, `cooks.distance()`, `rstudent()` and `rstandard()` to calculate Cook's distance, the studentized residuals and the standardized residuals, respectively.)

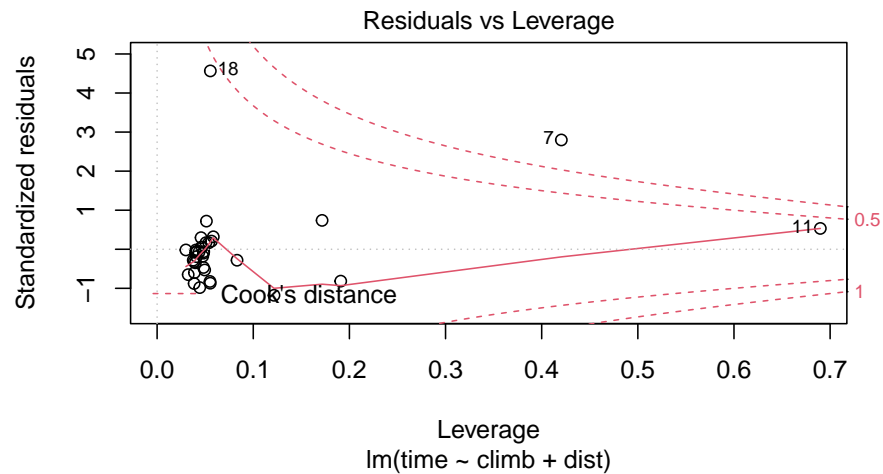
```
hills_lm <- lm(time ~ climb + dist, data = hills)
hills <-
  hills %>%
  add_column(
    r_leverage = hatvalues(hills_lm)
  )
hills
```

```
## # A tibble: 35 x 6
##   race      dist climb  time leverage r_leverage
##   <chr>    <dbl> <dbl> <dbl>    <dbl>    <dbl>
## 1 Greenmantle  2.5   650  16.1  0.0538  0.0538
## 2 Carnethy      6  2500  48.4  0.0495  0.0495
## 3 Craig Dunain  6   900  33.6  0.0384  0.0384
## 4 Ben Rha      7.5   800  45.6  0.0485  0.0485
## 5 Ben Lomond    8  3070  62.3  0.0553  0.0553
## 6 Goatfell     8  2866  73.2  0.0468  0.0468
## 7 Bens of Jura 16  7500 205.   0.420   0.420
## 8 Cairnpapple  6   800  36.4  0.0410  0.0410
## 9 Scolty       5   800  29.8  0.0403  0.0403
## 10 Traprain    6   650  39.8  0.0457  0.0457
## # ... with 25 more rows
```

- Obtain the usual sequence of diagnostic plots from R.

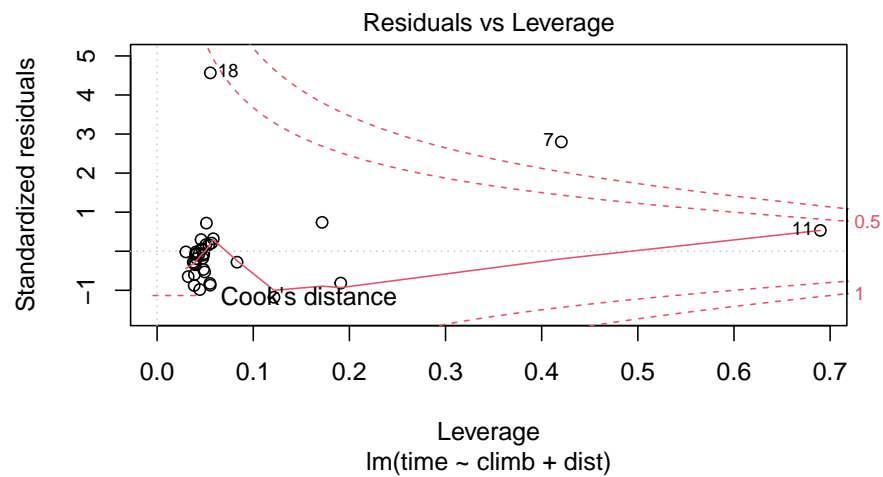
```
plot(hills_lm)
```





- Based on the residuals vs leverage plot, identify the most influential point.

```
plot(hills_lm, which = 5)
```



```
hills %>% slice(7)
```

```
## # A tibble: 1 x 6
##   race      dist climb  time leverage r_leverage
##   <chr>    <dbl> <dbl> <dbl>    <dbl>    <dbl>
## 1 Bens of Jura    16  7500  205.    0.420    0.420
```

- Identify the points with the largest residual and the highest leverage, and comment.

```
augment(hills_lm) %>%
  add_column(
    race = hills$race
  ) %>%
  filter(
    .resid == max(.resid) | .hat == max(.hat) | .cooks == max(.cooks)
  ) %>%
  select(race, .resid, .hat, .cooks)
```

```
## # A tibble: 3 x 4
##   race      .resid  .hat .cooks
##   <chr>      <dbl> <dbl> <dbl>
## 1 Bens of Jura 31.3  0.420  1.89
## 2 Lairig Ghru  4.36 0.690  0.211
## 3 Knock Hill  65.1 0.0554 0.407
```

So Knock Hill has largest residual, but leverage is small, while Lairig Ghru has the largest leverage, but small residual.

- Fit the same model to the data with the most influential point removed.

```
hills_lm2 <- lm(time ~ climb + dist,
  data = hills %>% filter(race != "Bens of Jura"))
```

- Calculate Cook's distance according to the formula

$$D_i^2 = \frac{\left(\hat{\beta} - \hat{\beta}^{(i)}\right)^T \left[\hat{\text{Var}}(\hat{\beta})\right]^{-1} \left(\hat{\beta} - \hat{\beta}^{(i)}\right)}{p}$$

$$= \frac{\left(\hat{\beta} - \hat{\beta}^{(i)}\right)^T (X^T X) \left(\hat{\beta} - \hat{\beta}^{(i)}\right)}{ps_e^2}.$$

Check that your value agrees with that produced by the built-in `cooks.distance` function.

```
beta <- coef(hills_lm)
beta_i <- coef(hills_lm2)
X <- model.matrix(hills_lm)
se2 <- glance(hills_lm)$sigma^2
p <- ncol(X)
t(beta - beta_i) %*% t(X) %*% X %*% (beta - beta_i) / (p*se2)
```

```
##           [,1]
## [1,] 1.893349
```