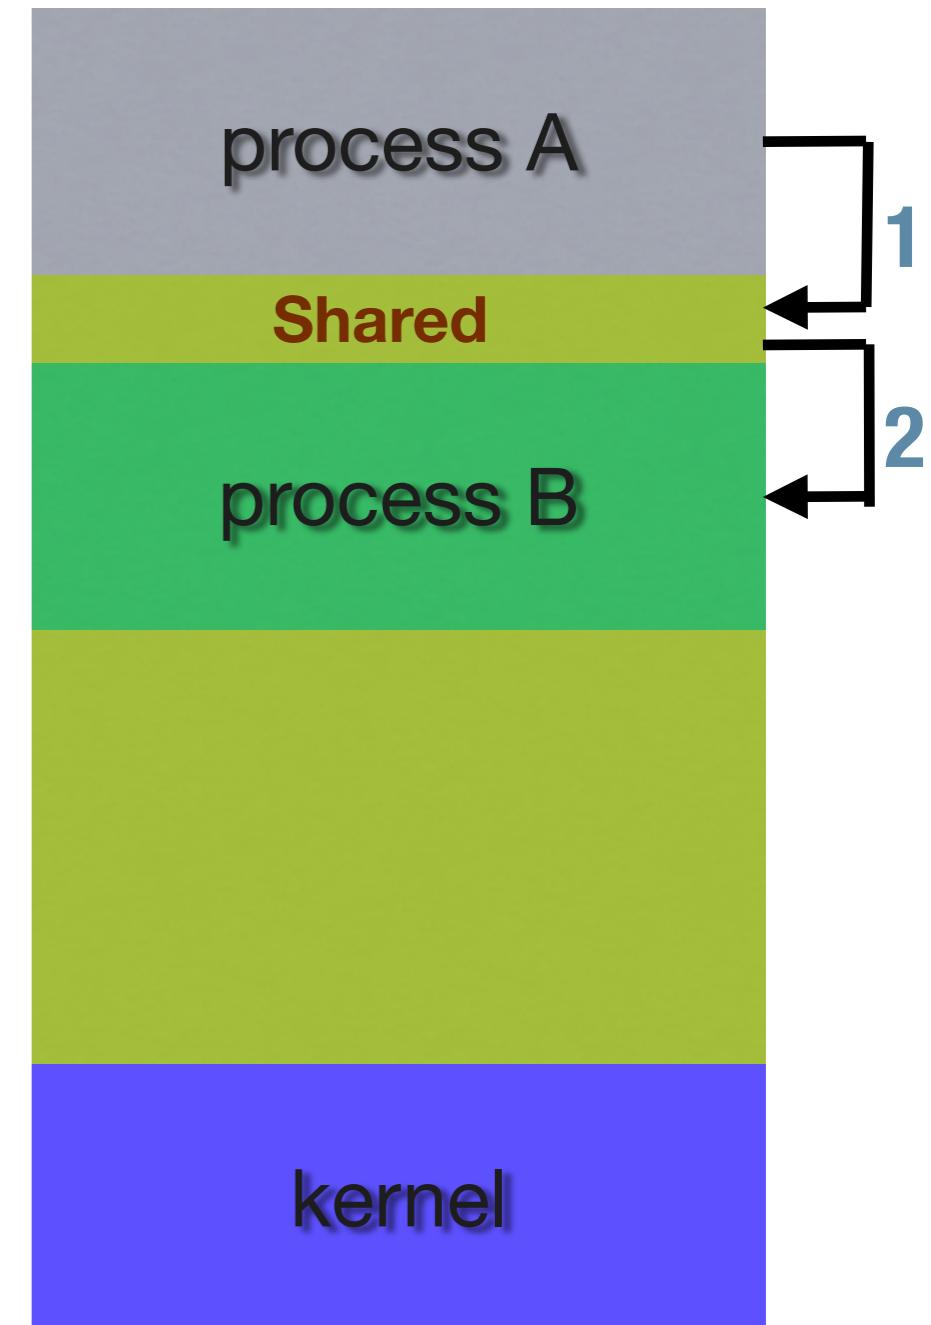
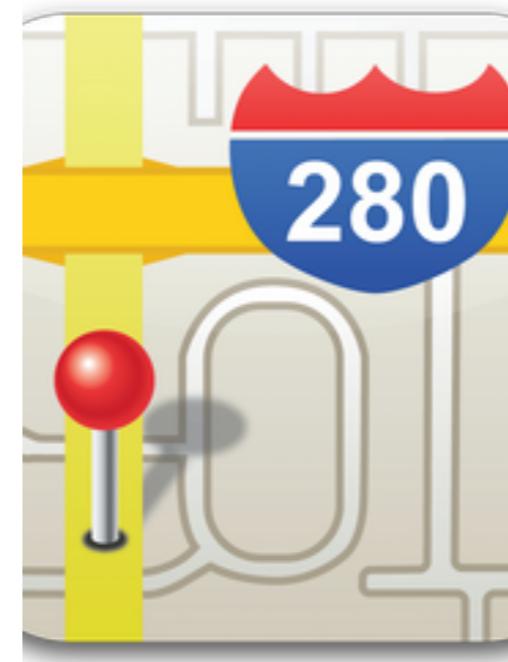
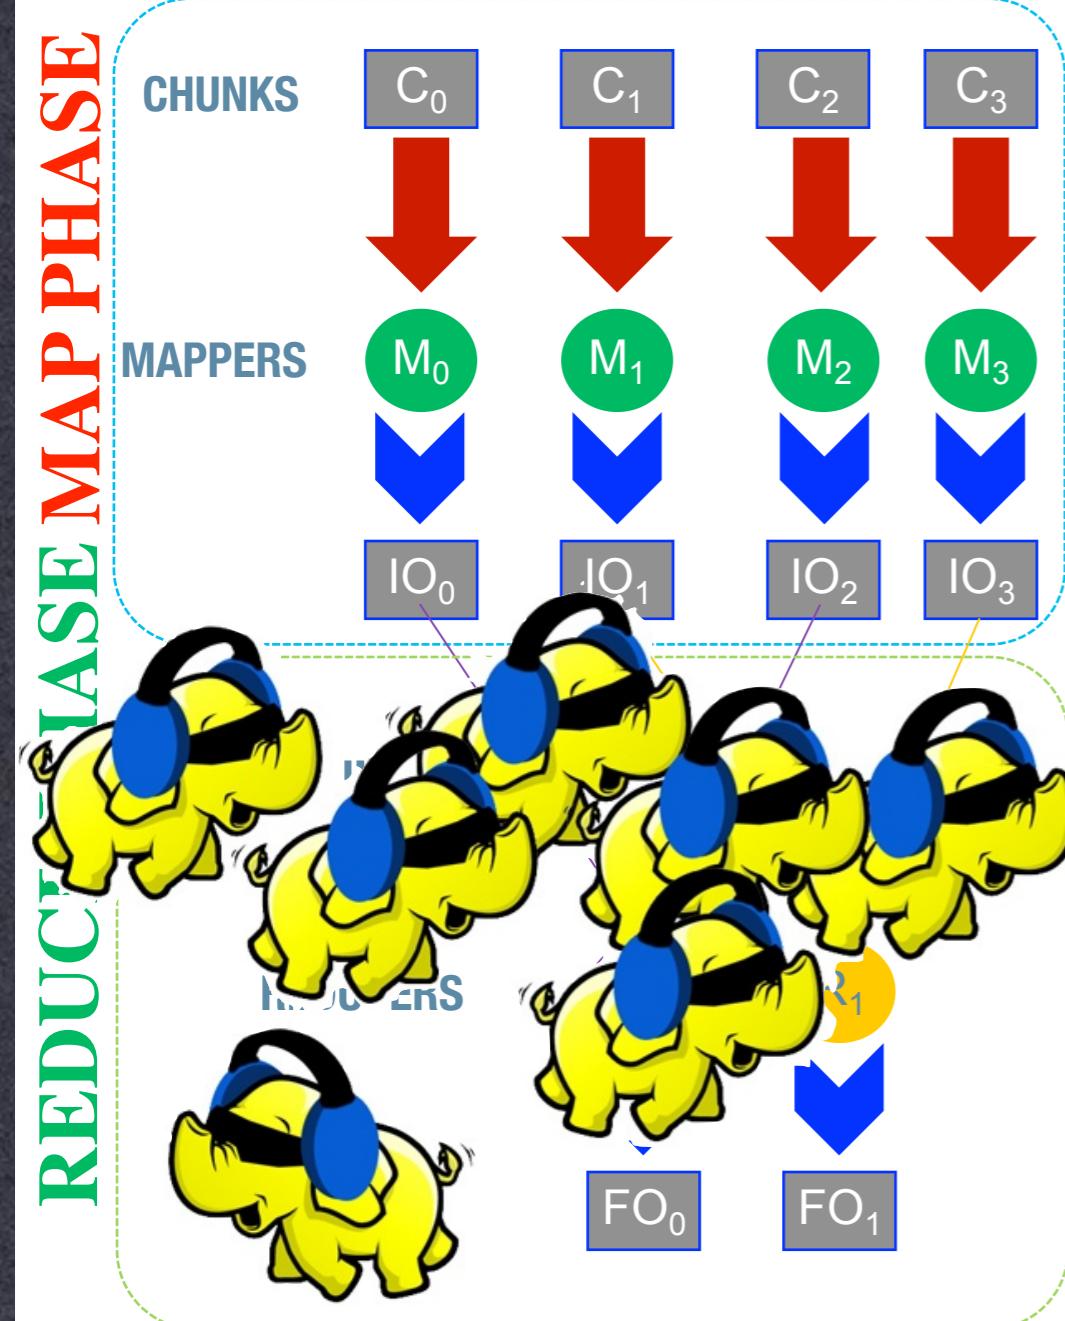


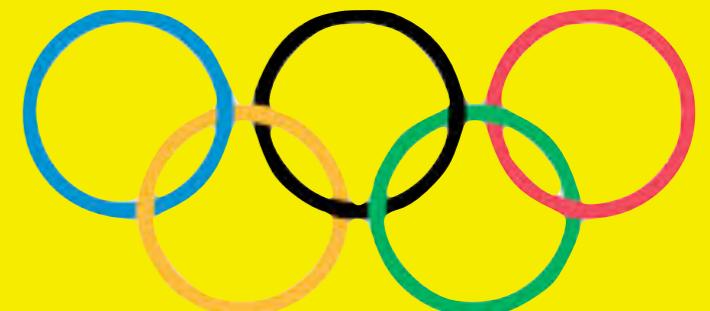
REDUCE PHASE MAP PHASE



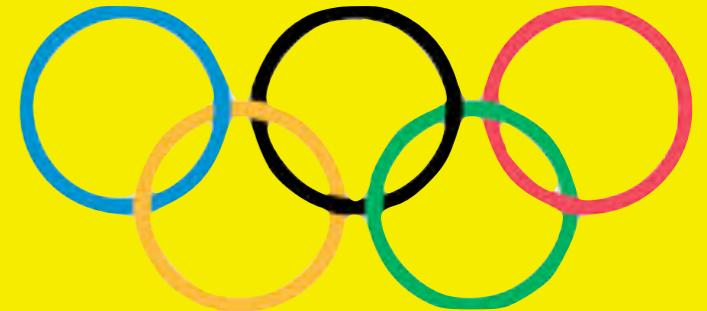
DISTRIBUTED SYSTEMS

PROCESSES, THREADS, LOCAL SYNCH

Activity: Distributed Systems Olympics



- * Form groups of three
- * Write down three examples of distributed systems
- * Decide as a group on your favorite and tell me why
- * Tell me what your favorite is
- * Give me the piece of paper – we will talk about your rejects at the end of the course



- * Asynchronous, lagging, communication
- * Communication is not instantaneous
 - * time matters
- * Failure is not an option
- * <http://www.youtube.com/watch?v=h2I8AoB1xgU>

Welcome (back to Earth)!

- ✳️ Last's week lecture introduced the course
- ✳️ Grand challenges in distributed systems
 - ✳️ synchronization
 - ✳️ transparency
 - ✳️ heterogeneity
 - ✳️ openness

Welcome (back to Earth)!

- ✳ Today's lecture will be all about processes, threads, local coordination & communication
- ✳ Is this an OS course?
 - ✳ No
 - ✳ Concepts are nevertheless essential

Program, process, threads

- * What is a program?
- * What is a process?

What is a process?

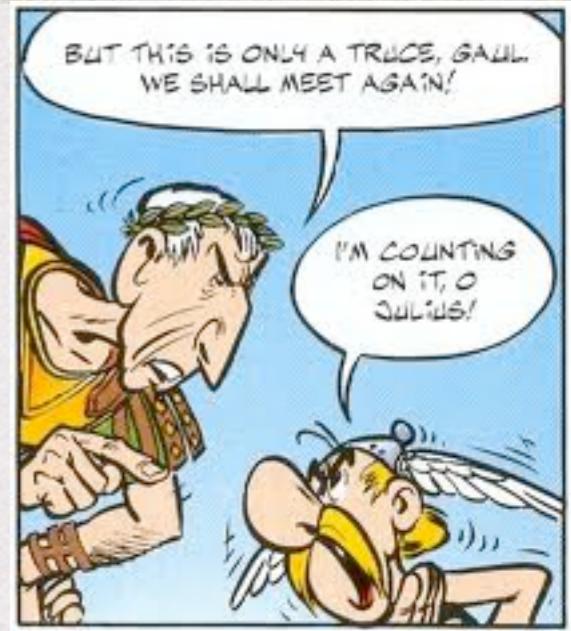
- ✳ Process: an execution stream (or program) in the context of a particular state
- ✳ Process state: determines effect of running code
 - ✳ Registers: general purpose, instruction pointer (program counter), floating point, ...
 - ✳ Memory: everything a process can address, code, data, stack, heap
 - ✳ I/O status: file descriptor table

Program vs process

- ✳ Program != process
 - ✳ Program: static code + static data
 - ✳ Process: dynamic instantiation of code + data + more
- ✳ Program : process: no 1:1 mapping
 - ✳ Process > program: more than code and data
 - ✳ Program > process: one program runs many processes
 - ✳ Process > program: many processes of same program

Why processes?

- ✳ Express concurrency
 - ✳ a web server can spawn multiple processes to take care of requests in parallel
 - ✳ managed by OS
- ✳ Follows divide and conquer
- ✳ Processes are isolated from each other
 - ✳ work sequentially
 - ✳ well-defined interfaces



Address Spaces

- ✳ All memory a process can address
- ✳ One process, one address space
- ✳ OS performs the isolation
 - ✳ one process cannot access other's address space
 - ✳ same pointer address in different process point to different memory

Context switching

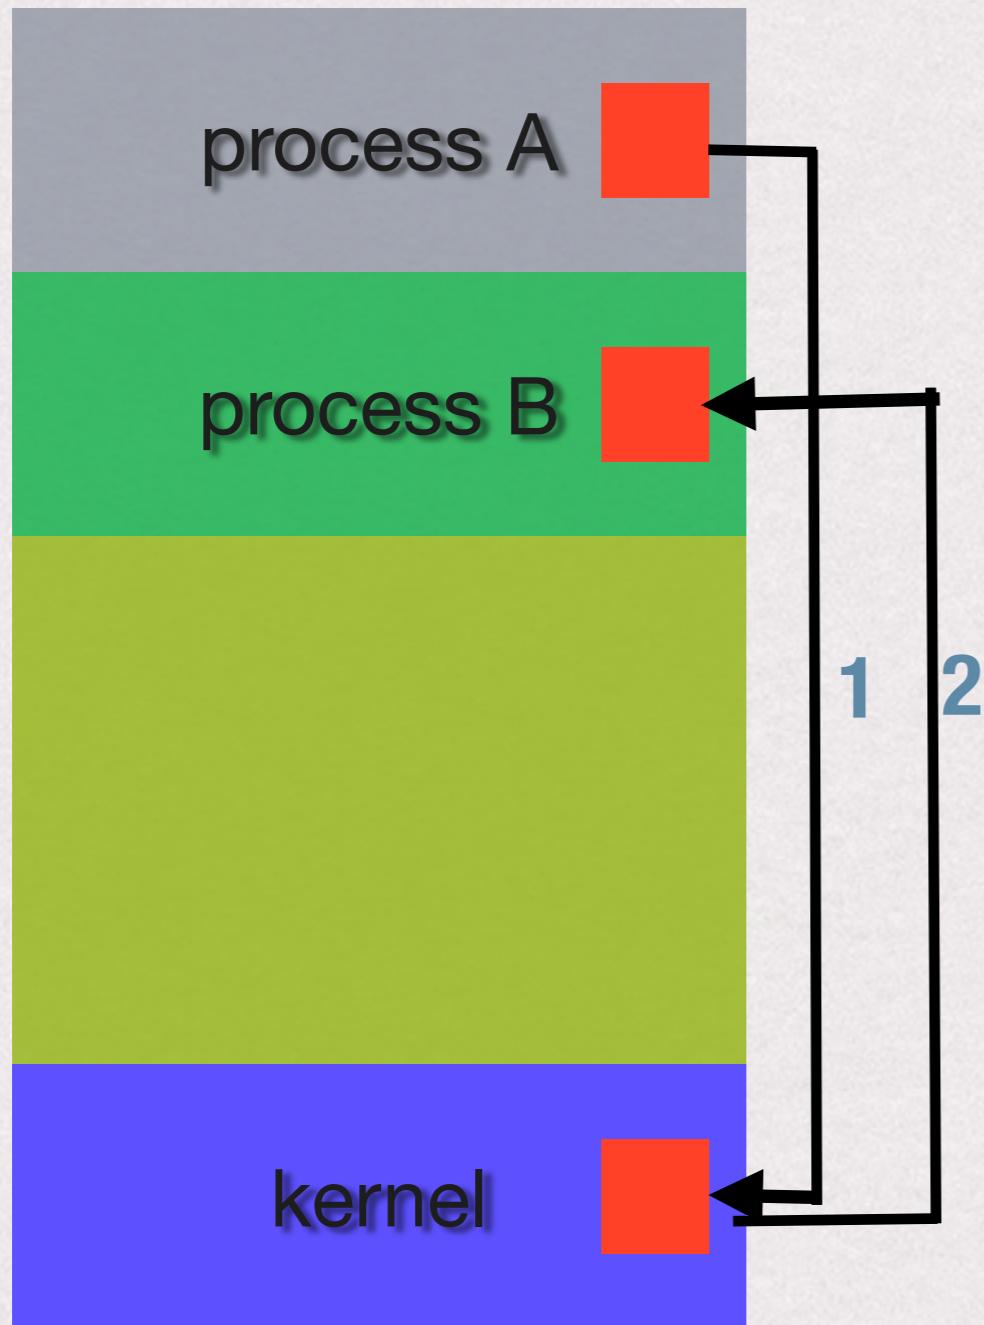
- ✳ Storing and restoring the state of a process such that execution can be resumed; crucial for multitasking
 - ✳ processes can wait for I/O or some other synchronization operation: other processes can run!
 - ✳ need to save and load register and memory maps, update tables, etc.
- ✳ Performance
 - ✳ Expensive
 - ✳ Needs to be limited

Inter-process Communication Example

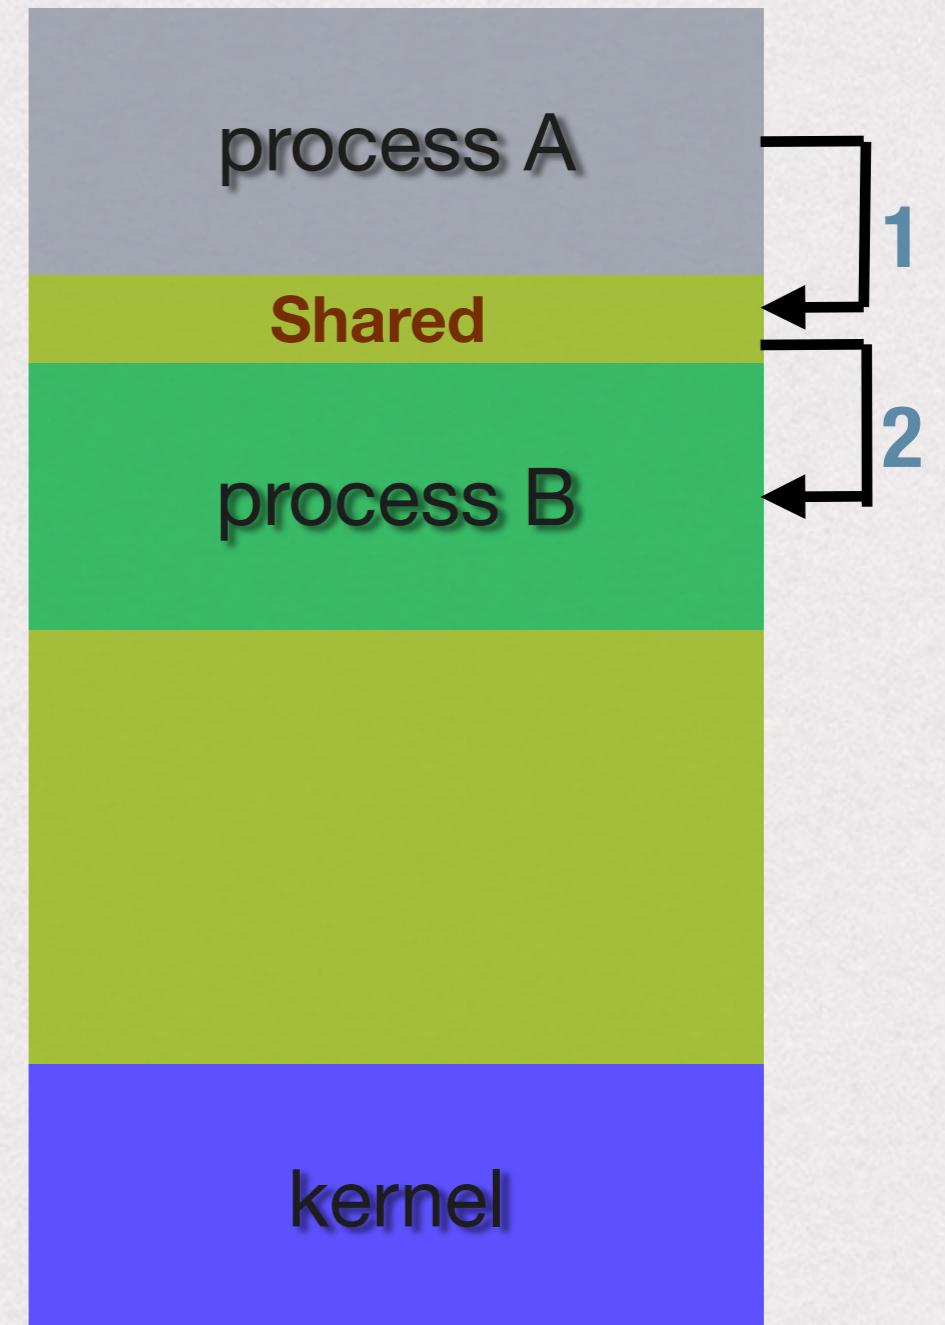
- ✳ Multiple processes are part of the same program
 - ✳ they need to coordinate
- ✳ Examples
 - ✳ # grep “DS is awesome” files | sort
 - ✳ servers like Apache spawn child processes to handle requests
 - ✳ producer-consumer - like communication

Fundamental Architectures

Message Passing



Shared Memory



Message Passing vs. Shared Memory

- ✳ Message passing
 - ✳ Good: all sharing is explicit
 - ✳ Bad: large overhead (copying data, context switching)
- ✳ Shared memory
 - ✳ Good: less overhead
 - ✳ Bad: no control about what changes on the shared space

Talk is cheap.

Show me the code. MPI

```
int argc;  
char *argv[];  
{  
    int rank, size;  
  
    MPI_Init (&argc, &argv); /* starts MPI */  
    MPI_Comm_rank (MPI_COMM_WORLD, &rank); /* get current process id */  
    MPI_Comm_size (MPI_COMM_WORLD, &size); /* get number of processes */  
    printf( "Hello world from process %d of %d\n", rank, size );  
    MPI_Finalize();  
    return 0;  
}
```

Talk is cheap.

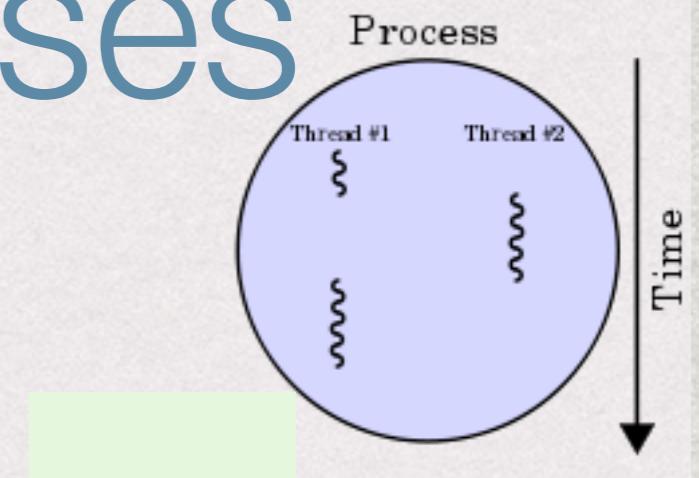
Show me the code. OpenMP

```
int main (int argc, char *argv[]) {  
    int nthreads, tid;  
  
    /* Fork a team of threads giving them their own copies of variables */  
    #pragma omp parallel private(nthreads, tid)  
  
    {  
        /* Obtain thread number */  
  
        tid = omp_get_thread_num();  
  
        printf("Hello World from thread = %d\n", tid);  
  
        /* Only master thread does this */  
  
        if (tid == 0)  
        {  
            nthreads = omp_get_num_threads();  
  
            printf("Number of threads = %d\n", nthreads);  
  
        }  
    } /* All threads join master thread and disband */
```

Threads

- ✳ Separate streams of execution that share one address space
- ✳ Thread state (not shared)
 - ✳ Program counter
 - ✳ Other registers
 - ✳ Stack
- ✳ Conceptually similar to processes, but different – Often called “lightweight processes”
- ✳ Threads in memory: each thread must have its own separate stack

Threads vs processes



- ✳️ Threads

- ✳️ allow running code concurrently within a single process
- ✳️ Switching among threads is lightweight
- ✳️ Sharing data among threads requires no IPC

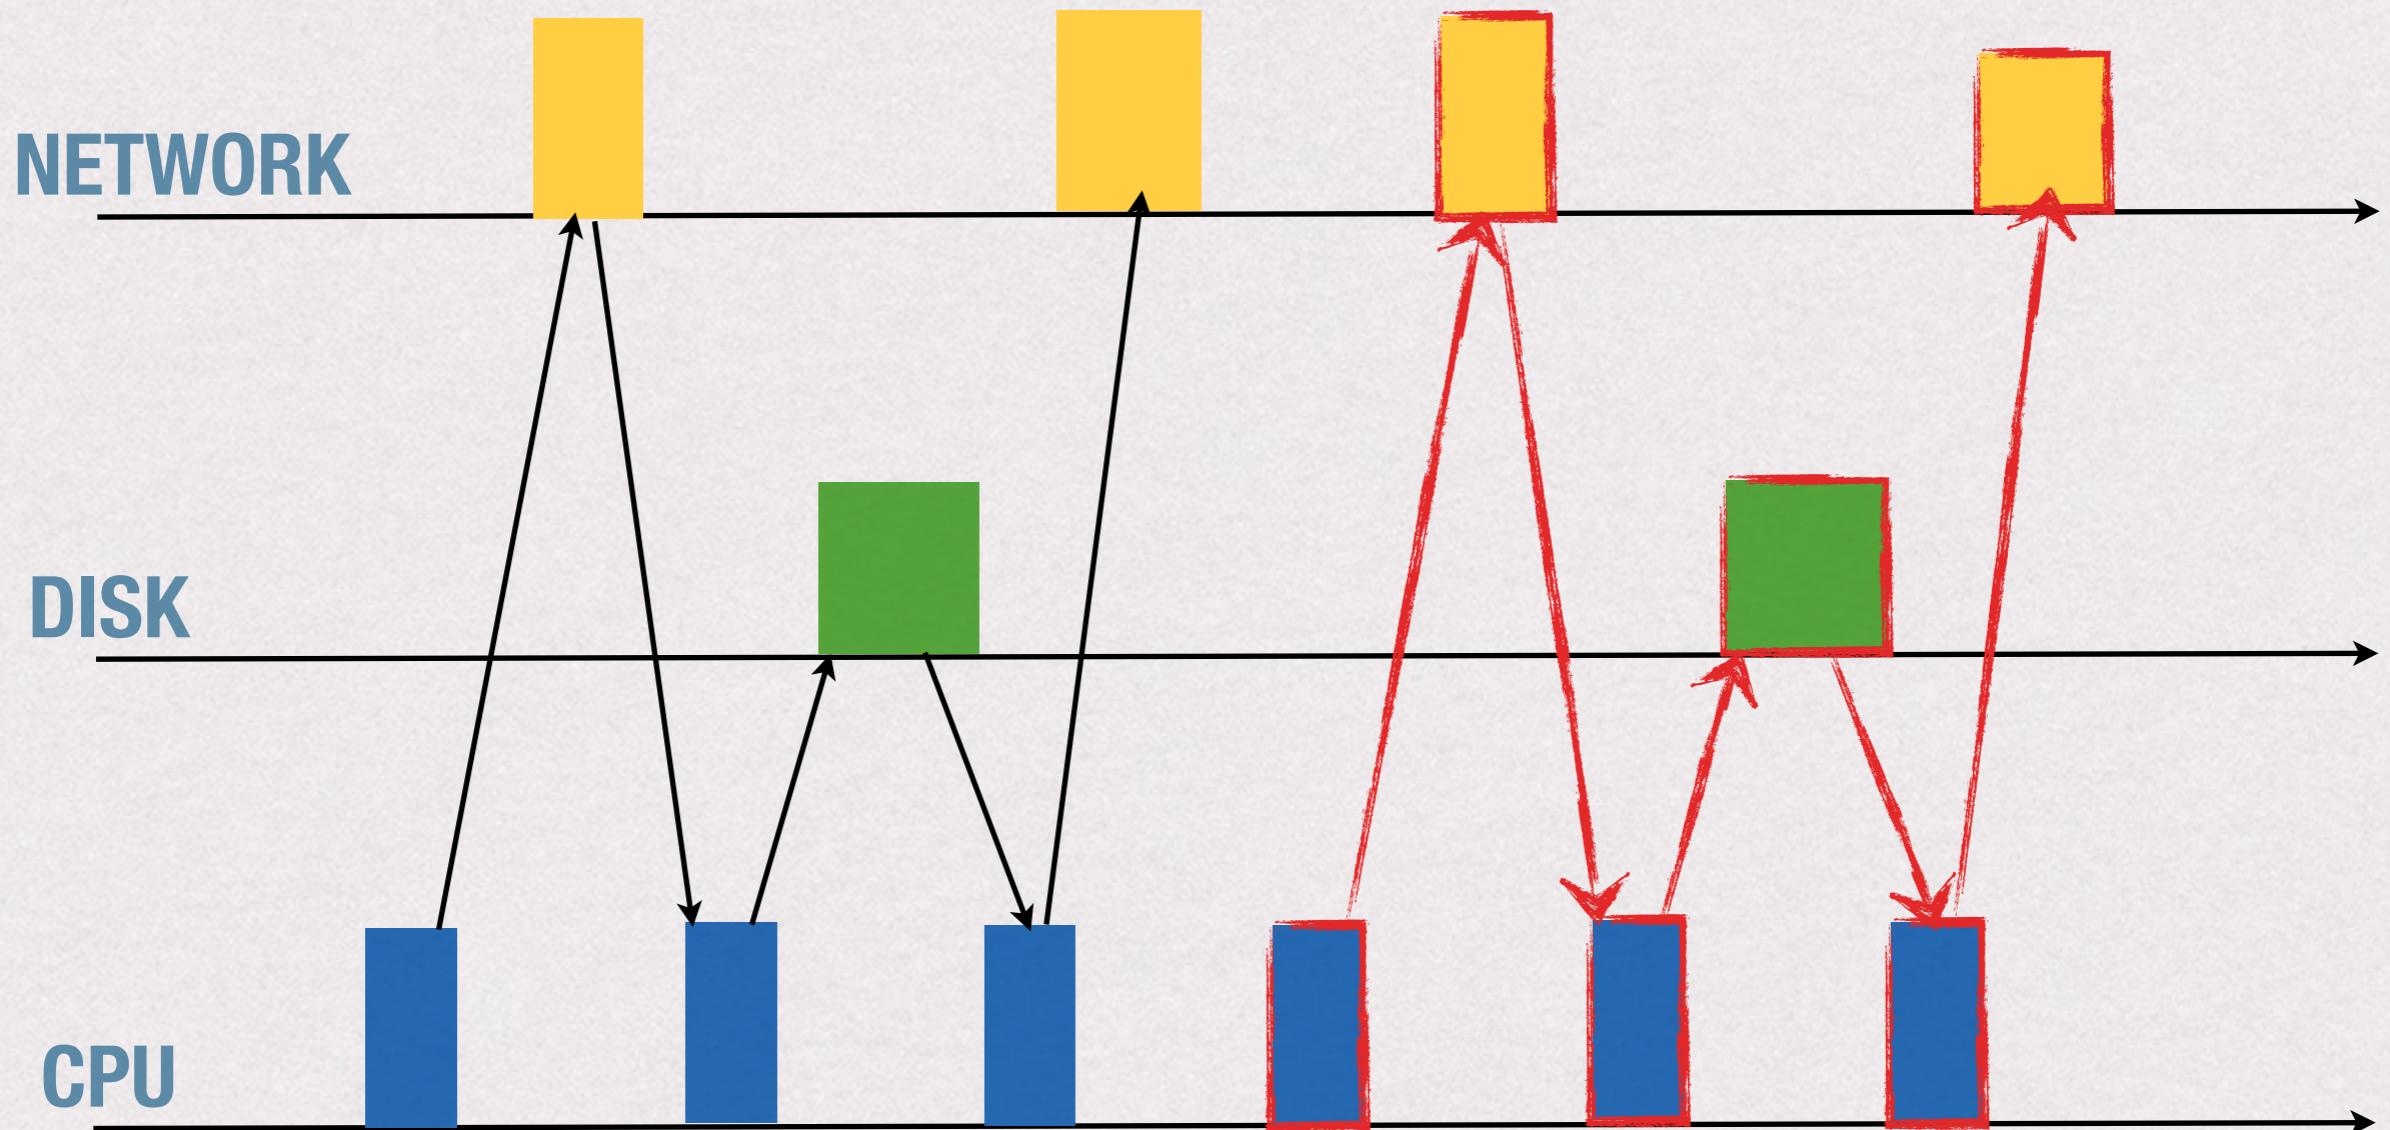
- ✳️ Processes

- ✳️ Fault isolation: One buggy process cannot crash others - see Chrome implementation

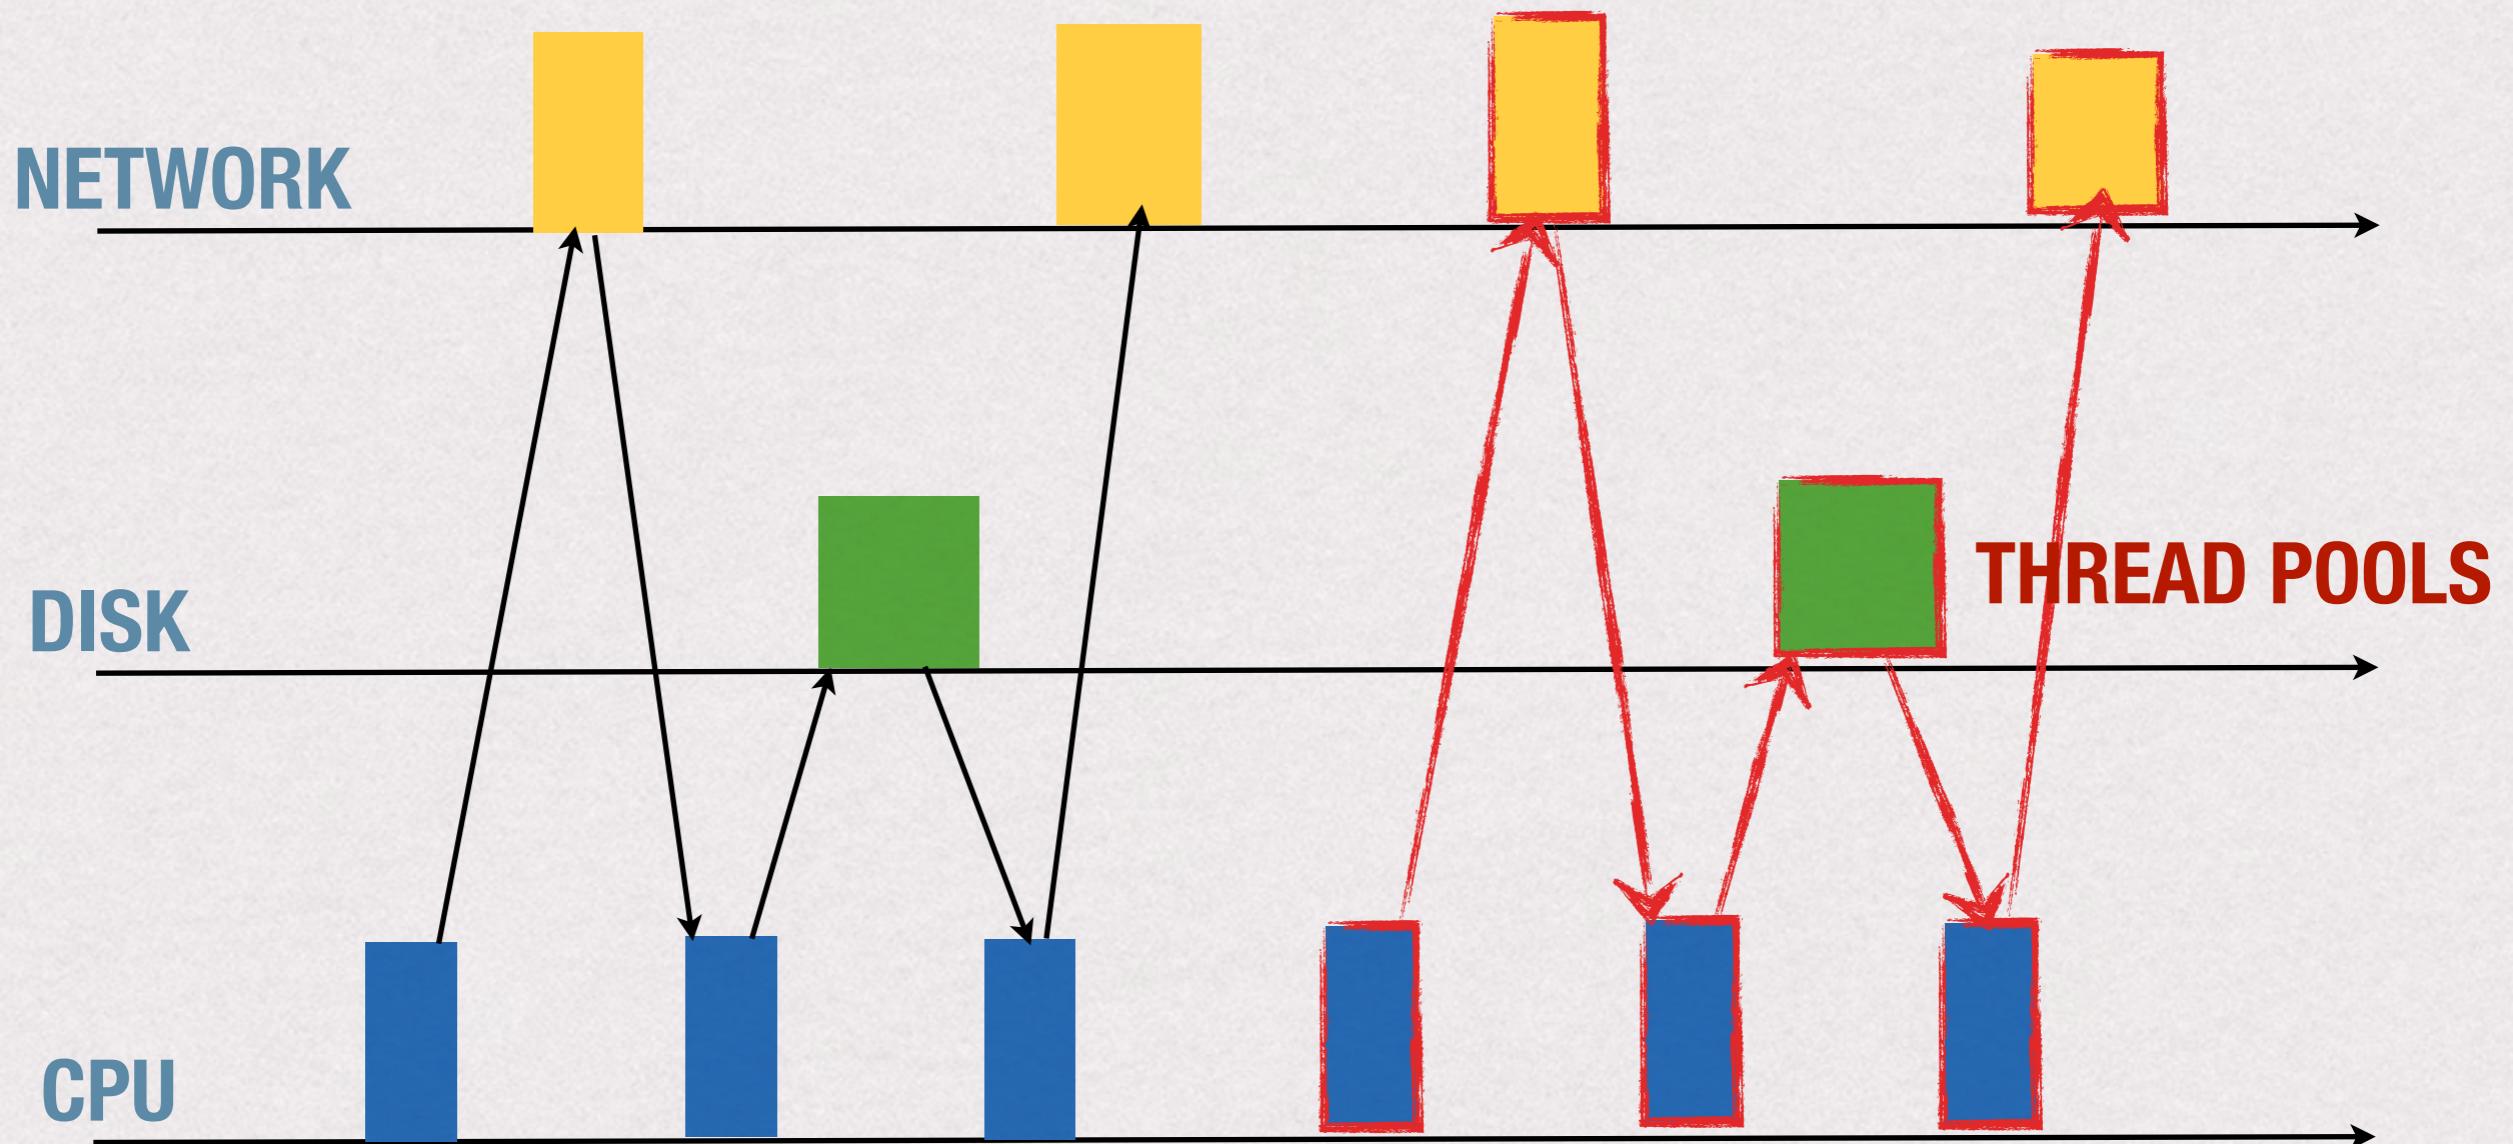
Multi-threaded programming

- ✳ Exploit multiple CPUs (multi-core) with little overhead
- ✳ Exploit I/O concurrency
 - ✳ Do some processing while waiting for disk, network, user
- ✳ Reduce latency of networked services
 - ✳ Servers serve multiple requests in parallel – Clients issue multiple requests in parallel
- ✳ Example:
 - ✳ threads in a web application: sending email after registration

Single-threaded server



Multi-threaded server



How?

```
while (true) {  
  
    try {  
  
        clientSocket = serverSocket.accept();  
  
        int i = 0;  
  
        for (i = 0; i < maxClientsCount; i++) {  
  
            if (threads[i] == null) {  
  
                (threads[i] = new clientThread(clientSocket, threads)).start();  
  
                break;  
            }  
        }  
  
        if (i == maxClientsCount) {  
  
            PrintStream os = new PrintStream(clientSocket.getOutputStream());  
  
            os.println("Server too busy. Try later.");  
  
            os.close();  
  
            clientSocket.close();  
        }  
  
    } catch (IOException e) { }  
}
```

```
public clientThread(Socket clientSocket, clientThread[] threads)
{
    this.clientSocket = clientSocket;
    this.threads = threads;
    maxClientsCount = threads.length;
}

public void run() {
    int maxClientsCount = this.maxClientsCount;
    clientThread[] threads = this.threads;
    //do magick!!!
}
```

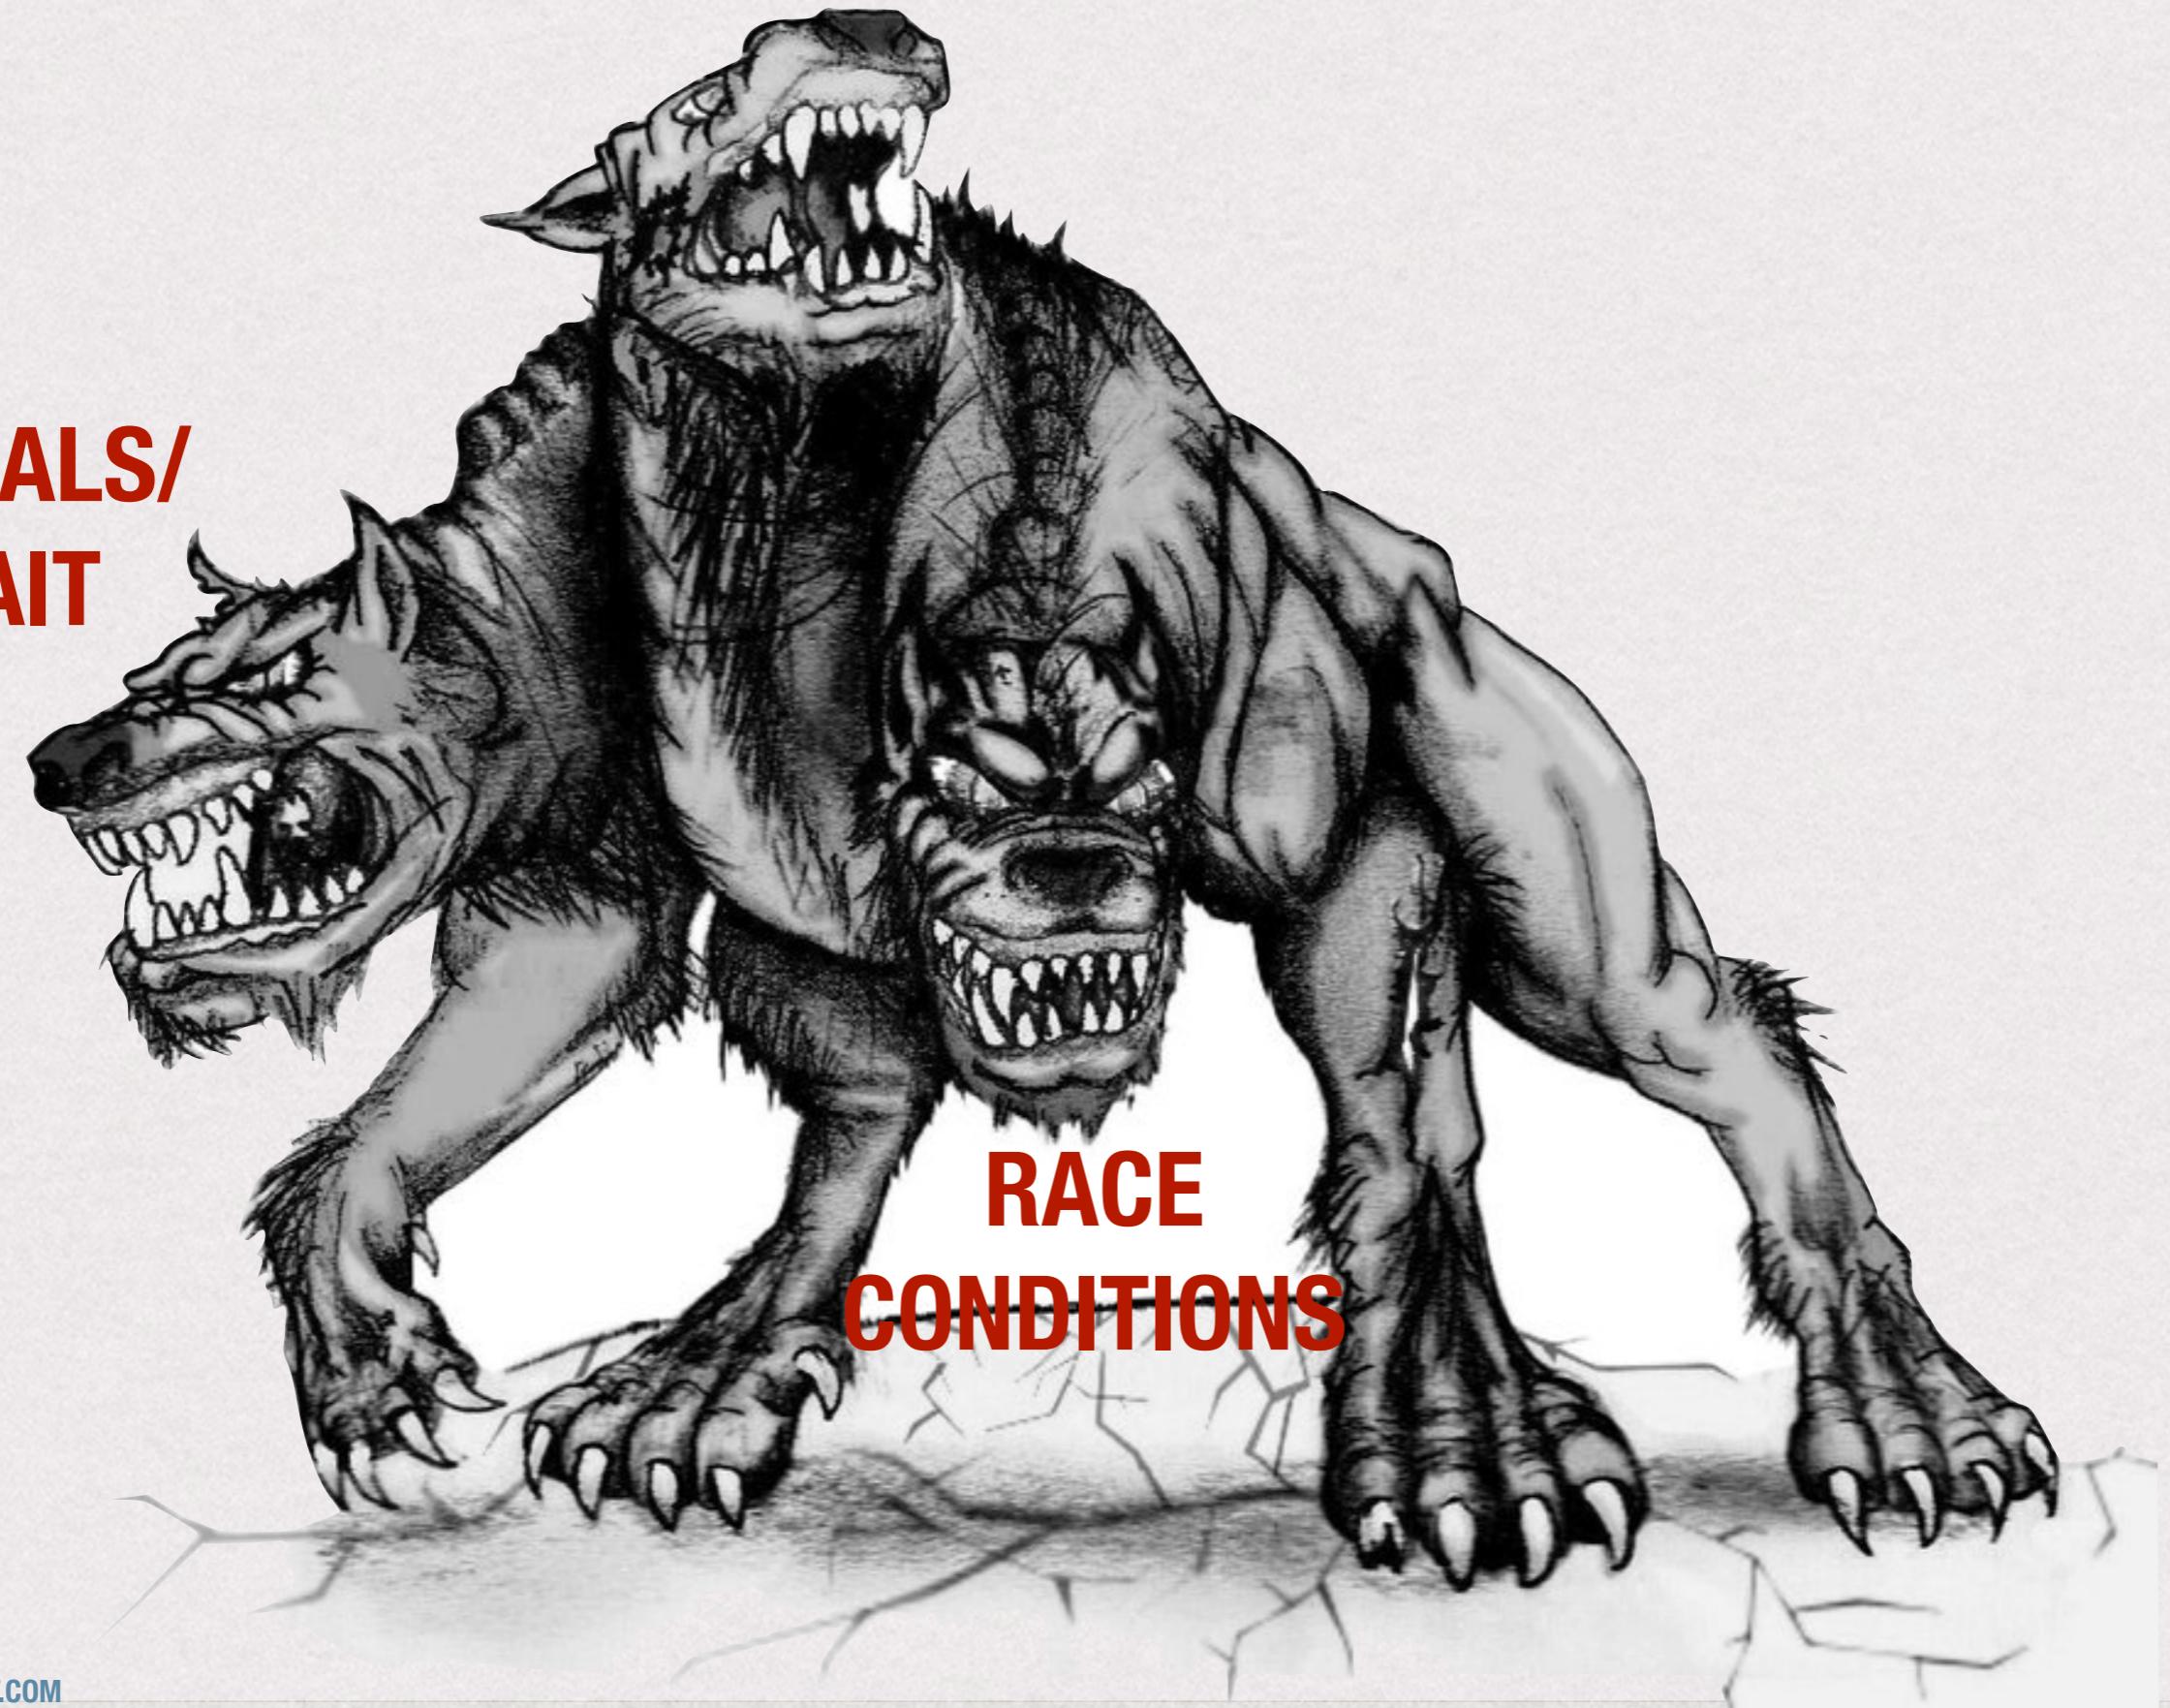
Thread synchronization

- * Memory is shared across all threads
- * Coordination is needed!

DEADLOCK

SIGNALS/
WAIT

RACE
CONDITIONS



Example

```
int balance = 1000;

void* withdraw(void *arg){

    int amount = (int) arg;

    if (balance >= amount) {

        balance -= amount;

        printf("ATM gives user $%d\n", amount); } }

int main() {

    pthread_t t1, t2;

    pthread_create(&t1, NULL, withdraw, (void*)800);

    pthread_create(&t2, NULL, withdraw, (void*)800);

    pthread_join(t1, NULL);

    pthread_join(t2, NULL);

    printf("All done: balance is $%d\n", balance);

    return 0; }
```

Discuss

- a) 200
- b) -600
- c) 1000

Result? Vote!

- a) 200
- b) -600
- c) 1000
- e) [https://www.mentimeter.com/app/presentation/
alh3h6g8fmeuek7yv5sxm9py7ewhy7bh/
9nwggkzz5hhk](https://www.mentimeter.com/app/presentation/alh3h6g8fmeuek7yv5sxm9py7ewhy7bh/9nwggkzz5hhk)

Schedule

```
matterhorn:resources gecko$ ./bank
ATM gives user $800
All done: balance is $200
```

THREAD 1

```
if (balance >= amount) {
    balance -= amount;
```

THREAD 2

```
if (balance >= amount) {
    balance -= amount;
```



Schedule

THREAD 1

```
if (balance >= amount) {  
    balance -= amount;  
  
    printf ...
```

```
matterhorn:resources gecko$ ./bank  
  
ATM gives user $800  
  
ATM gives user $800  
  
All done: balance is $-600
```

THREAD 2

```
if (balance >= amount) {  
  
    balance -= amount;  
  
    printf ...
```



Schedule

THREAD 1

```
if (balance >= amount) {  
  
    balance -= amount;  
  
    printf ...
```

matterhorn:resources gecko\$./bank

ATM gives user \$800

ATM gives user \$800

All done: balance is \$200

THREAD 2

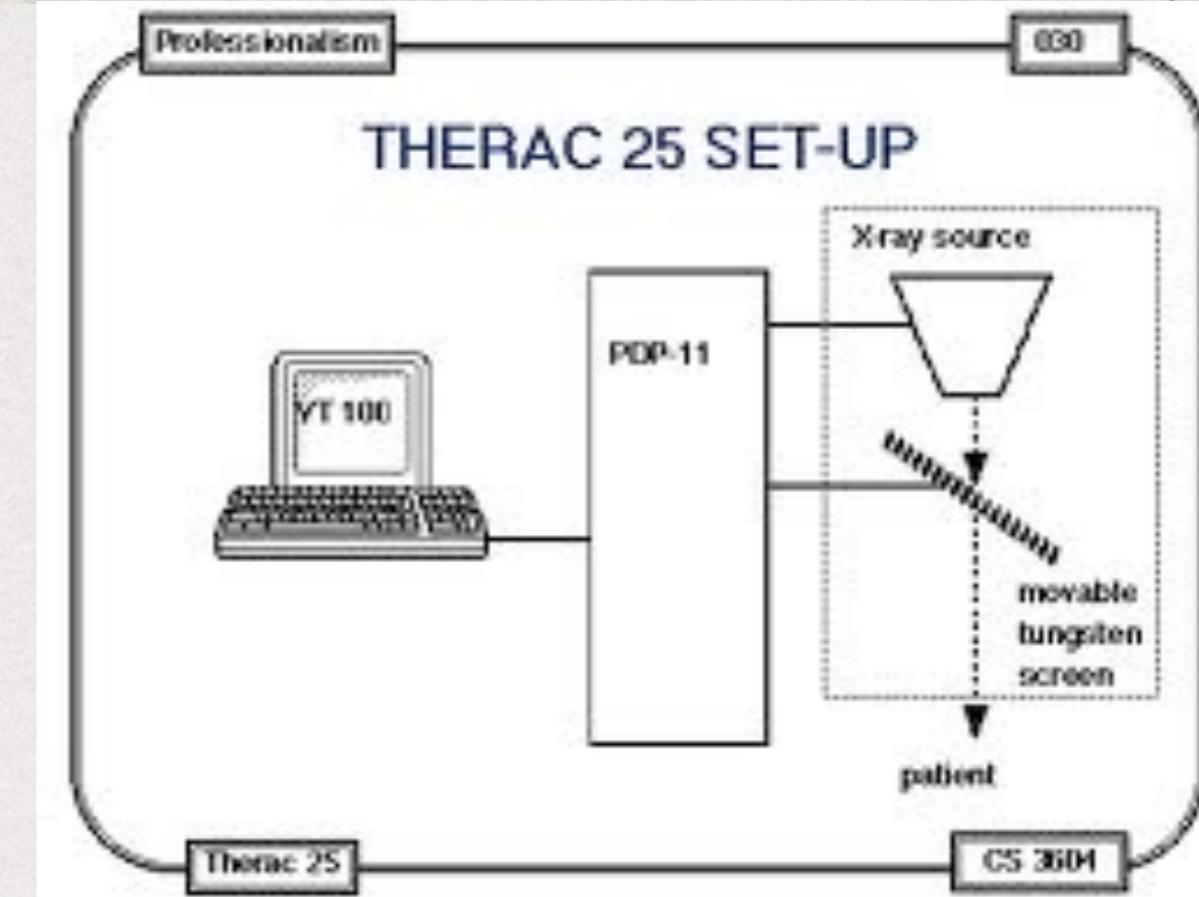
```
if (balance >= amount) {  
  
    balance -= amount;  
  
    printf ...
```



Race conditions

- ✳ Definition: a timing dependent error involving shared state
- ✳ Can be disastrous
 - ✳ Non-deterministic: don't know what the output will be, and it is likely to be different across runs
 - ✳ Hard to detect: too many possible schedules
 - ✳ Hard to debug: debugging changes timing so hides bugs ("heisenbug")
 - ✳ W. Heisenberg: the act of observing a system fundamentally changes its state

Therac 25



- * Radiation therapy machine
- * Race condition gave patients 100 times intended radiation dose

Northeast Blackout

- * August 14, 2003 - 50 million people, 8 states + Canada
- * blackout bug: <http://www.securityfocus.com/news/8412>
- * <http://www.securityfocus.com/news/8016>

I S A T G e o S t a r 4 5
2 3 : 1 5 E S T 1 4 A u g . 2 0 0 3

Discuss

a) How would you solve synchronisation problems?

Synchronization Mechanisms

- ✳ Used extensively in distributed systems
- ✳ Multiple mechanisms, each solving a different problem
 - ✳ Locks
 - ✳ Condition variables
 - ✳ Semaphores
 - ✳ Monitors
 - ✳ Barriers

Locks

- * Allow only one thread to pass through a critical section at any time
 - * lock: acquire lock exclusively; wait if not available
 - * unlock: allow other threads to have access to lock

Locks in Java

```
public class SynchronizedCounter {  
  
    private int c = 0;  
  
    public synchronized void increment() {  
  
        c++;  
  
    }  
  
    public synchronized void decrement() {  
  
        c--;  
  
    }  
  
    public synchronized int value() {  
  
        return c;  
  
    } }  
  
http://docs.oracle.com/javase/tutorial/essential/concurrency/  
syncmeth.html
```

Common pitfalls

- ✳ Wrong lock granularity
 - ✳ Too small: leads to races - why?
 - ✳ Too large: leads to bad performance - why?
- ✳ Deadlocks
 - ✳ Better bugs than race
- ✳ Starvation

What does this have to do with DS?

- ✳ Every server is multi-threaded
 - ✳ need to coordinate: similar to IPC, shared memory, locking, barriers, etc.
- ✳ When testing, we will use multi-threaded programs to start multiple clients