

BOINC server

# DISTRIBUTED SYSTEMS

SYSTEMS TALK | RPC (CONTINUED)

# Activity

- \* What is the most powerful computer that any of you has ever used?
  - \* number of cores
  - \* frequency
  - \* memory
  - \* hard disk
- \* Answer here: <https://www.mentimeter.com/app/presentation/alsnfrkirk4y9p3kynim7anx47fjqey5/va1pwcros3wz>

SHOULD I ASK?

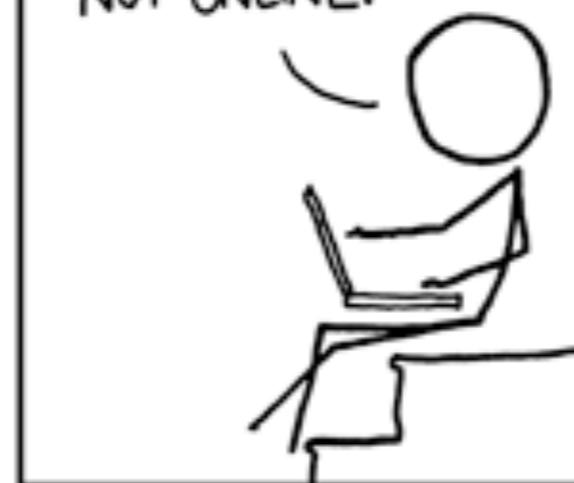
I'M LOCKED OUT,  
AND TRYING TO GET  
MY ROOMMATE TO  
LET ME IN.



FIRST I TRIED  
HER CELL PHONE,  
BUT IT'S OFF.



THEN I TRIED  
IRC, BUT SHE'S  
NOT ONLINE.



I COULDN'T FIND  
ANYTHING TO THROW  
AT HER WINDOW,



SO I SSH'D INTO THE MAC  
MINI IN THE LIVING ROOM  
AND GOT THE SPEECH SYNTH  
TO YELL TO HER FOR ME.

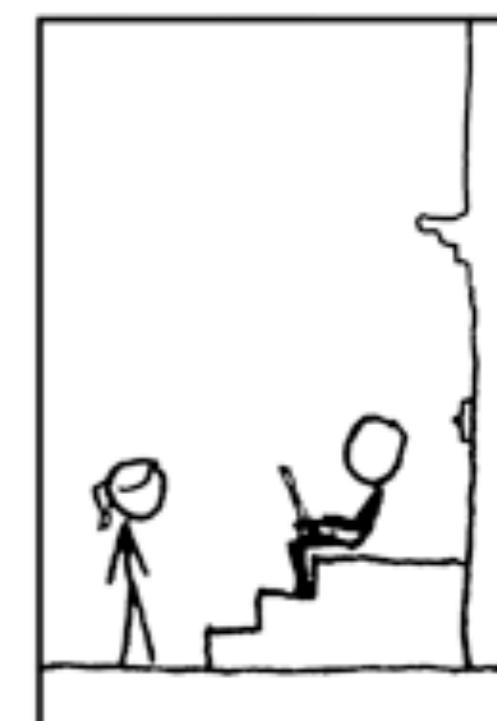


BUT I THINK I LEFT THE  
VOLUME WAY DOWN,  
SO I'M READING THE OS X  
DOCS TO LEARN TO SET THE  
VOLUME VIA COMMAND LINE.



AH.

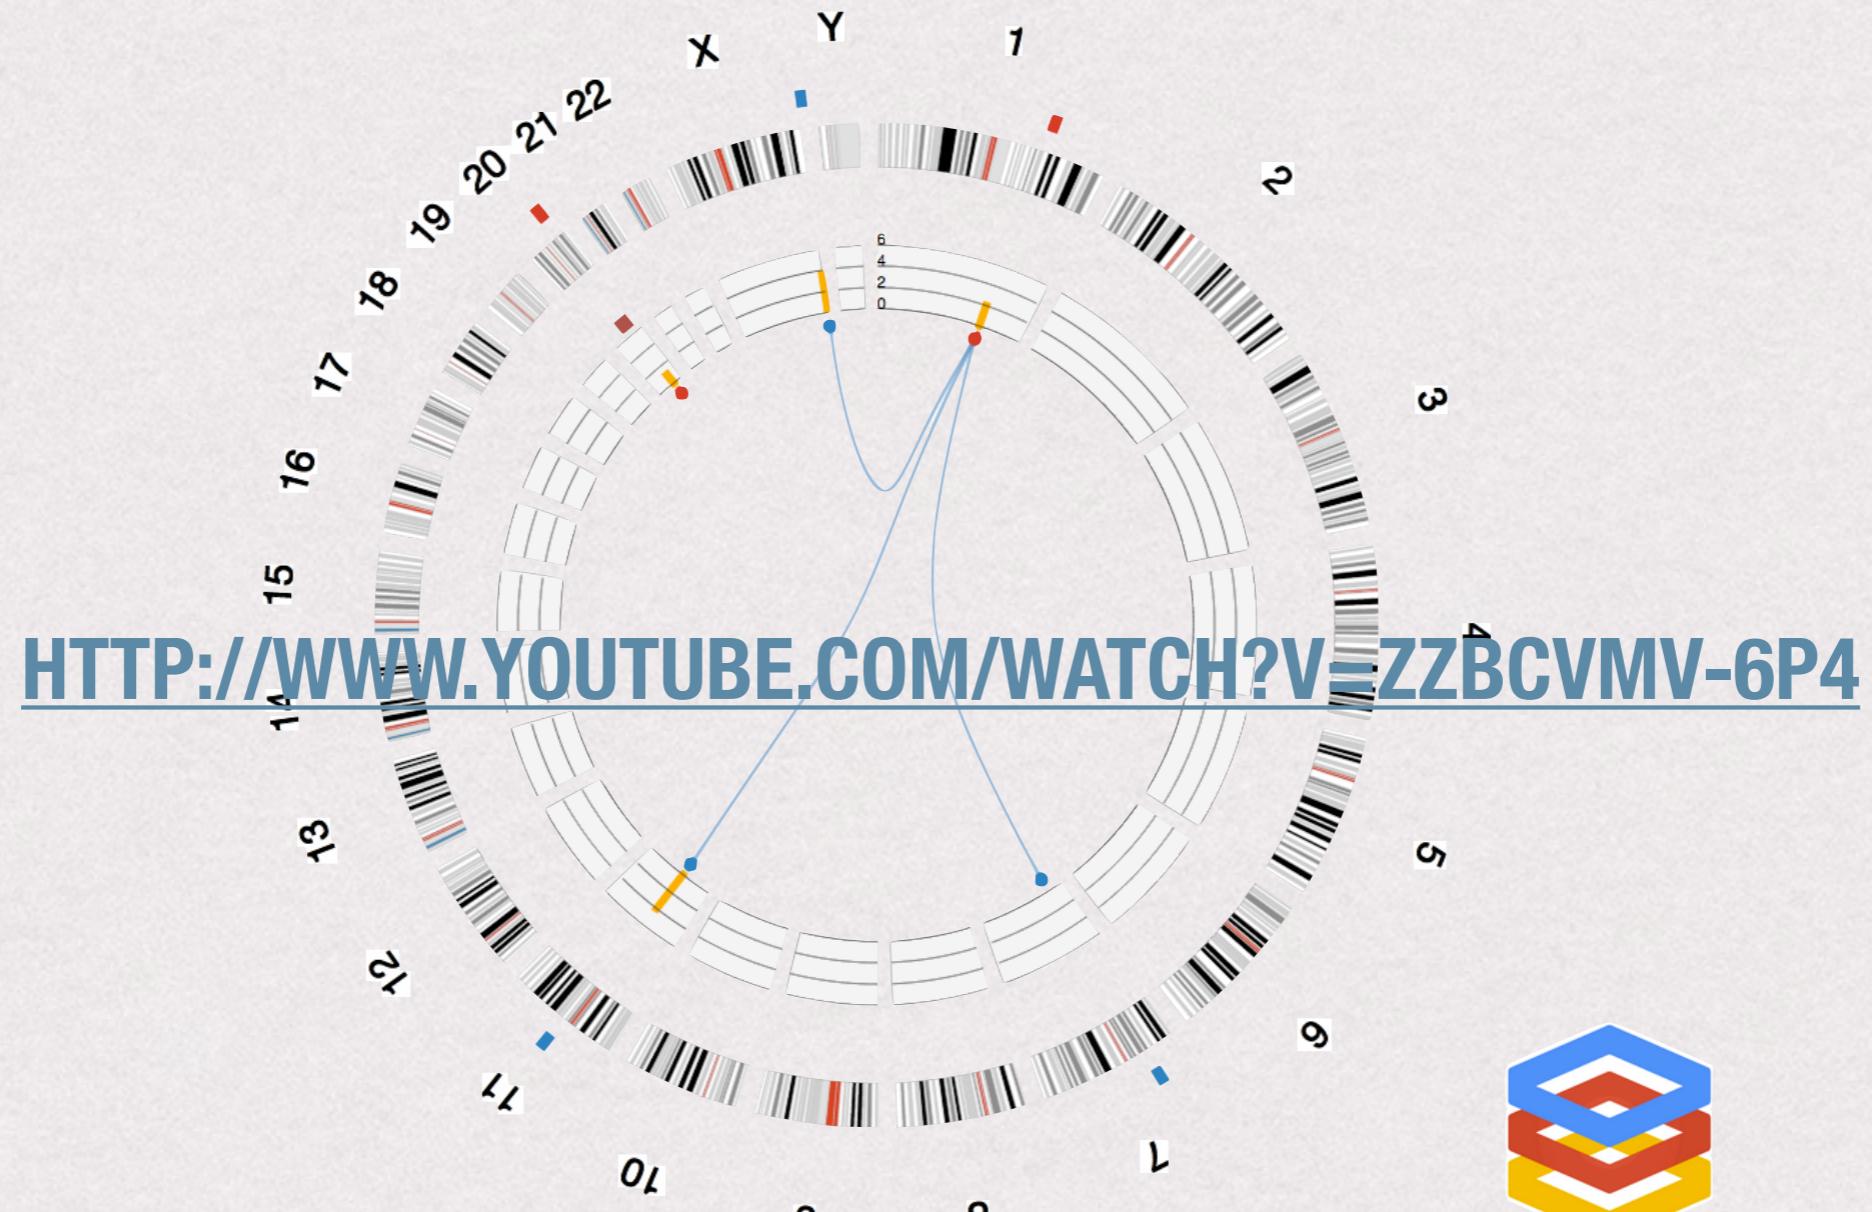
I  
I TAKE IT  
THE DOORBELL  
DOESN'T WORK?



## SYSTEMS TALK (XKCD 530)

# Google Compute Engine

Genome Explorer - powered by **Google** and **Institute for Systems Biology**

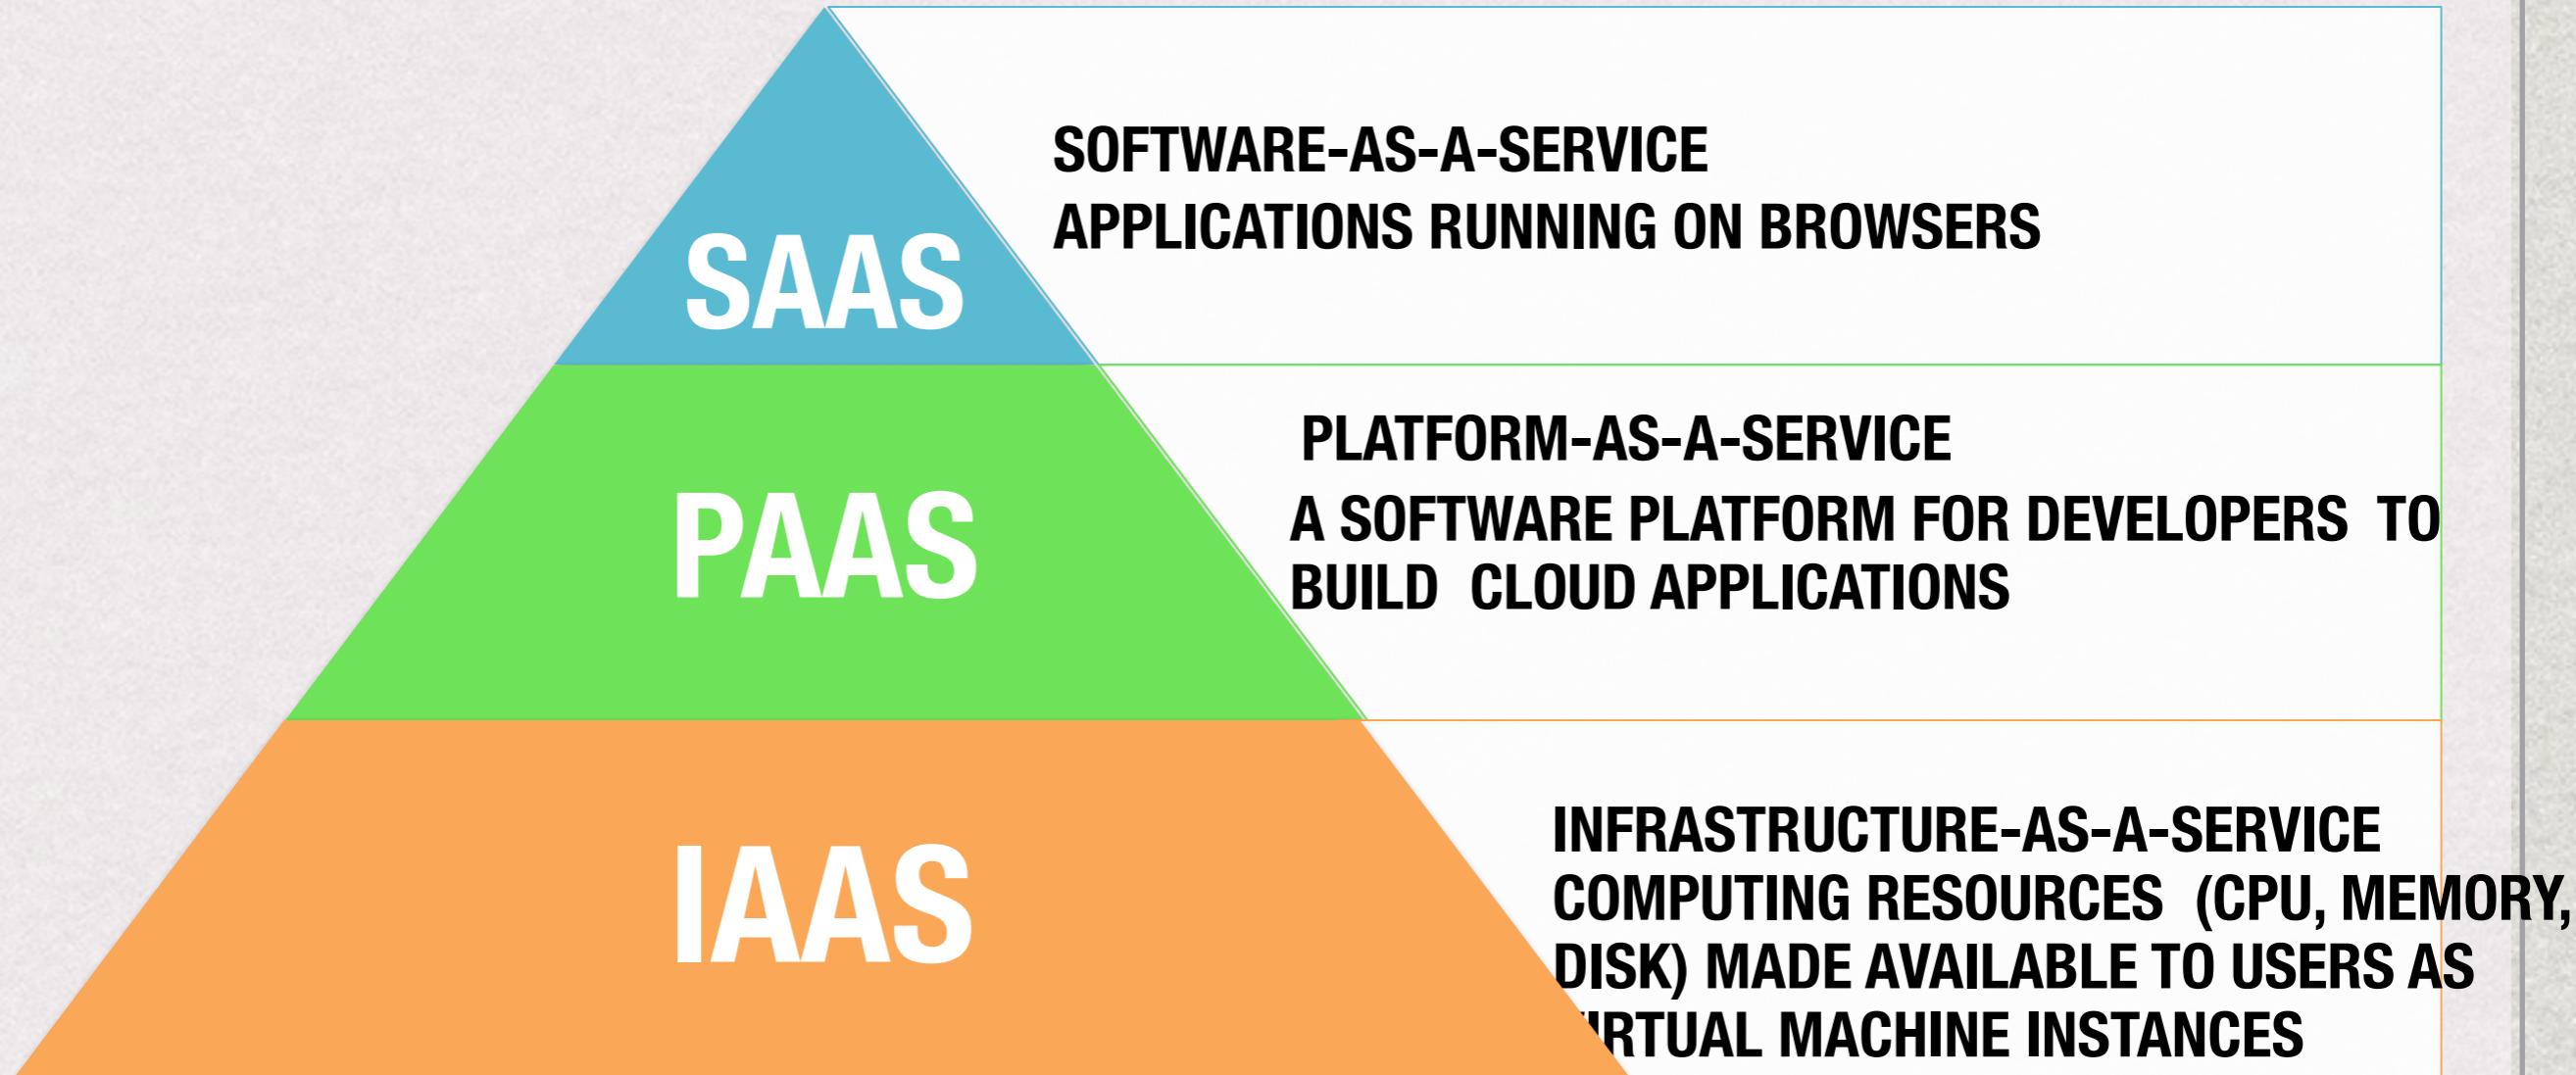


Google Compute Engine

# Key Cloud Characteristics

- ✳ On-demand self-service through a service portal
- ✳ Rapid elasticity – time to market / fast deployment
- ✳ Pay per use
- ✳ Ubiquitous access
- ✳ Location-independent resource pooling

# Cloud Service Models



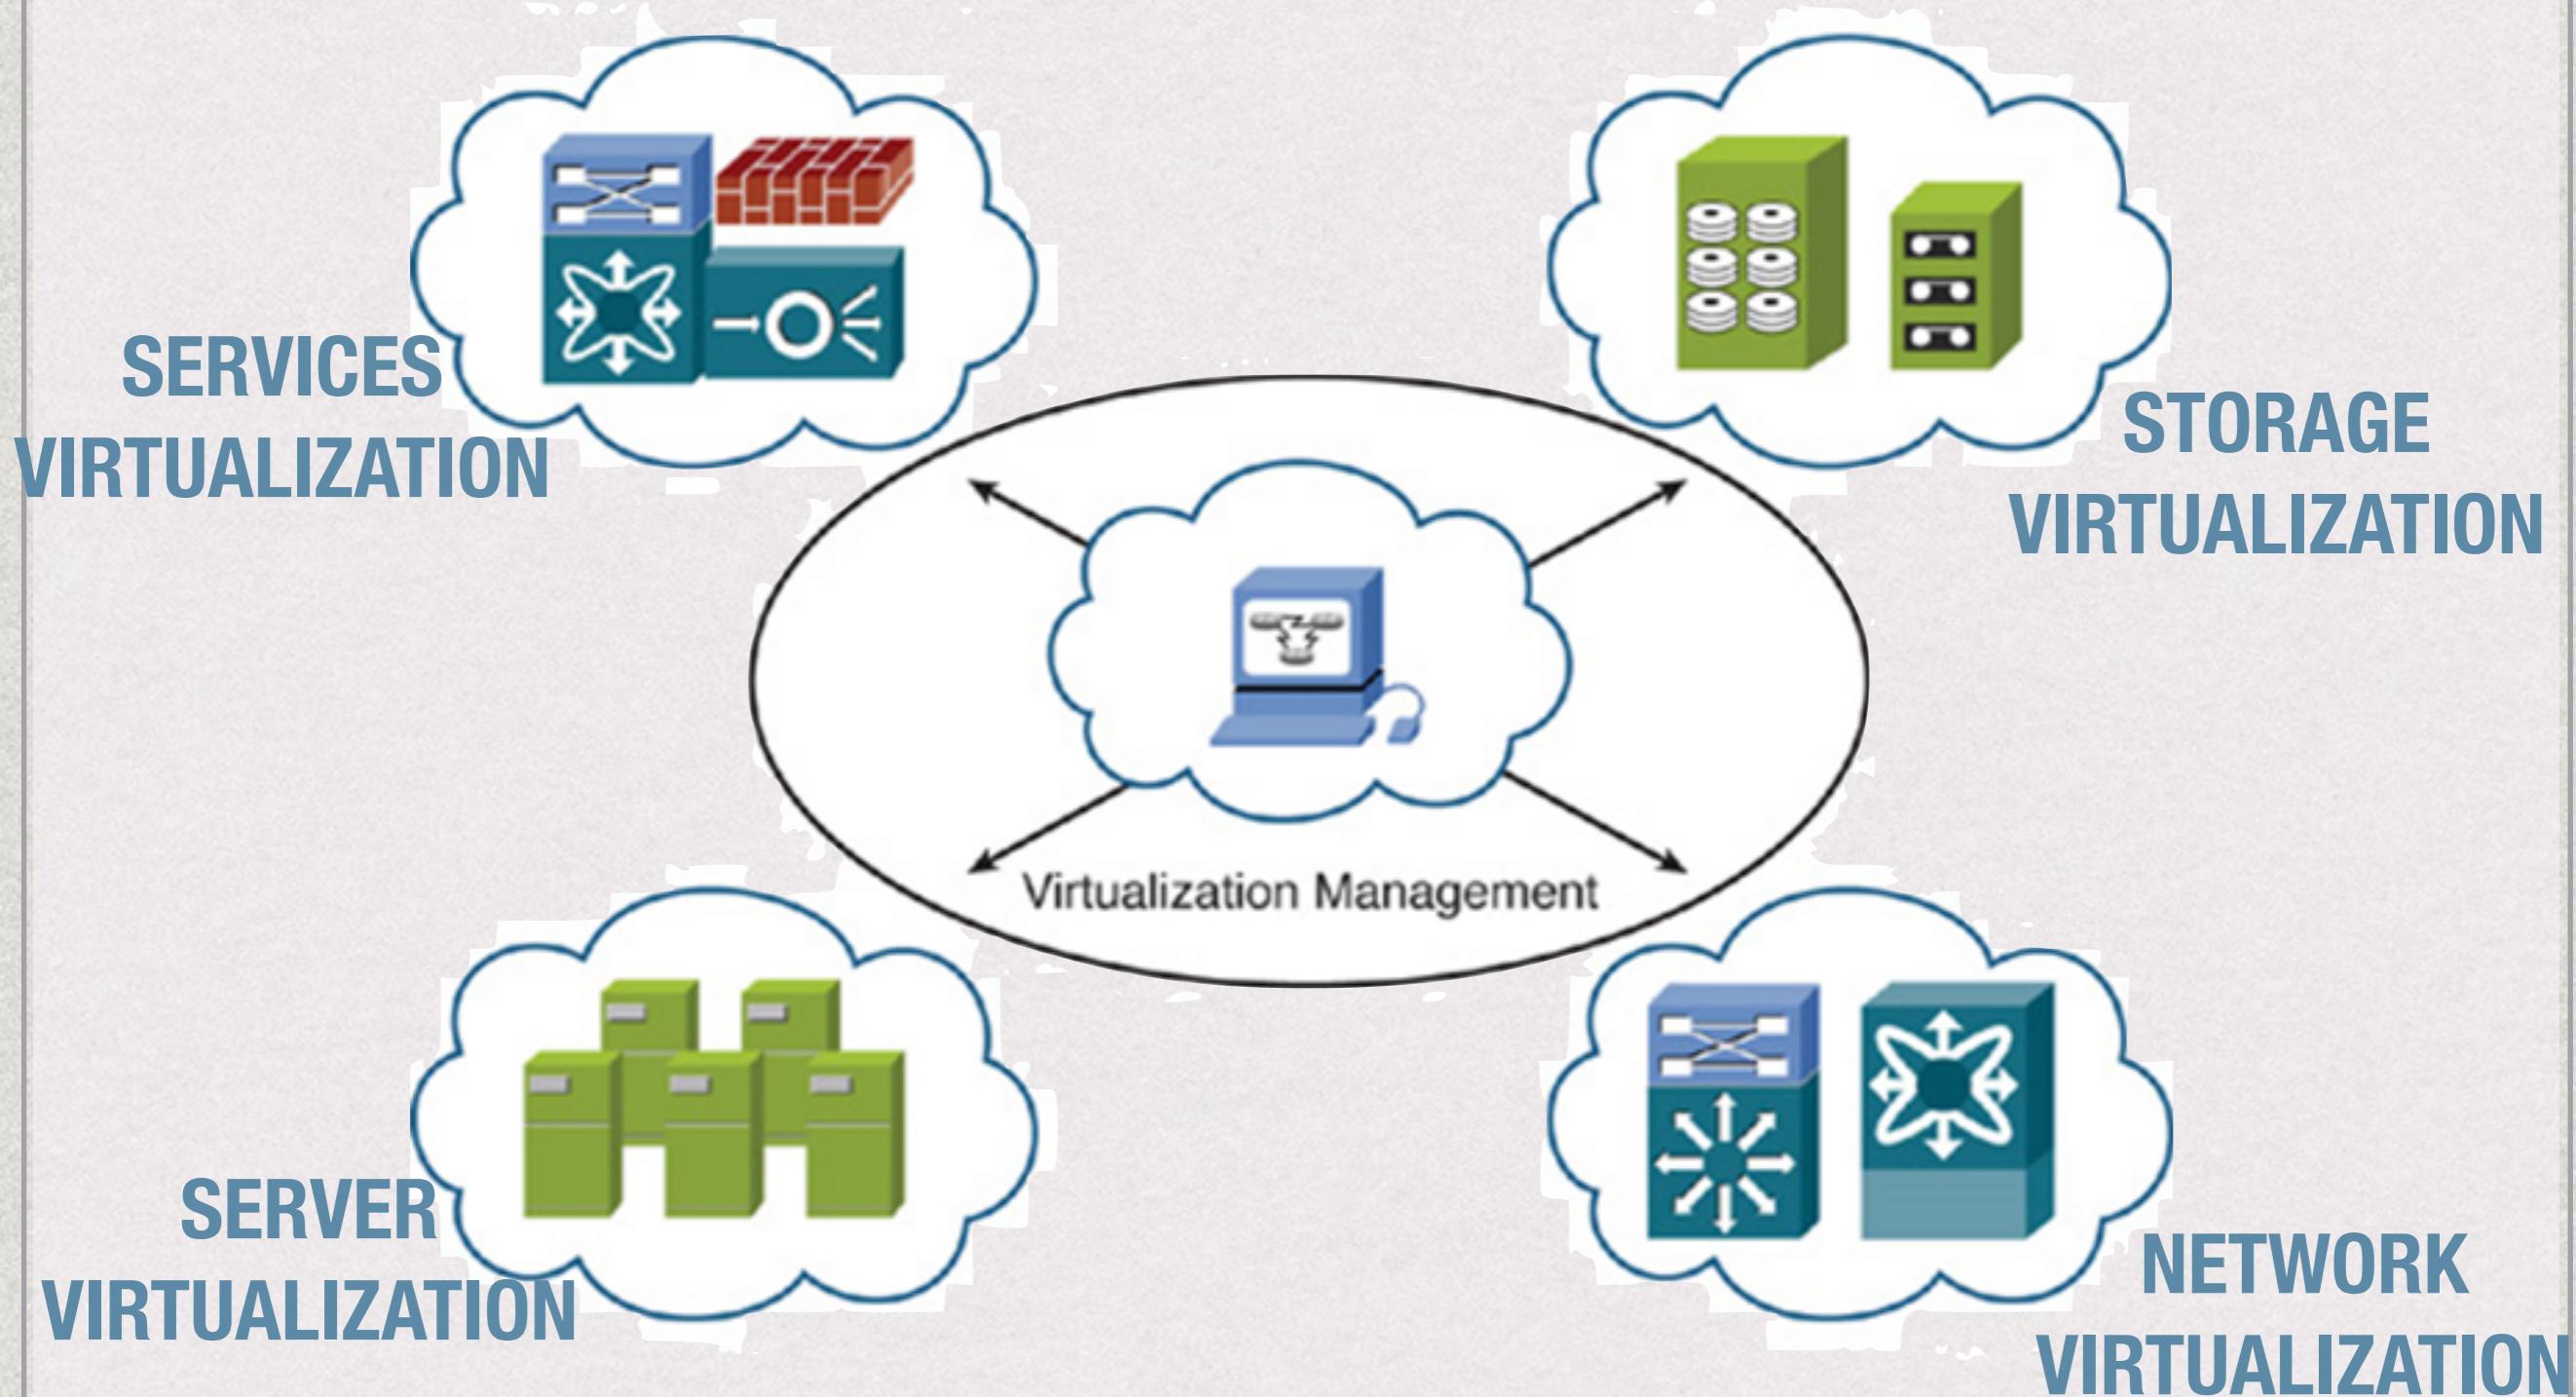
# Pros and Cons of Service Models

<b>Service Models</b>	<b>Pros</b>	<b>Cons</b>
Traditional	highest flexibility	time-to-market
IaaS	scalability, no hw procure	privacy
PaaS	DB, Frameworks, middleware ready	vendor lock in
SaaS	time-to-market	lowest flexibility

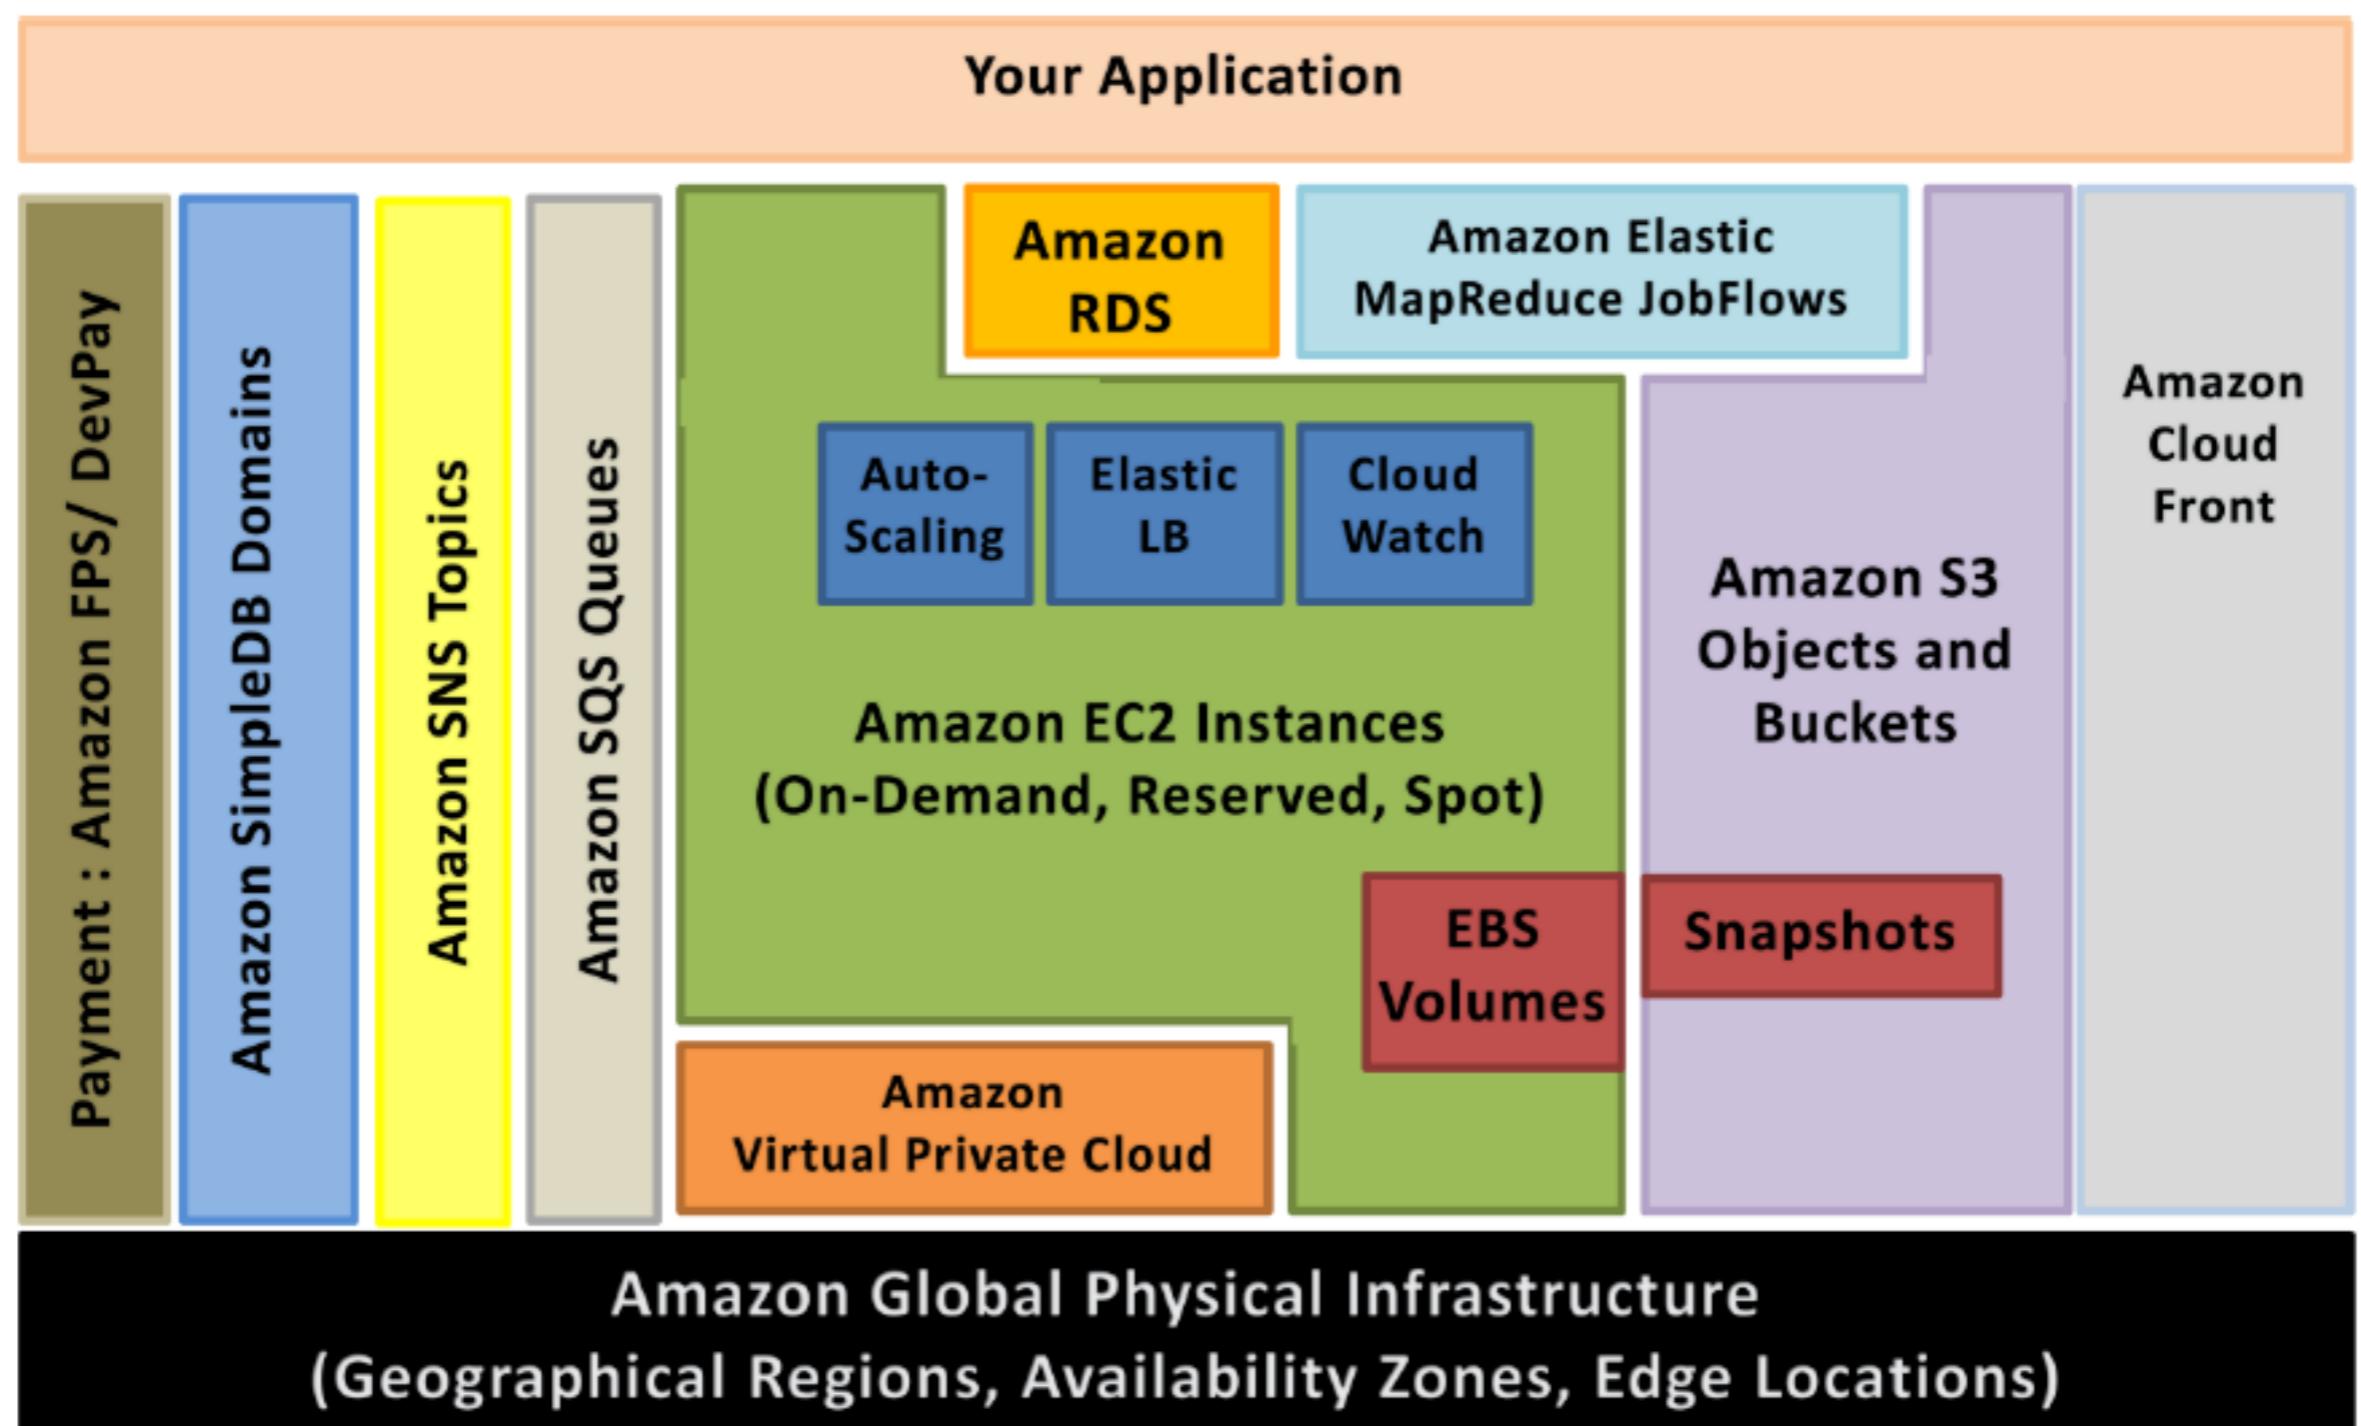
# Virtualization

- ✳ Key technology in cloud computing
- ✳ What does it mean?
- ✳ Creation of a virtual (rather than actual) version of something, such as an operating system, computing devices (server), storage devices or network devices

# Types of Virtualization

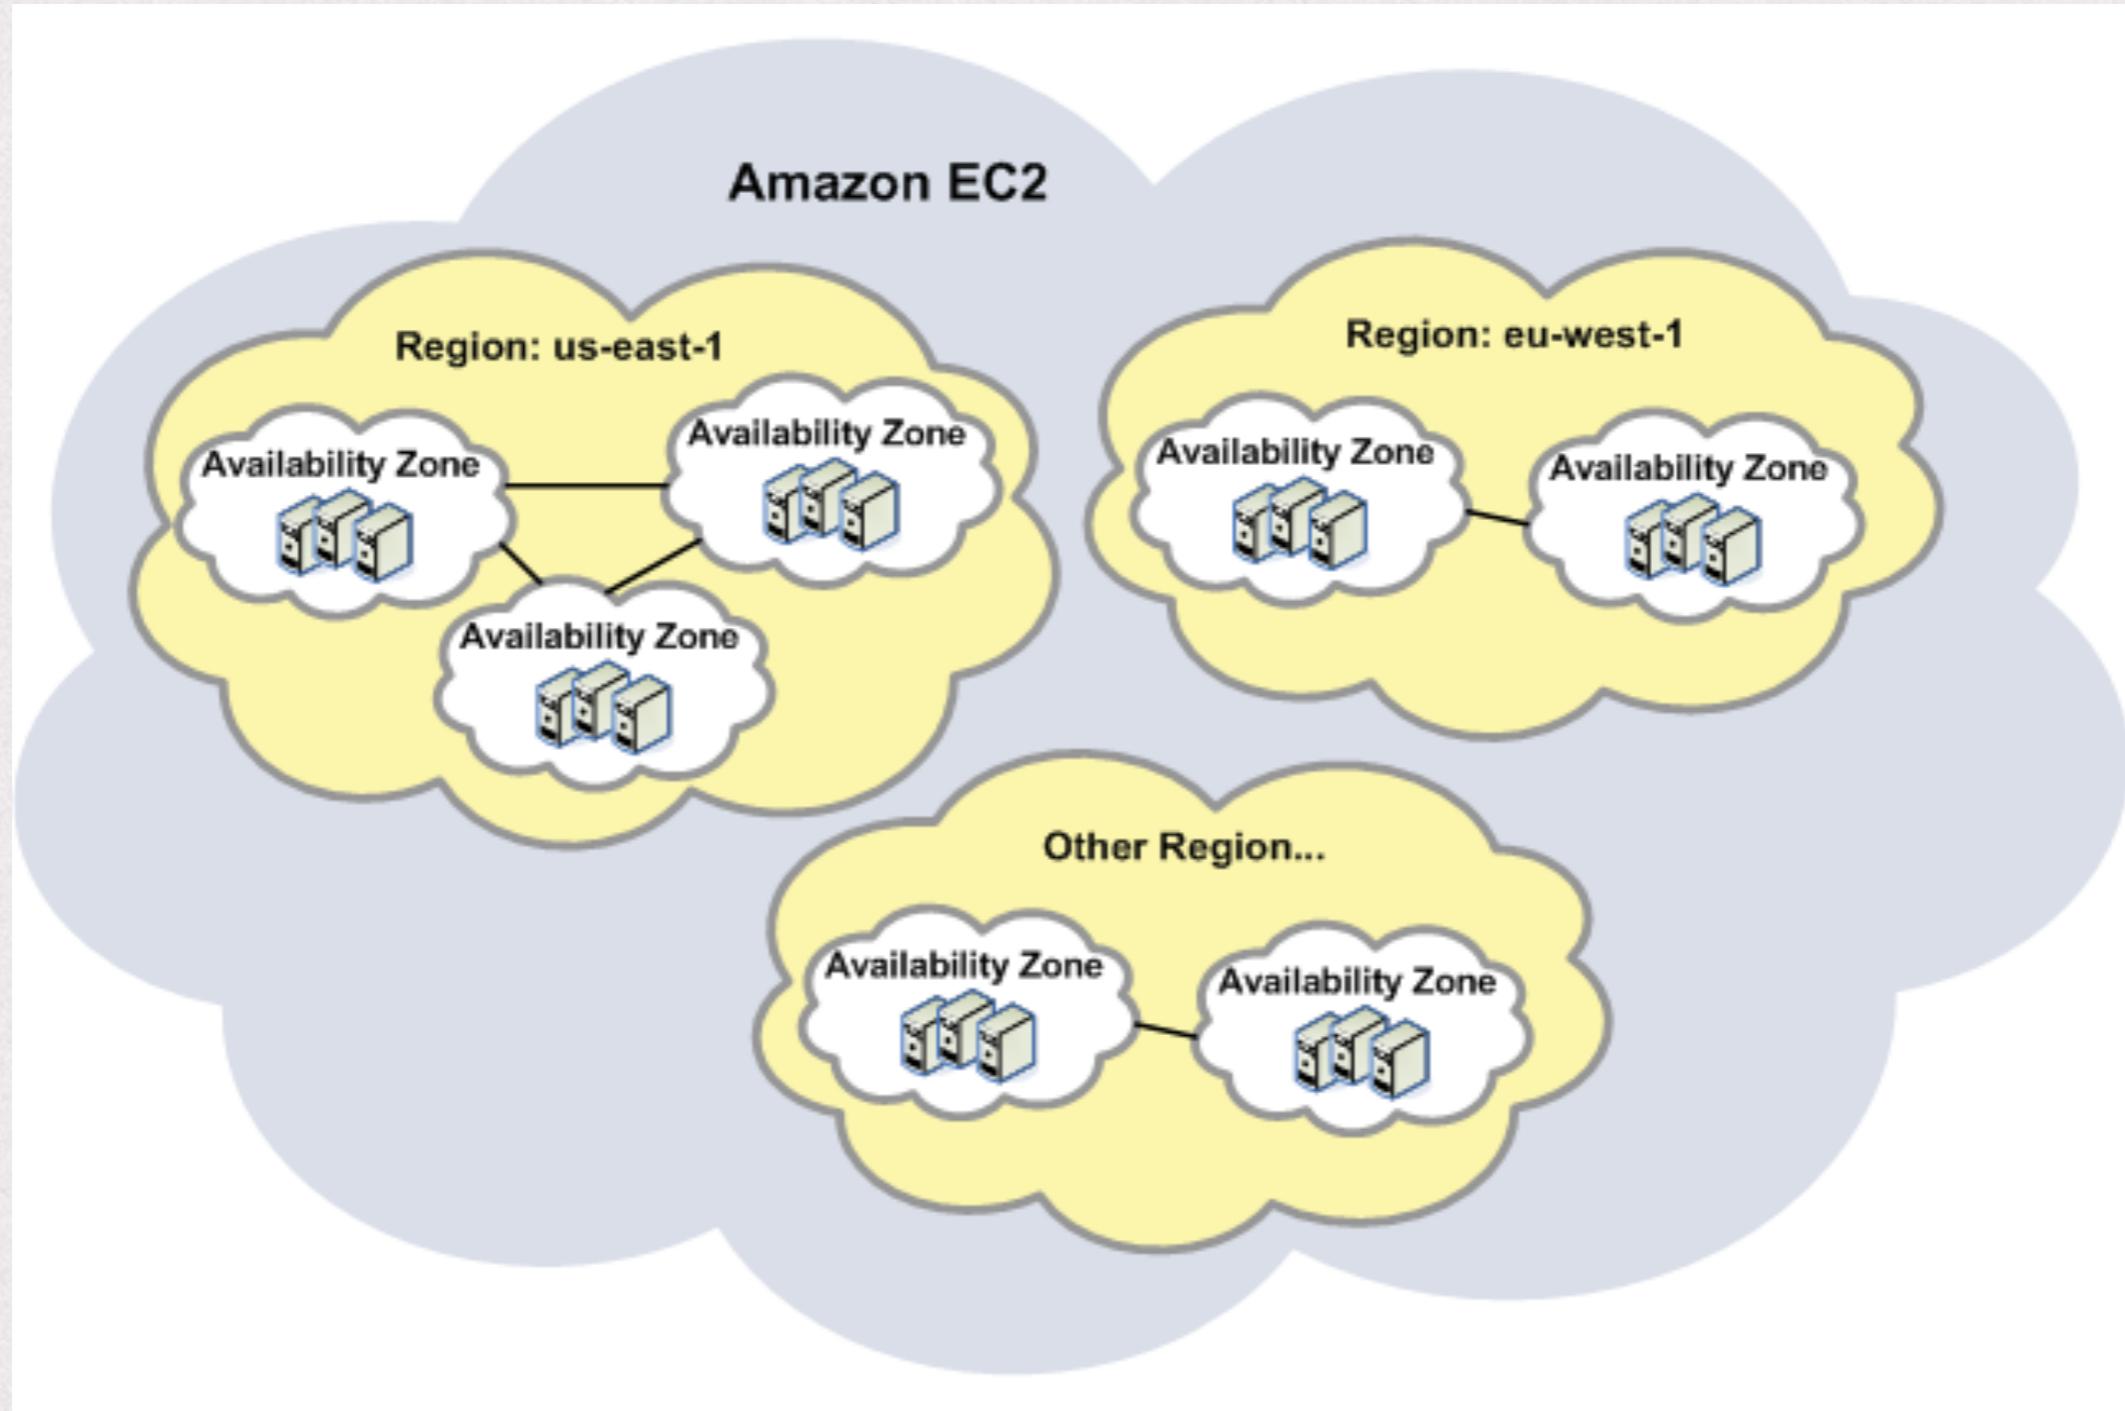


# AWS Ecosystem



J. VARIA, ARCHITECTING FOR THE CLOUD: BEST PRACTICES, AMAZON  
WEB SERVICES, MAY 2010.

# Regions and Availability Zones



# Regions and Availability Zones

- \* Amazon EC2 provides the ability to place instances in multiple locations
- \* Amazon EC2 locations are composed of Availability Zones and Regions
- \* Regions are distributed and located in separate geographic areas (US, EU, etc.)
  - \* You can design your application to be closer to specific customers or to meet legal or other requirements
- \* Availability Zones are distinct locations within a Region
  - \* You can protect your applications from the failure of a single location

# Amazon's Global Datacenters



- Regions
- Coming Soon

**AMAZON EC2 IS CURRENTLY AVAILABLE IN EIGHT REGIONS: US EAST (NORTHERN VIRGINIA),  
US WEST (OREGON), US WEST (NORTHERN CALIFORNIA), EU (IRELAND), ASIA PACIFIC (SINGAPORE),  
ASIA PACIFIC (TOKYO), SOUTH AMERICA (SAO PAULO), AND AWS GOVCLOUD**

# Remote Procedure Call

- \* Systems such as Java's RMI (Remote Method Invocation) are part of Remote Procedure Call (RPC) systems:
  - \* Java RMI is an RPC facility for objects, for calls of the form: O.m(params...)
  - \* The original RPC systems are formulated in terms of procedures, with calls of the form: P(params...)

# Group Activity

- \* Java's garbage collection can (simplistically) be implemented using reference counting
- \* How would an RMI distributed garbage collection algorithm work?
  - \* what different kinds of objects are there in RMI?
  - \* how would you deal with them?
  - \* where would your DGC be (on the client? on the server?)
  - \* Please do not use any internet-capable devices

# Group Activity

- \* When the client creates (unmarshalls) a remote reference, it calls dirty() on the server-side DGC. After the client has finished with the remote reference, it calls the corresponding clean() method.
- \* A reference to a remote object is leased for a time by the client holding the reference.
- \* Lease starts when the dirty() call is received.
- \* The client must renew the leases by making additional dirty() calls on the remote references it holds before such leases expire.
- \* If the client does not renew the lease before it expires, the distributed garbage collector assumes that the remote object is no longer referenced by that client.

# Last week

- \* RPC
- \* Reducing remote operation latency
  - \* Latency hiding
  - \* Latency reduction

# Latency Hiding – Readings

- ✳ Futures and promises paper:
  - ✳ B. Liskov and L. Shrira. Promises: Linguistic support for efficient asynchronous procedure calls in distributed systems. In Proceedings of the ACM SIGPLAN '88 Conference on Programming Languages Design and Implementation. ACM, June 1988. 63

# Latency Hiding

- \* Two camps:
  - \* system level - multi-threading OS, language support for general threading
  - \* language-level - futures and promises; superior?

# Recall ...

- \* The ***latency reduction*** strategy:
  - \* Tries to reduce the average latency of remote calls by running multiple calls at the same time.
  - \* However, this is only possible if calls can be sent off before the results of all previous calls are available.

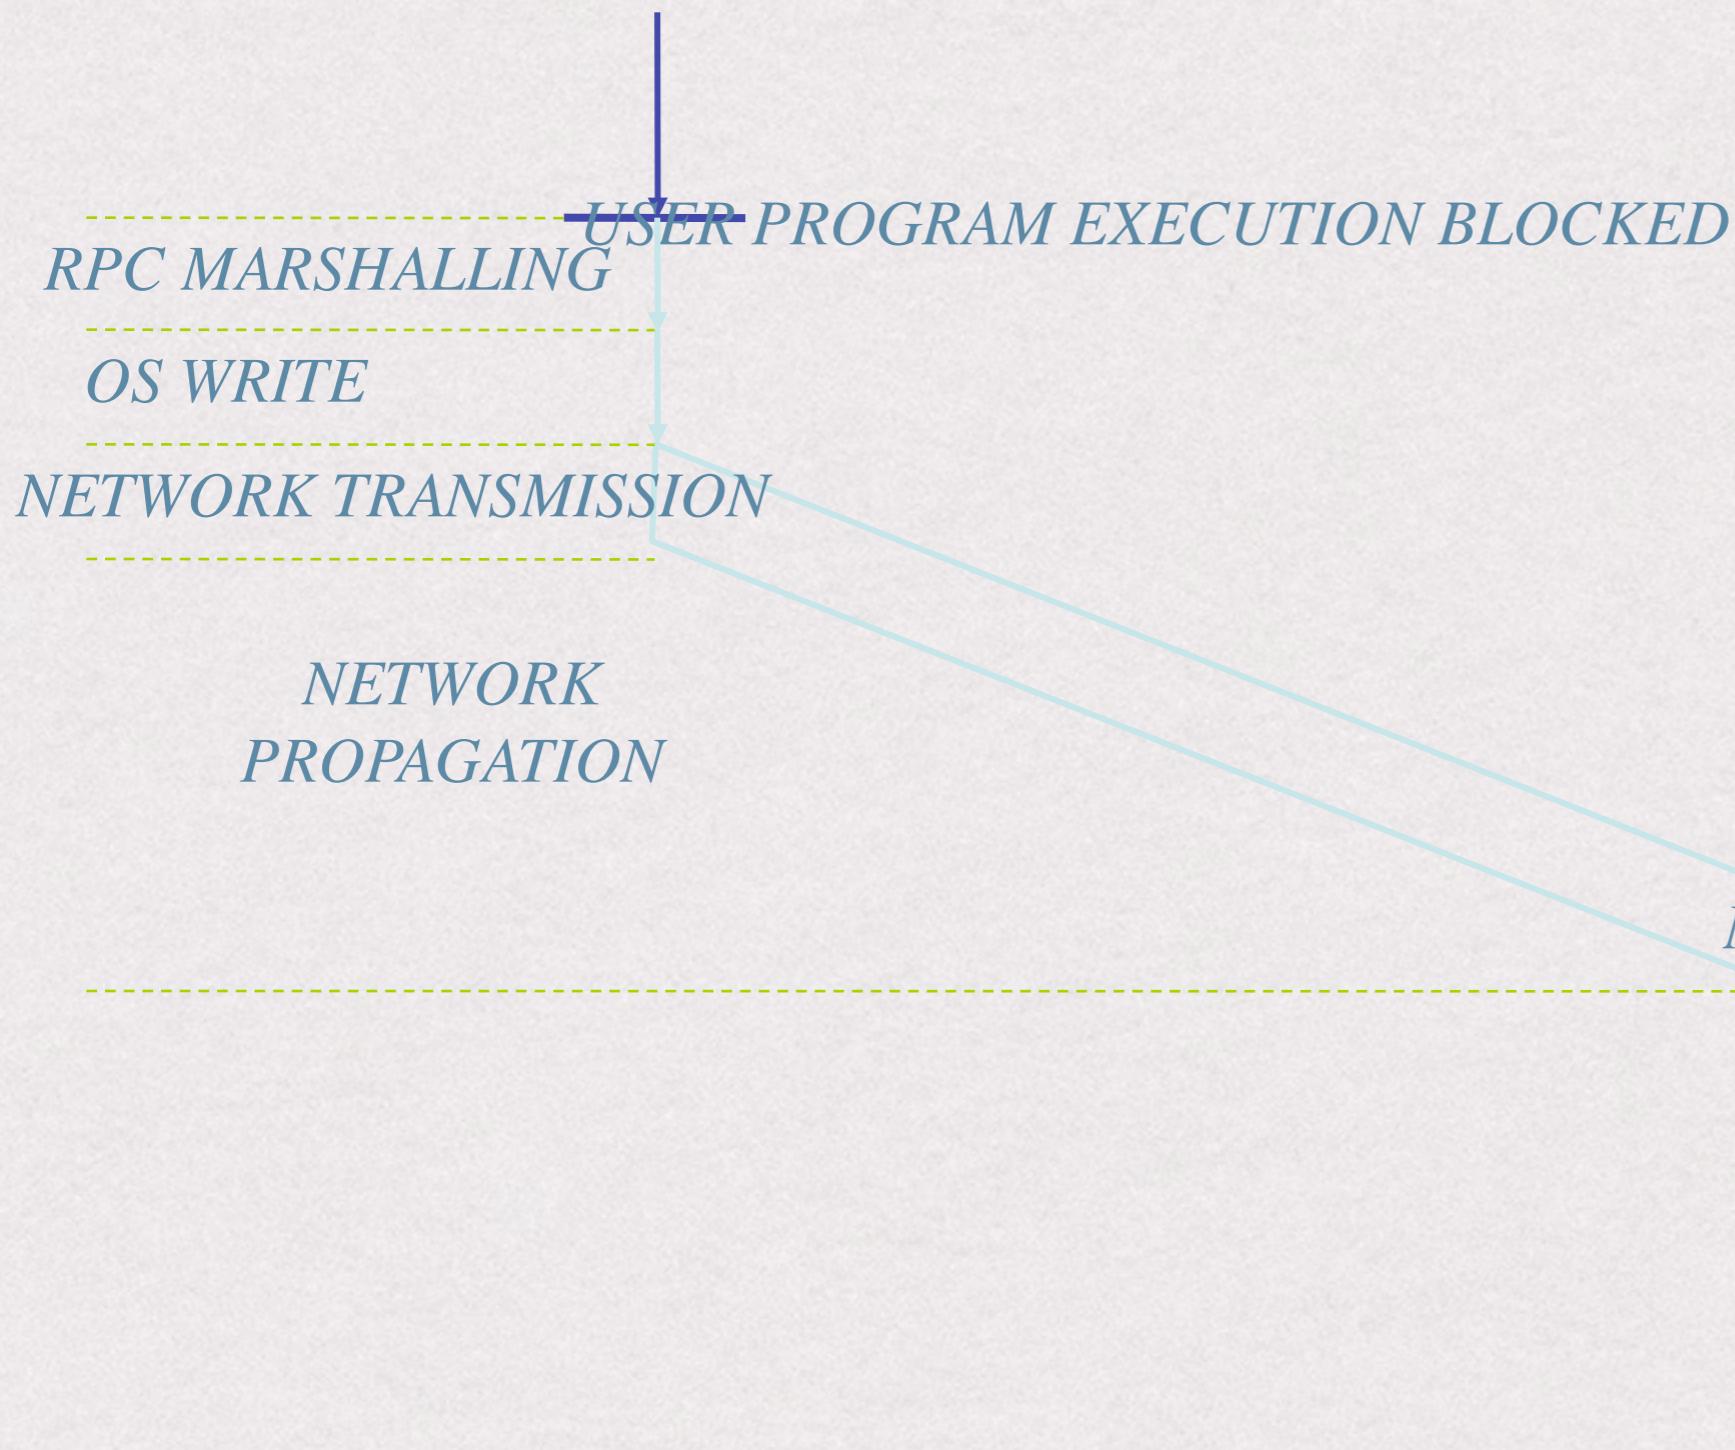
# Recall ... Average Latency

Total time taken for client divided by the number of operations

vs

Average time taken for each operation individually

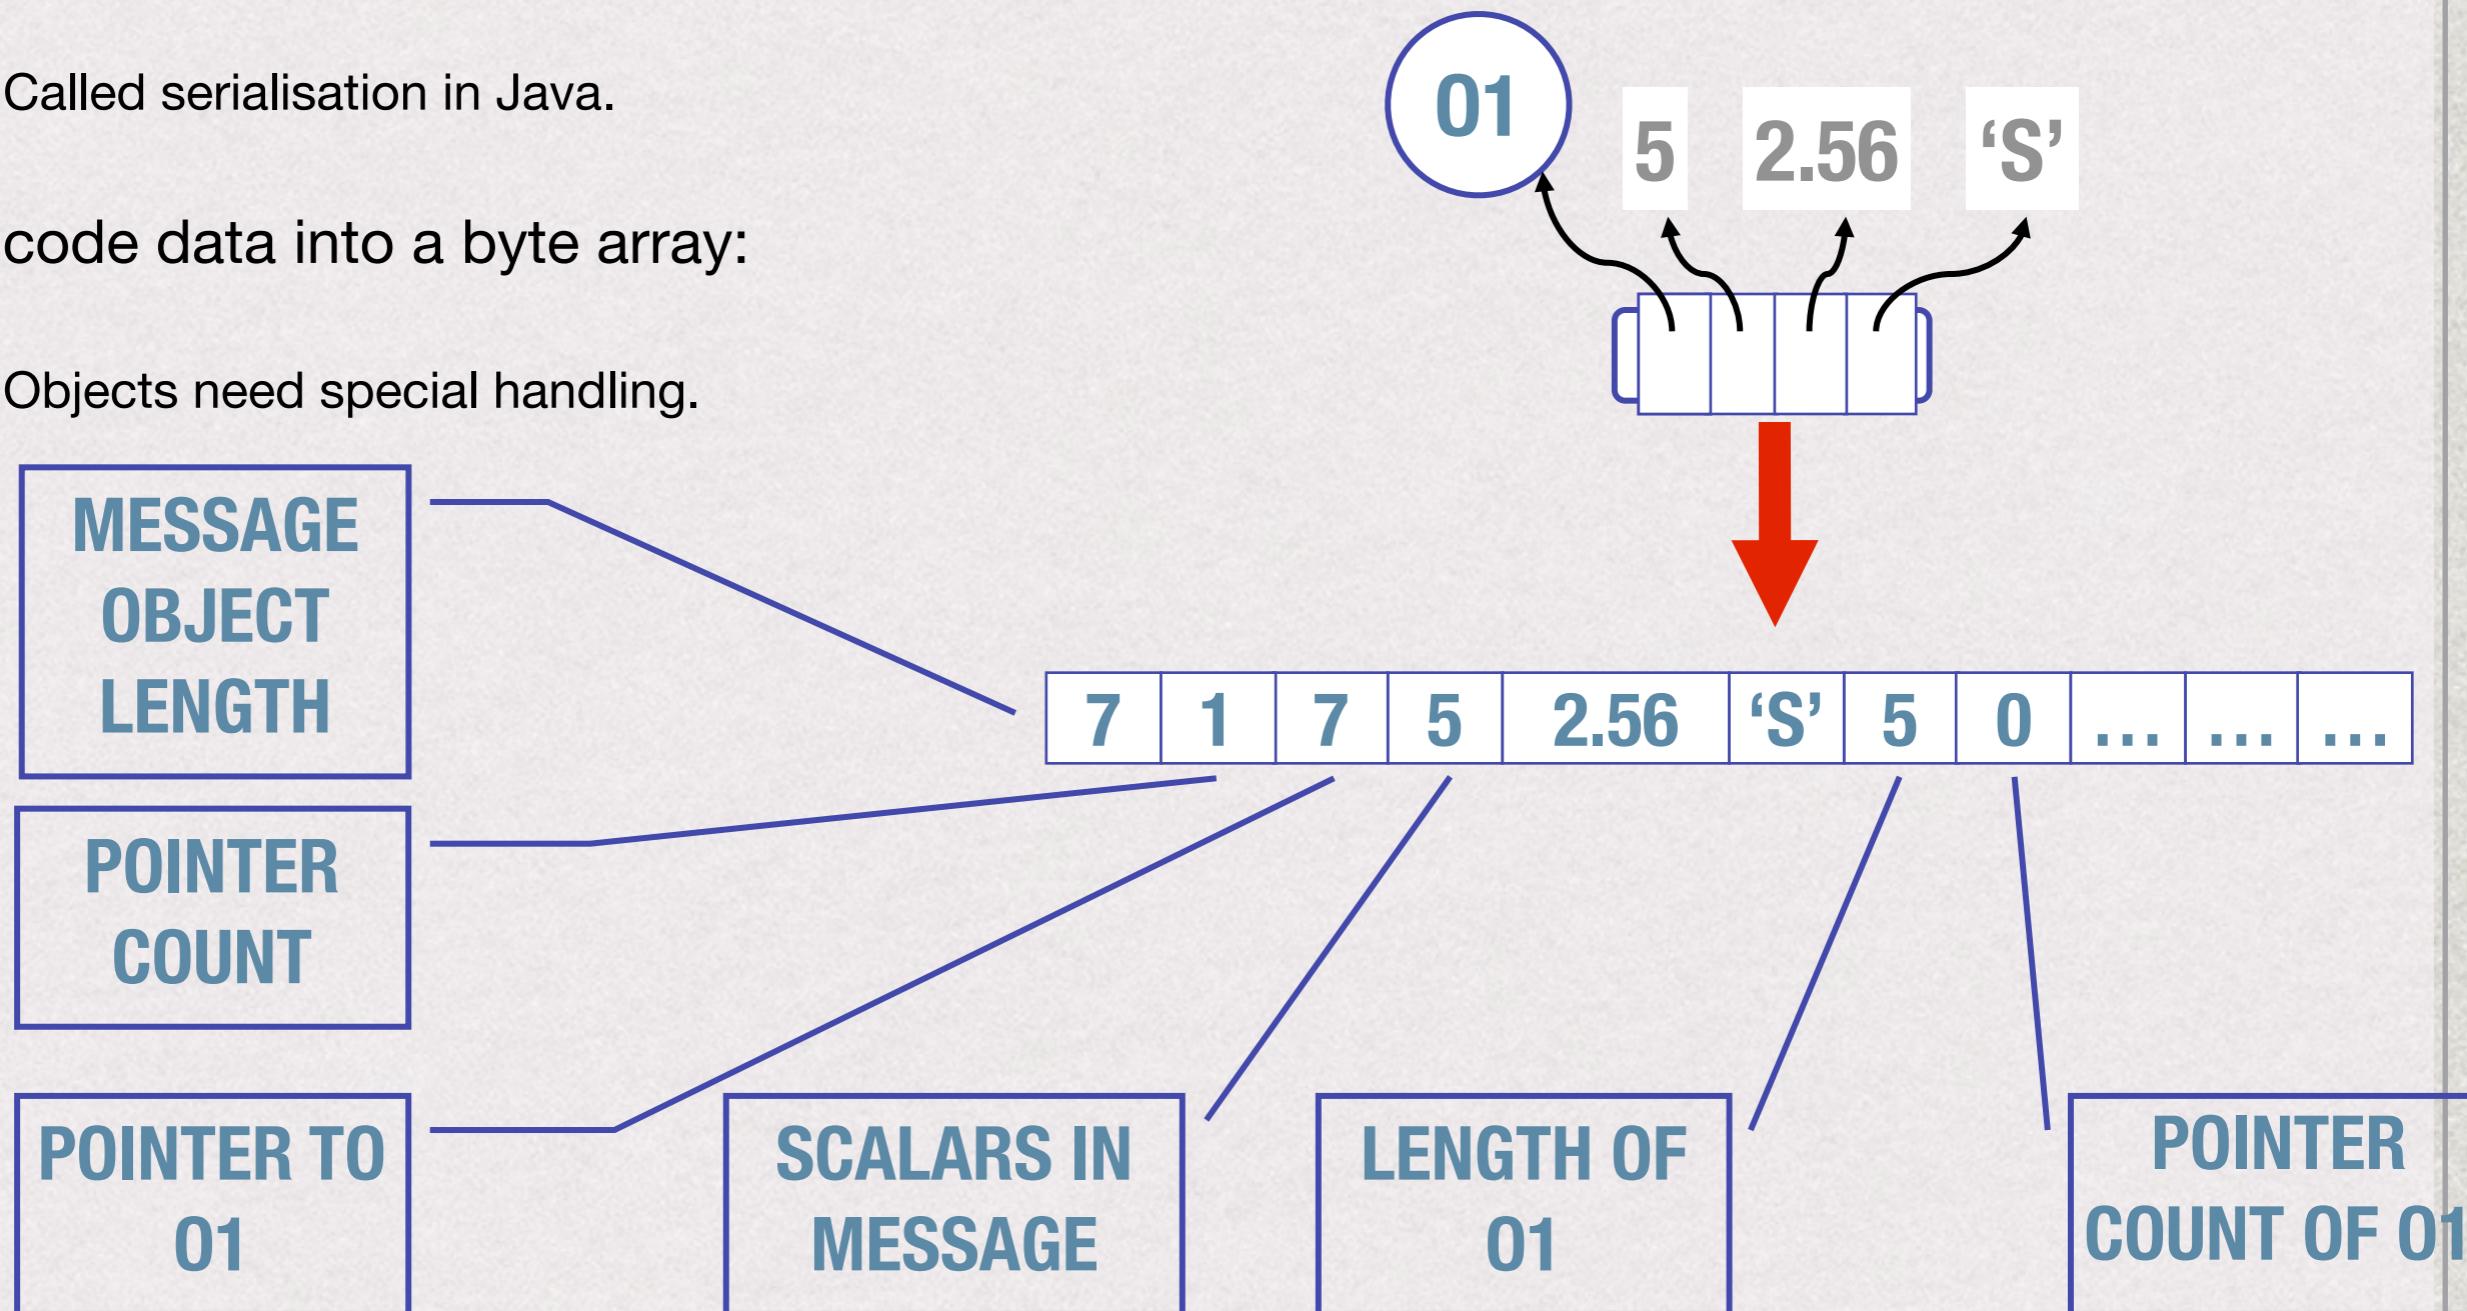
## CLIENT EXECUTION



## TRACKING A CALL FROM CLIENT TO SERVER

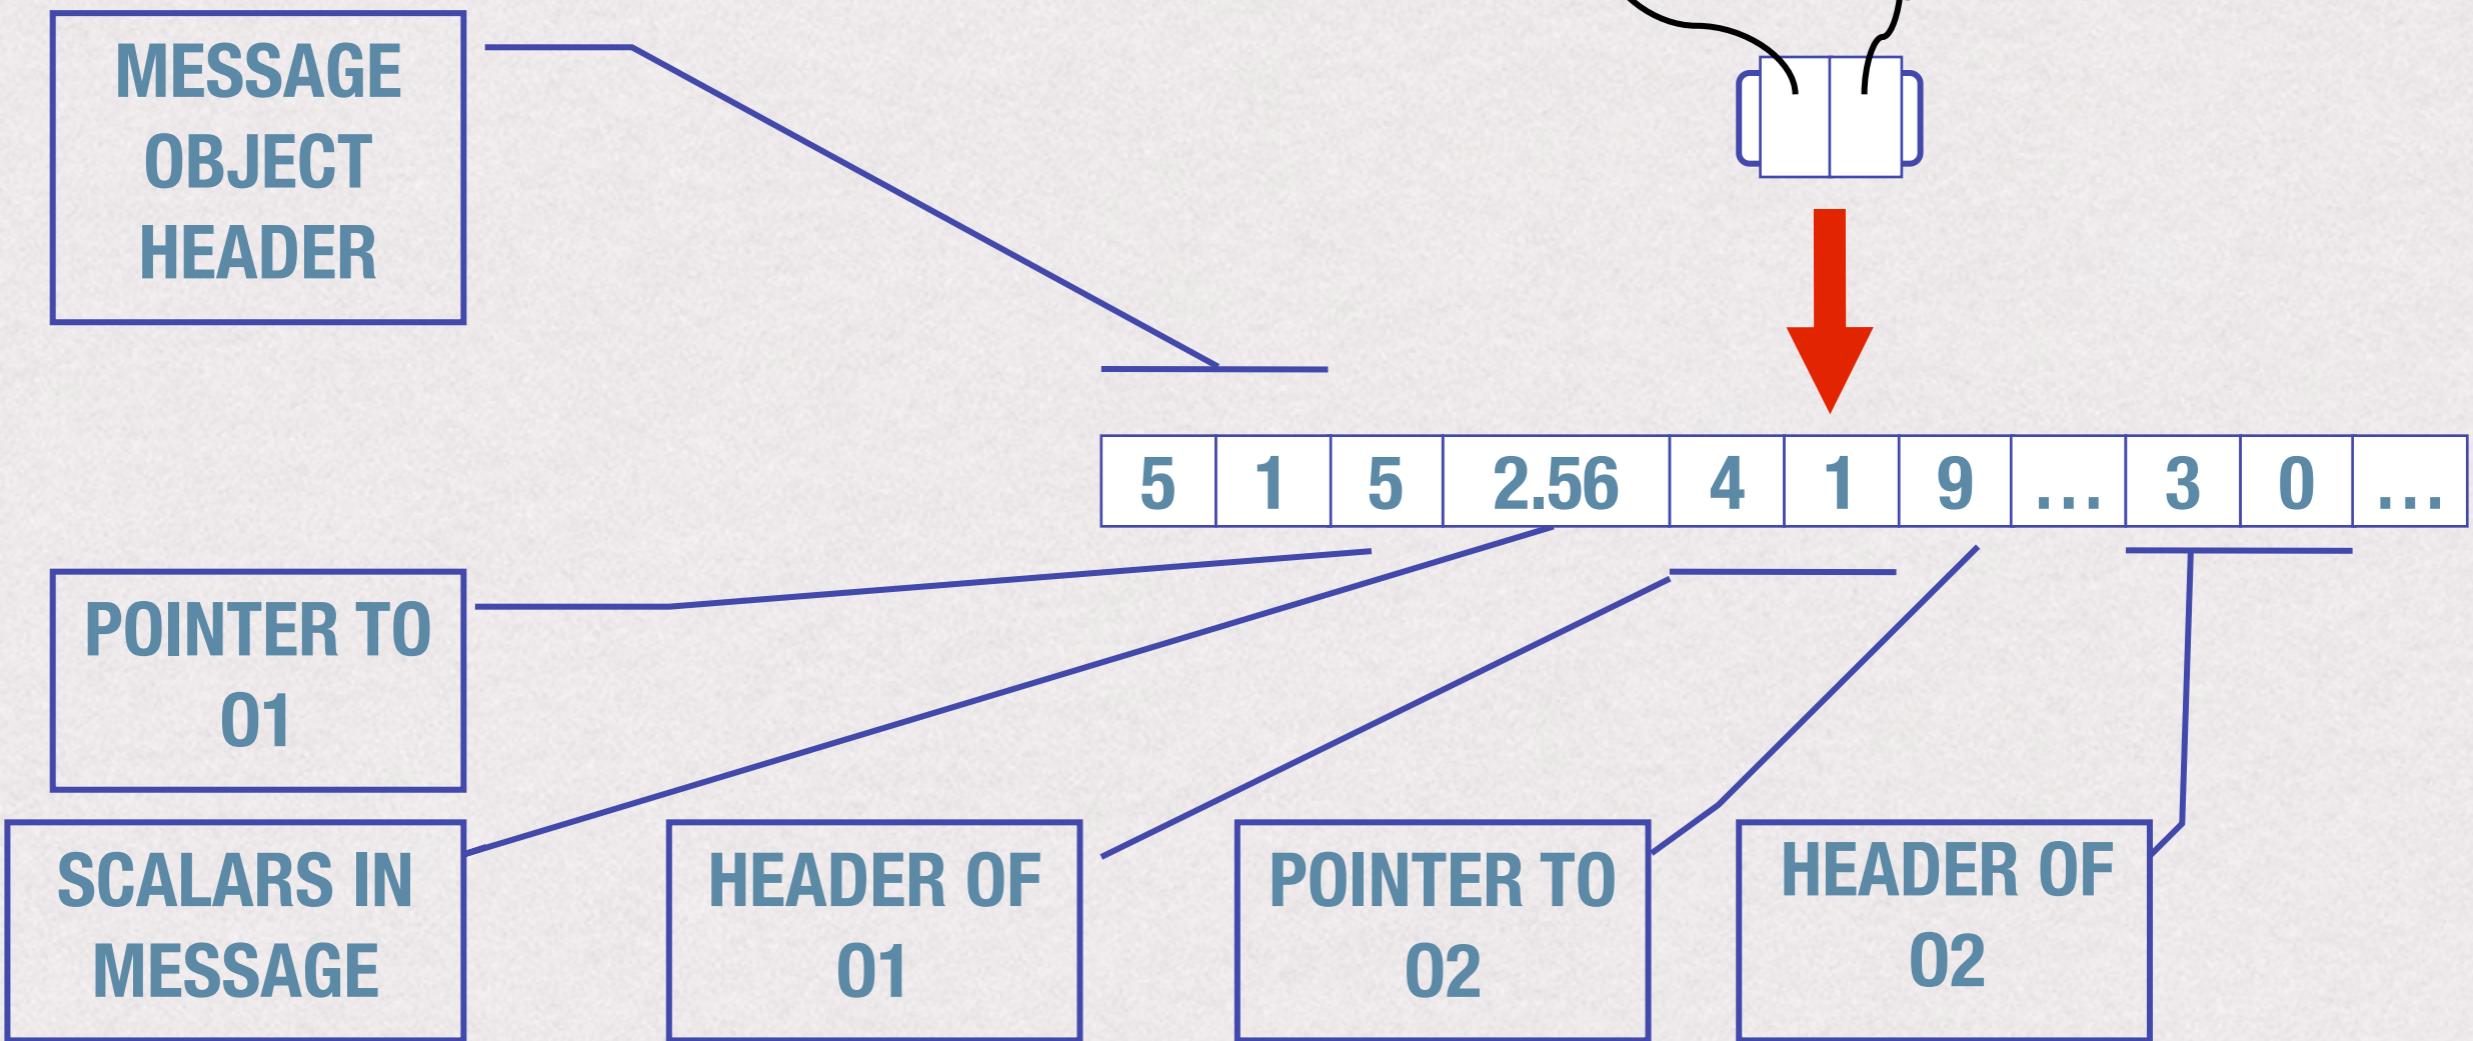
# Marshalling

- The process of flattening data to send a copy of it over a network:
  - Called serialisation in Java.
- Encode data into a byte array:
  - Objects need special handling.

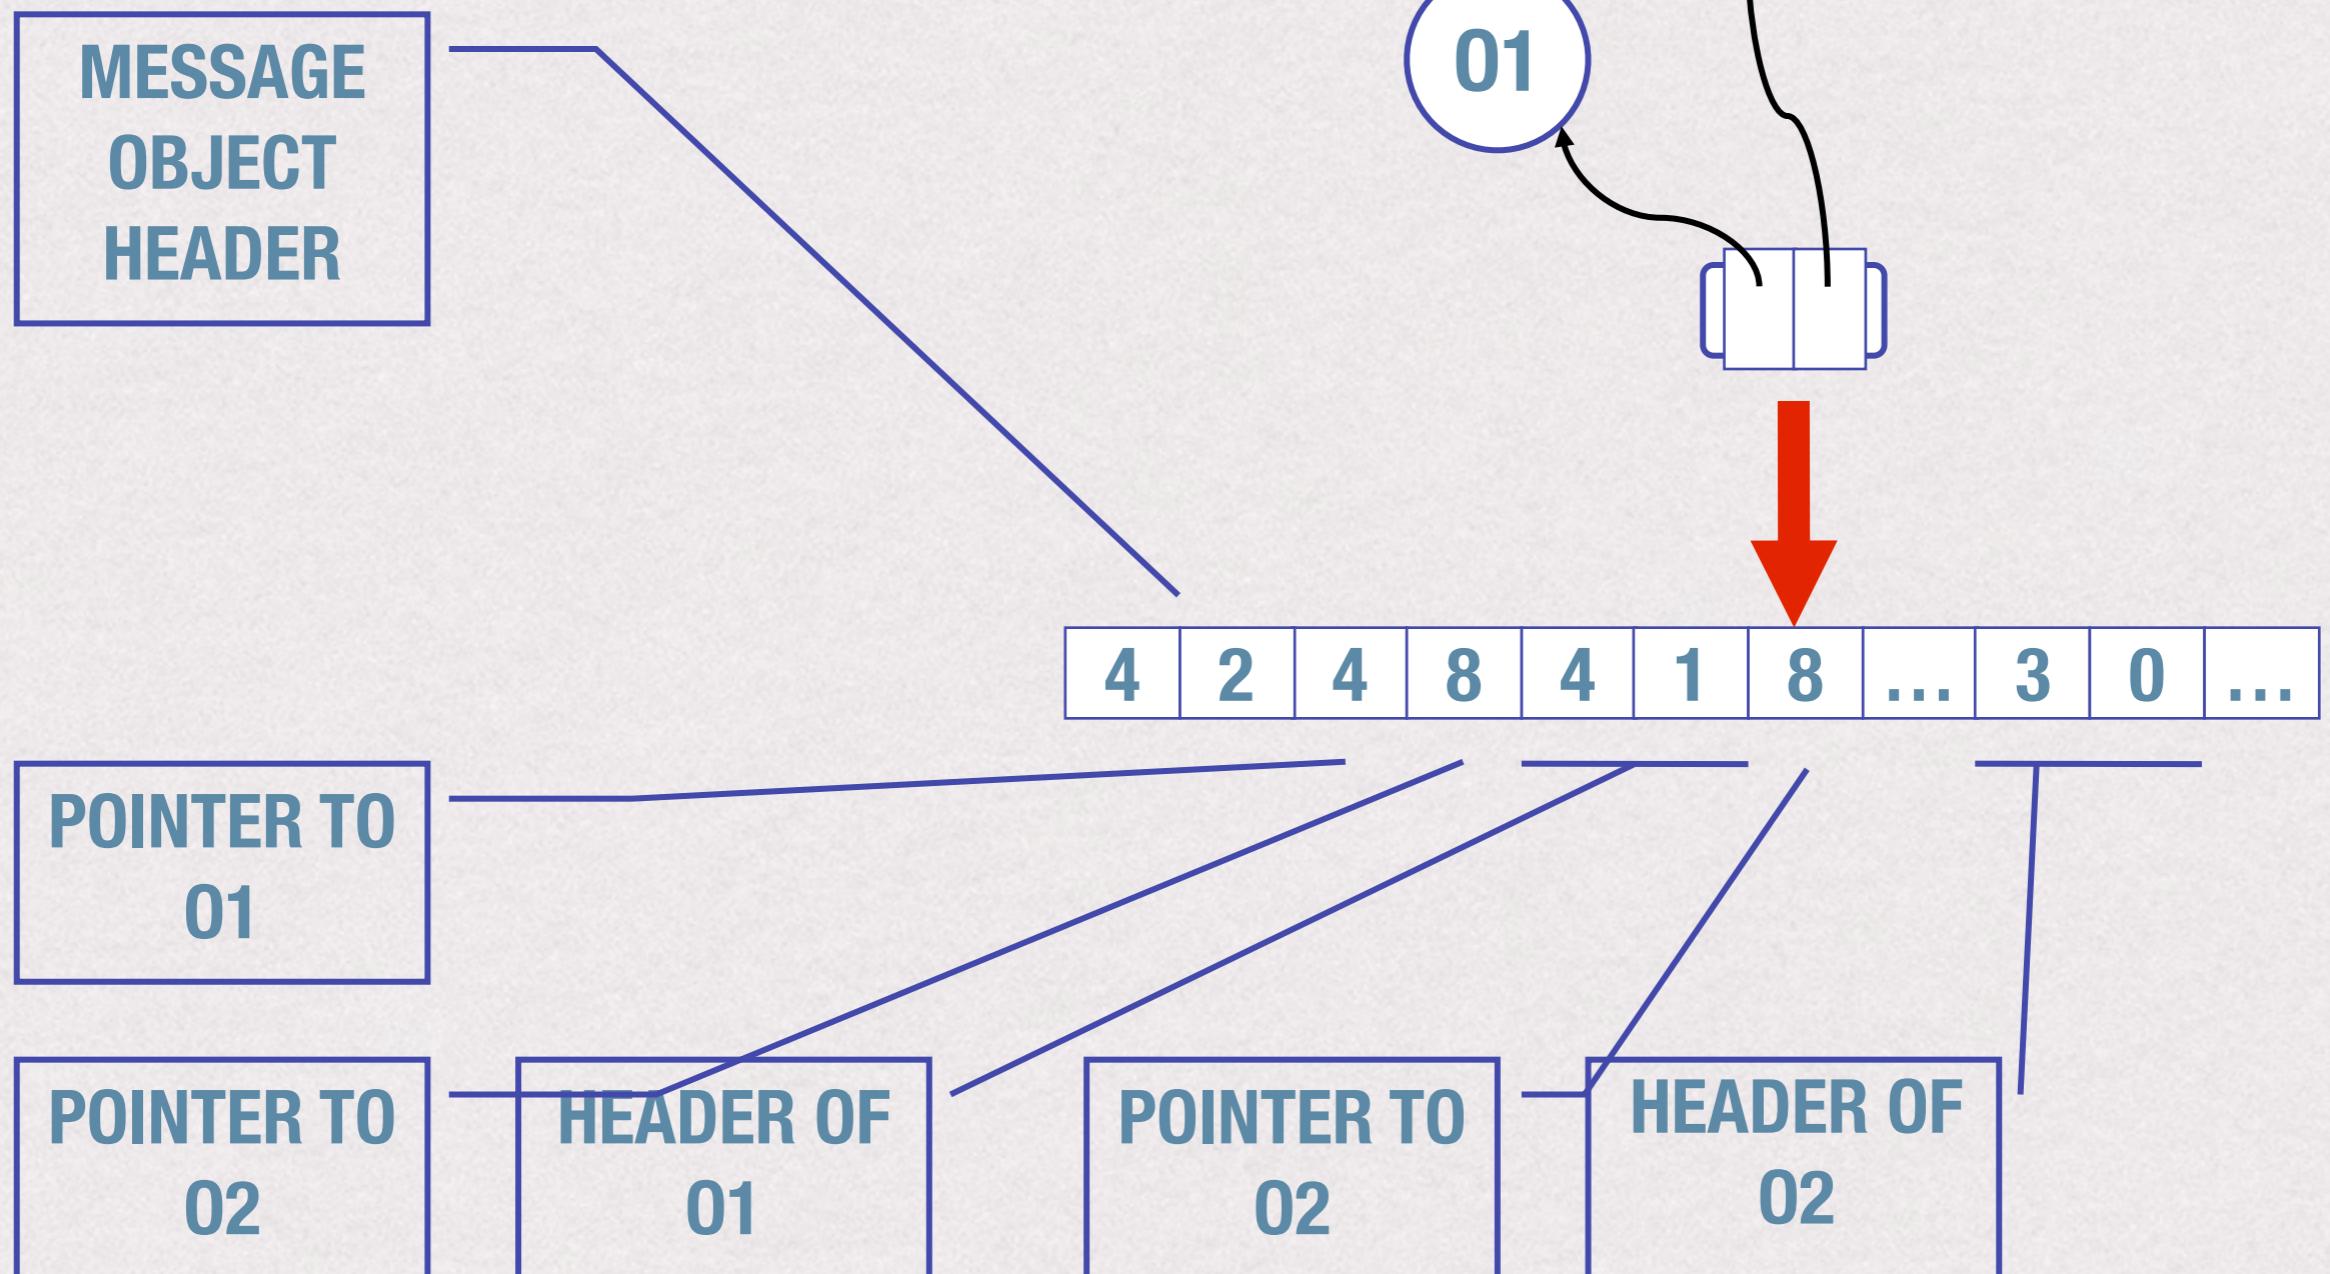


# Marshalling Object Containing Pointers

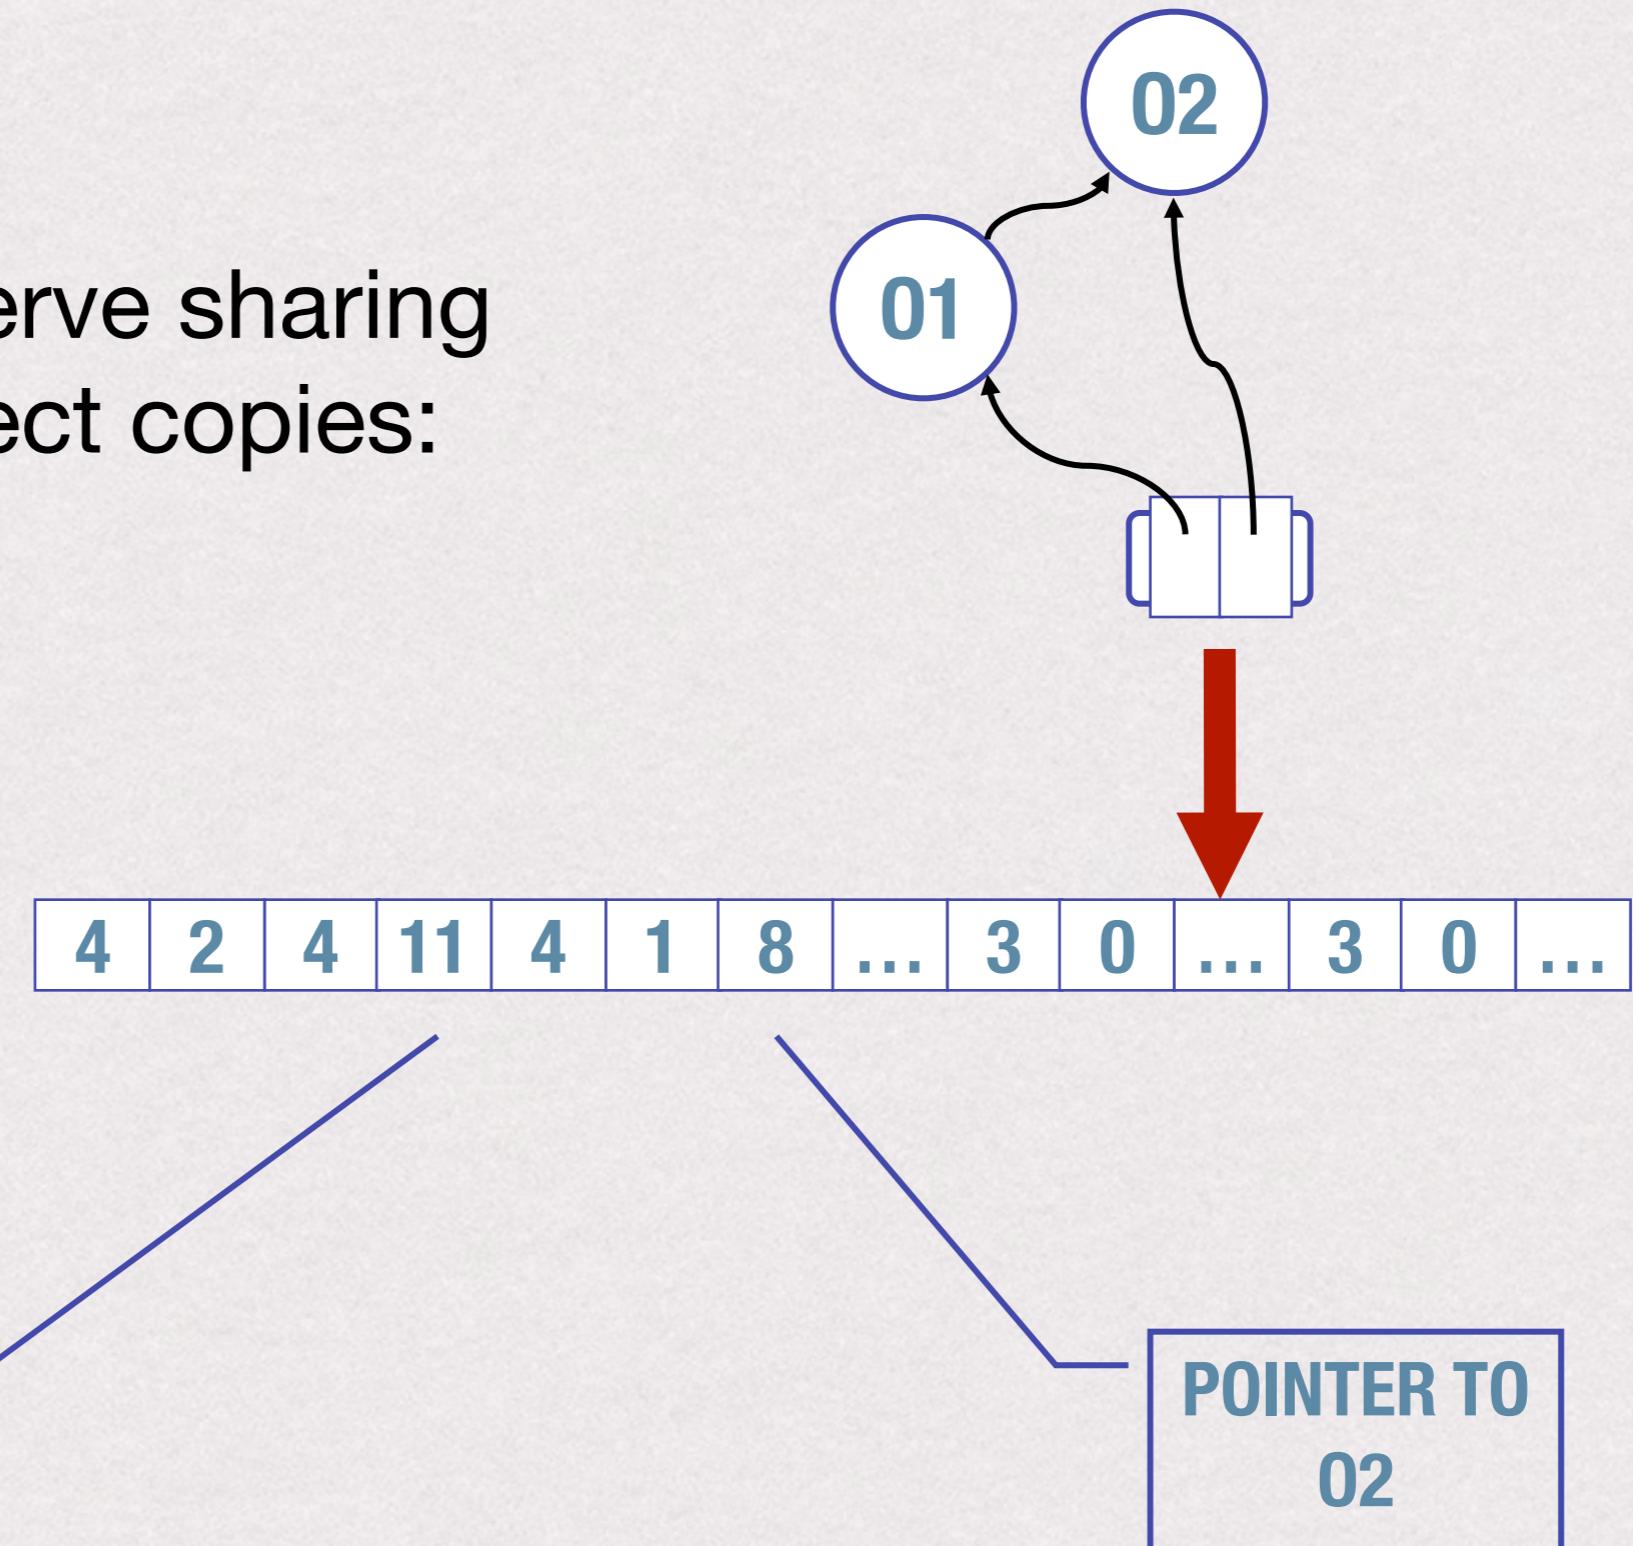
- Marshalling needs to follow pointers inside objects.



- Marshalling needs to preserve sharing within data structures.

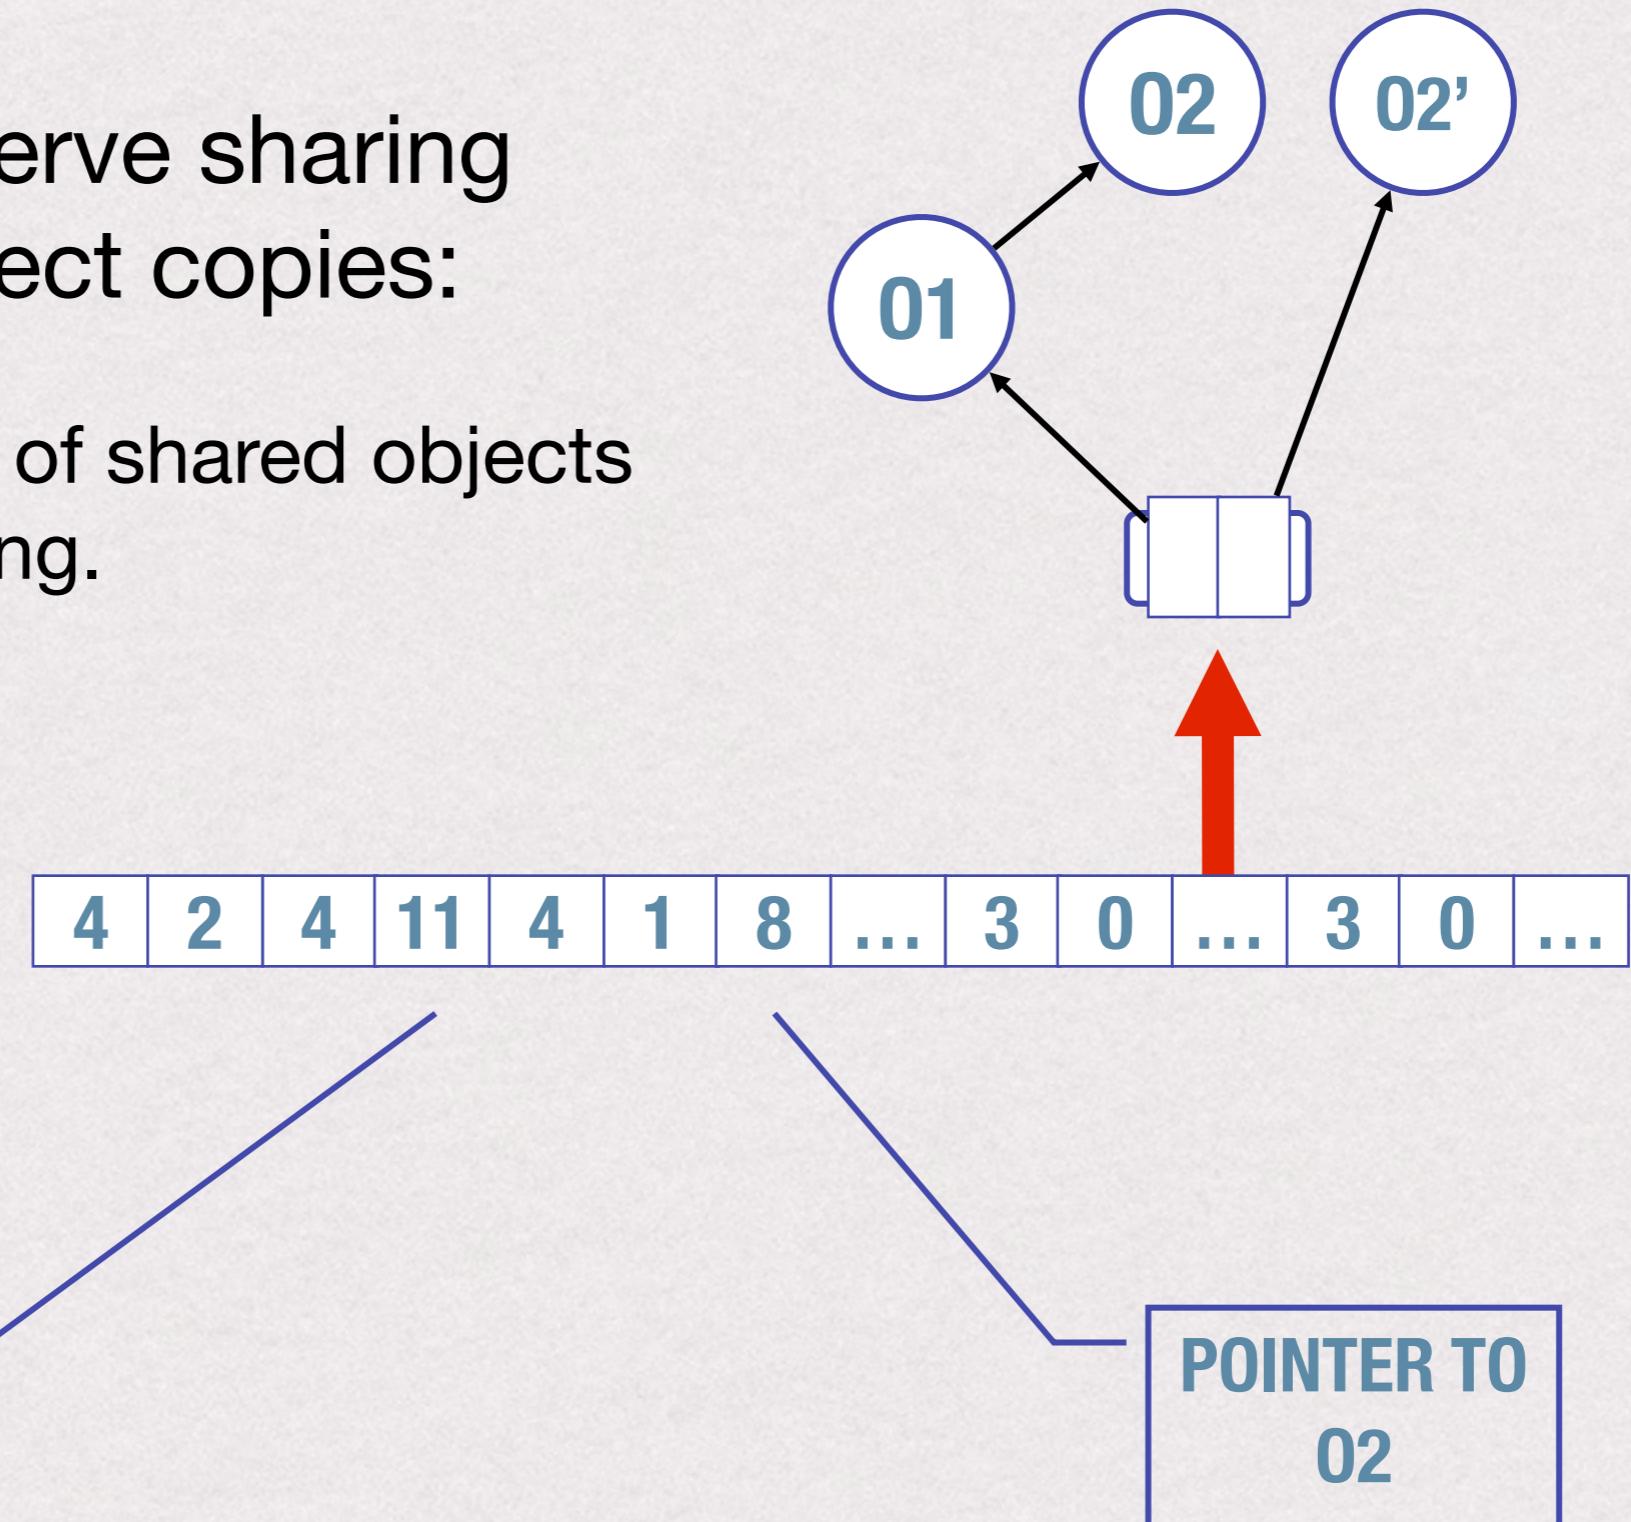


- Failure to preserve sharing leads to incorrect copies:

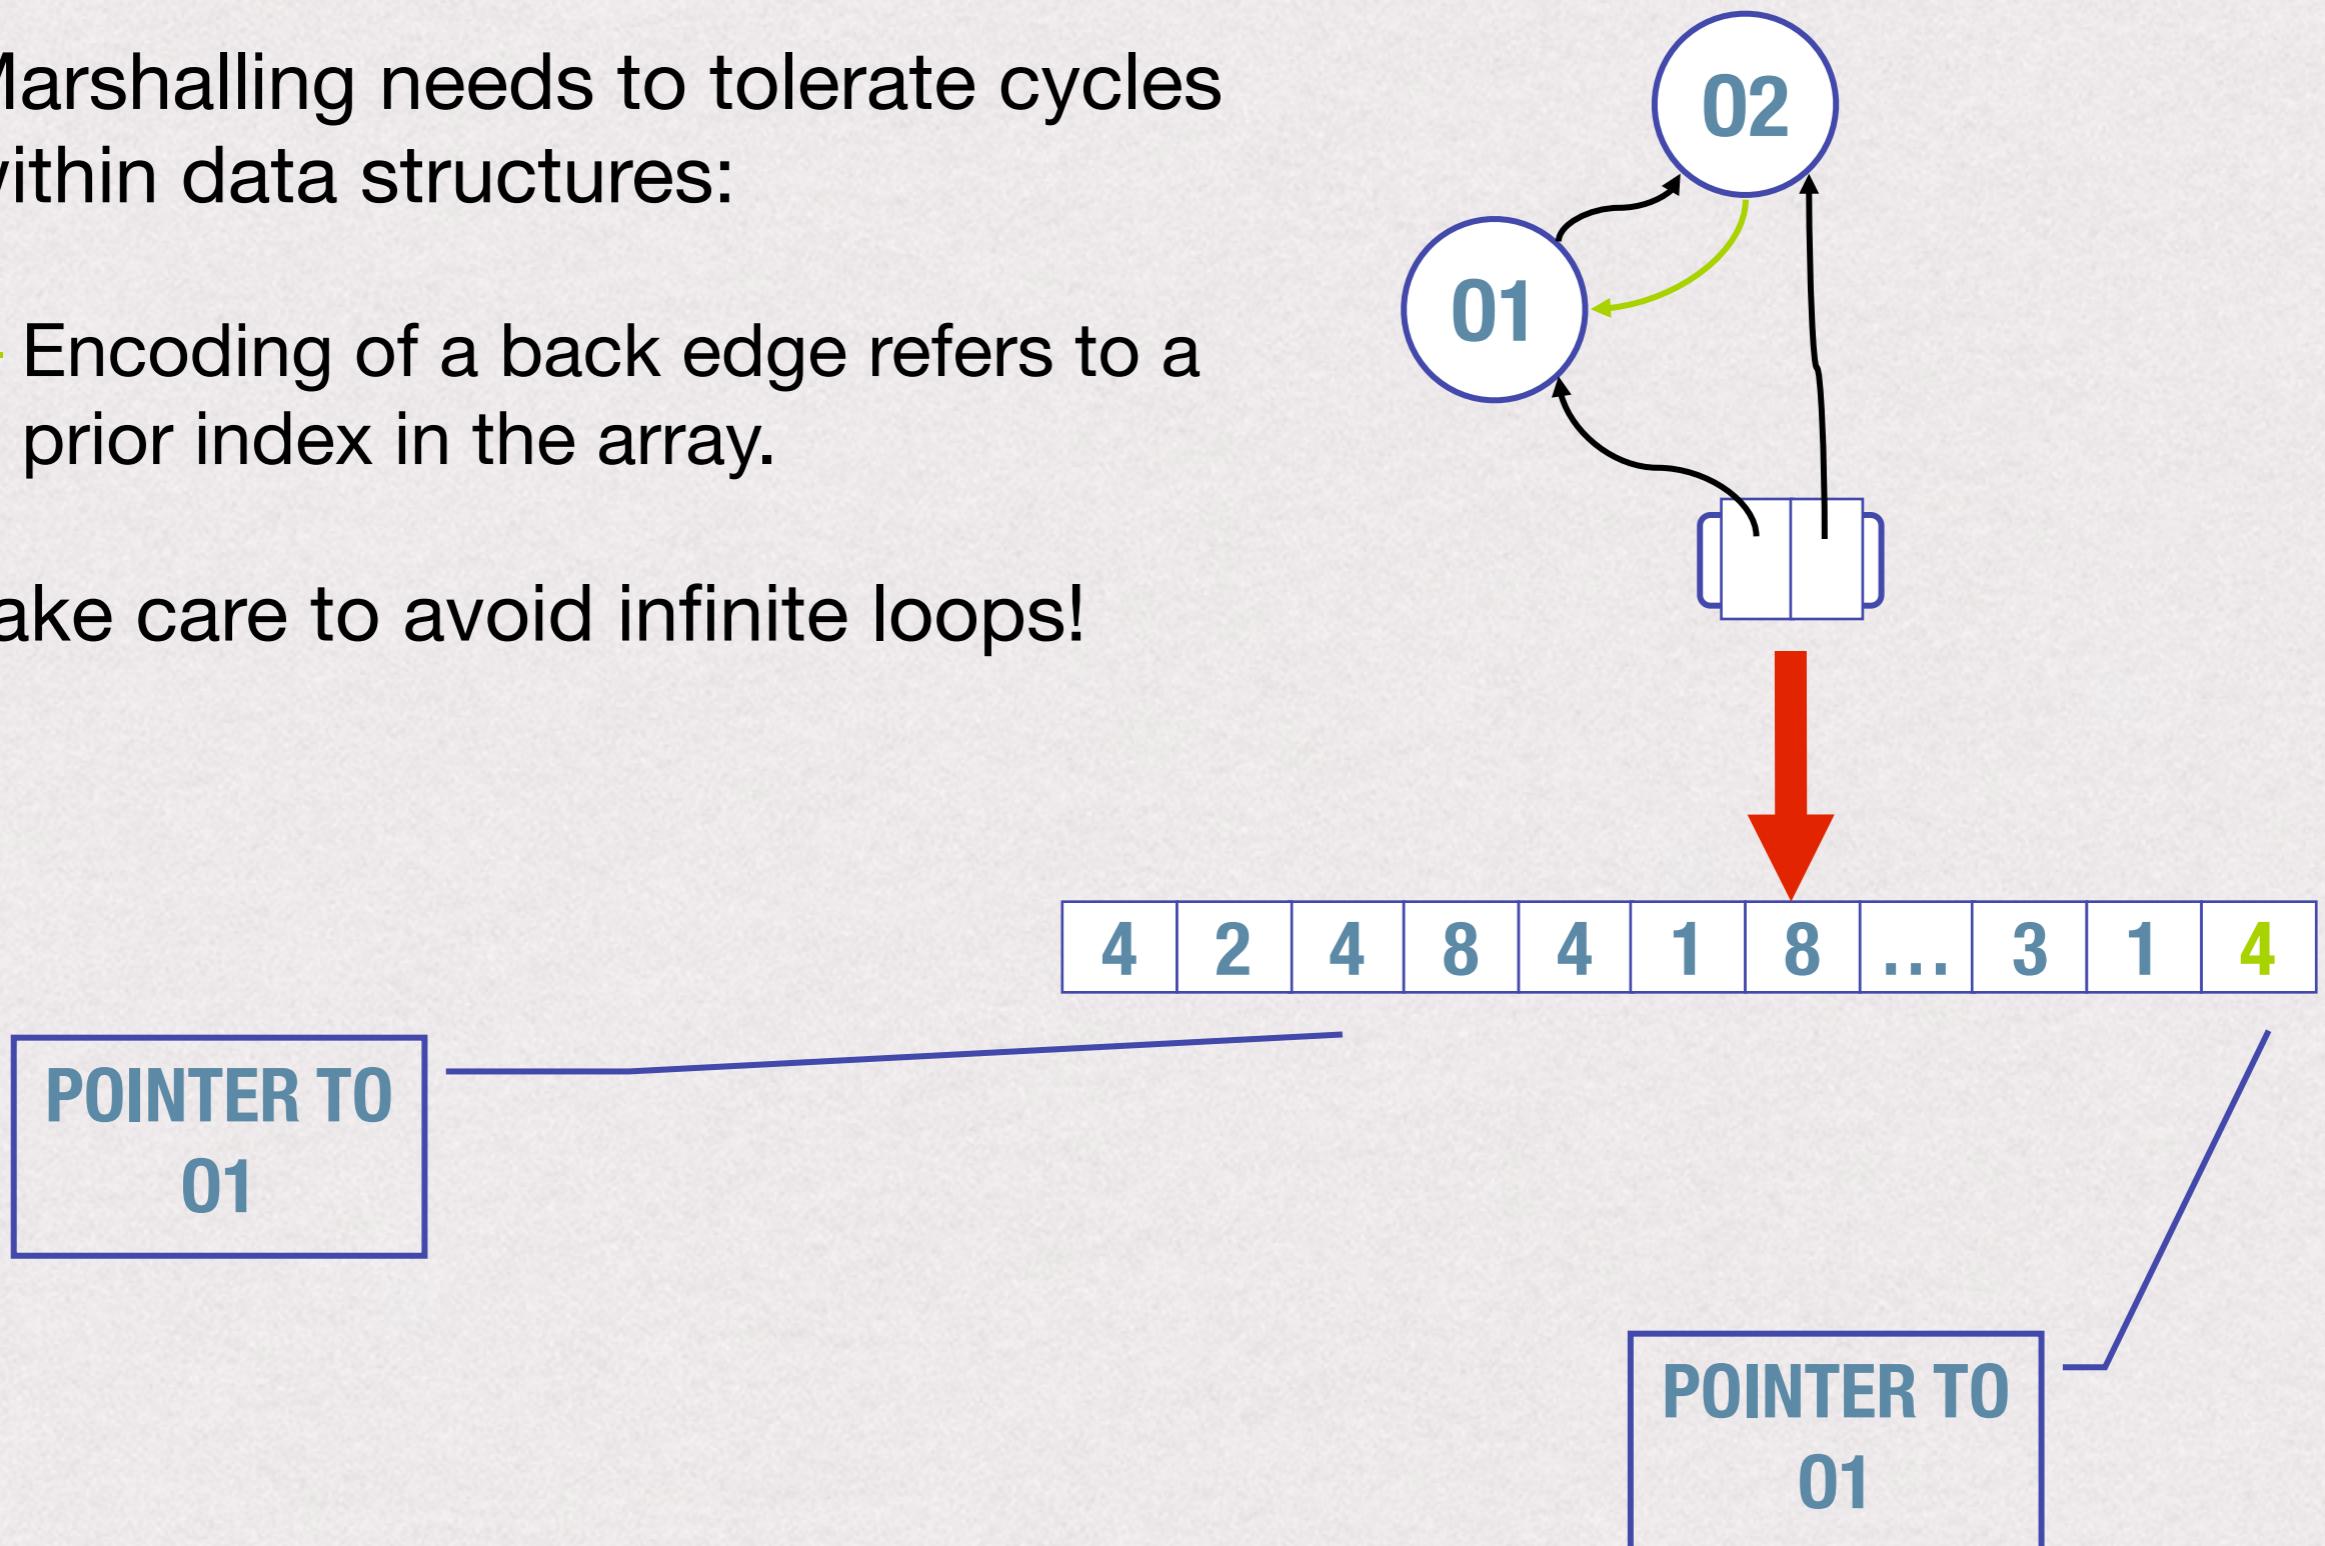


- Failure to preserve sharing leads to incorrect copies:

- Multiple copies of shared objects are unmarshalling.



- Marshalling needs to tolerate cycles within data structures:
  - Encoding of a back edge refers to a prior index in the array.
- Take care to avoid infinite loops!



# Latency Reduction Issues

- \* Network propagation delay is usually dominant
  - \* Other terms reducible by increasing network bandwidth (transmission delay) or processing power (all else).
  - \* Lower bound on propagation delay due to speed of light .
- \* Latency reduction aims to reduce the mean count of propagation delays per remote call:
  - \* For a single RPC, the number of propagation delays per call is 2/1, i.e. two – one for the call, one for the result.
  - \* Alternative terminology is to talk in terms of “round trips” a term referring to a pair of propagation delays.

# Focus

- \* Dependencies constraints
  - \* Classify different types of dependencies
- \* Abstractions to reduce propagation delays

# Latency Reduction with Parallel RPC

```
pcall {  
    res1 = O1.m1(paramsA ...)  
    res2 = O2.m2(paramsB ...)  
    res3 = O3.m3(paramsC ...)  
}
```

- \* All three (or however many) calls are sent off together, in any order, then the client blocks until the results of all calls have arrived back.
- \* *What is the overall latency for the calls in a parallel RPC?*
- \* *What is the average latency for the calls in a parallel RPC?*
- \* *What dependencies are permitted between the calls in a parallel RPC?*

# Latency Reduction with One-way RPC

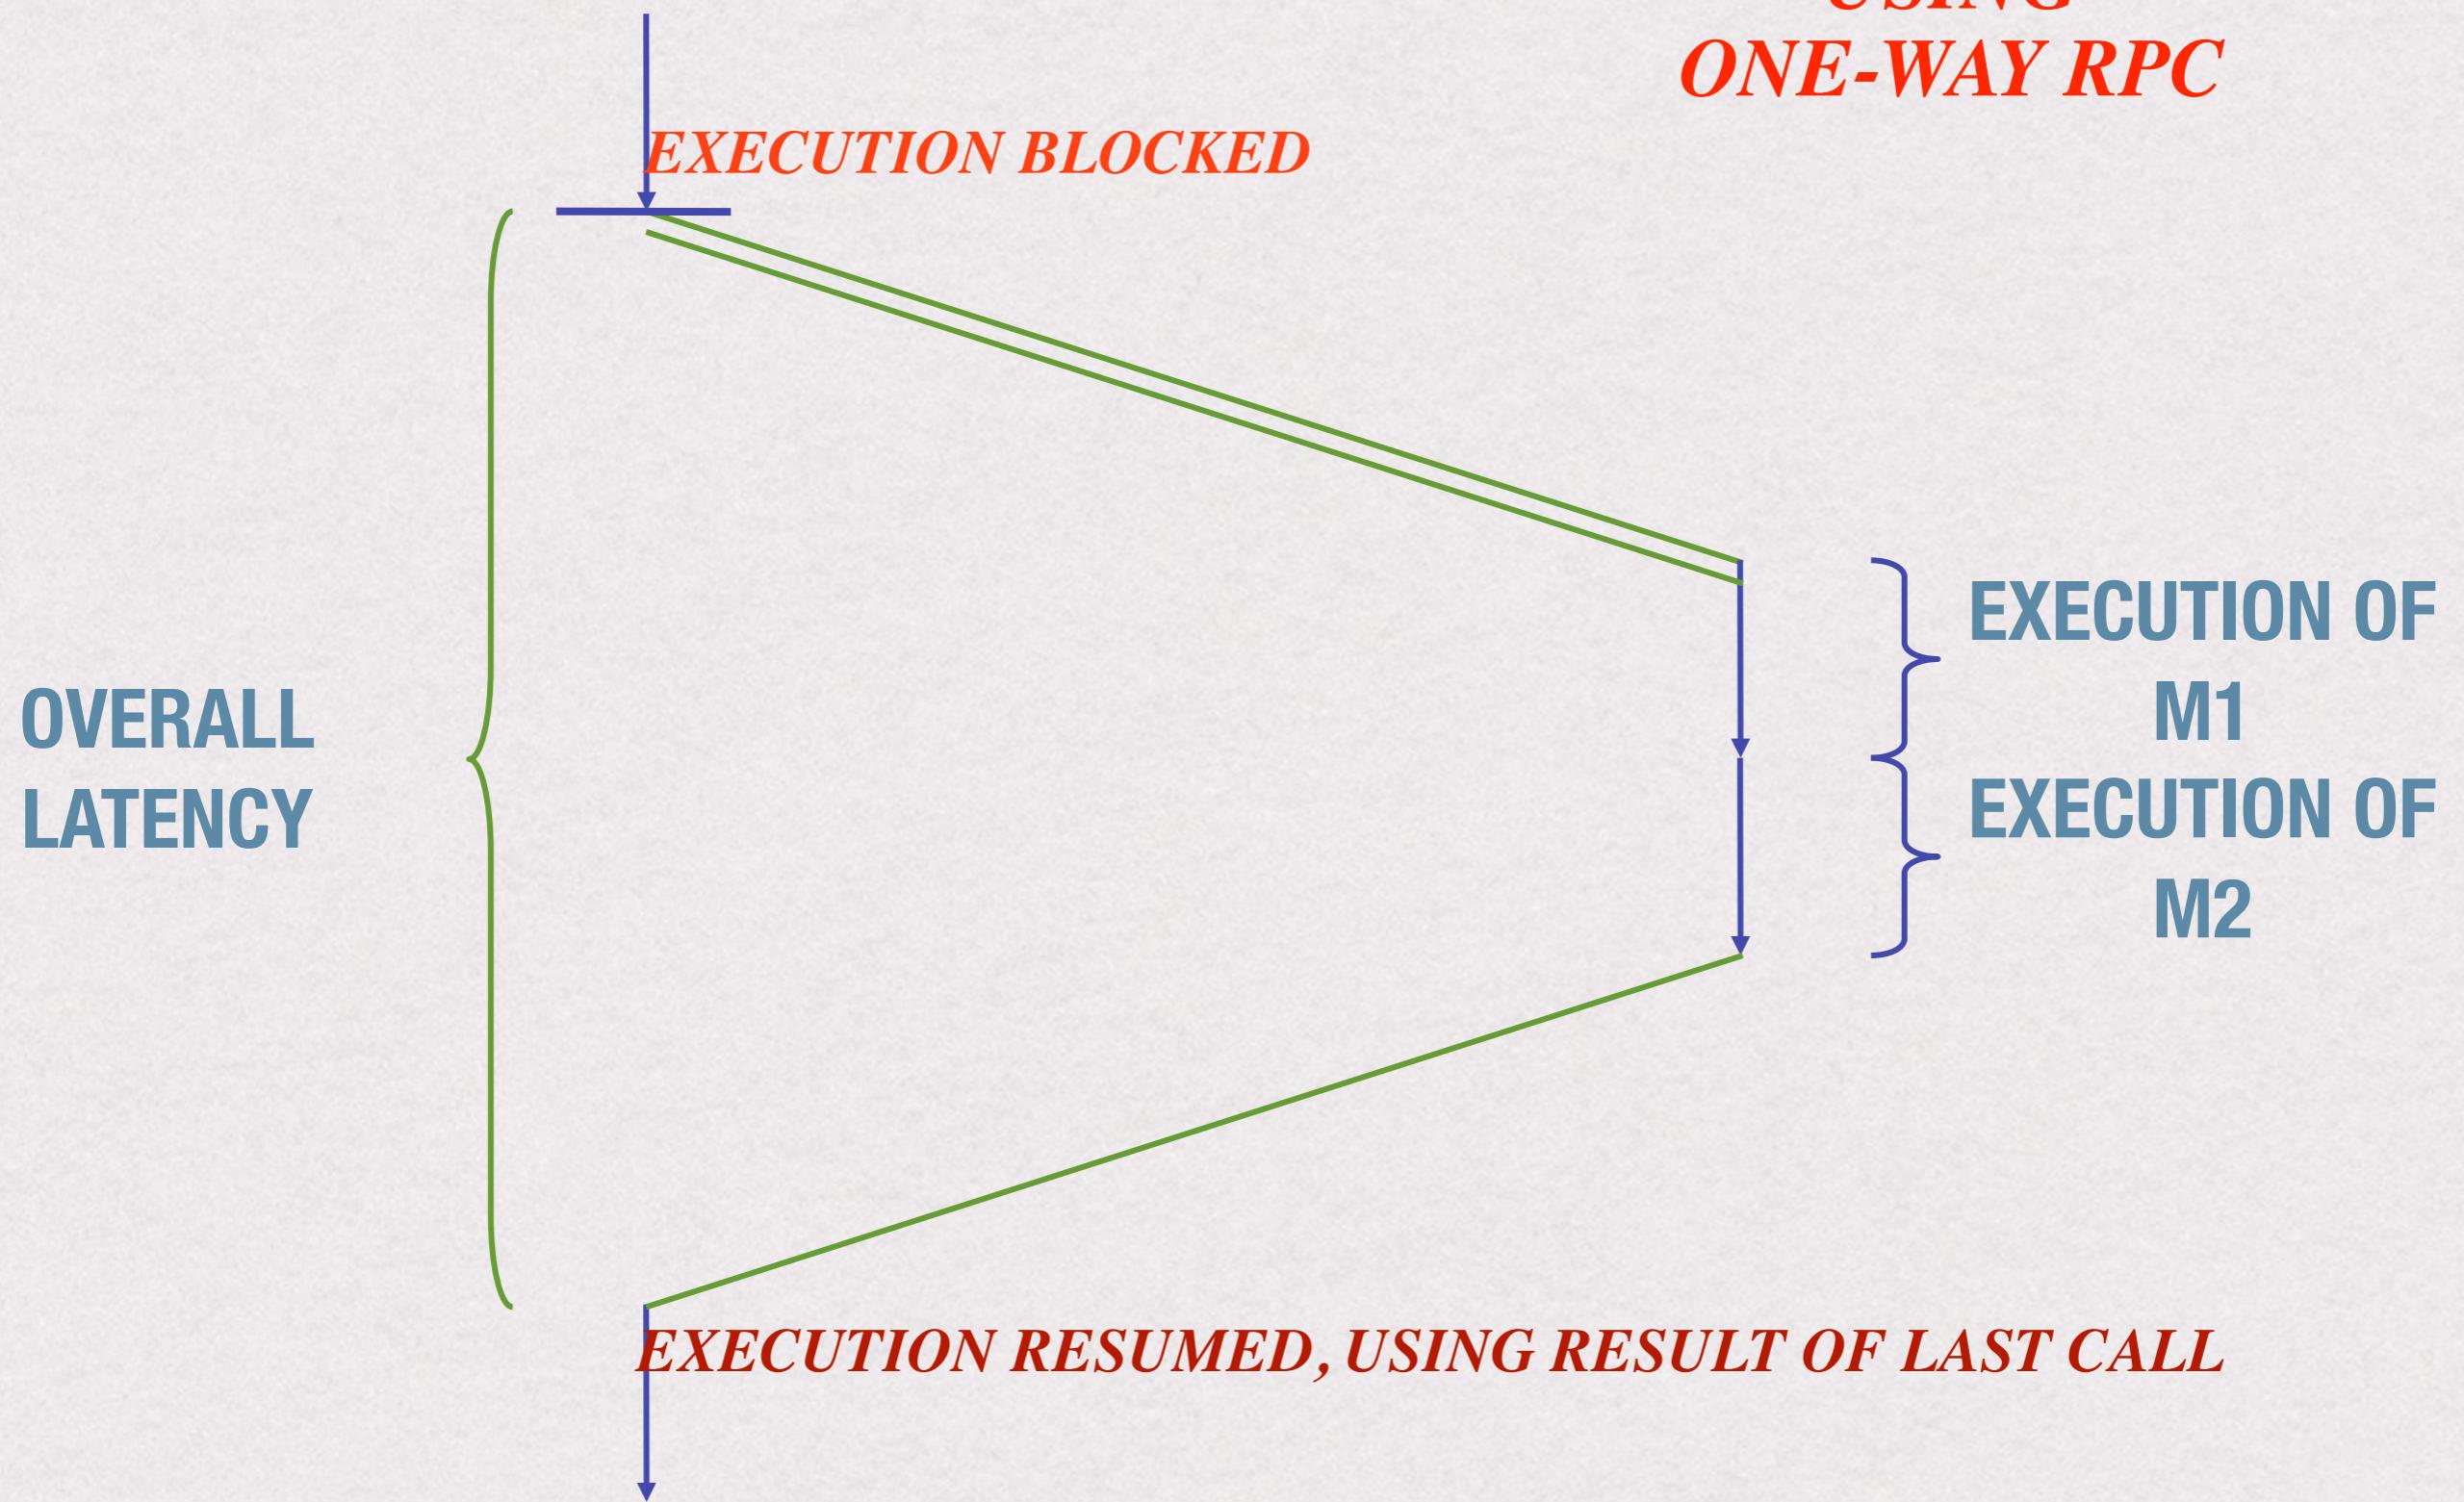
2

- The ***One-way RPC*** idea is:

```
O1.m1(paramsA ...)
res2 = O1.m2(paramsB ...)
```

- All except the last of this group of calls are one-way RPCs – they produce no result.

## CLIENT EXECUTION



# One-way RPC

- The **One-way RPC** idea is:

O1.m1(paramsA ...)

res2 = O1.m2(paramsB ...)

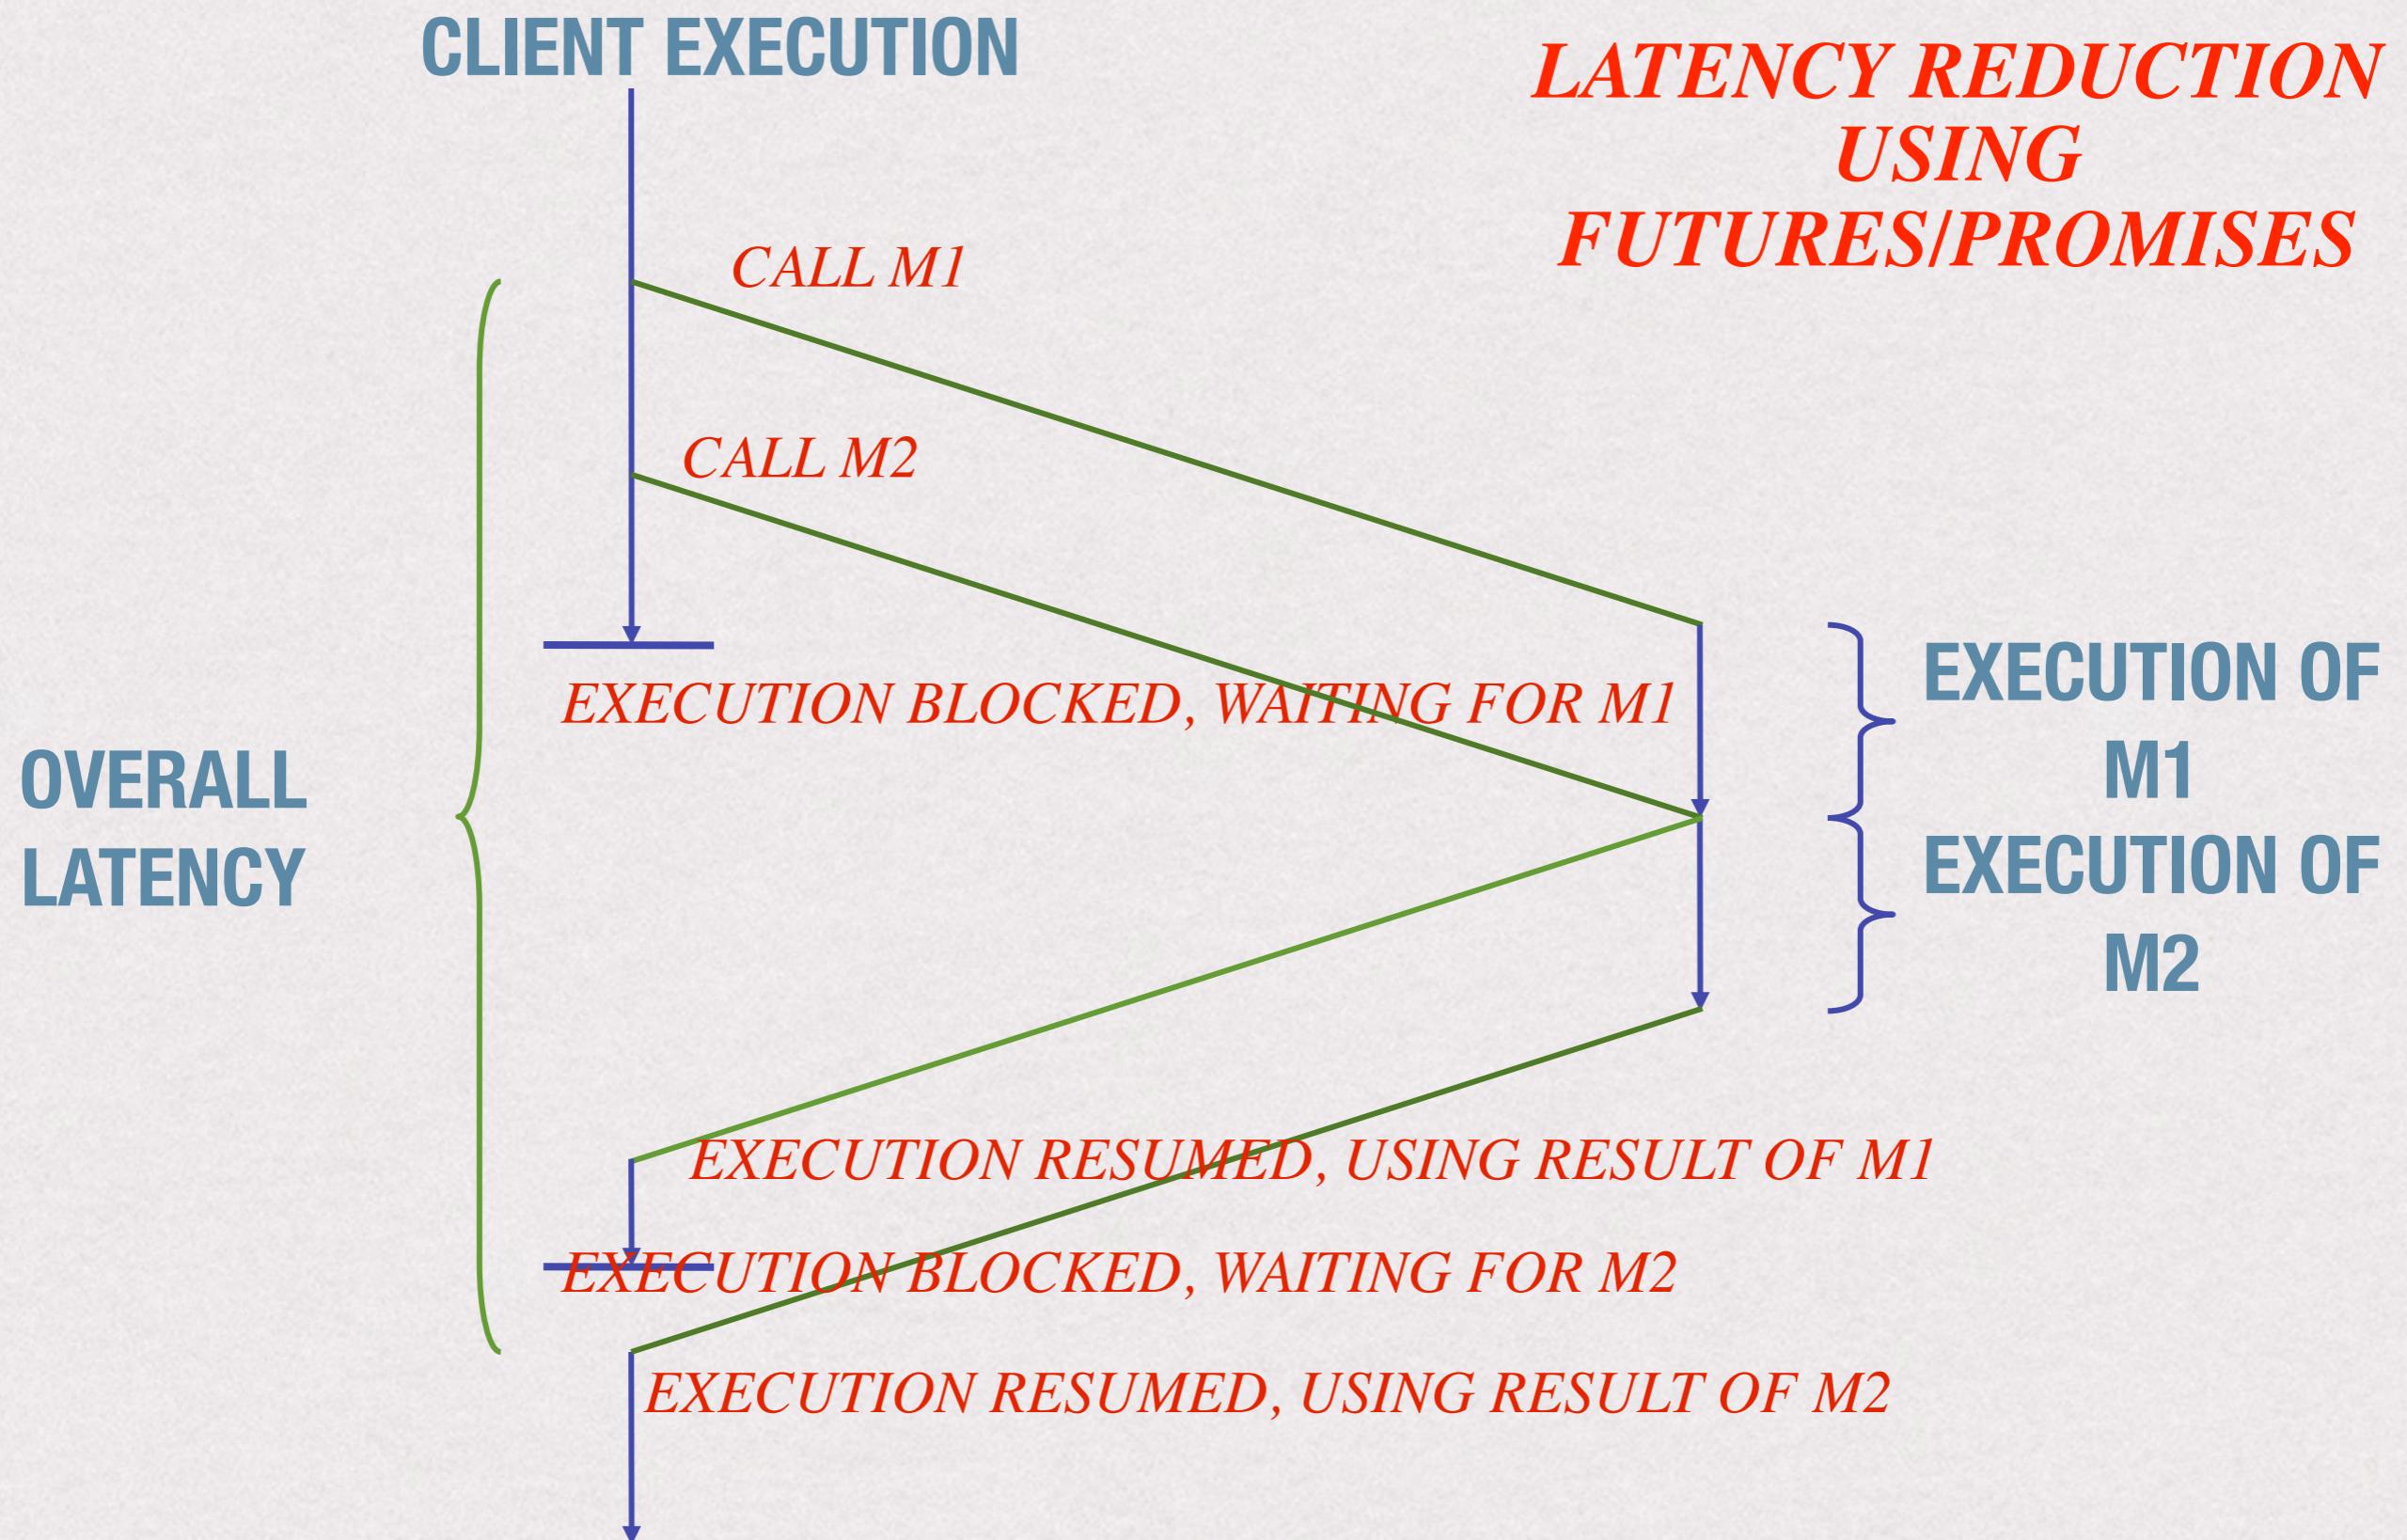
- All except the last of this group of calls are one-way RPCs – they produce no result.
- \* *What is the overall latency for a series of calls using one-way RPC?*
- \* *What is the mean latency for a series of calls using one-way RPC?*
- \* *What can we say about success or failure of one-way RPC calls?*
- \* *What can be said about the server used to process the calls in such a group?*
- \* *What dependencies are permitted between the calls using one-way RPC?*
- \* *There is a variant of one-way RPC call **maybe RPC**, i.e. the call may be executed, but there is no way to find out whether it did*

# Latency Reduction with Futures and Promises

- \* An example of the Future/Promises idea is:

```
res1 = O1.m1(paramsA ...)  
... computation ...  
res2 = O1.m2(paramsB ...)  
... computation ...  
... computation using res1 ...  
... computation using res2 ...  
res3 = O2.m3(paramsC ...)
```

- \* Much more flexible than Parallel RPC!



# Futures and Promises

- \* An example of the Future/Promises idea is:

```
res1 = O1.m1(paramsA ...)  
... computation ...  
res2 = O1.m2(paramsB ...)  
... computation ...  
... computation using res1 ...  
... computation using res2 ...  
res3 = O2.m3(paramsC ...)
```

- \* What is the best case for total latency of m1 and m2 considered together?
- \* What is the best case for mean latency of m1 and m2 considered together?
- \* What can we say about the dependencies between m1 and m2?
- \* What can we say about the dependencies between m2 and m3, assuming the objects called are on different servers?

# Recall ...

- \* Latency reduction strategy:
  - \* It became increasingly difficult to analyse as the mechanisms became more complex, more application dependent
  - \* However, as mechanisms required less blocking they tended to achieve better latency reduction.
- \* In this part of the lecture, we will explore this tendency further.

# Use of Results and Blocking

- \* Consider the following:

```
res1 = O.m1(...)  
res2 = O.m2(..., res1, ...)
```

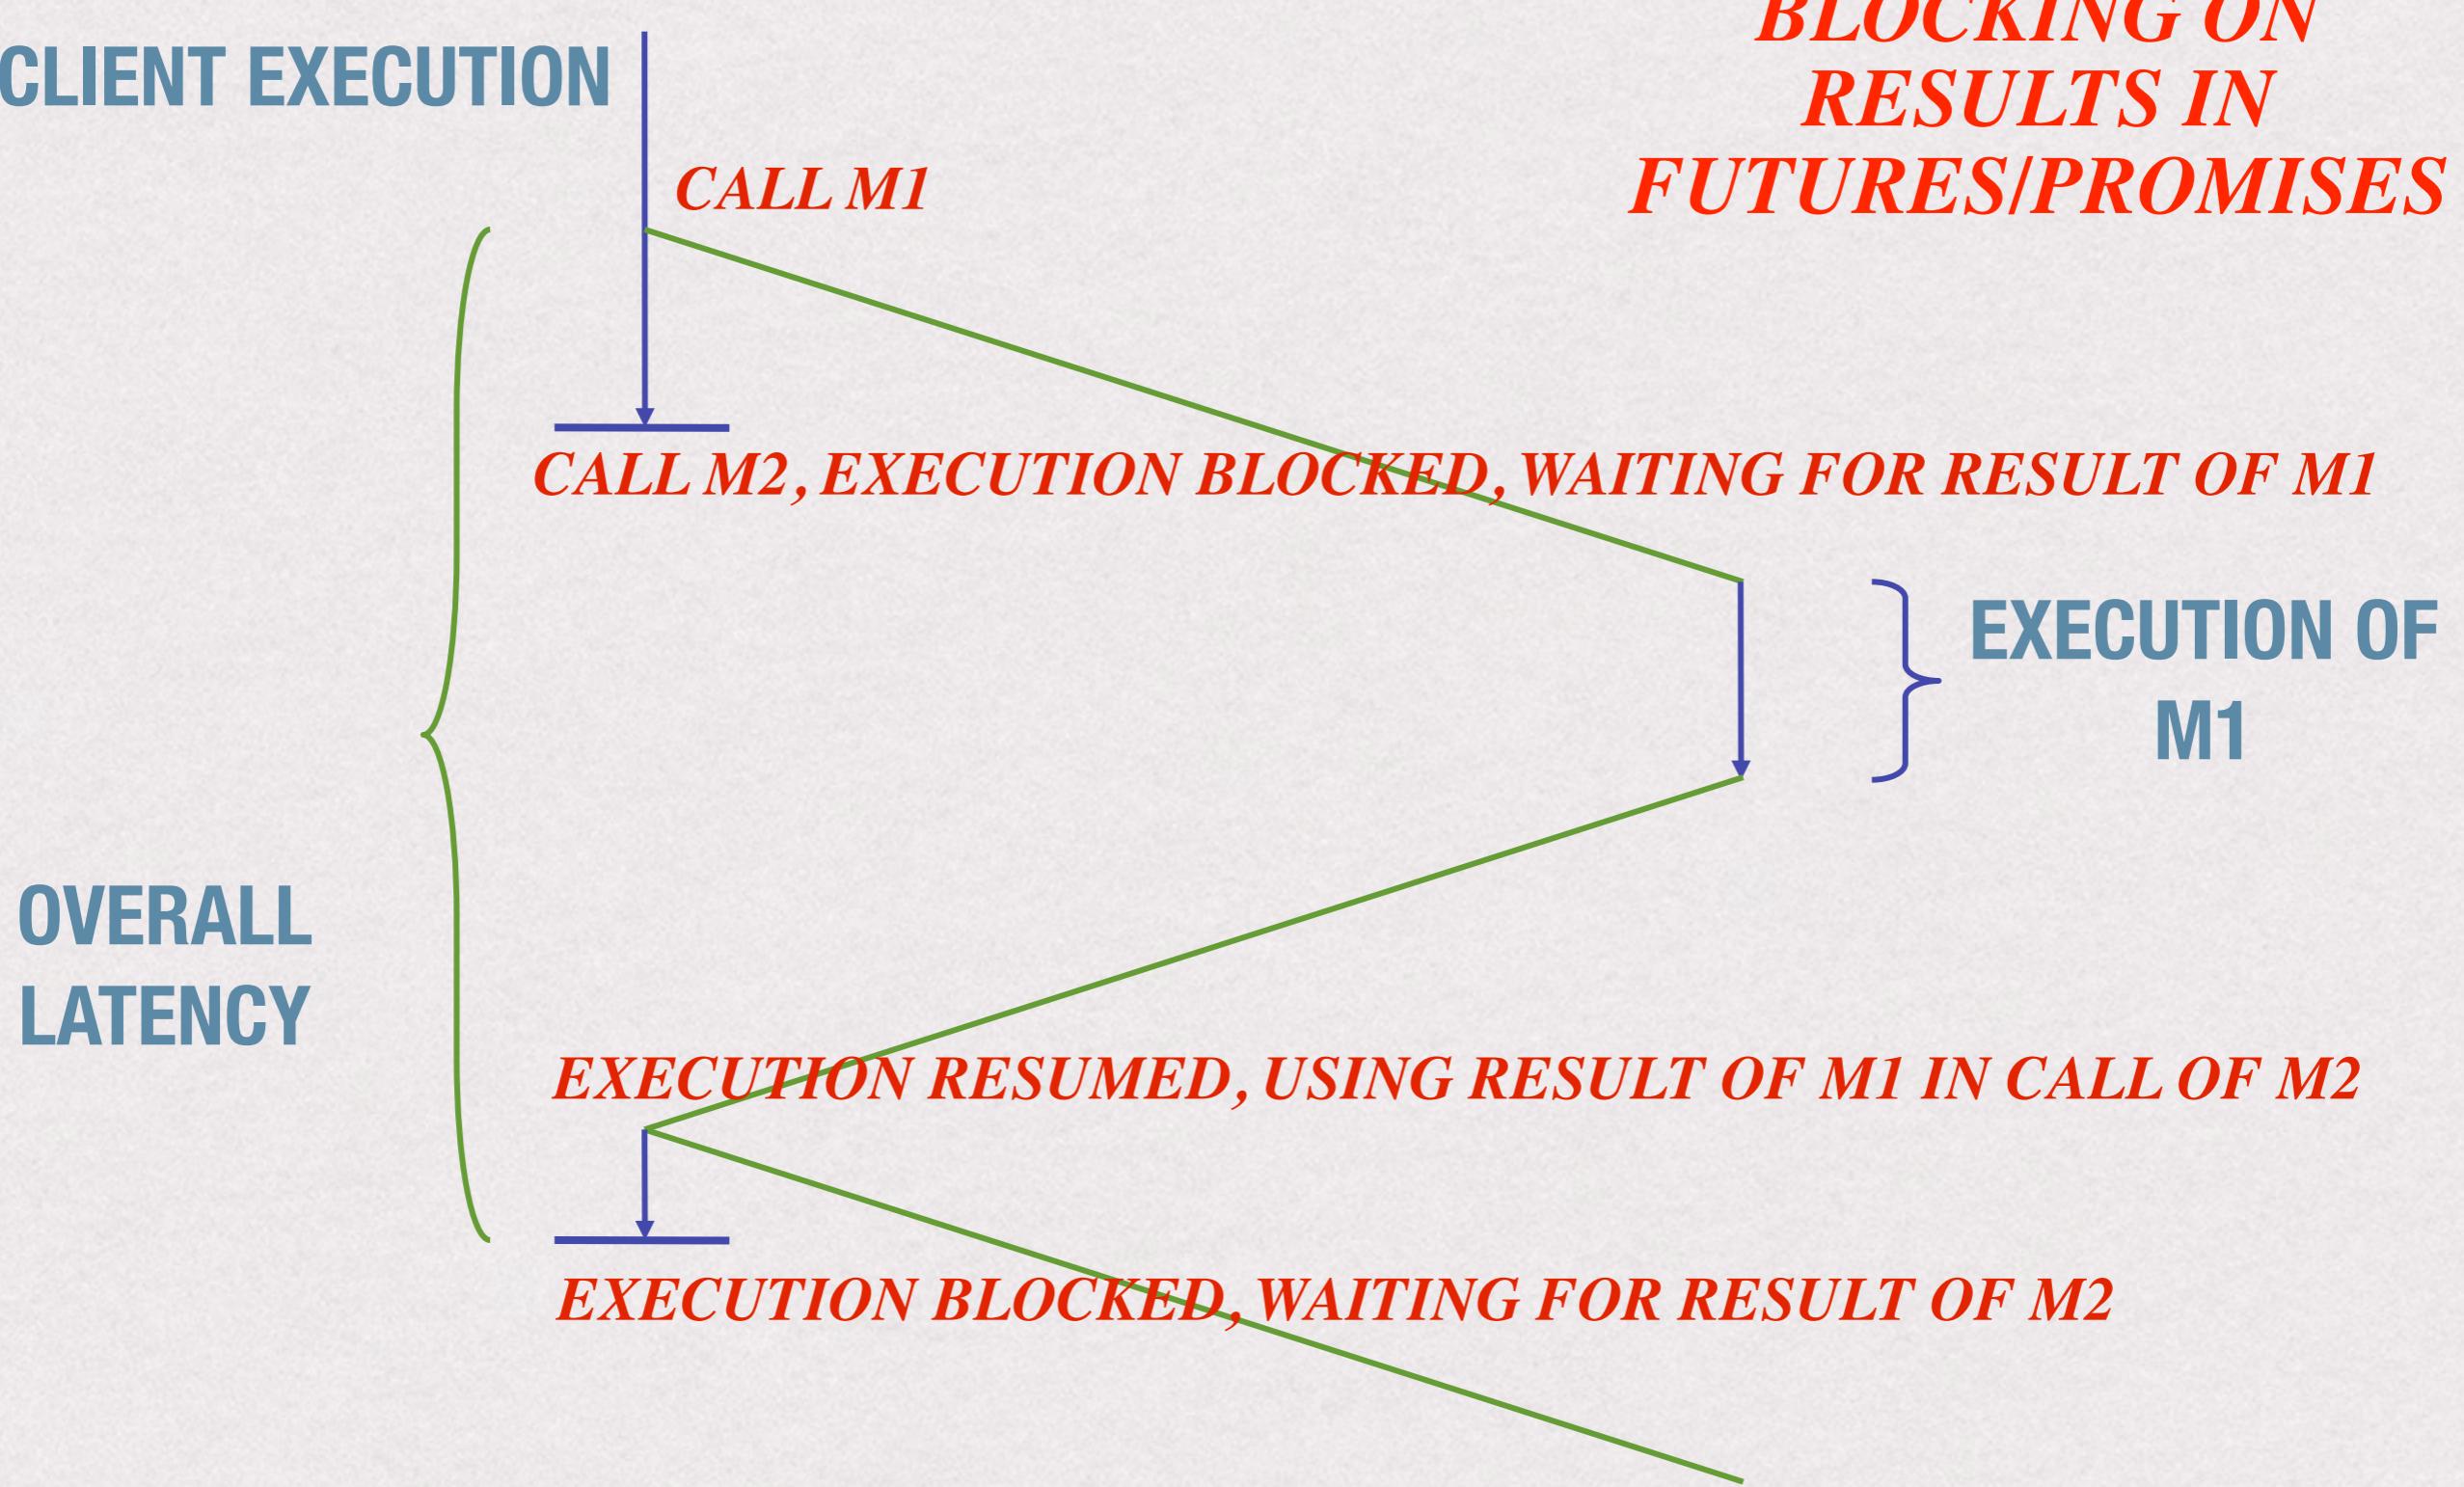
- \* With explicit claims, this must be re-written:

```
res1 = O.m1(...)  
res2 = O.m2(..., claim(res1), ...)
```

- \* With implicit claims, also known as Wait by Necessity, the code can be used as is.
- \* In either case, the claim operation (whether implicit or explicit) may block the client:
  - \* it will block if the result has not arrived, and until it does arrive.

## CLIENT EXECUTION

## BLOCKING ON RESULTS IN FUTURES/PROMISES



# Can we do better than this?



# Latency Reduction with Batched Futures

- \* An example of the Batched Futures idea is:

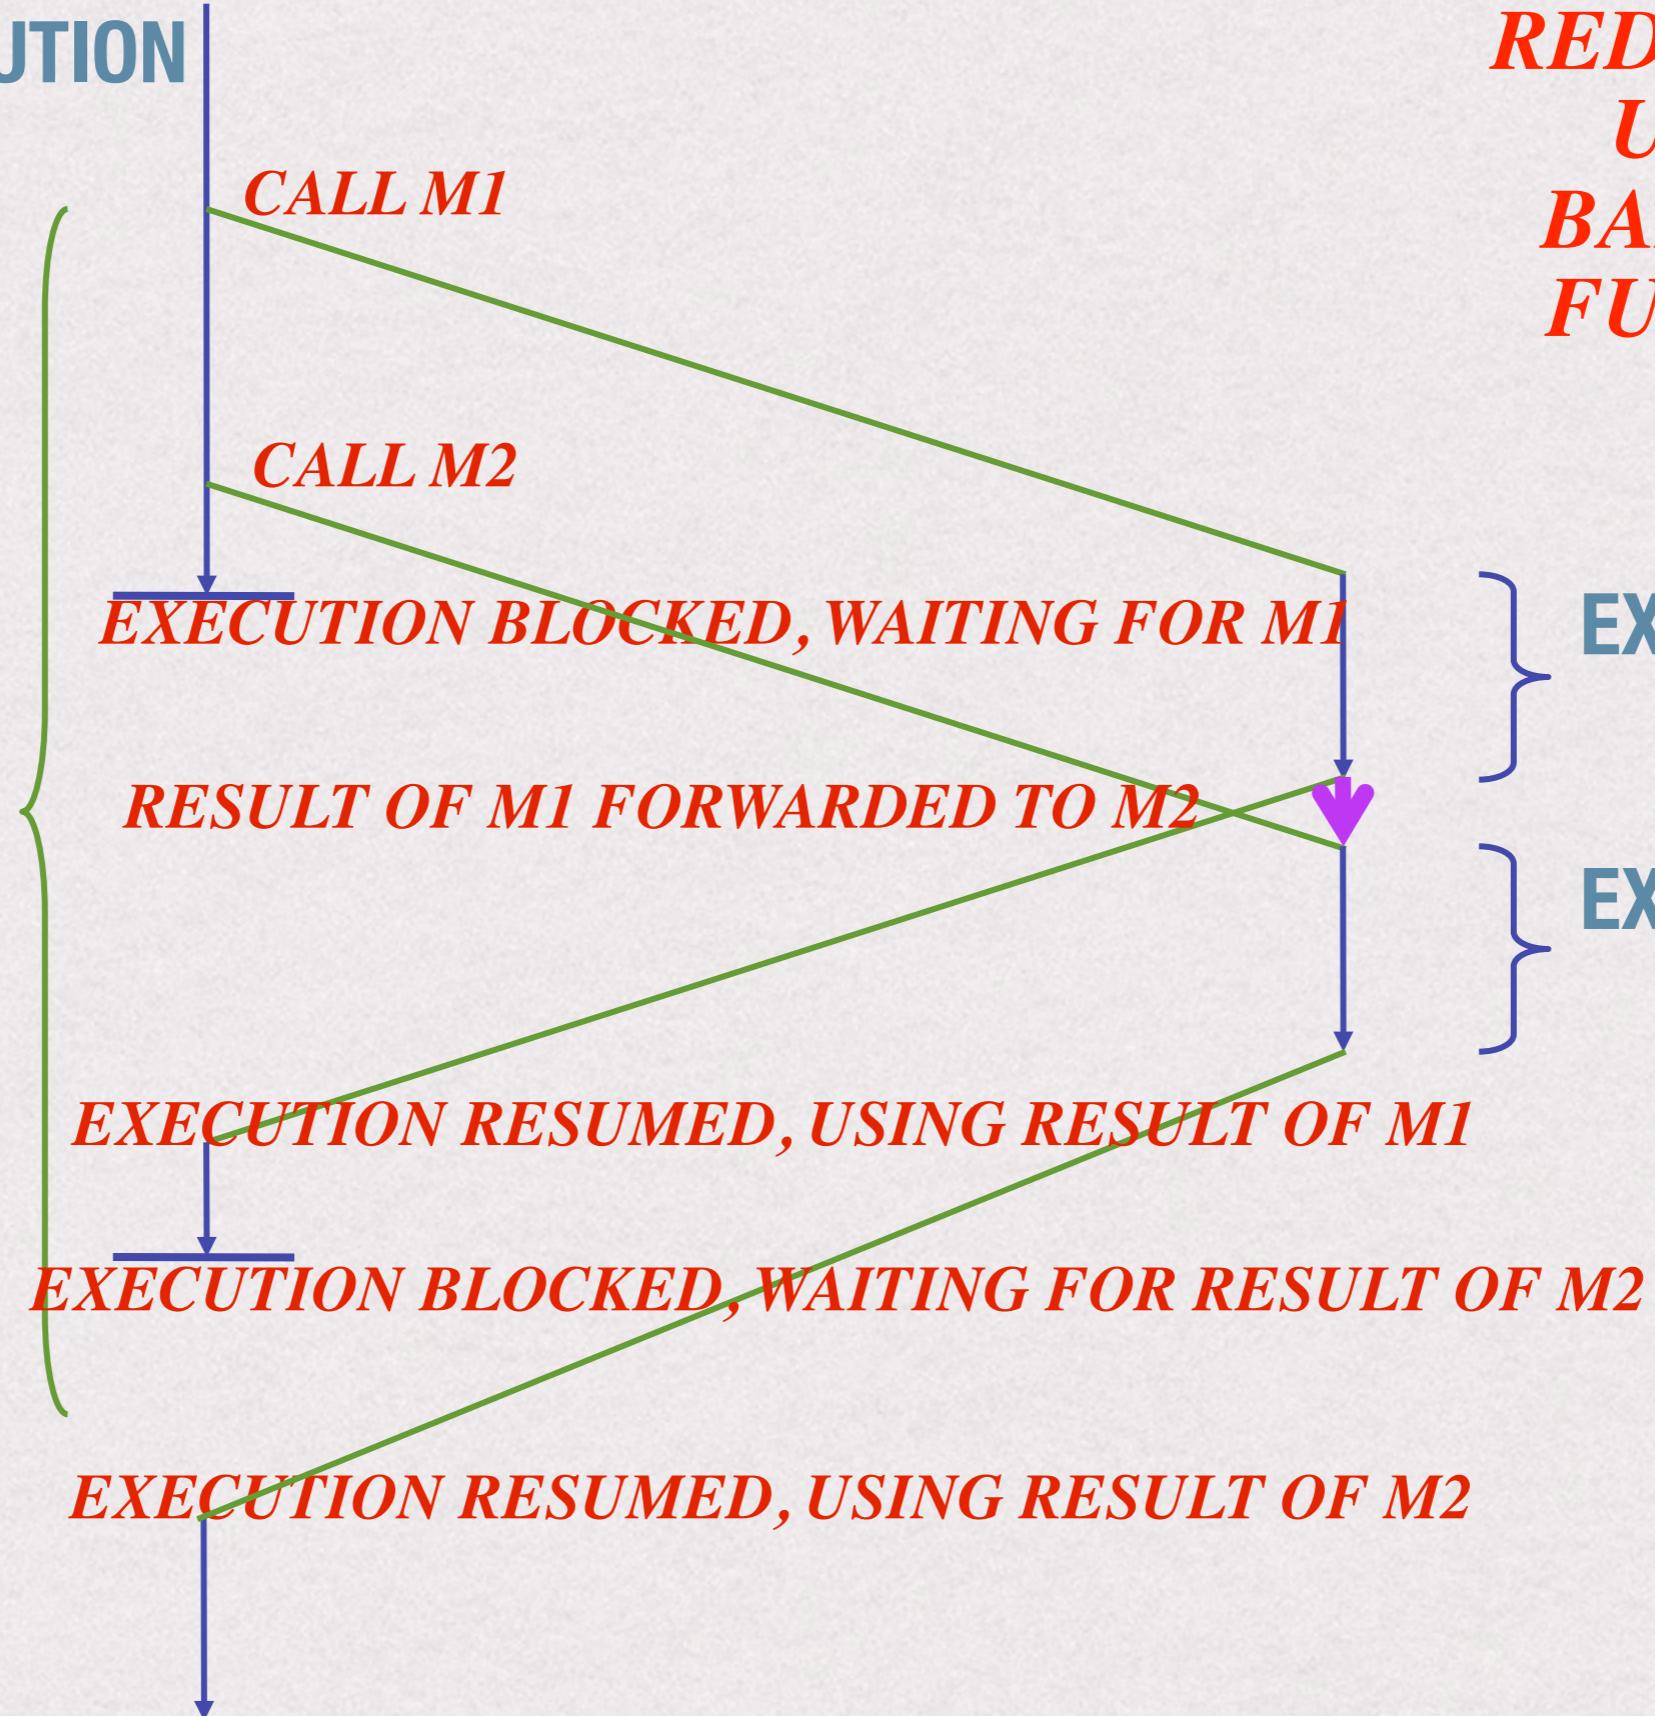
```
res1 = 01.m1(paramsA ...)  
... computation ...  
res2 = 01.m2(..., res1, ...)
```

- \* The result of the first call is able to be used as a parameter of the second call, without the possibility of blocking.

## LATENCY REDUCTION USING BATCHED FUTURES

### CLIENT EXECUTION

### OVERALL LATENCY



# Batched Futures

```
res1 = o1.m1(paramsA ...)  
... computation ...  
res2 = o1.m2(..., res1, ...)
```

- \* What is the longest path (i.e critical path) length in the network propagation graph?
- \* What is the best case for total latency of m1 and m2 considered together?
- \* What is the best case for mean latency of m1 and m2 considered together?
- \* What can we say about the dependencies between m1 and m2?
- \* Using batched futures, must m1 and m2 be calls to object on the same server?

# Latency Reduction with Responsibilities

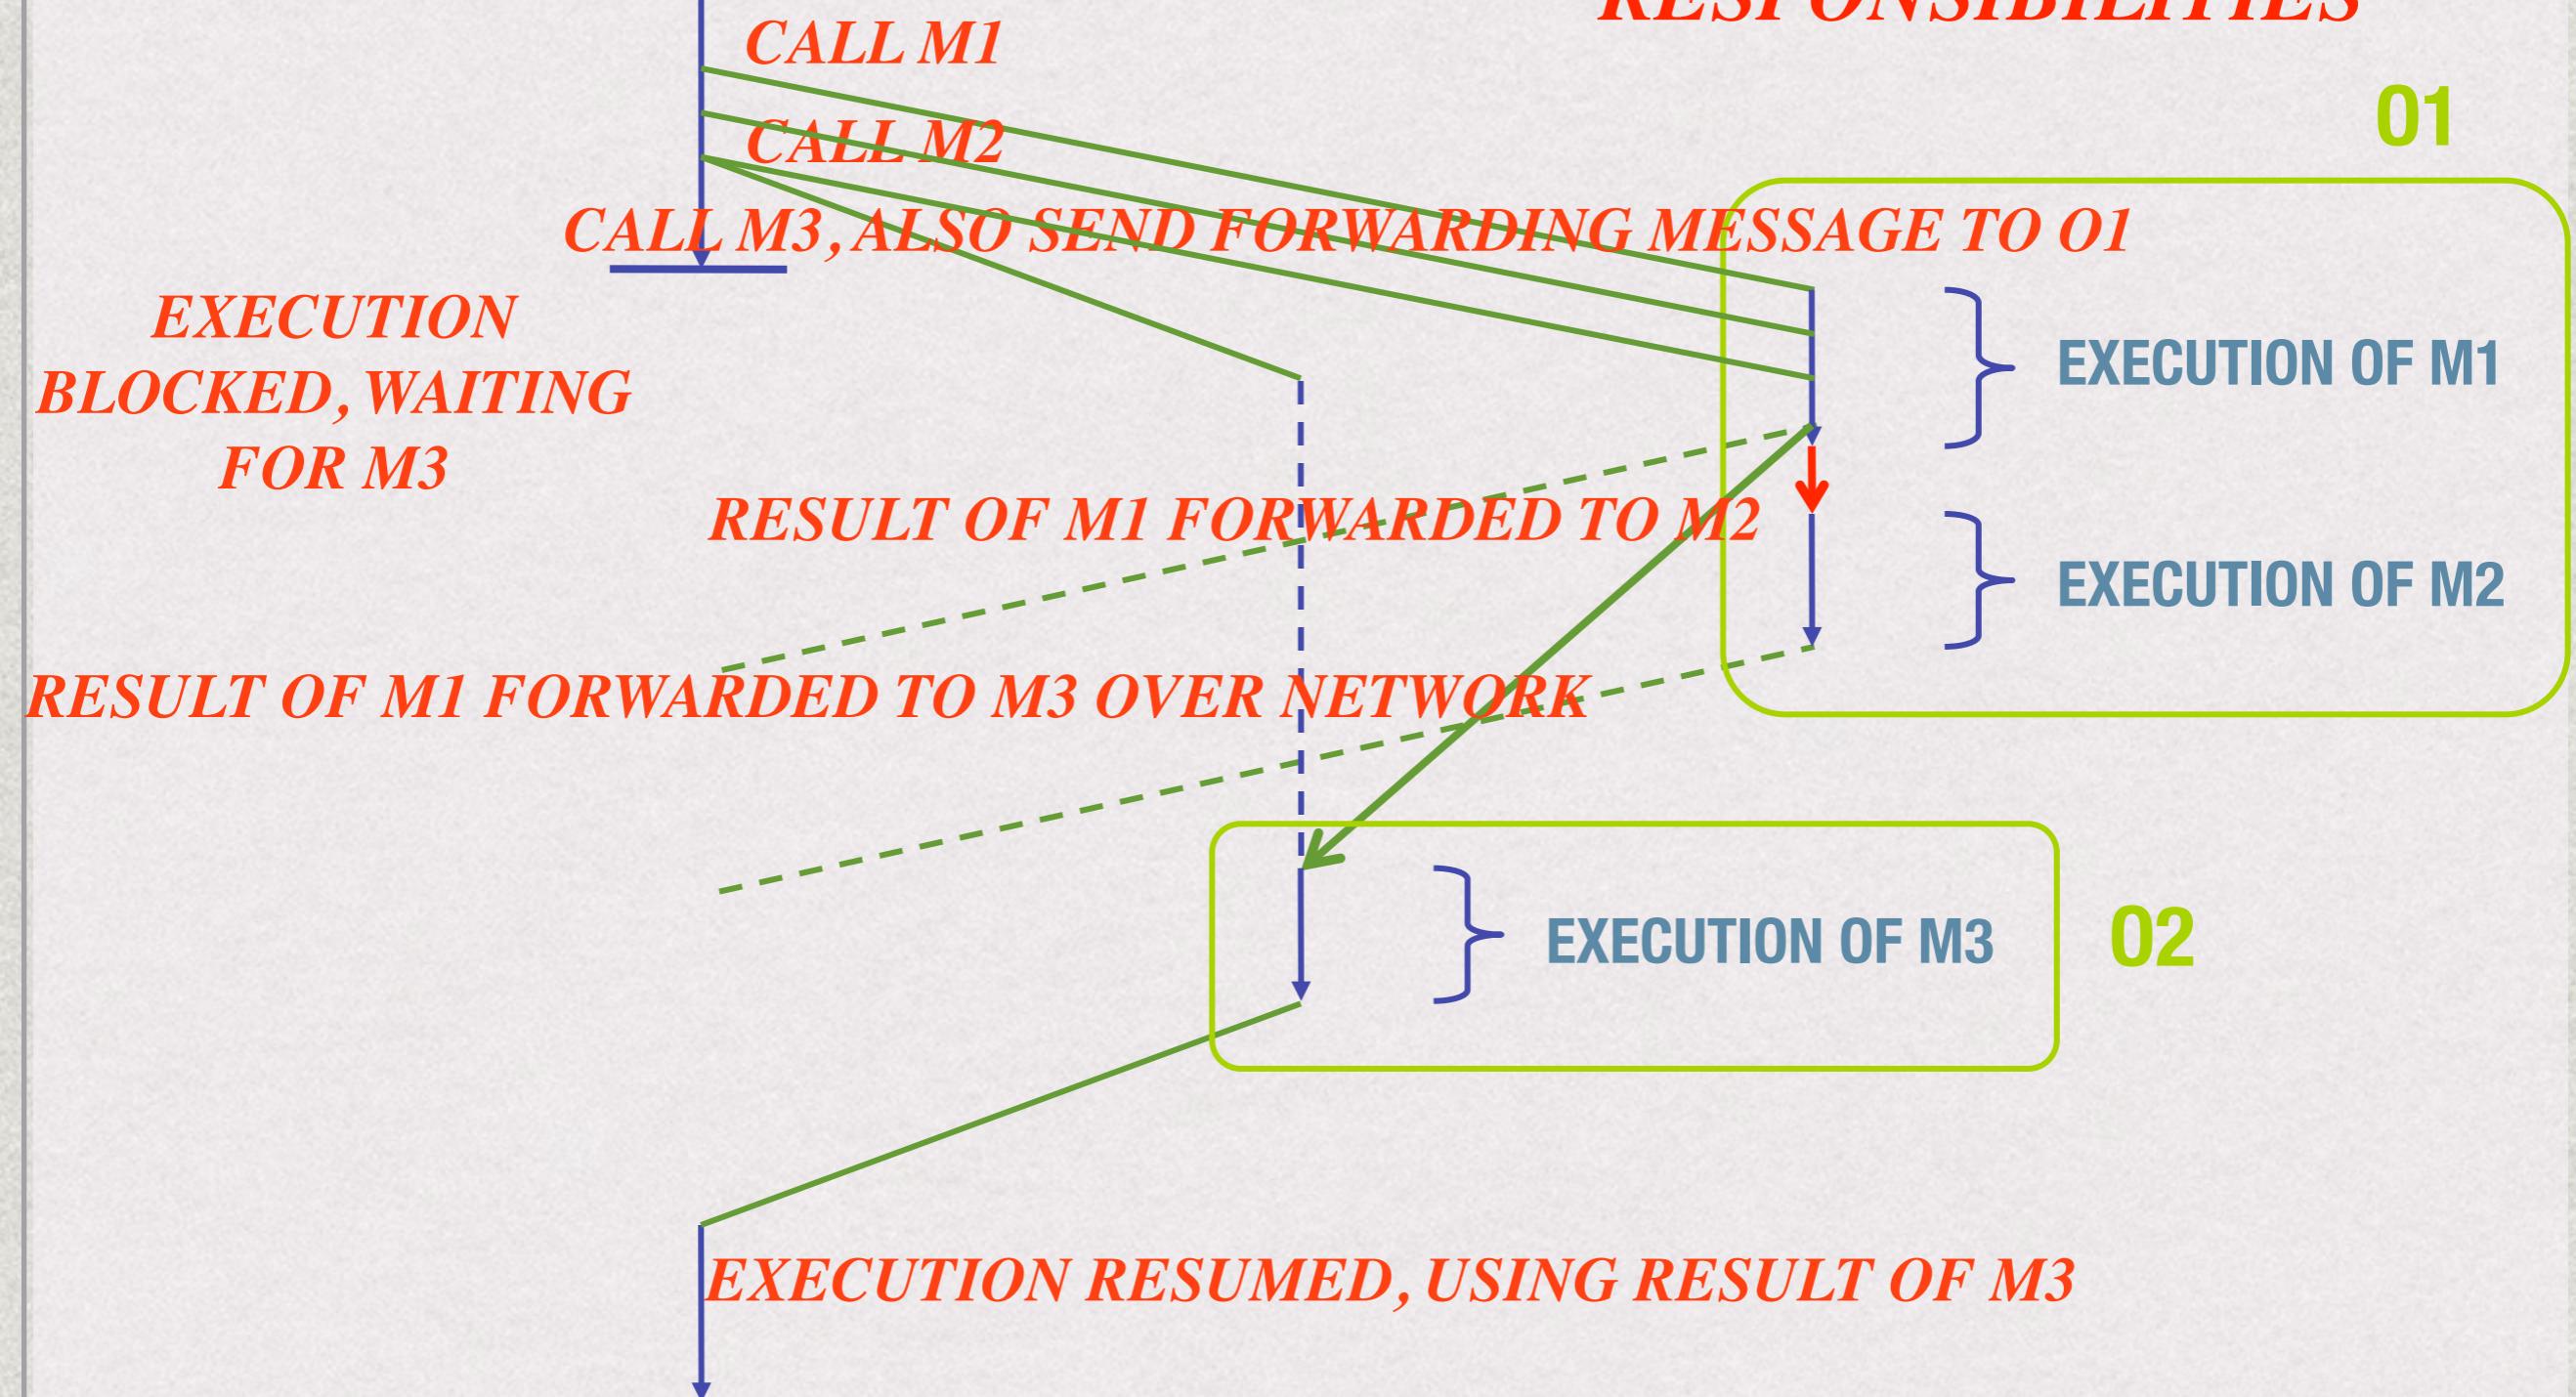
- \* An example of the Responsibilities idea is:

```
res1 = 01.m1(paramsA ...)  
res2 = 01.m2(..., res1, ...)  
res3 = 02.m3(..., res1, ...)  
x = res3 + 5
```

- \* With Batched Futures, the third call will block the client until res1 arrives, since the call is being sent to a different server (object).
- \* With Responsibilities, the third call is also non-blocking.

## CLIENT EXECUTION

## LATENCY REDUCTION USING RESPONSIBILITIES



# Responsibilities

```
res1 = 01.m1(paramsA ...)  
res2 = 01.m2(..., res1, ...)  
res3 = 02.m3(..., res1, ...)  
x = res3 + 5
```

- \* What is the longest path length through the network propagation graph?
- \* What is the best case for total latency of m1, m2 and m3 considered together?
- \* What would it be for batched futures?
- \* What is the best case for mean latency of m1, m2 and m3 considered together?
- \* What would it be for batched futures?

# Analyse That!

- \* One factor in requiring less blocking is the kinds of dependencies between calls that can be traversed without blocking:
  - \* Result dependent < Order dependent < Independent
- \* Another factor is whether dependencies between calls to different servers can be traversed without blocking:
  - \* Multi server < Single server

# Types of Dependencies

- \* Calls m1 and m2:
- \* ***Independent*** – m1 and m2 can execute in any order and the result of either is not used in the other.
- \* ***Order Dependent*** – m2 must begin execution after the end of the execution of m1, but m2 does not use the result of m1.
- \* ***Result Dependent*** – m2 takes as parameters one or more results of m1, and hence must begin execution after the end of the execution of m1:
  - \* Res1 = O.m1(...)
  - Res2 = O.m2(..., res1, ...)

# Types of Dependencies

- \* ***Functionally Dependent*** – m2 takes as parameters one or more non-identity functions of one or more results of m1,
- \* must begin execution after the end of the execution of m1:
- \*  $\text{Res1} = \text{O.m1}(\dots)$   
 $\text{Res2} = \text{O.m2}(\dots, \text{res1} + 3, \dots)$