

DISTRIBUTED SYSTEMS

INTRODUCTION TO FAULT TOLERANCE

Past lectures ...

- * Remote operations
 - * Latency reduction
 - * Pcall, one-way RPC, futures and promises, batched futures, responsibilities, ...
 - * Linguistic heterogeneity
 - * CORBA IDL
 - * XML

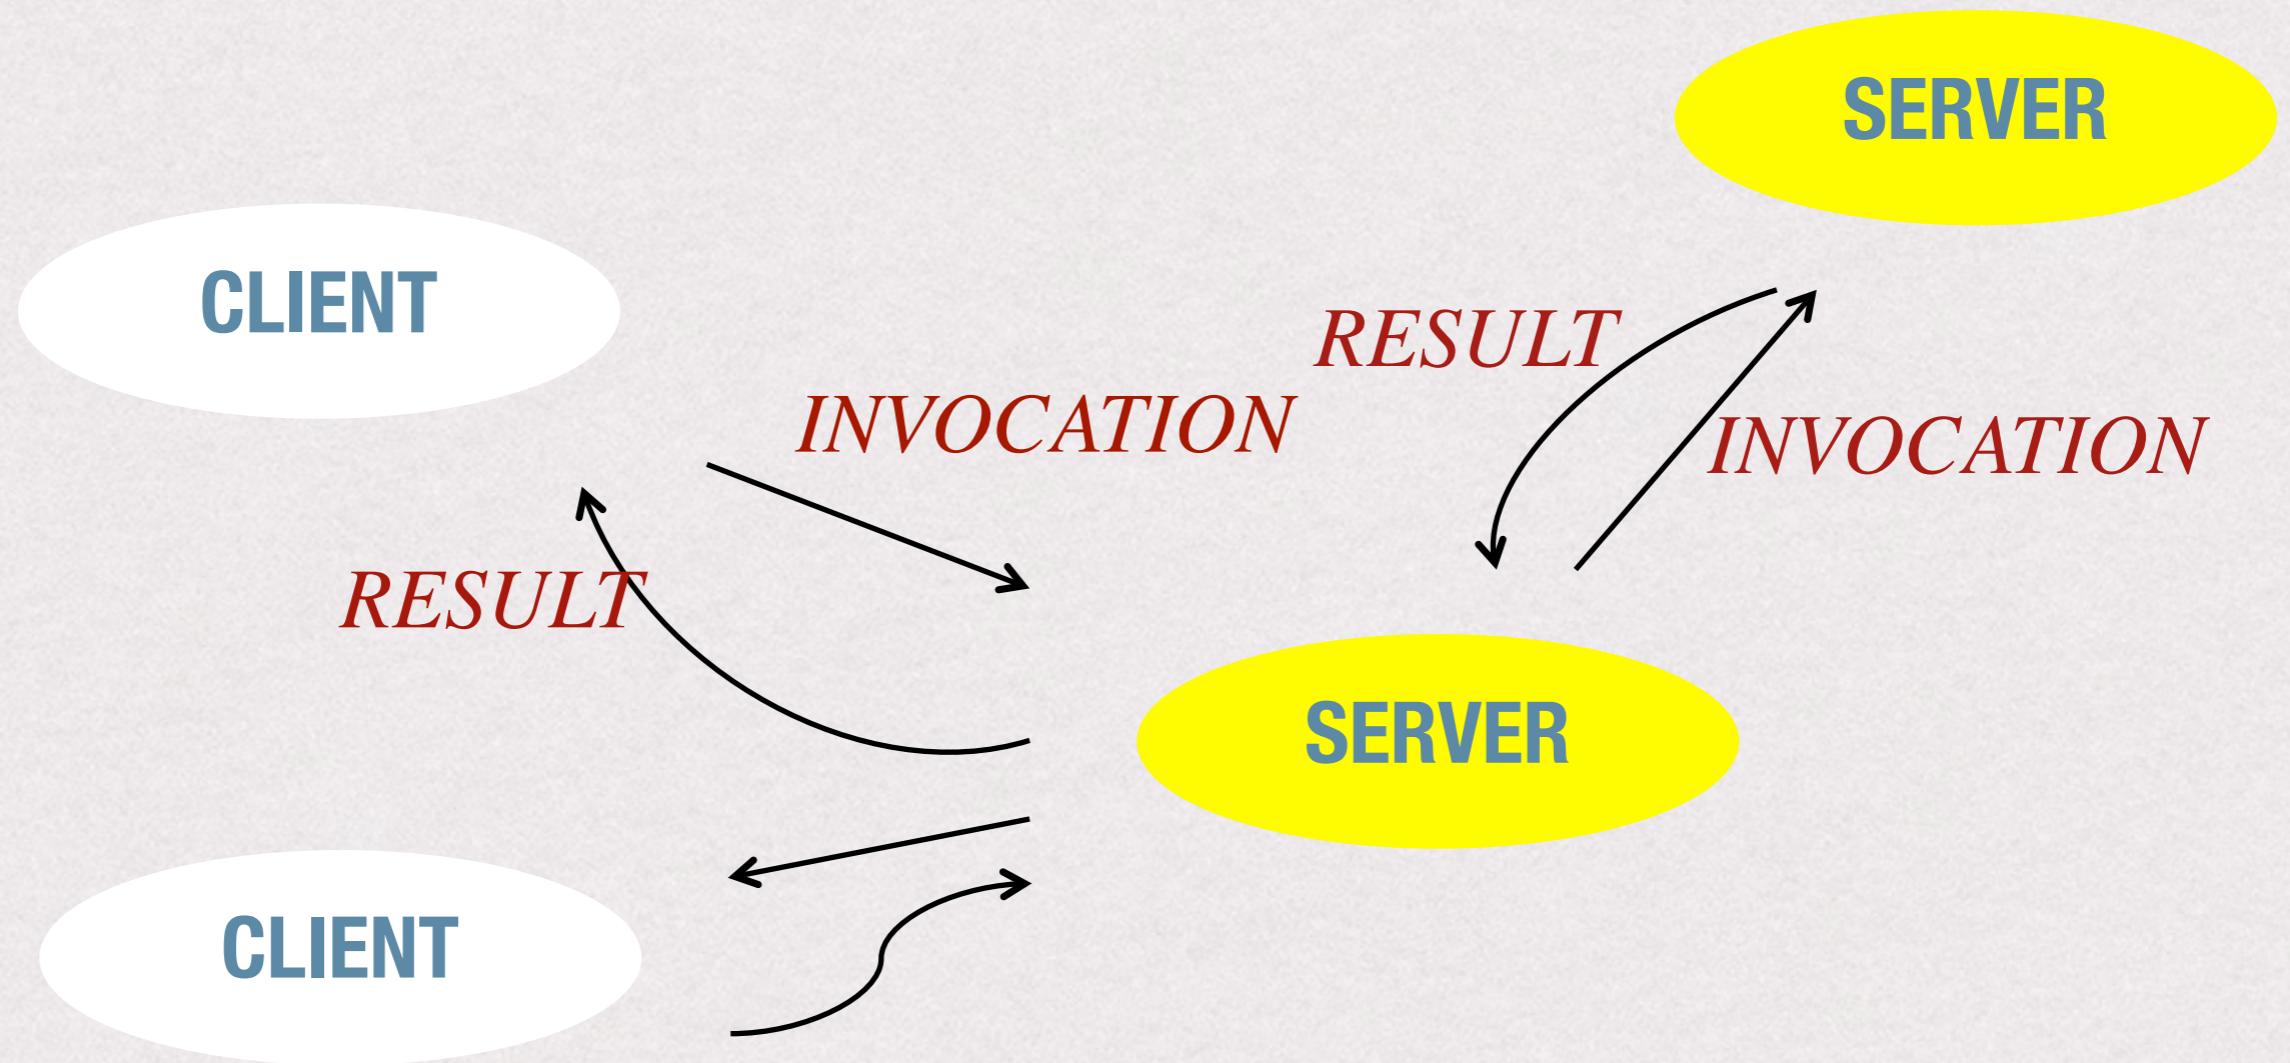
Architectural models

- ✳ Distribution of responsibility
 - ✳ Client-server
 - ✳ Peer-to-peer

Client - Server

- * Servers manage a set of resources or provide a set of services
- * Clients interact with servers to obtain access from resources or result from services
- * Servers can be clients; many clients can interact with many servers

Client - Server



Peer-to-Peer

- * All of the processes play similar roles, interacting cooperatively as equal peers
- * No distinction between peers
- * Scales well as peers are responsible for parts of a service
- * Sometimes, provides better service if there is a large number of peers
- * E.g. BT

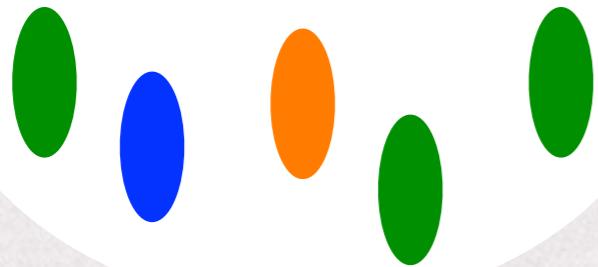
Peer-to-Peer

PEER 1

PEER 3

APPLICATION

APPLICATION



PEER 2

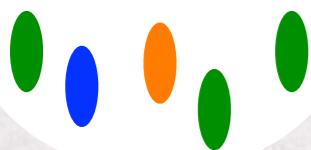
APPLICATION

PEER 4

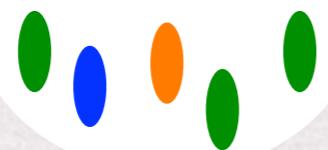
PEER 5

PEER 6

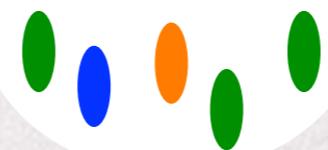
APPLICATION



APPLICATION



APPLICATION



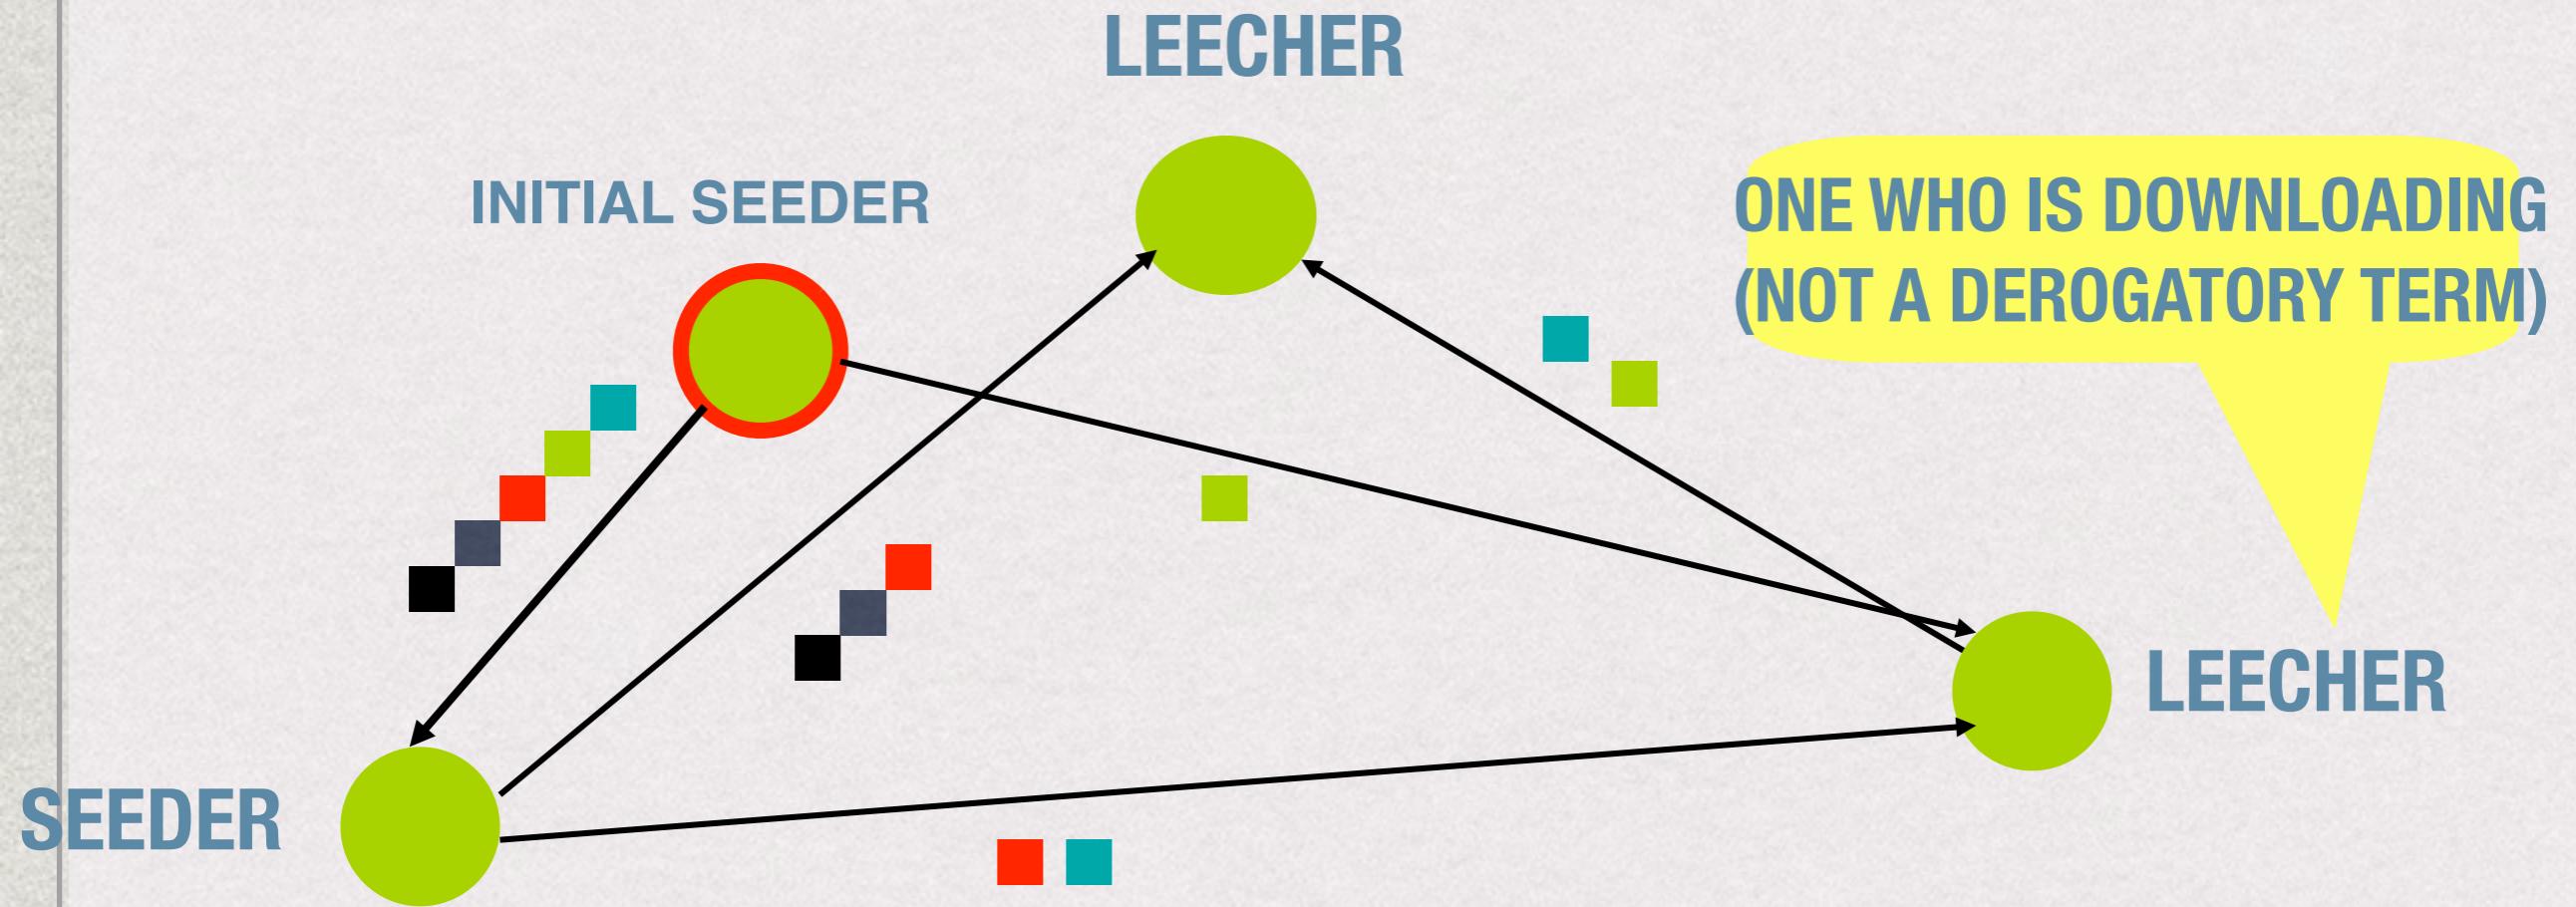
BitTorrent

- * Protocol for transmitting large files
- * User that wants to distribute a file creates a torrent
- * User makes file available to a seed
- * File is divided into segments: pieces
- * As each peer receives a piece for a file, it becomes a seed for that piece
- * Pieces are downloaded non-sequentially

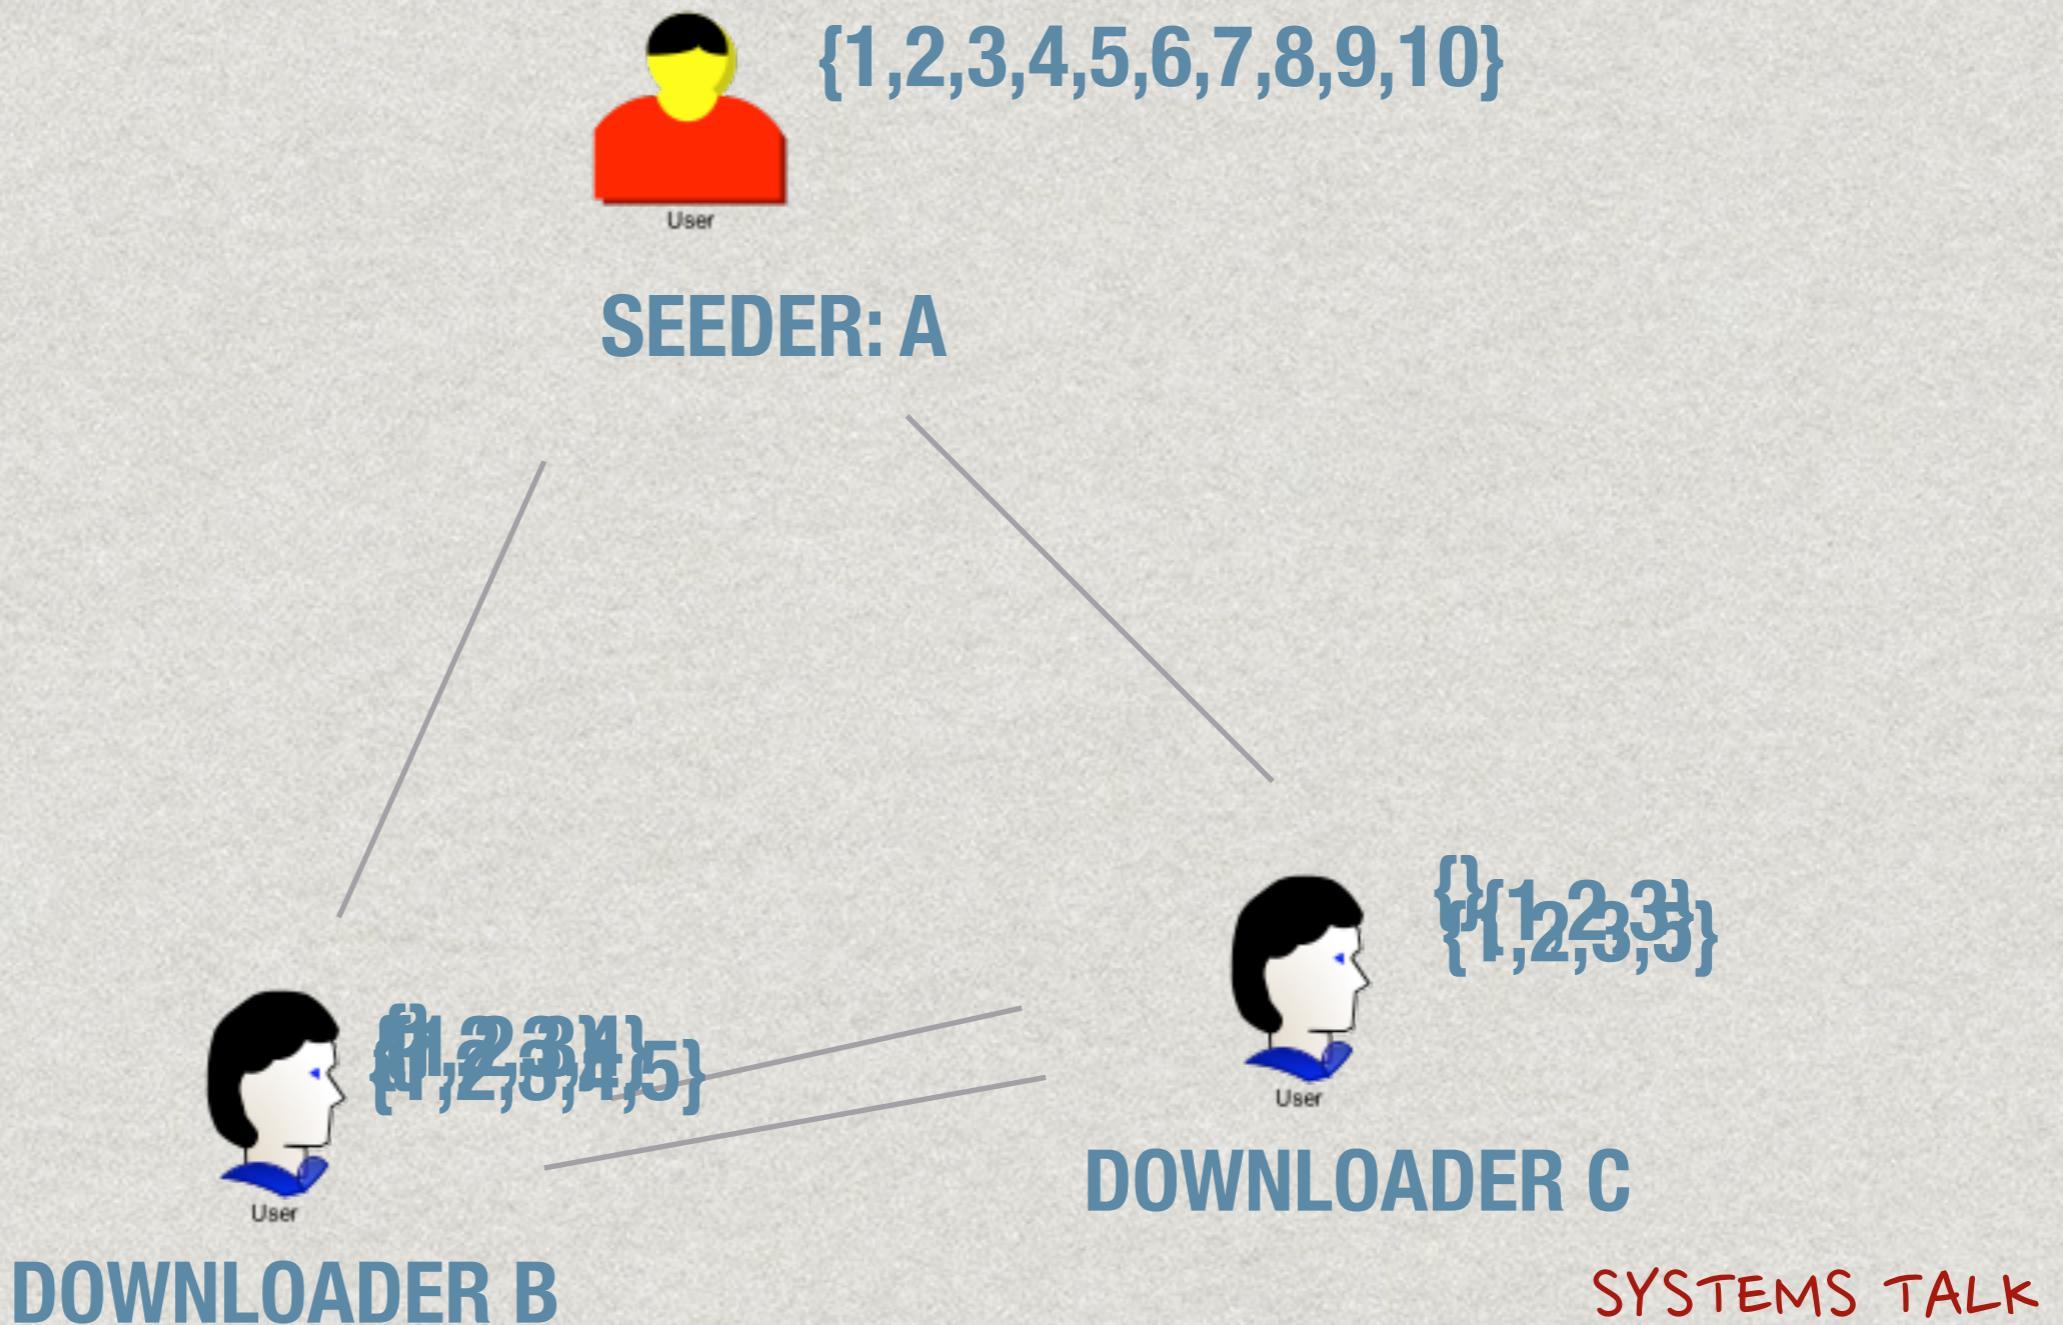
BitTorrent Lingo

SEEDER = A PEER THAT PROVIDES THE COMPLETE FILE.

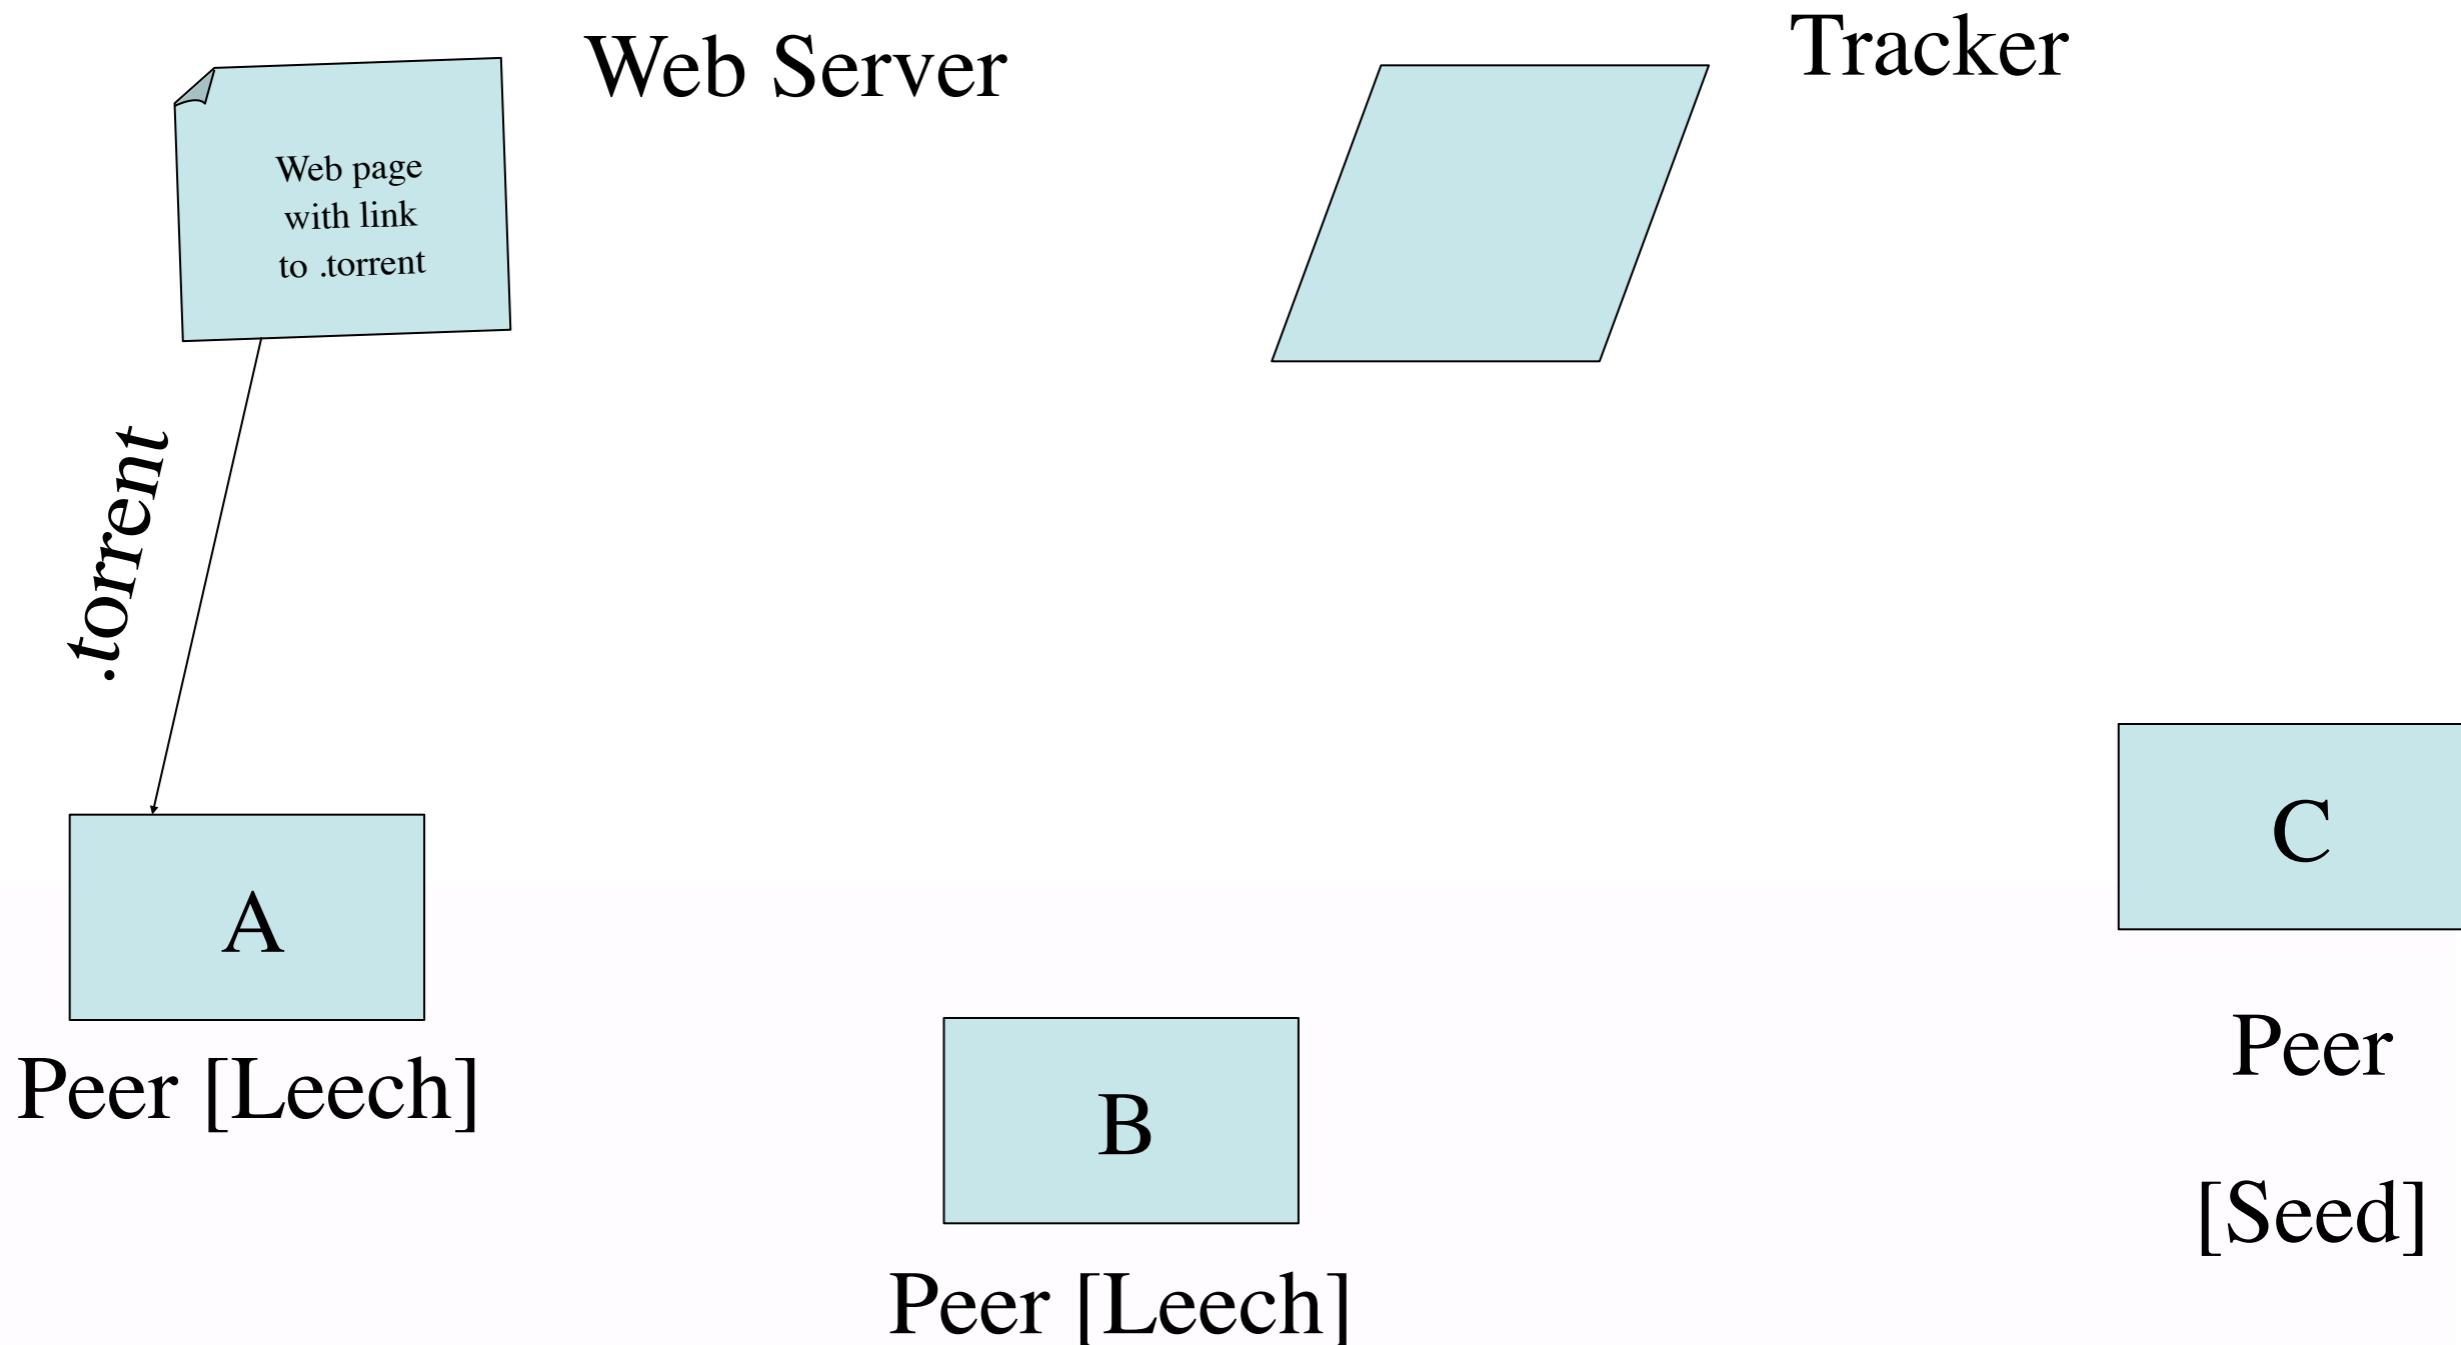
INITIAL SEEDER = A PEER THAT PROVIDES THE INITIAL COPY.



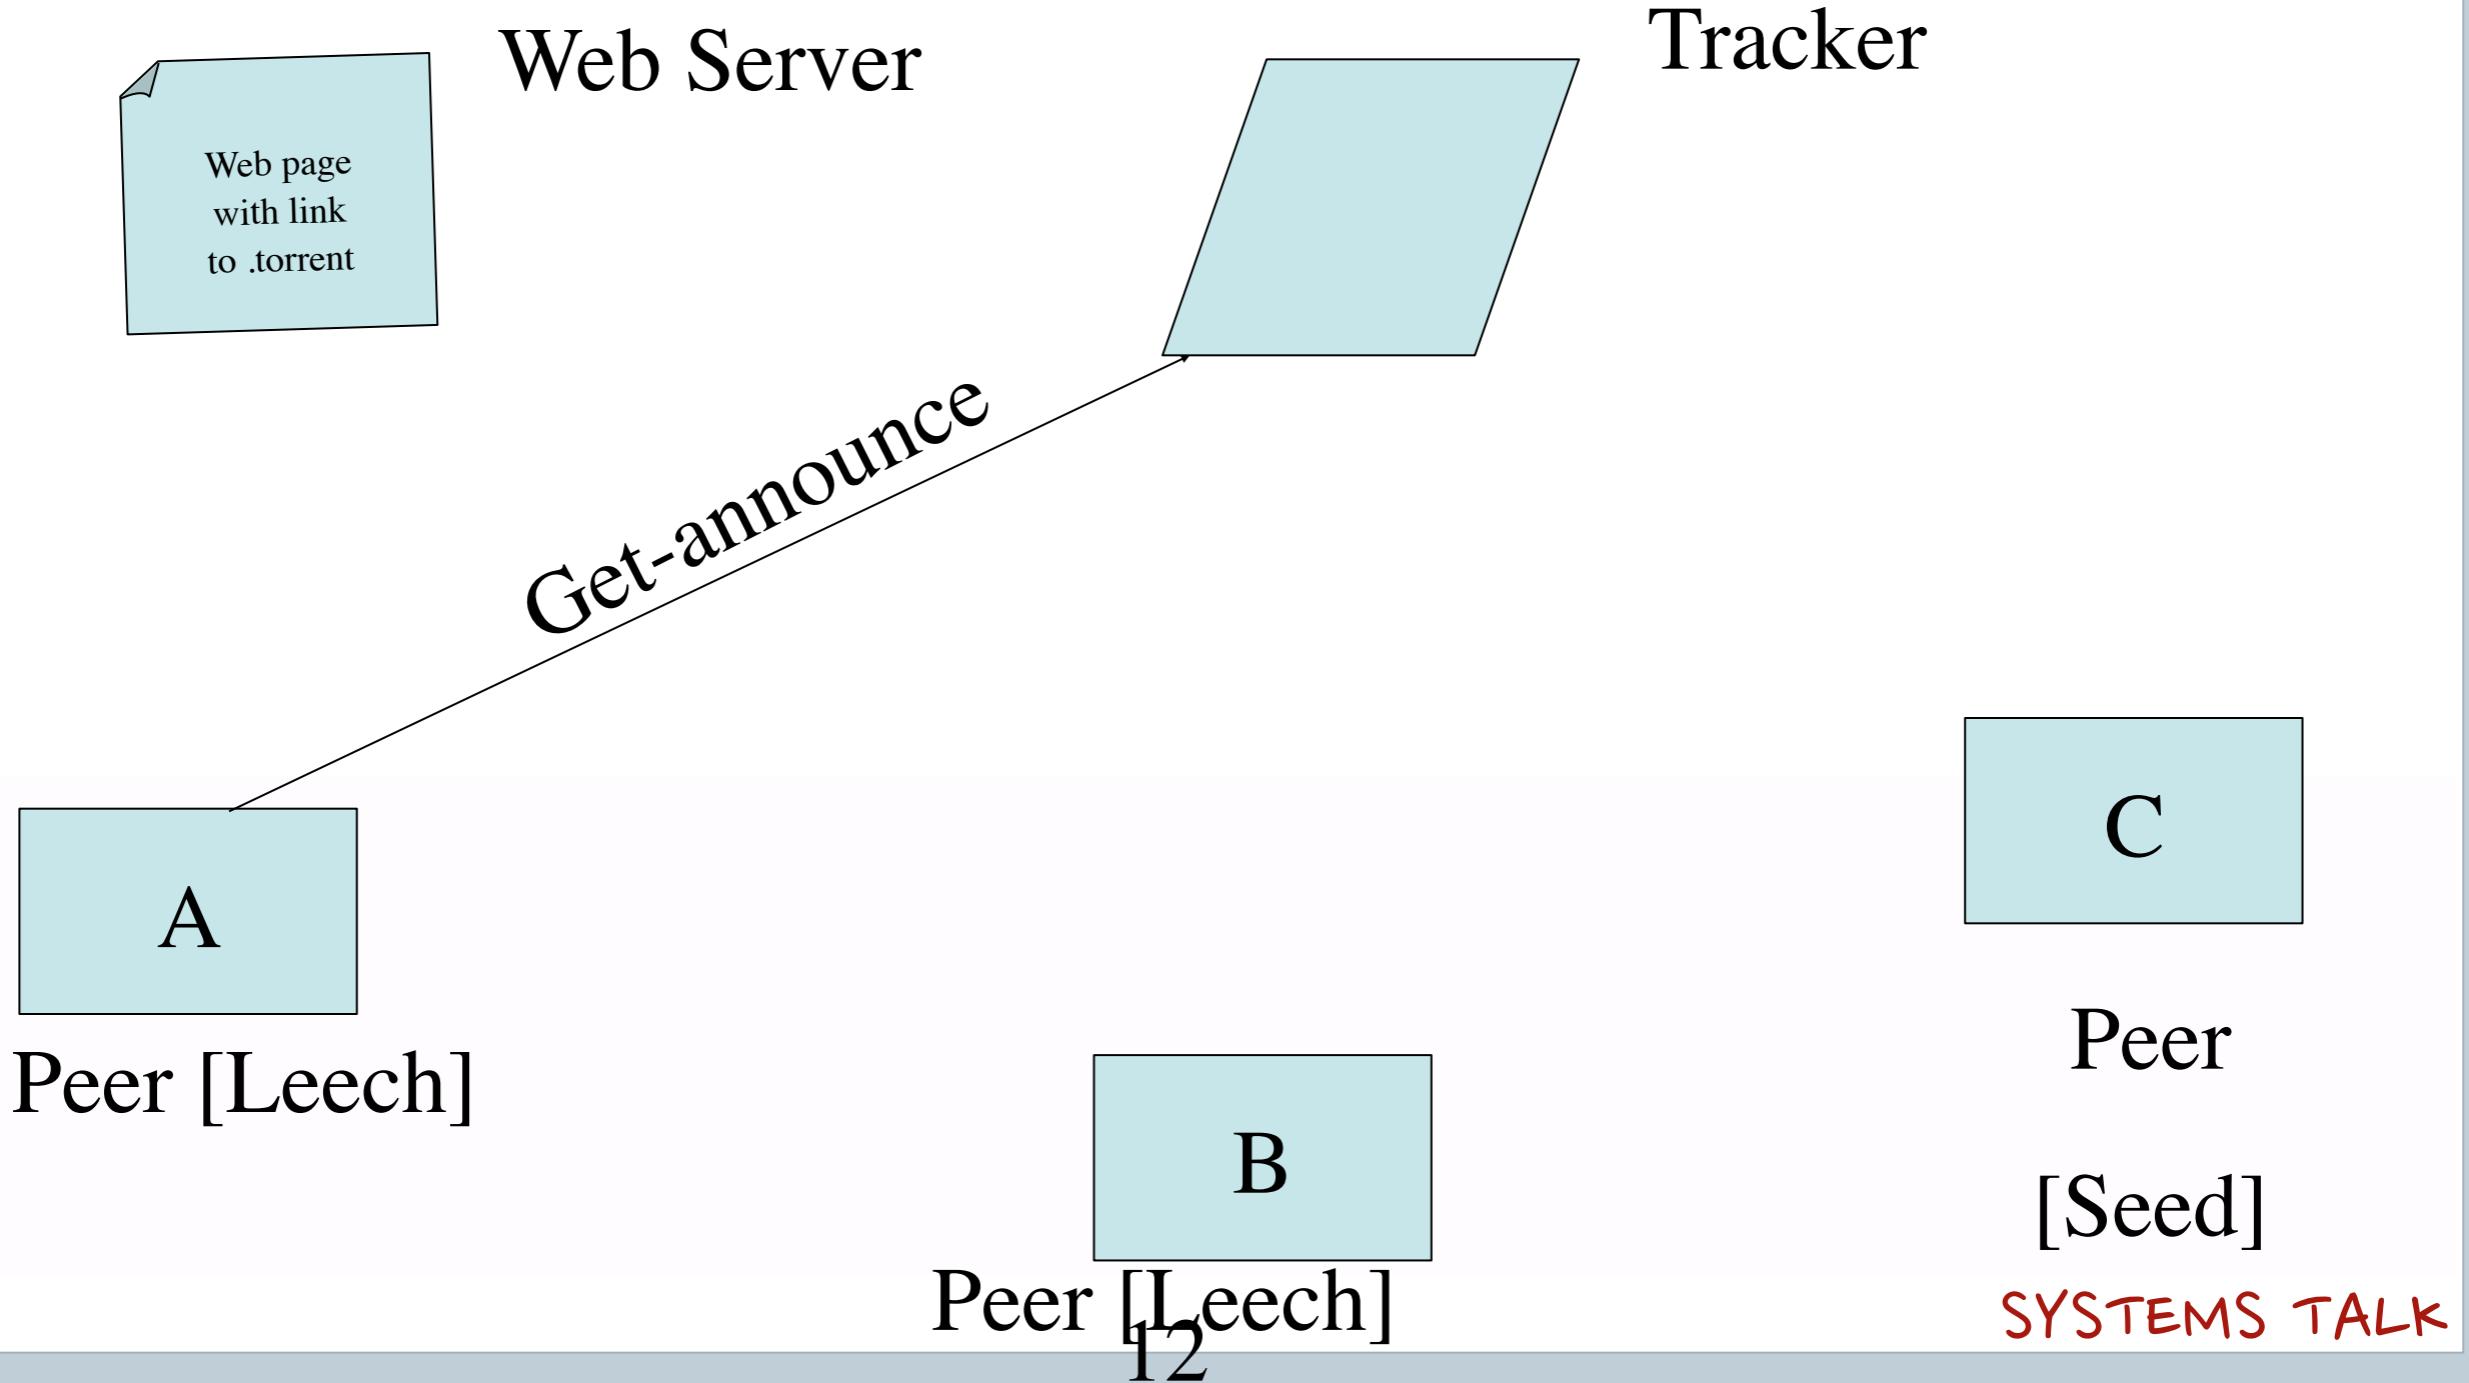
Simple example



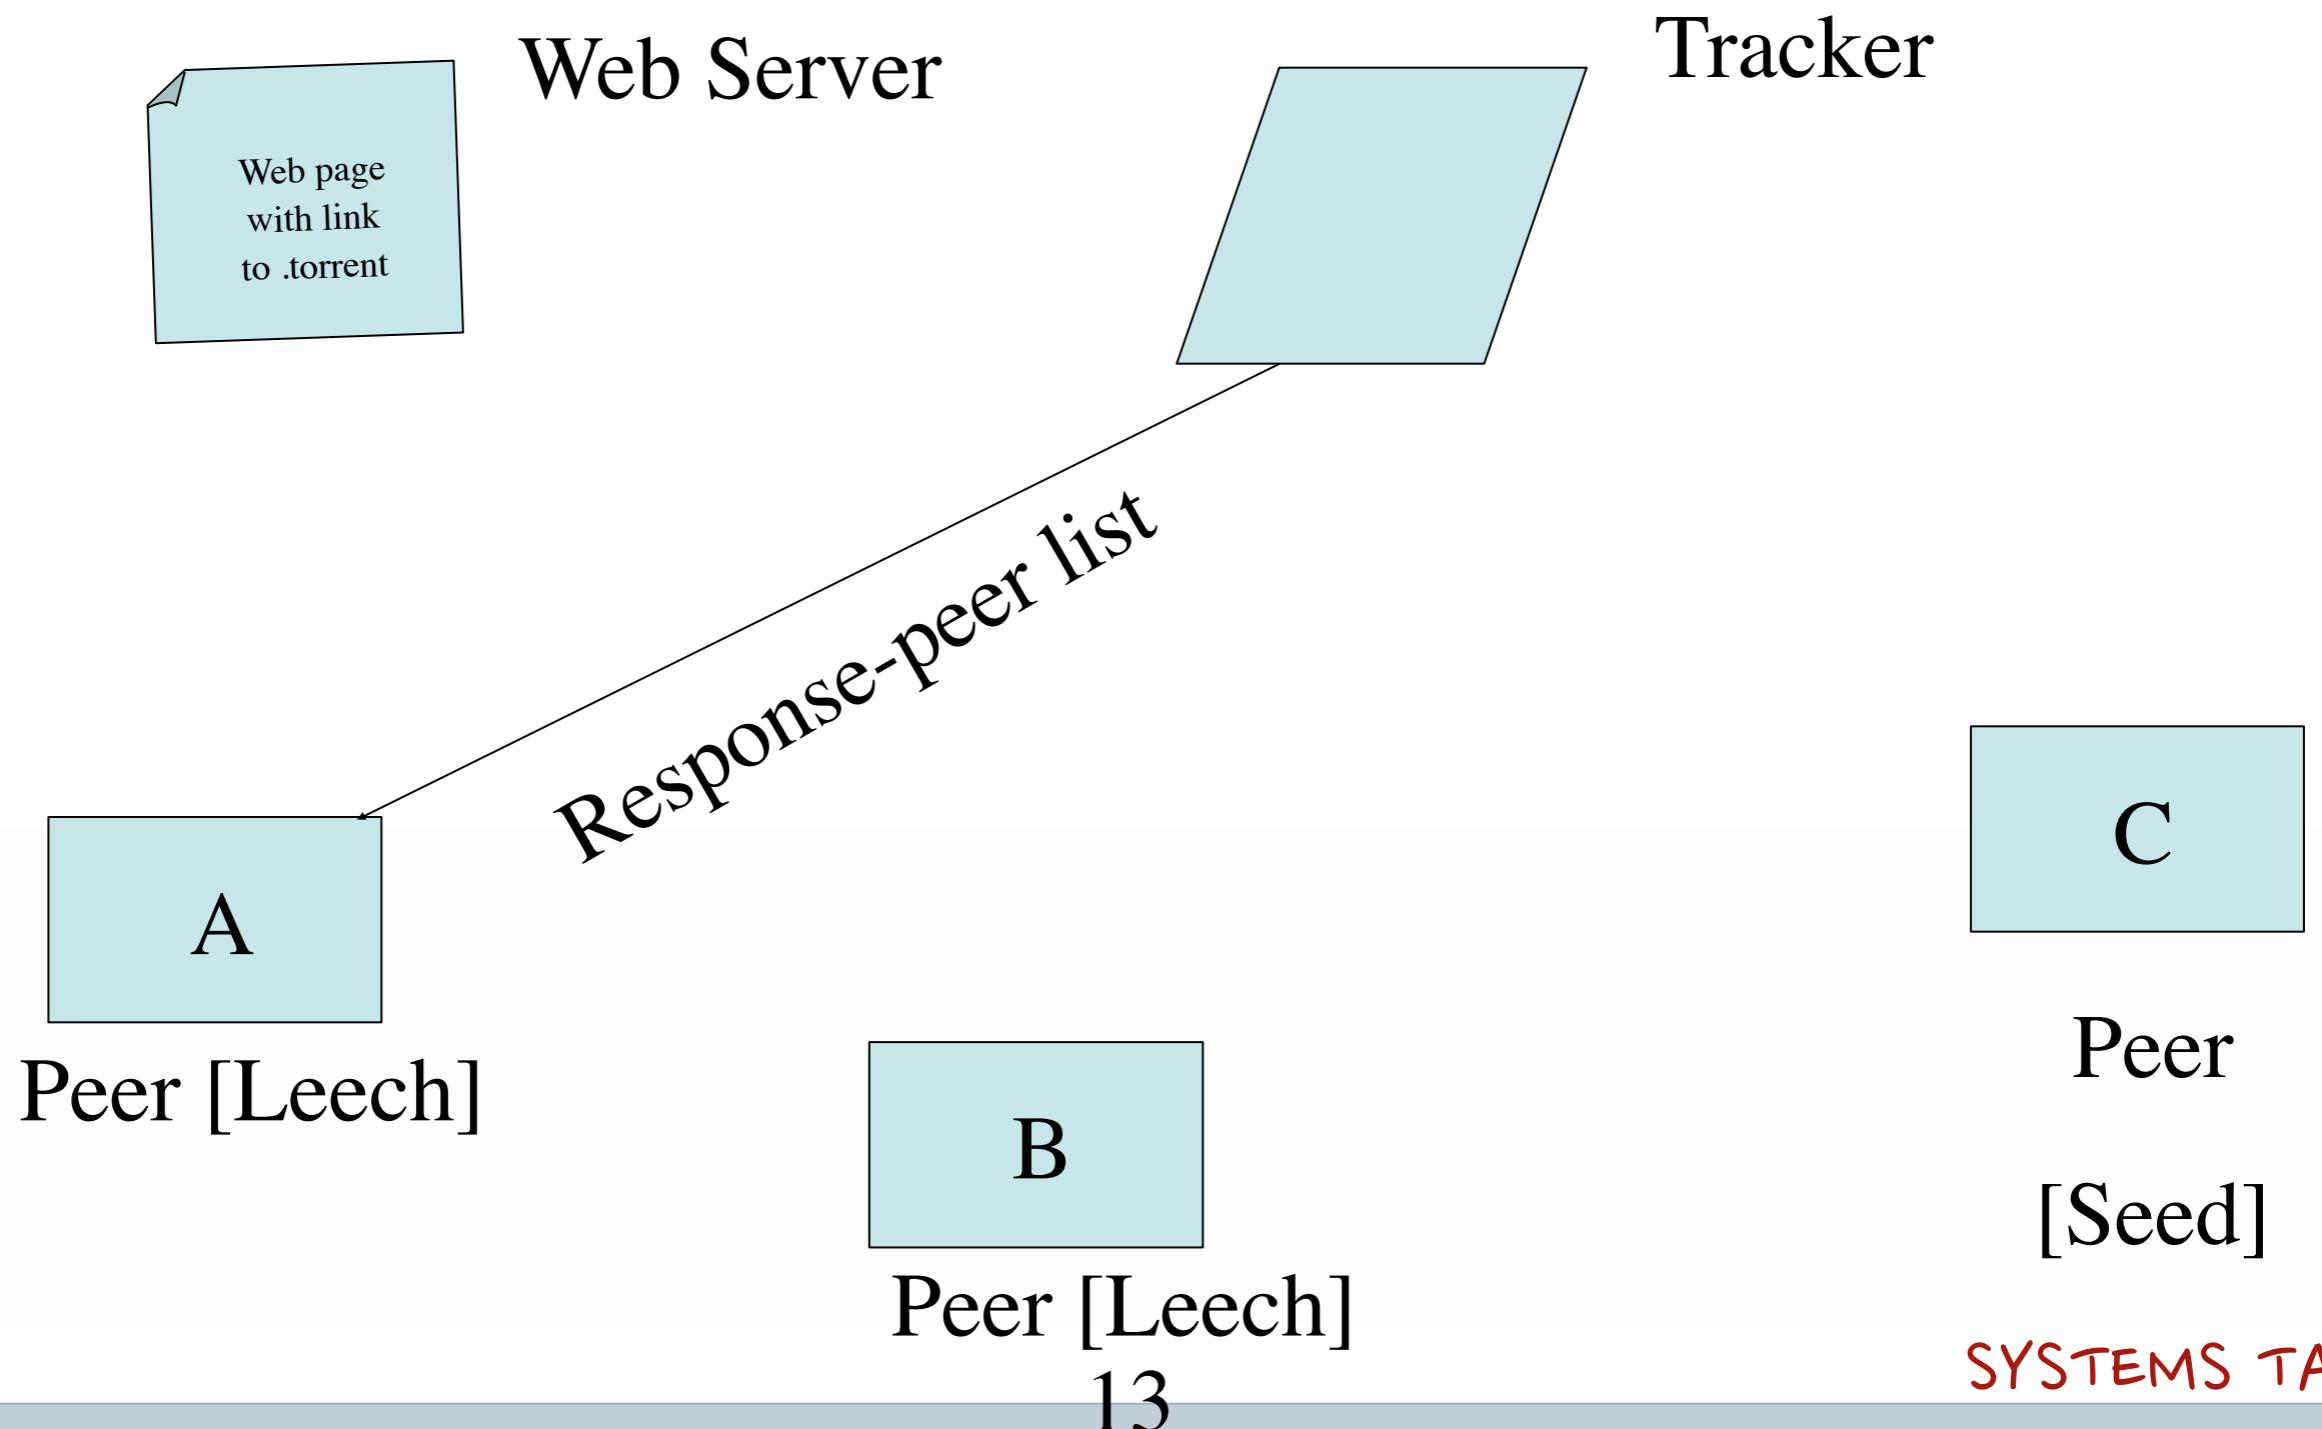
Overview



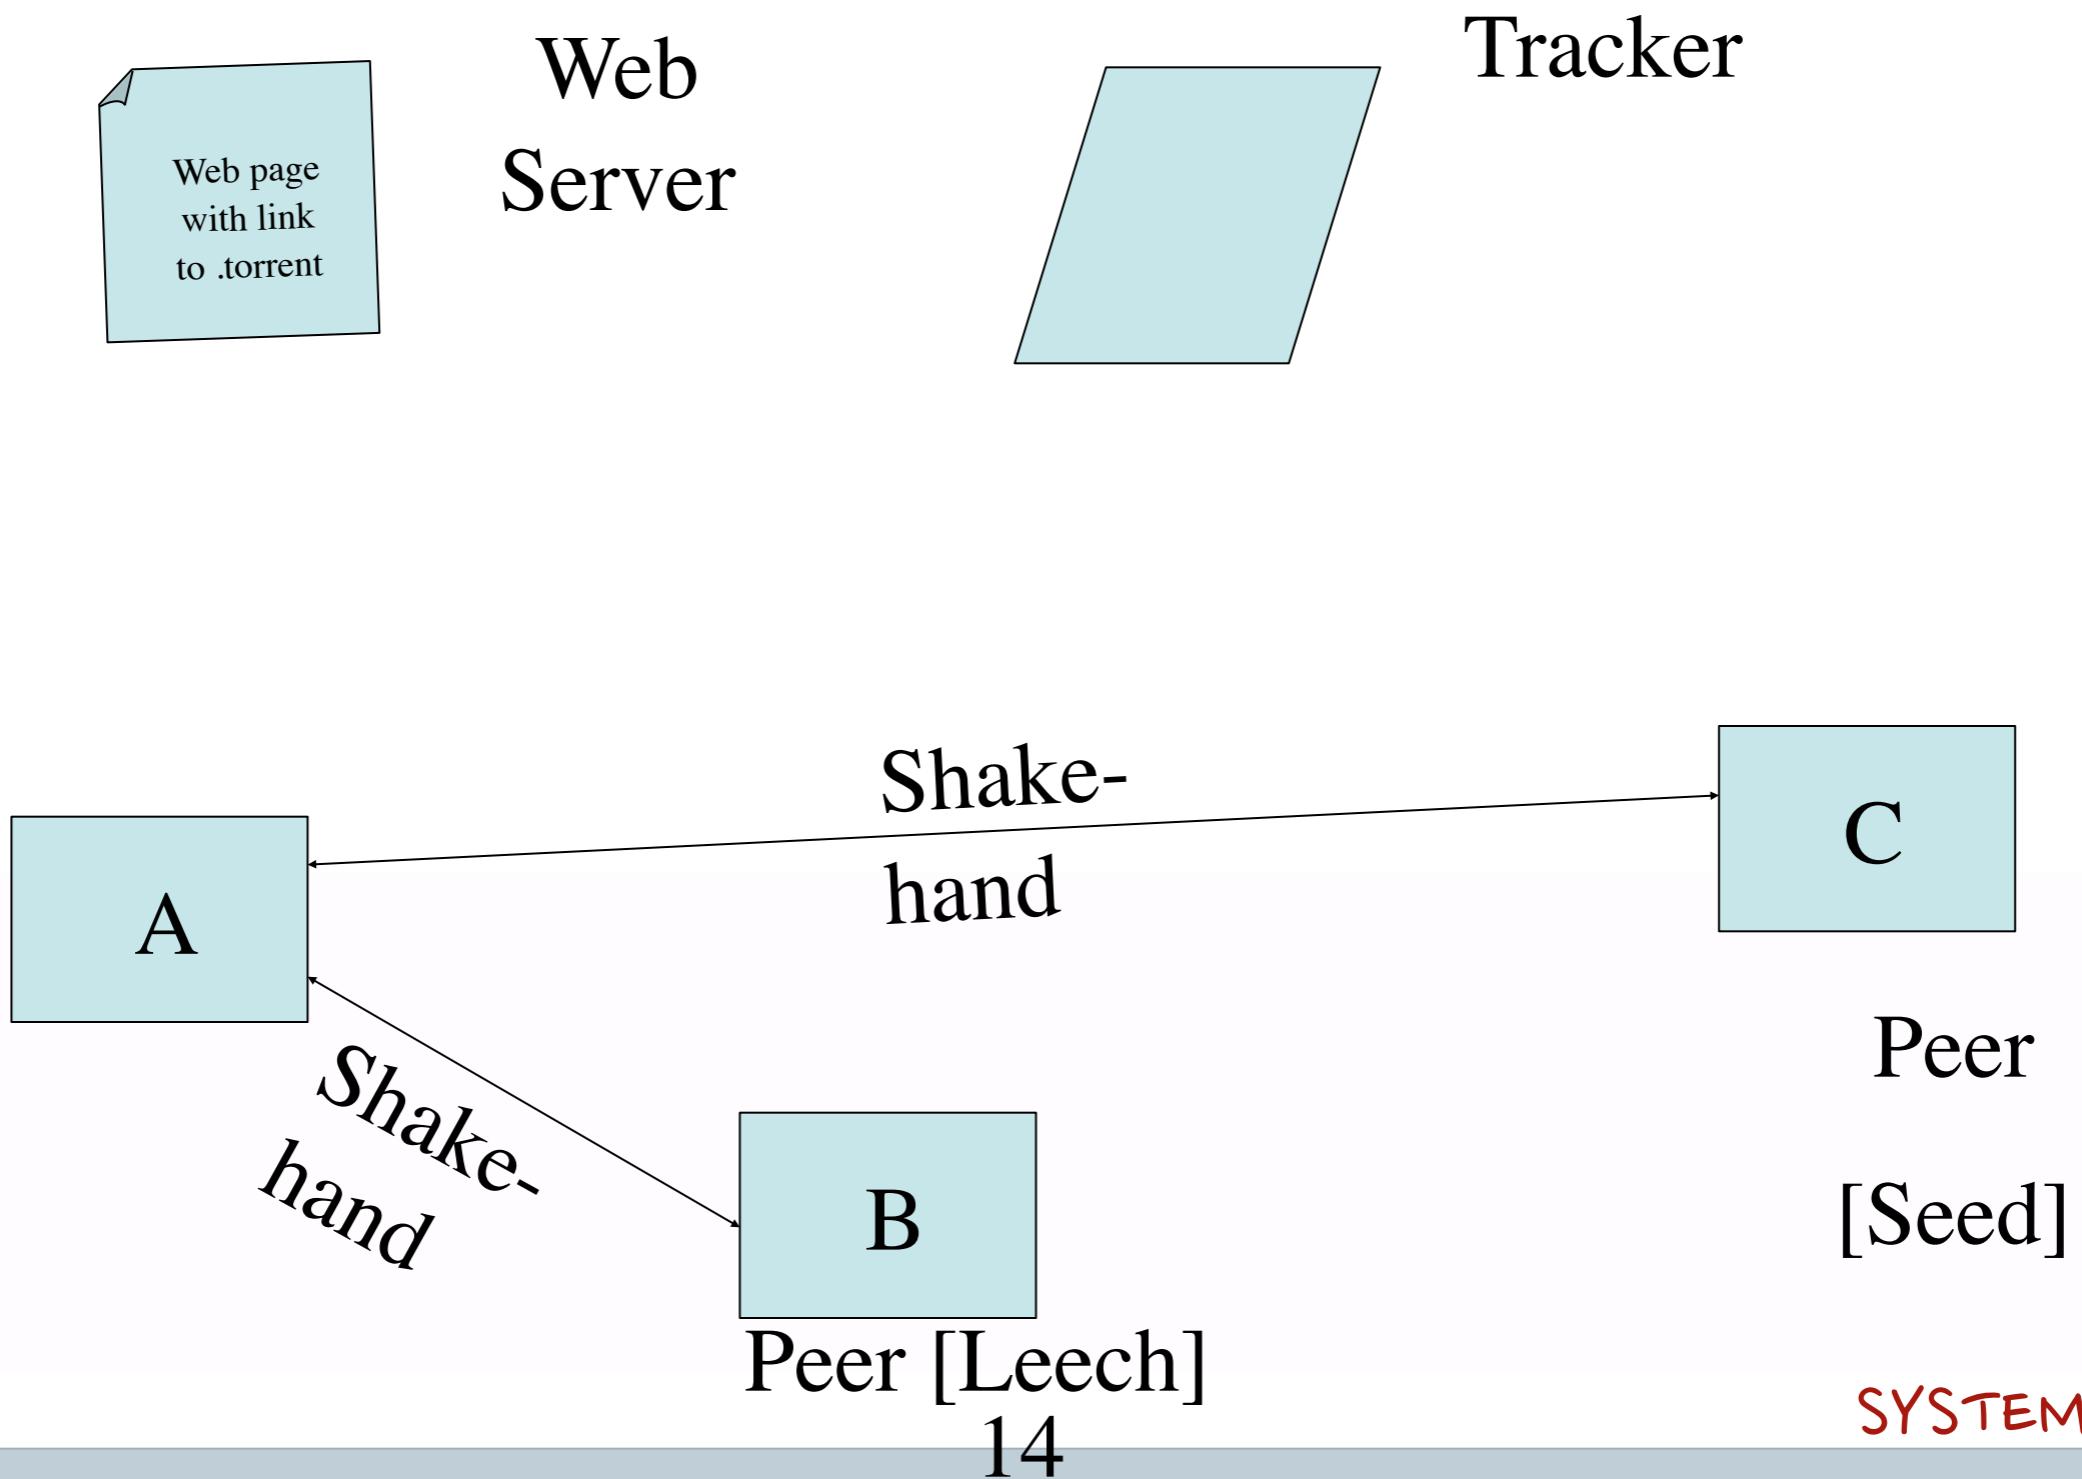
Overview



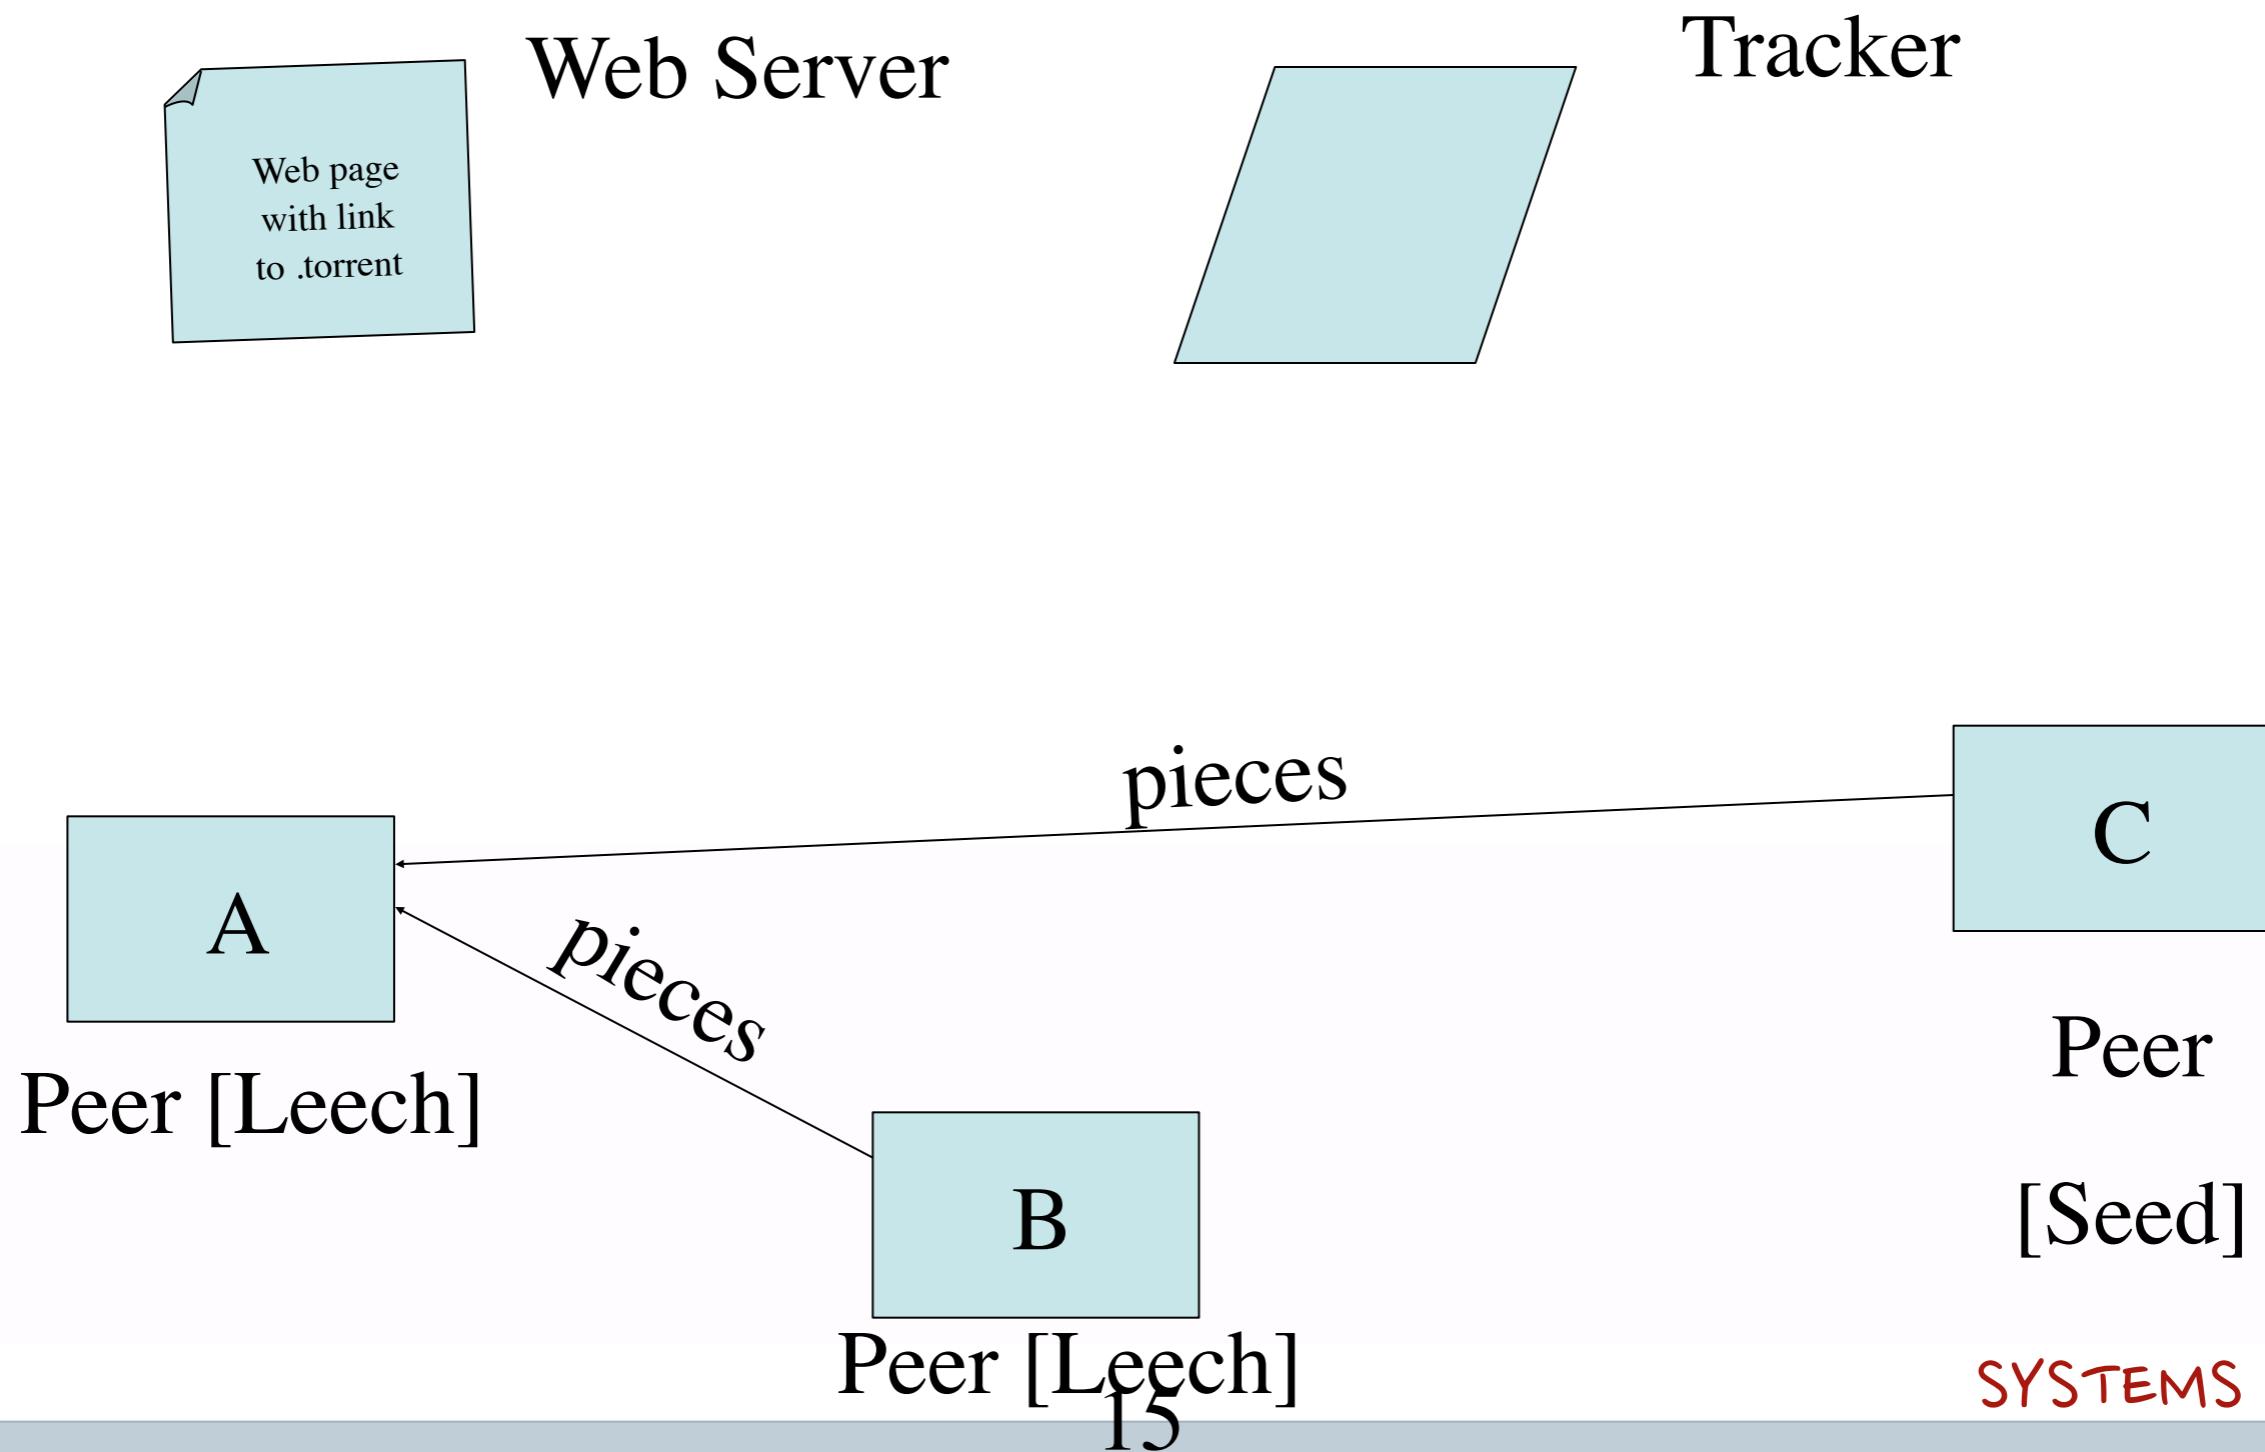
Overview



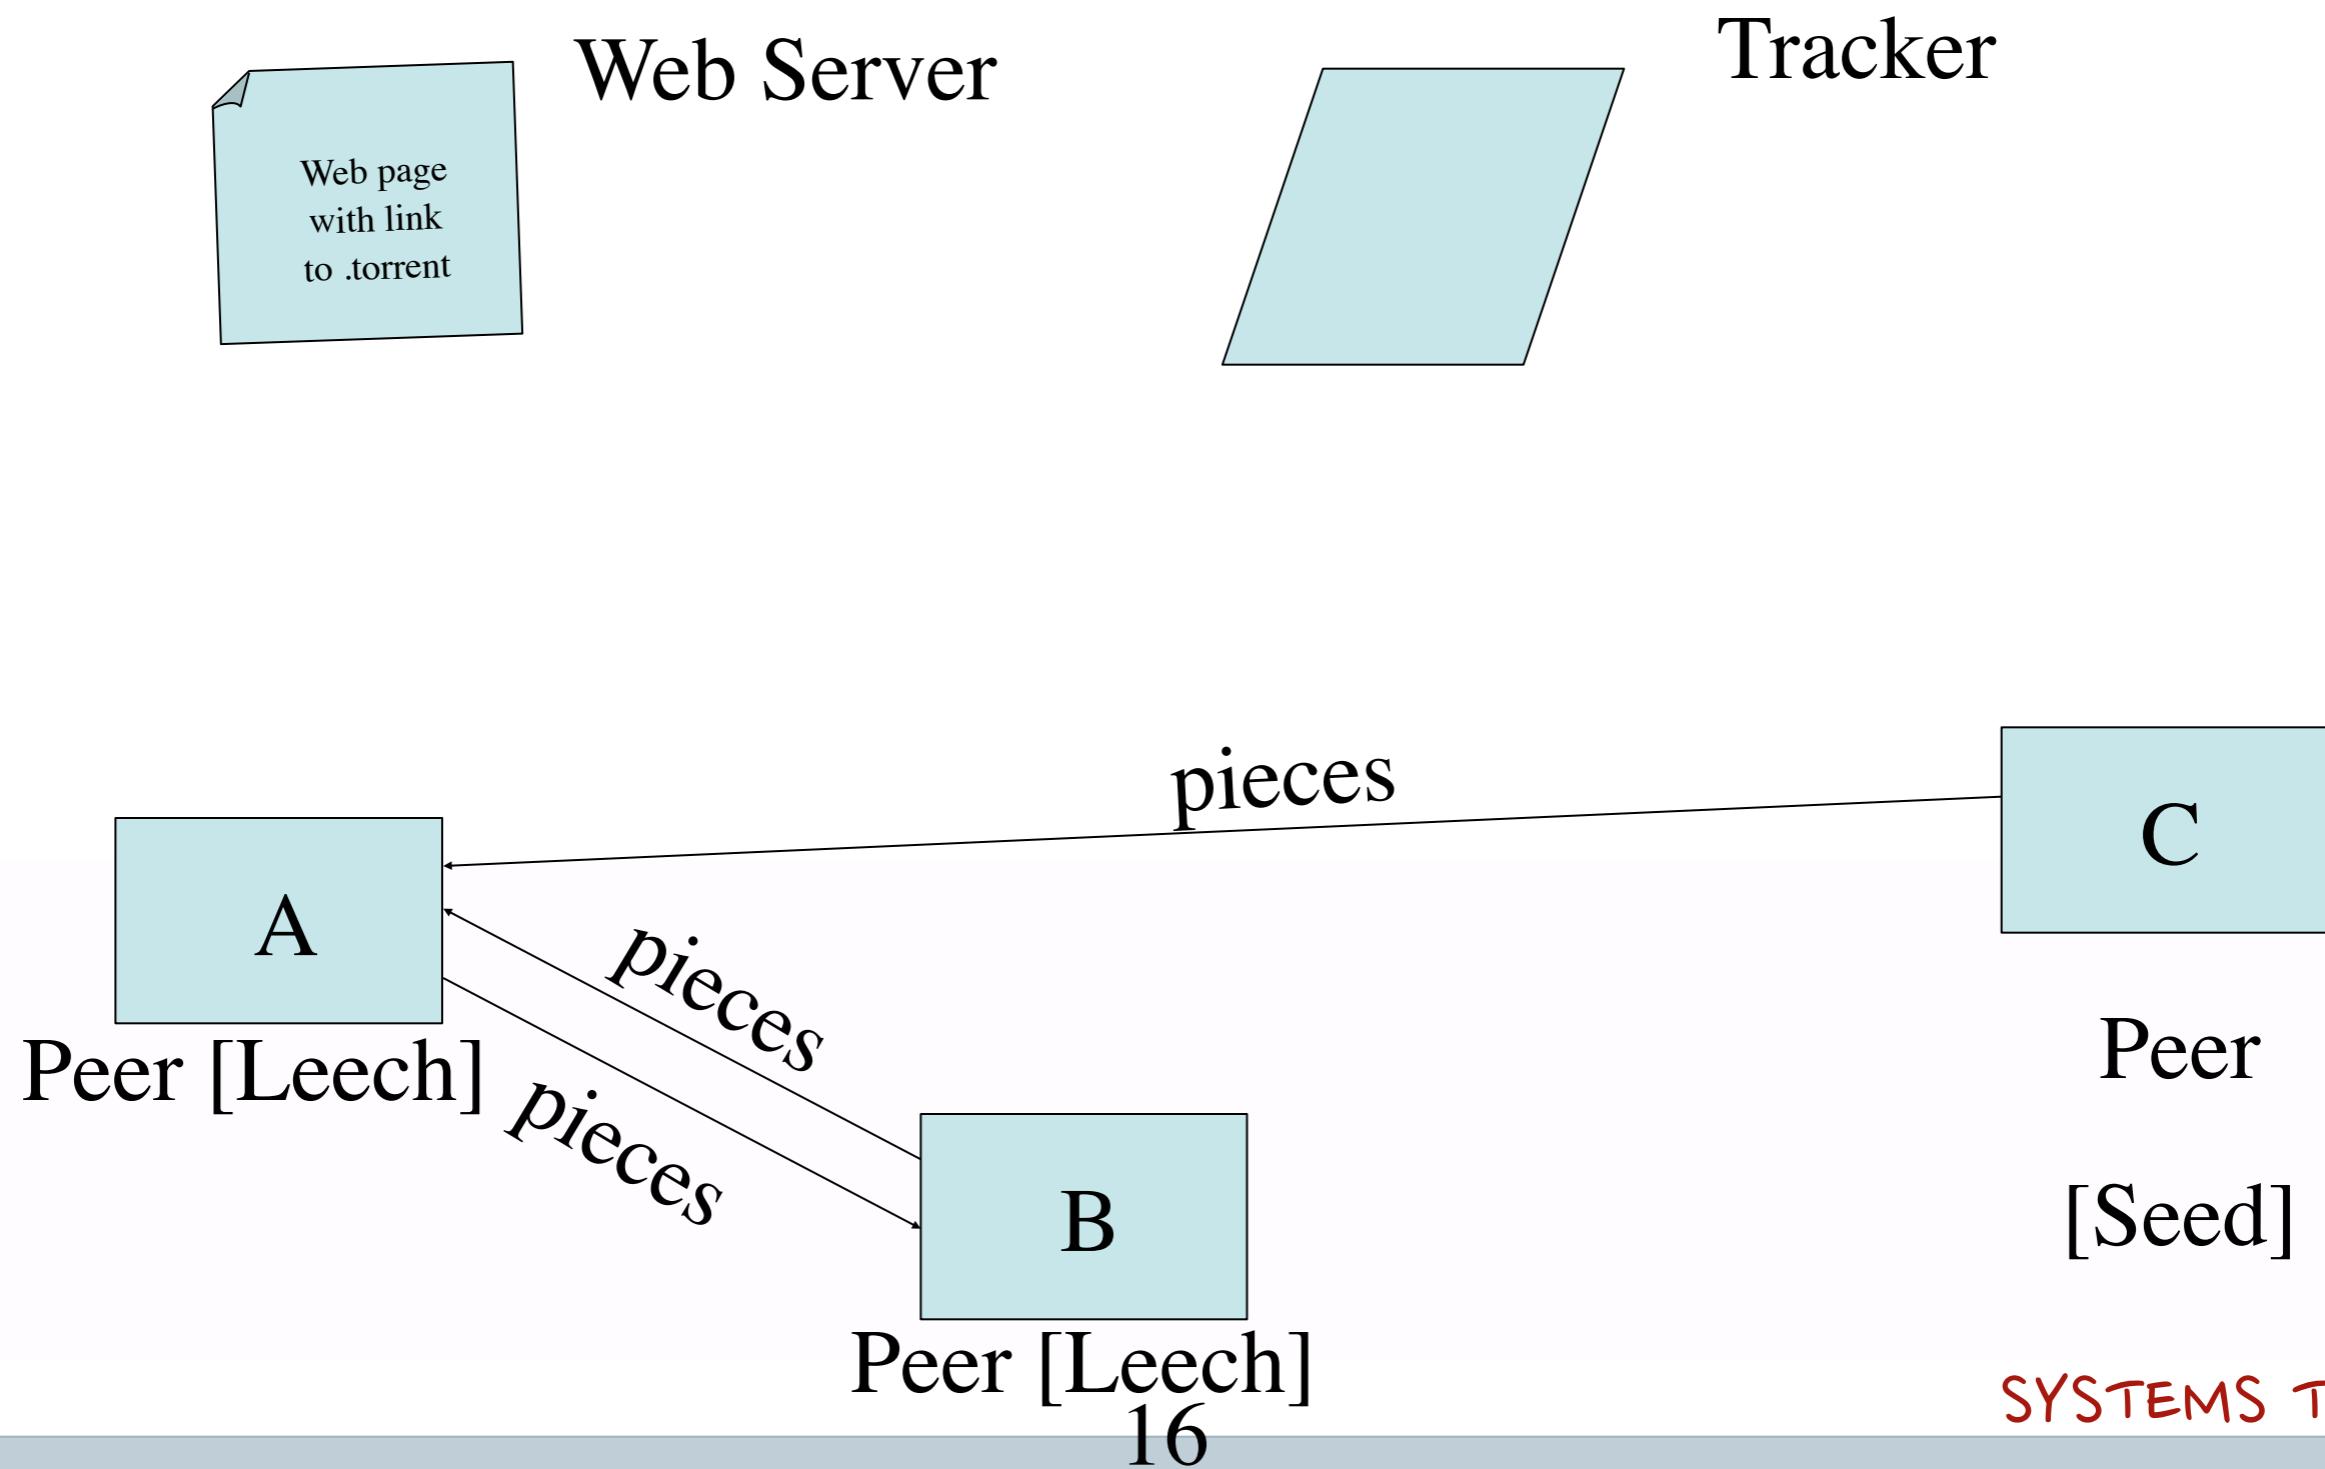
Overview



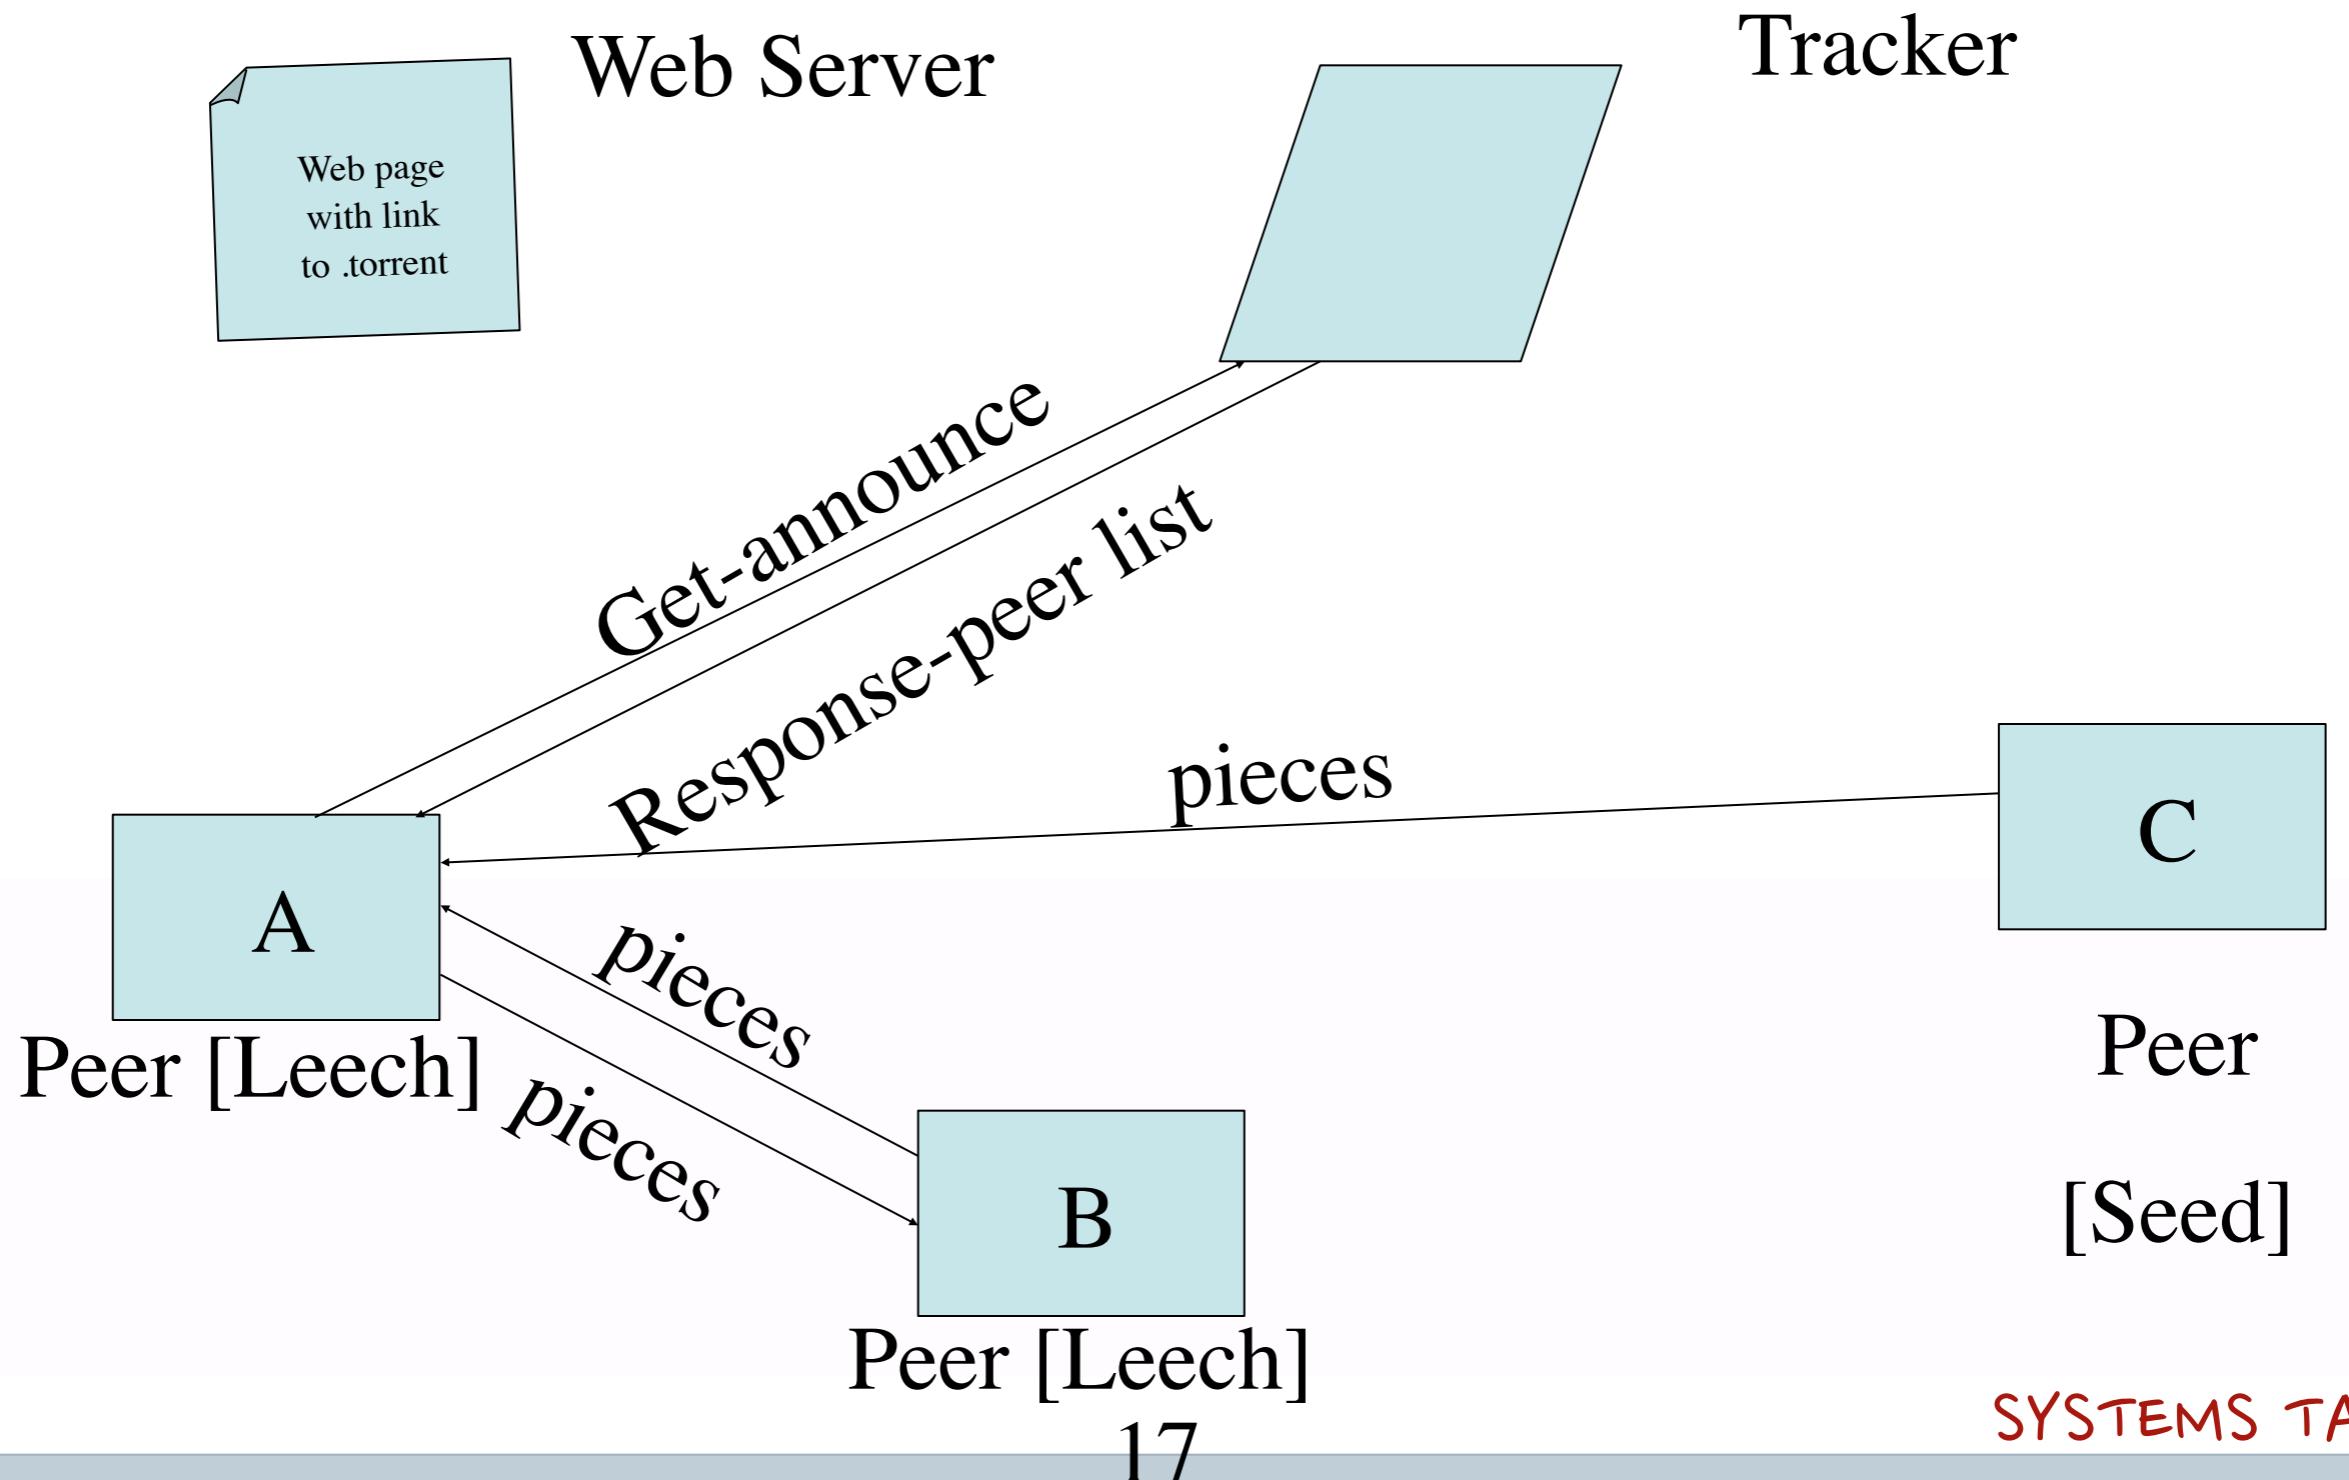
Overview



Overview



Overview



Cool stuff

- Distributed trackers
 - Trackerless torrents;
 - DHT implementation
- Decentralized keyword search
- Choking
- Optimistic unchoking
 - Clients try to periodically sending to choked connections

Chocking

- Ensures that nodes cooperate and eliminates the free-rider problem
- Cooperation involves uploaded sub-pieces that you have to your peer
- Choking is a temporary refusal to upload; downloading occurs as normal
- Connection is kept open so that setup costs are not borne again and again
- Based on game-theoretic concepts
 - Tit-for-tat strategy in Repeated Games

Piece Selection

- The order in which pieces are selected by different peers is critical for good performance
- If a bad algorithm is used, we could end up in a situation where every peer has all the pieces that are currently available and none of the missing ones
- If the original seed is taken down, the file cannot be completely downloaded!

Random First Piece

- Initially, a peer has nothing to trade
- Important to get a complete piece ASAP
- Rare pieces are typically available at fewer peers, so downloading a rare piece initially is not a good idea
- Policy: Select a random piece of the file and download it

Rarest Piece First

- Policy: Determine the pieces that are most rare among your peers and download those first
- This ensures that the most common pieces are left till the end to download
- Rarest first also ensures that a large variety of pieces are downloaded from the seed

Chocking Algorithm

- Goal is to have several bidirectional connections running continuously
- Upload to peers who have uploaded to you recently
- Unutilized connections are uploaded to on a trial basis to see if better transfer rates could be found using them

How BitTorrent works

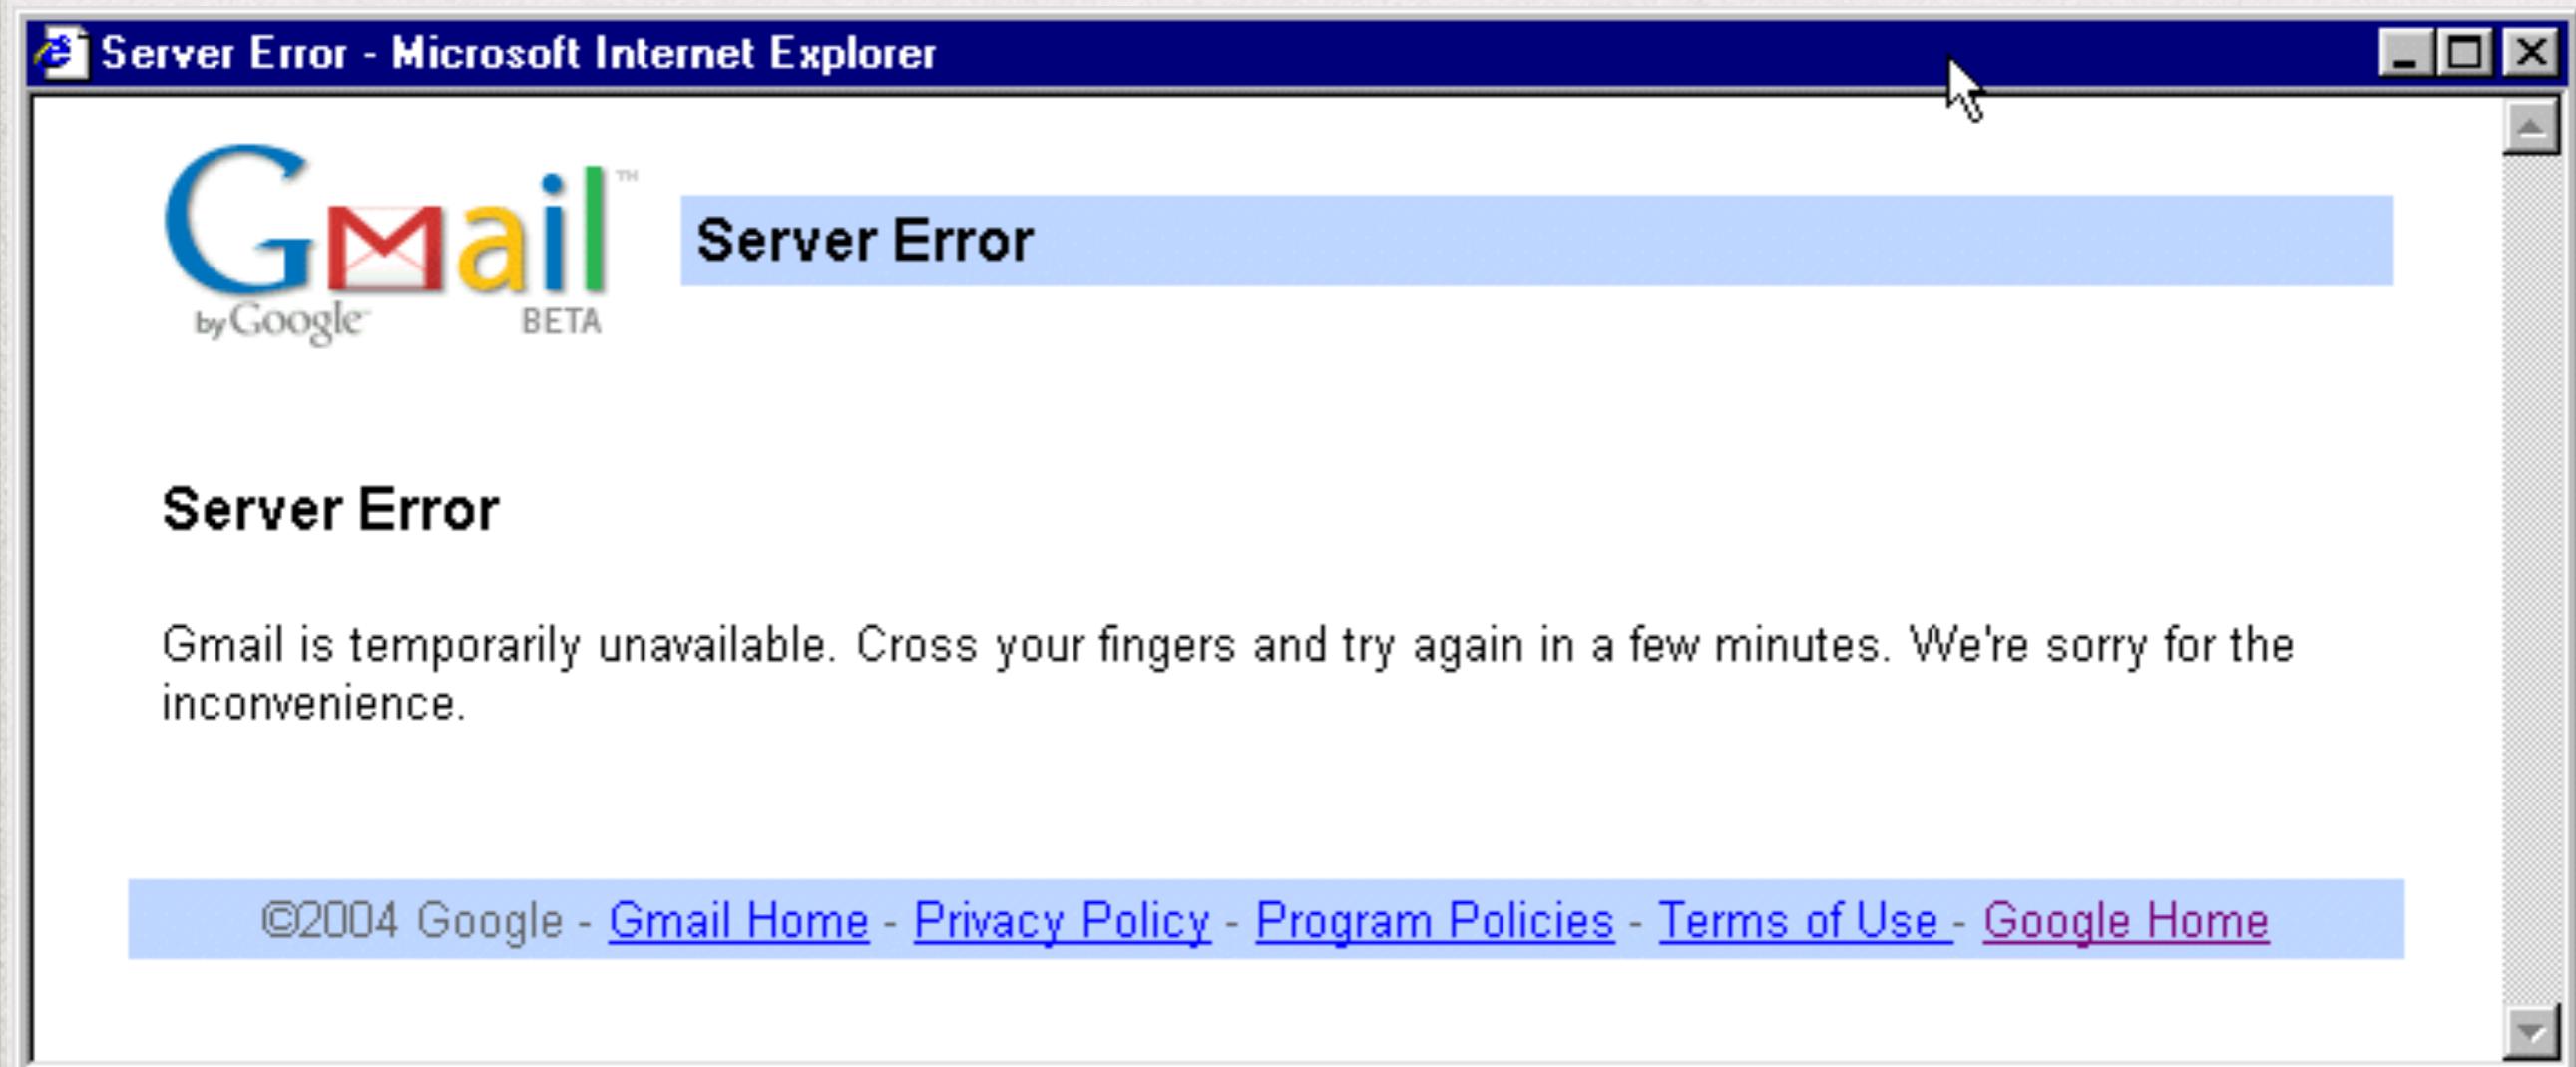
- * <http://mg8.org/processing/bt.html>

This lecture ...

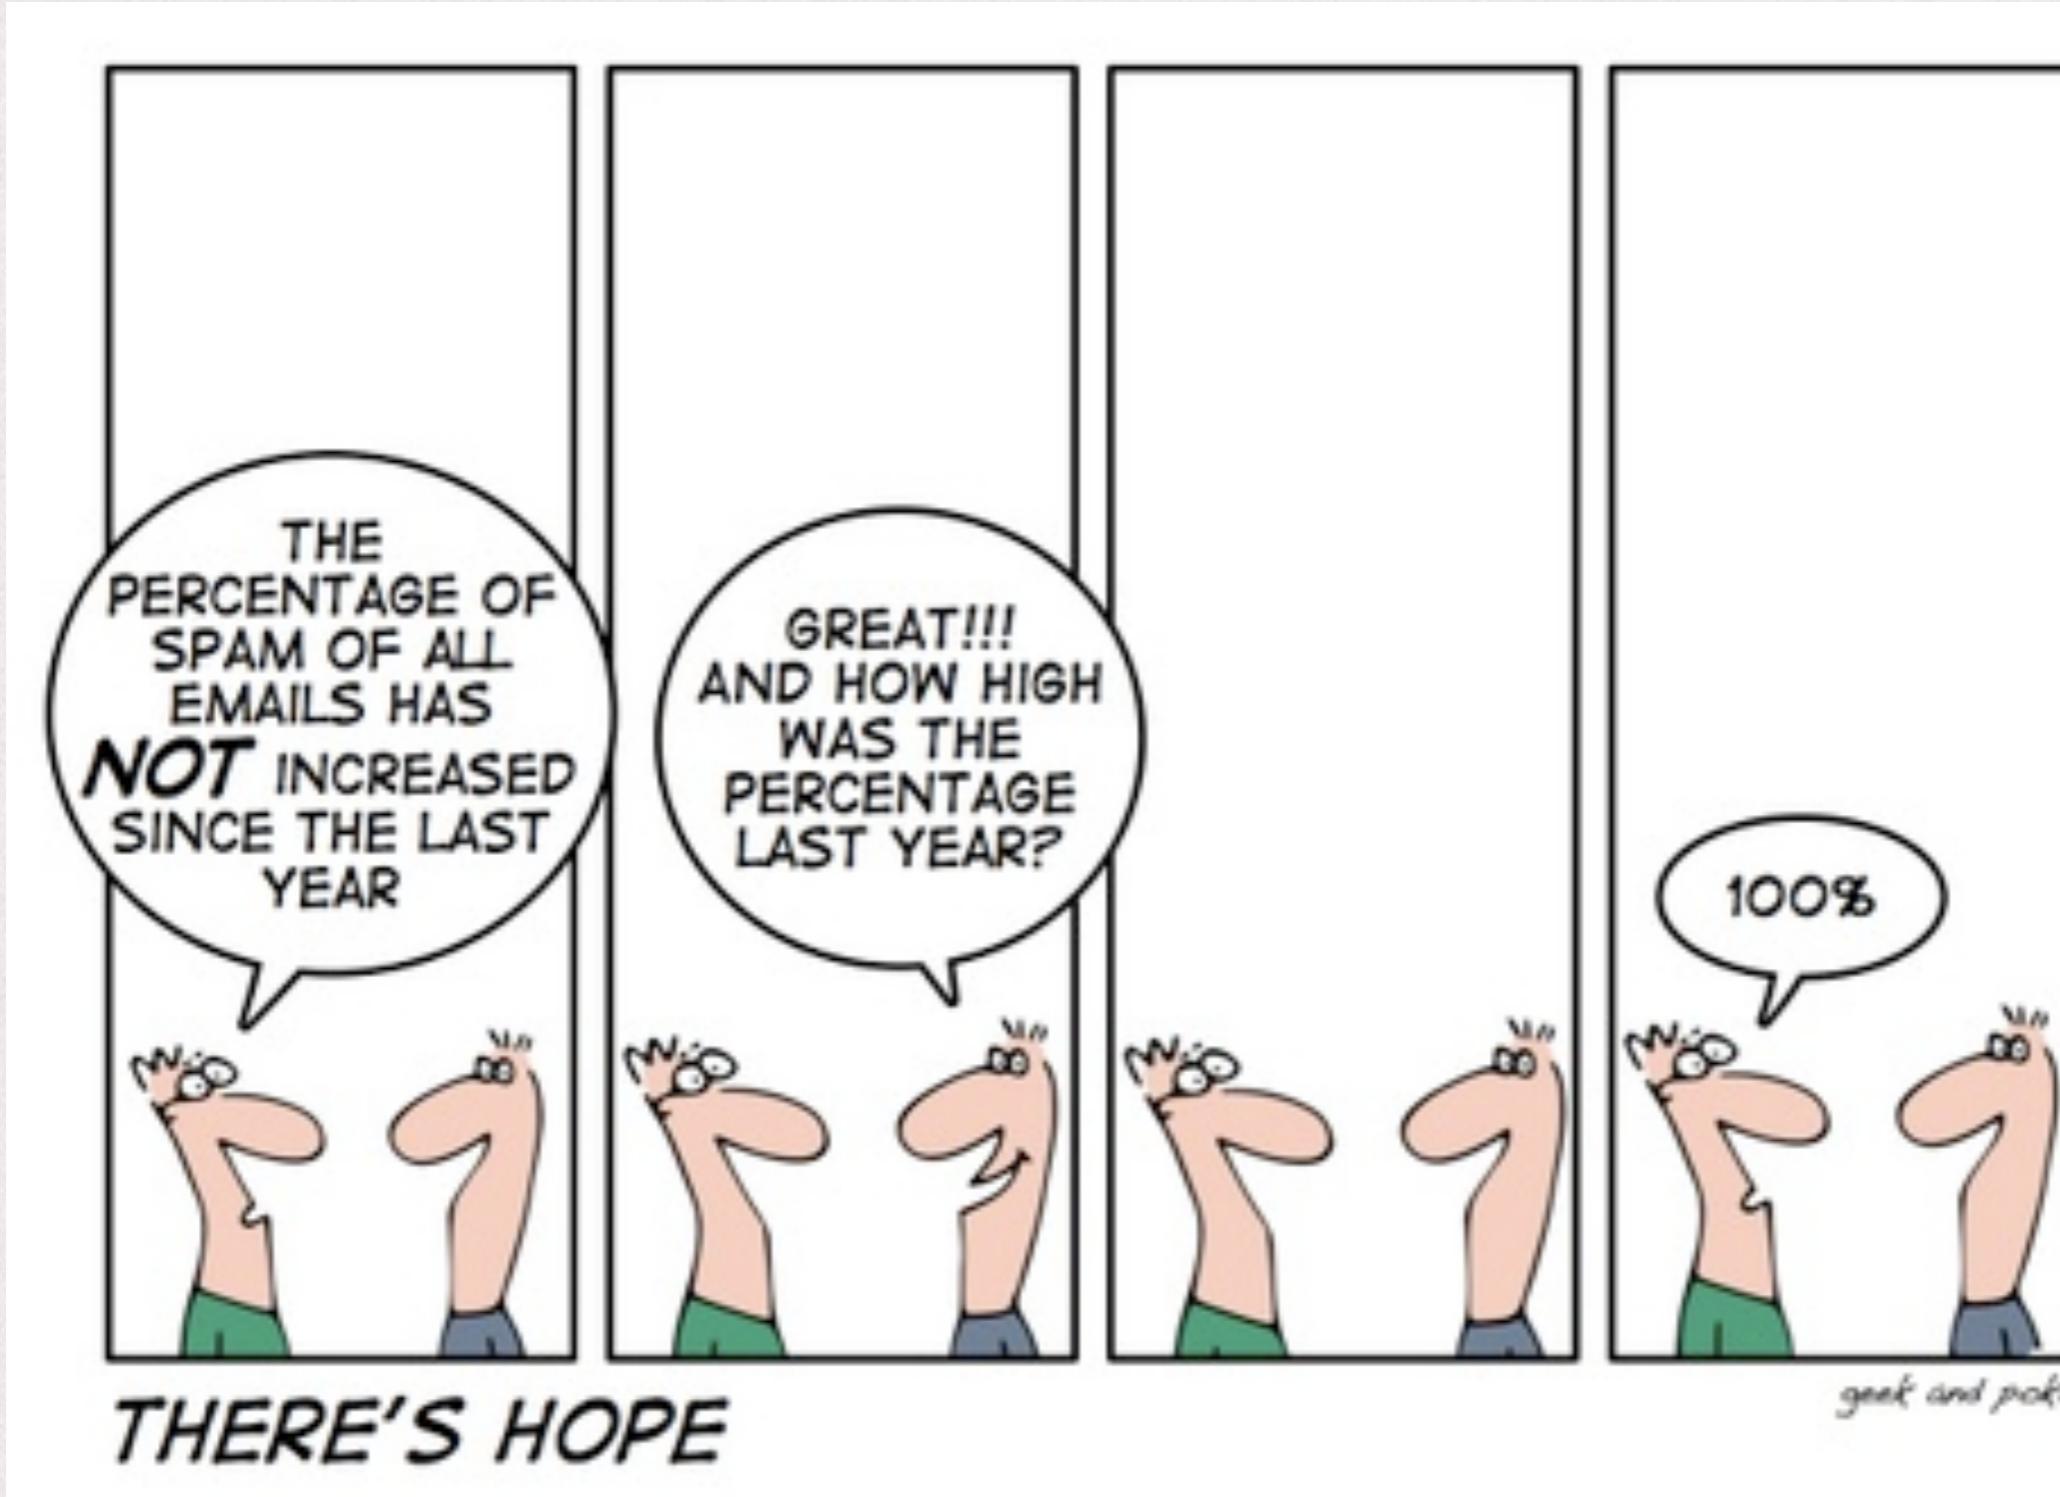
The image shows a screenshot of a Twitter error page. At the top left is the Twitter logo. To its right are navigation links: Home, Public Timeline, and Help. The main message is "Twitter is over capacity." followed by the subtext "Too many tweets! Please wait a moment and try again." Below the text is a cartoon illustration of a white whale swimming in blue water. Numerous orange Twitter birds are shown flying upwards from the whale's back, with arcs above them indicating their flight paths. The bottom of the image features a red-orange wavy line representing the horizon or ocean floor.

© 2008 Twitter [About Us](#) [Contact](#) [Blog](#) [Status API](#) [Help](#) [Jobs](#) [TOS](#) [Privacy](#)

This lecture ...



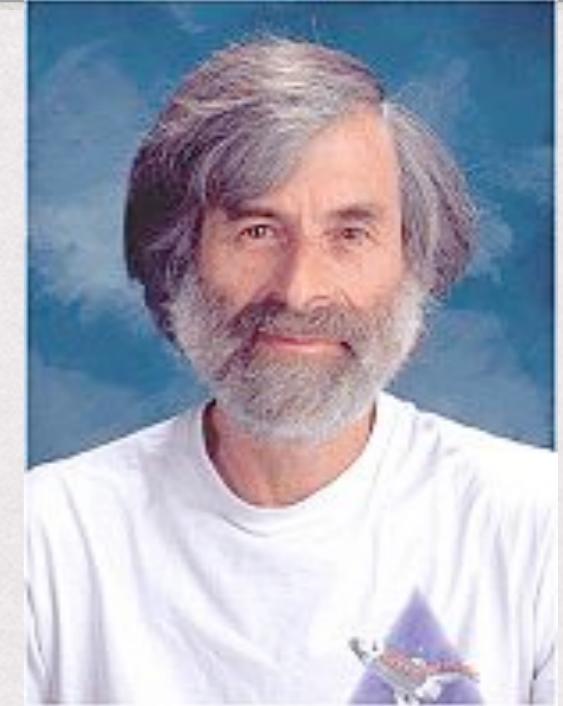
This lecture ...



Theme

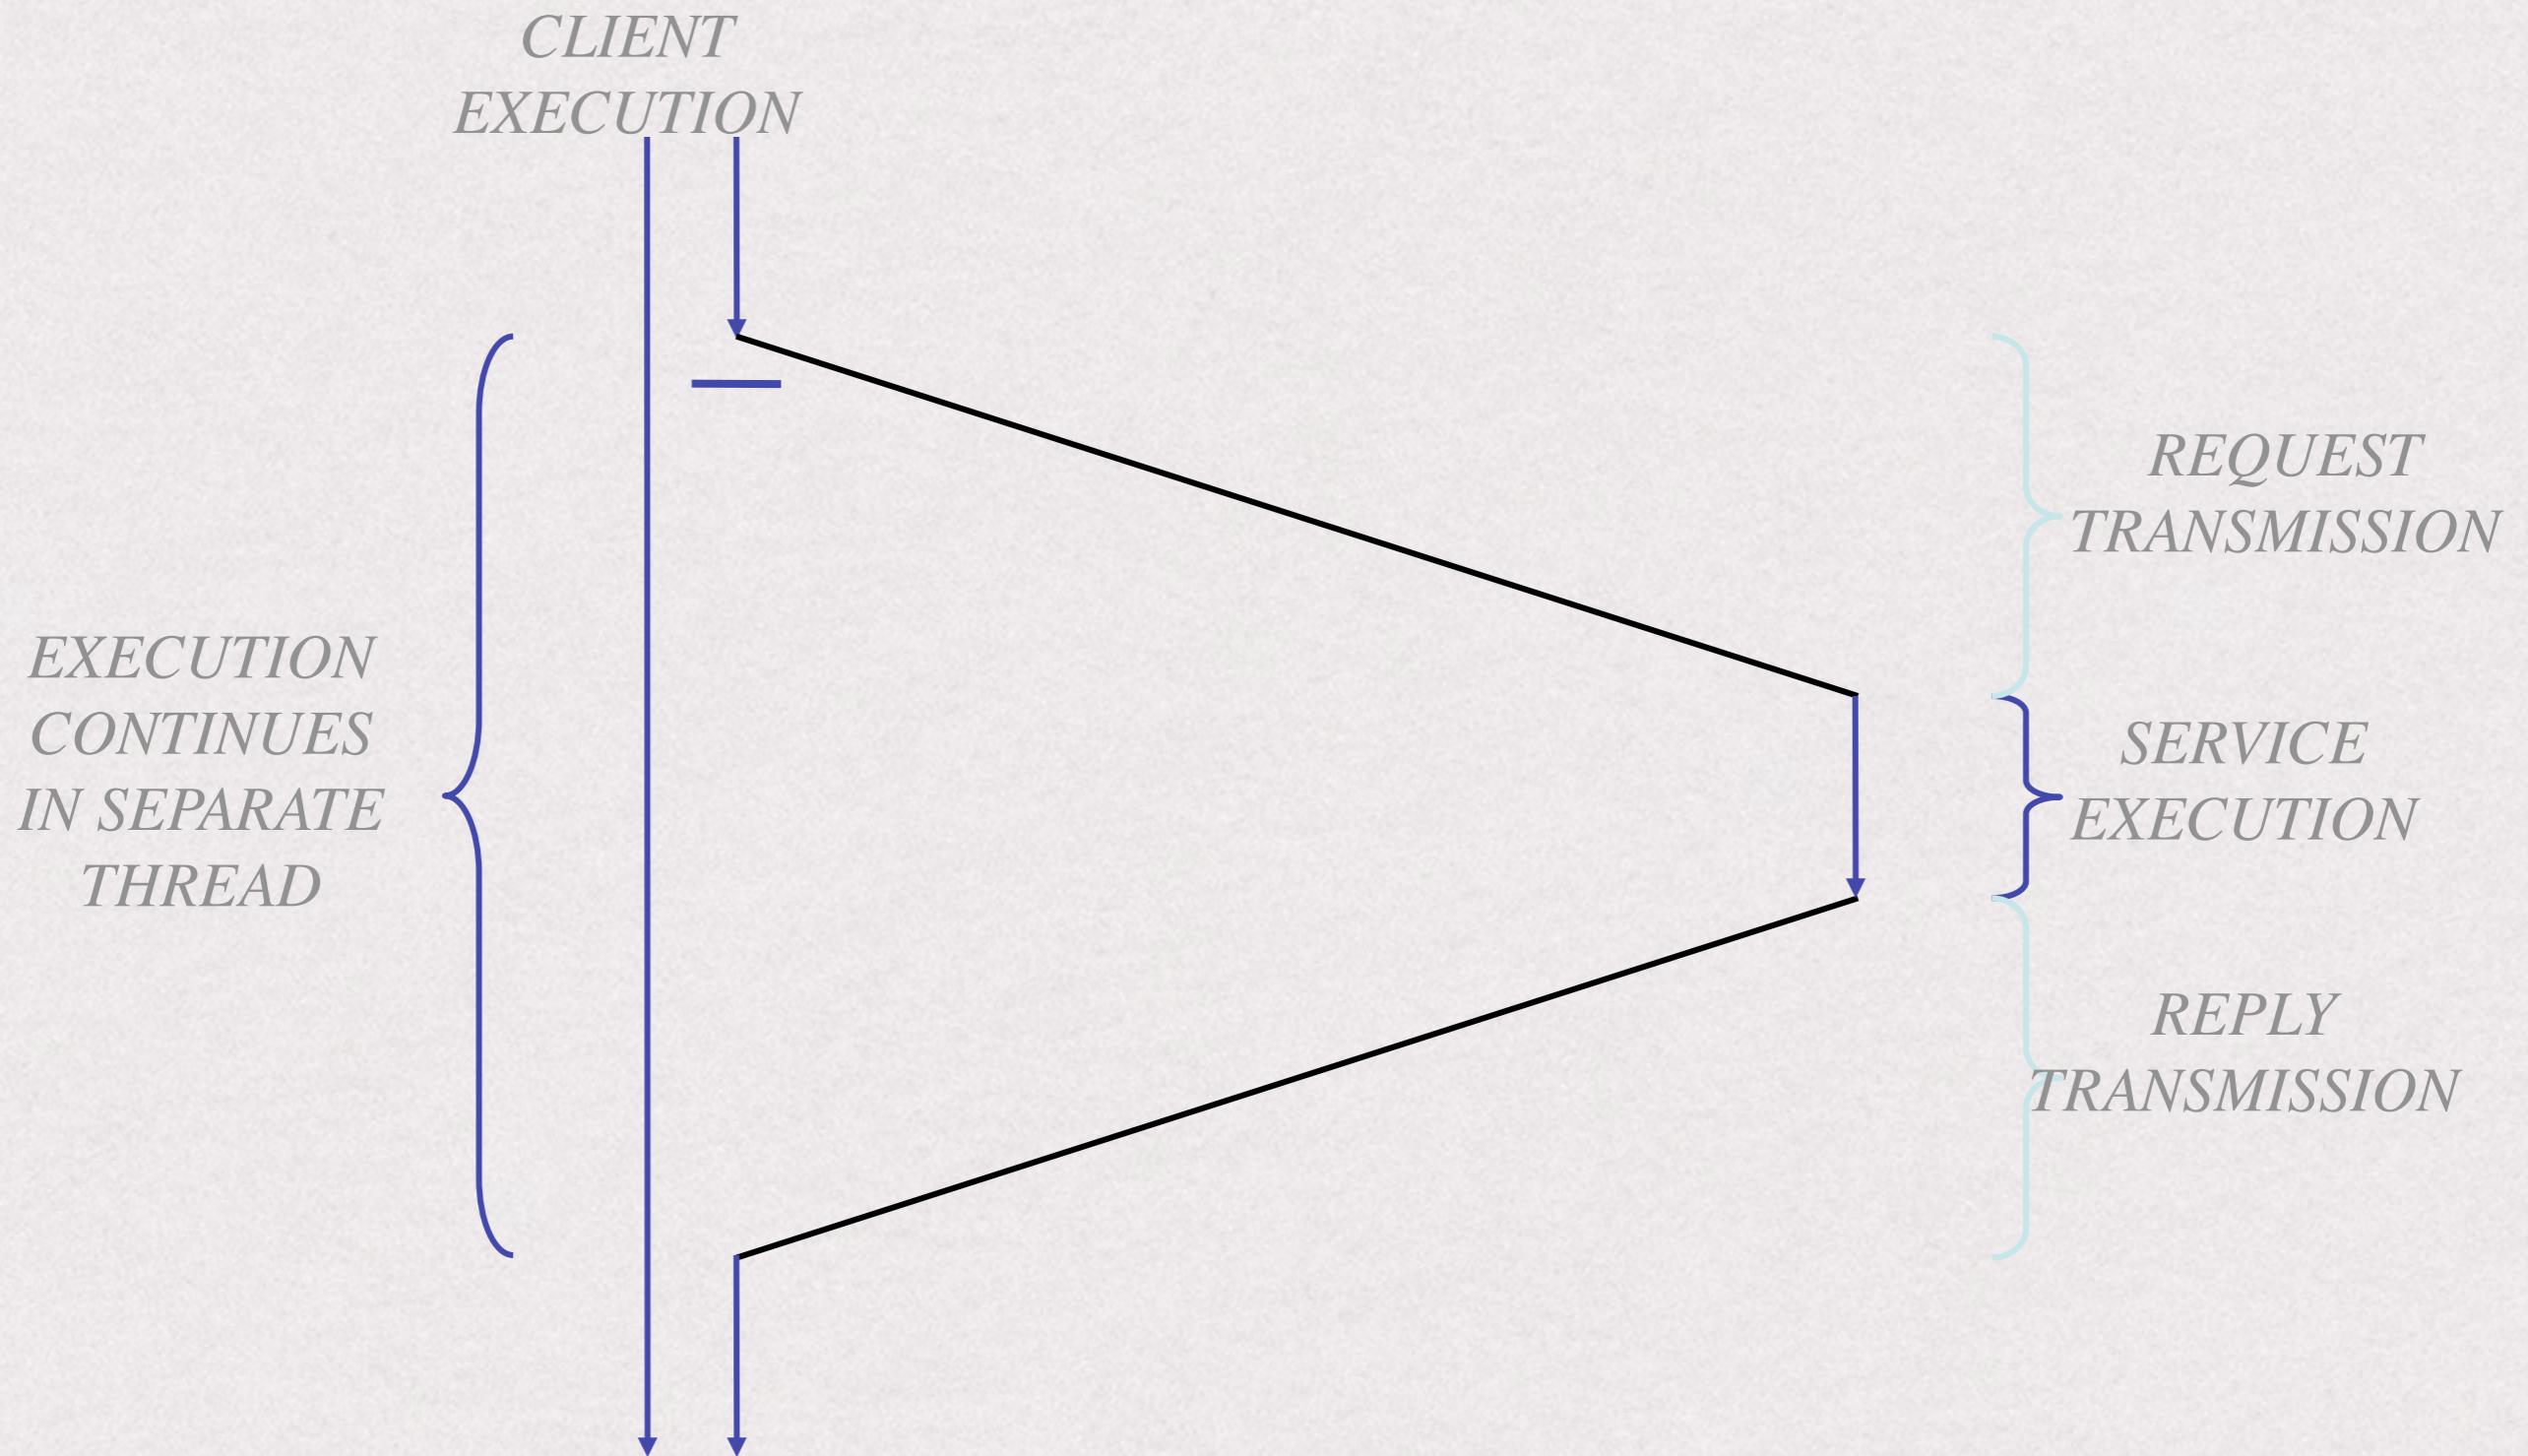
- * A distributed system is, “*One on which I cannot get any work done because some machine I have never heard of has crashed*”
- * A well-known quotation from Leslie Lamport
- * Concerned with the (increased) propensity of distributed systems to failure, and the complex consequences of those failures.

Leslie Lamport



- * Seminal work in distributed systems
 - * Time, clocks and the order of events
 - * The byzantine general's problem
 - * Reaching agreement in the presence of faults
 - * Mutual exclusion using the bakery algorithm
 - * Snapshot algorithms for determining consistent global states
- * The initial developer of LaTeX

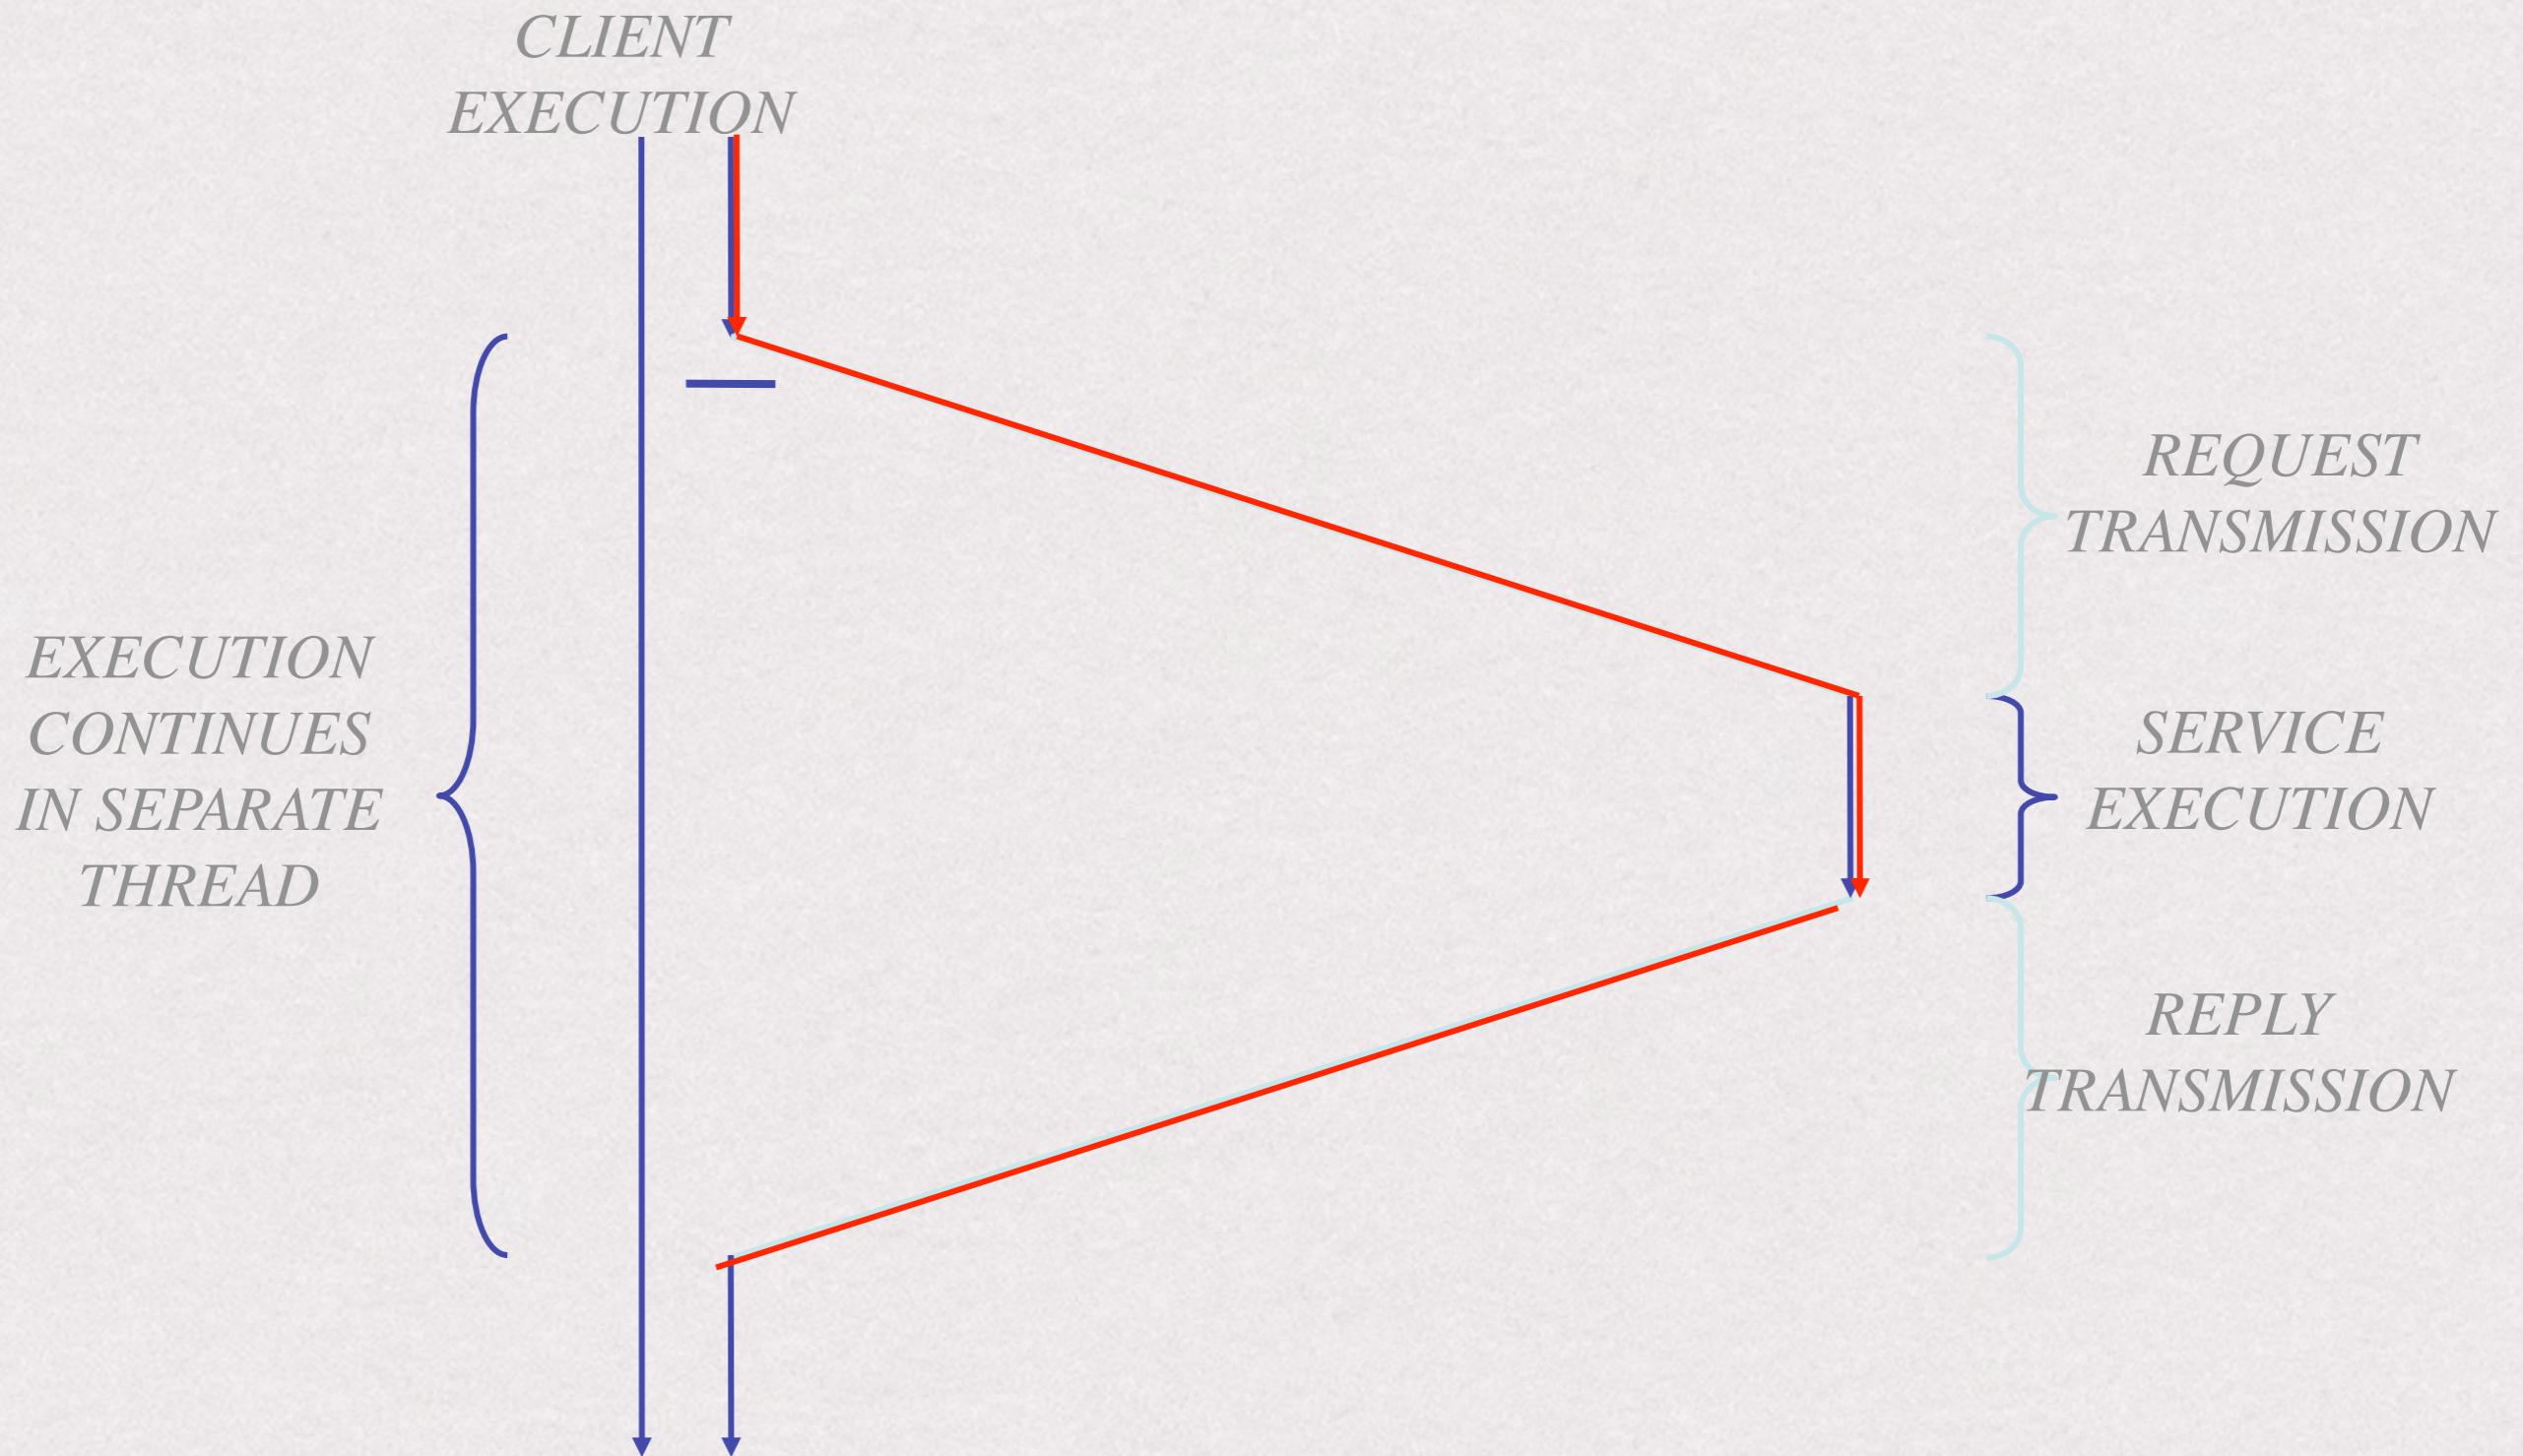
Remote Calls



Questions

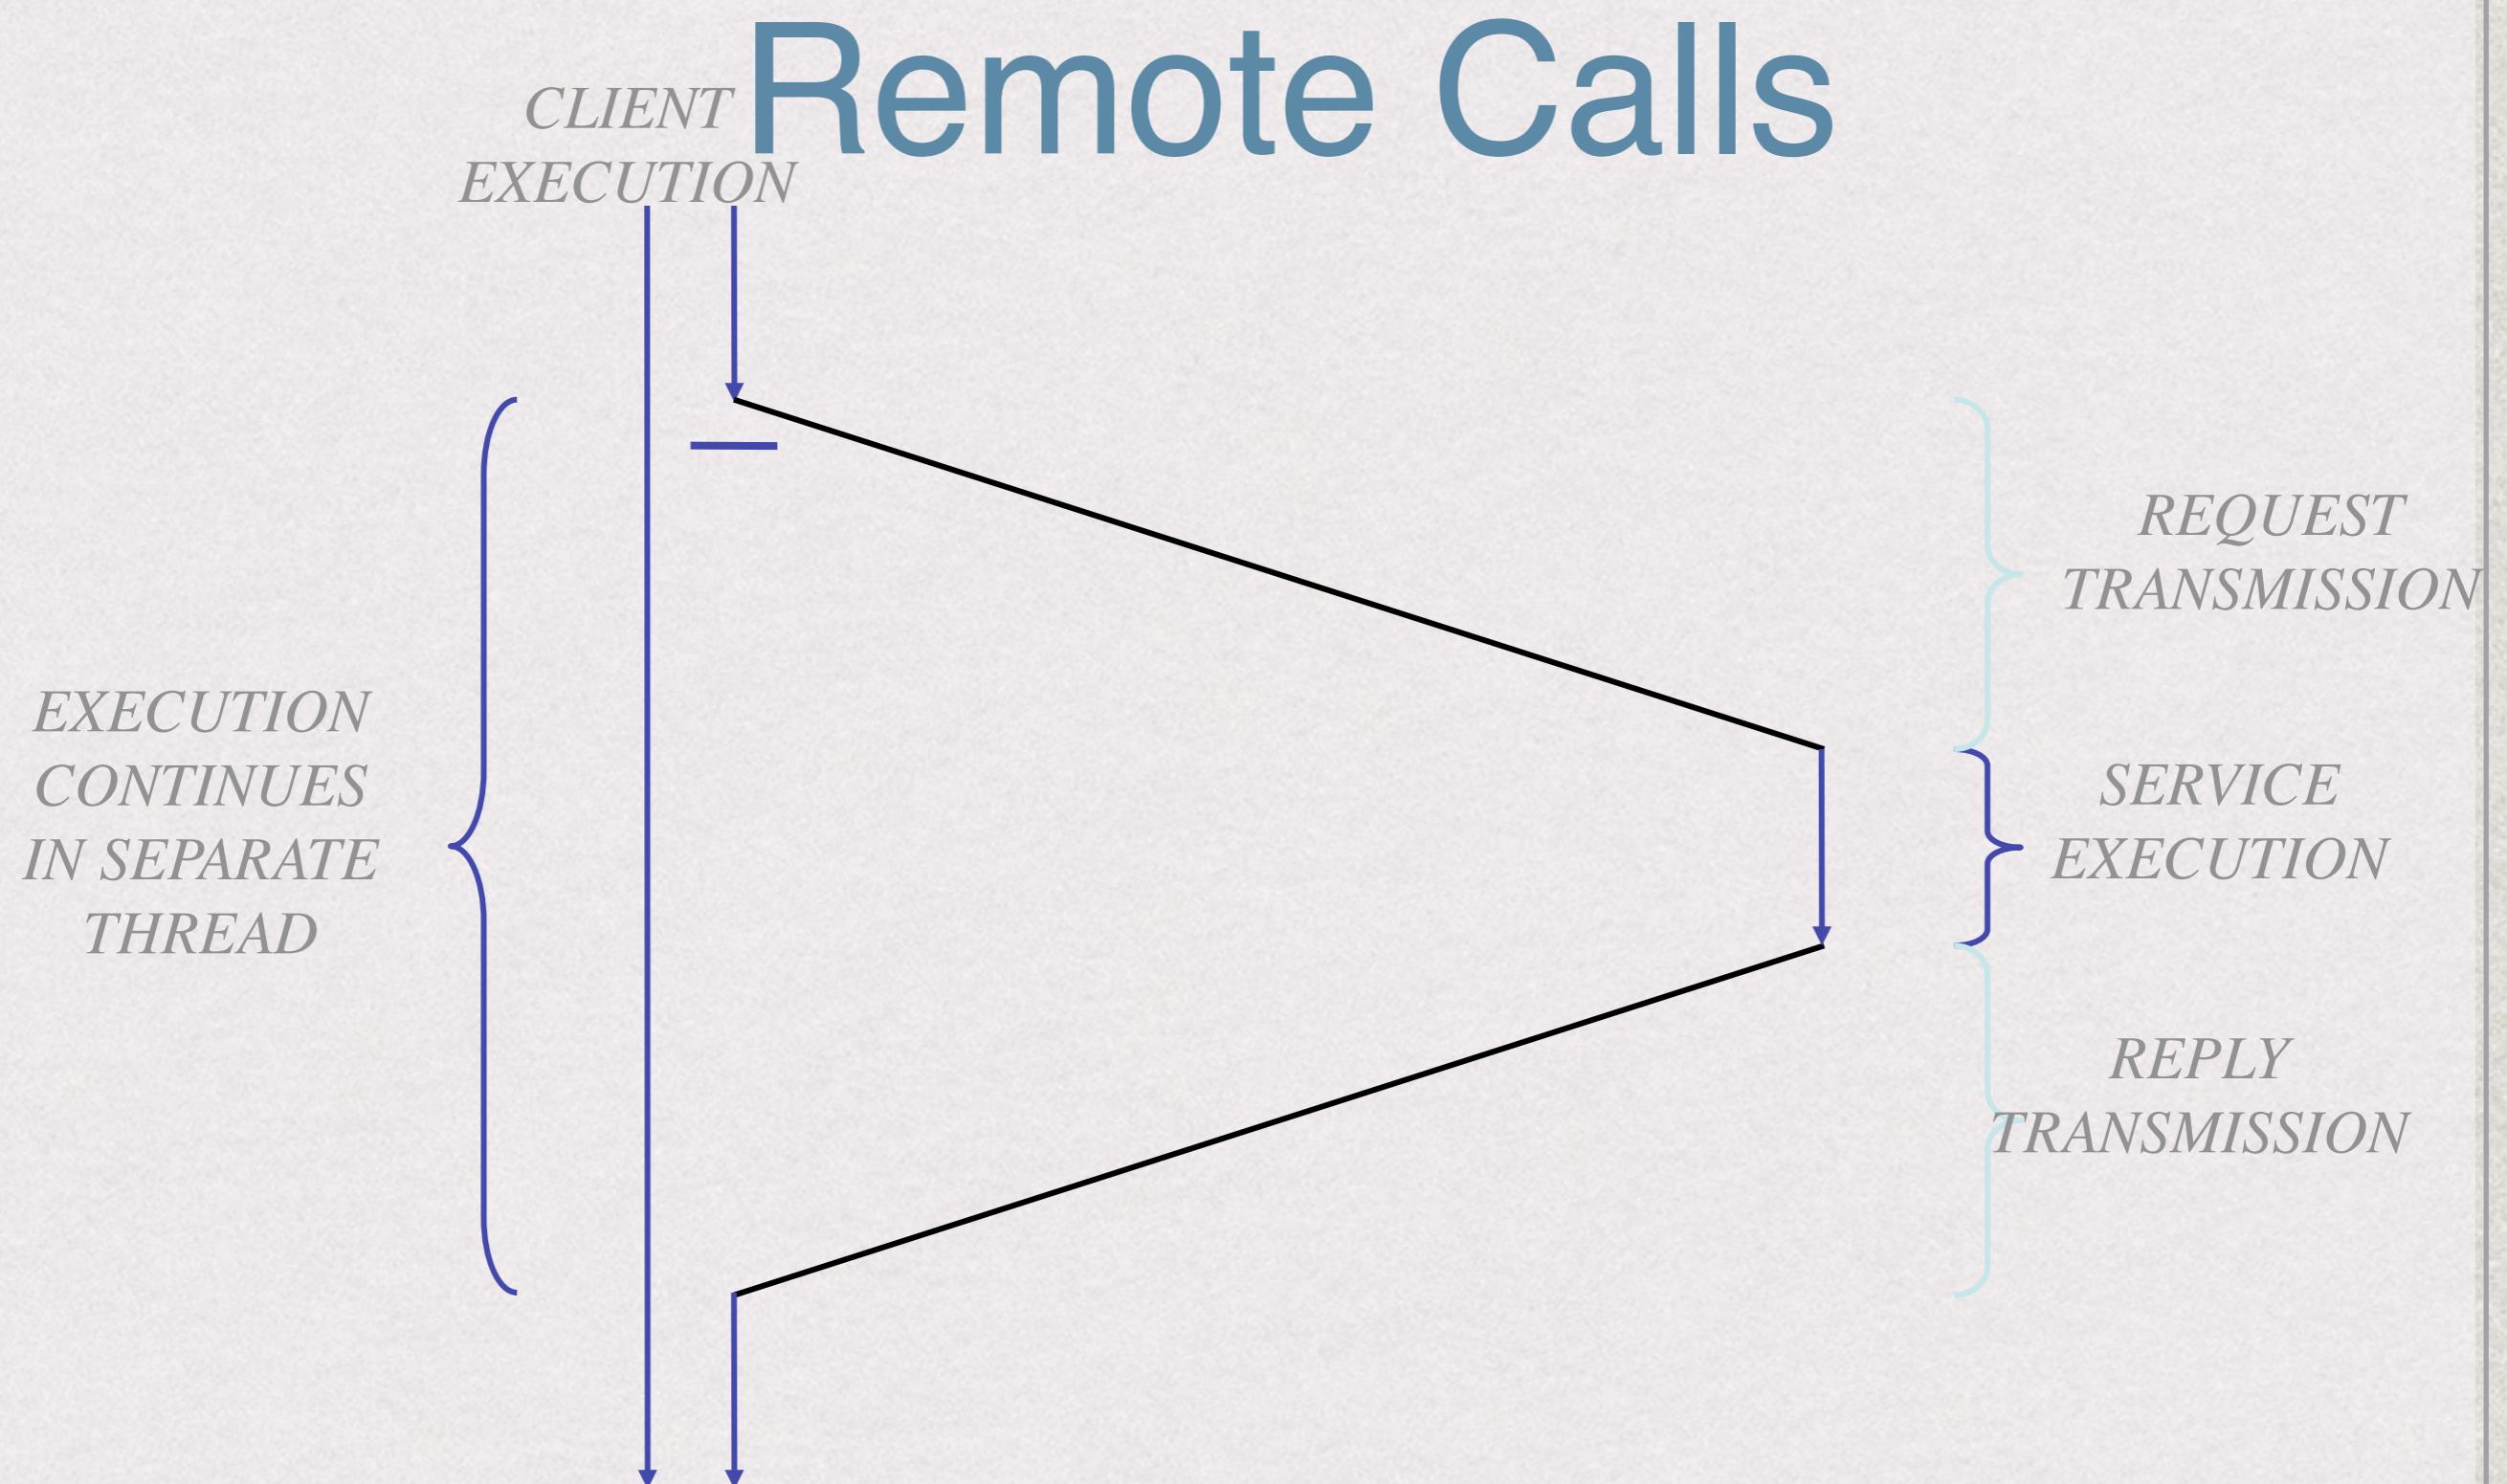
- * What can fail between the client starting an RPC and the end of that RPC?

Remote Calls



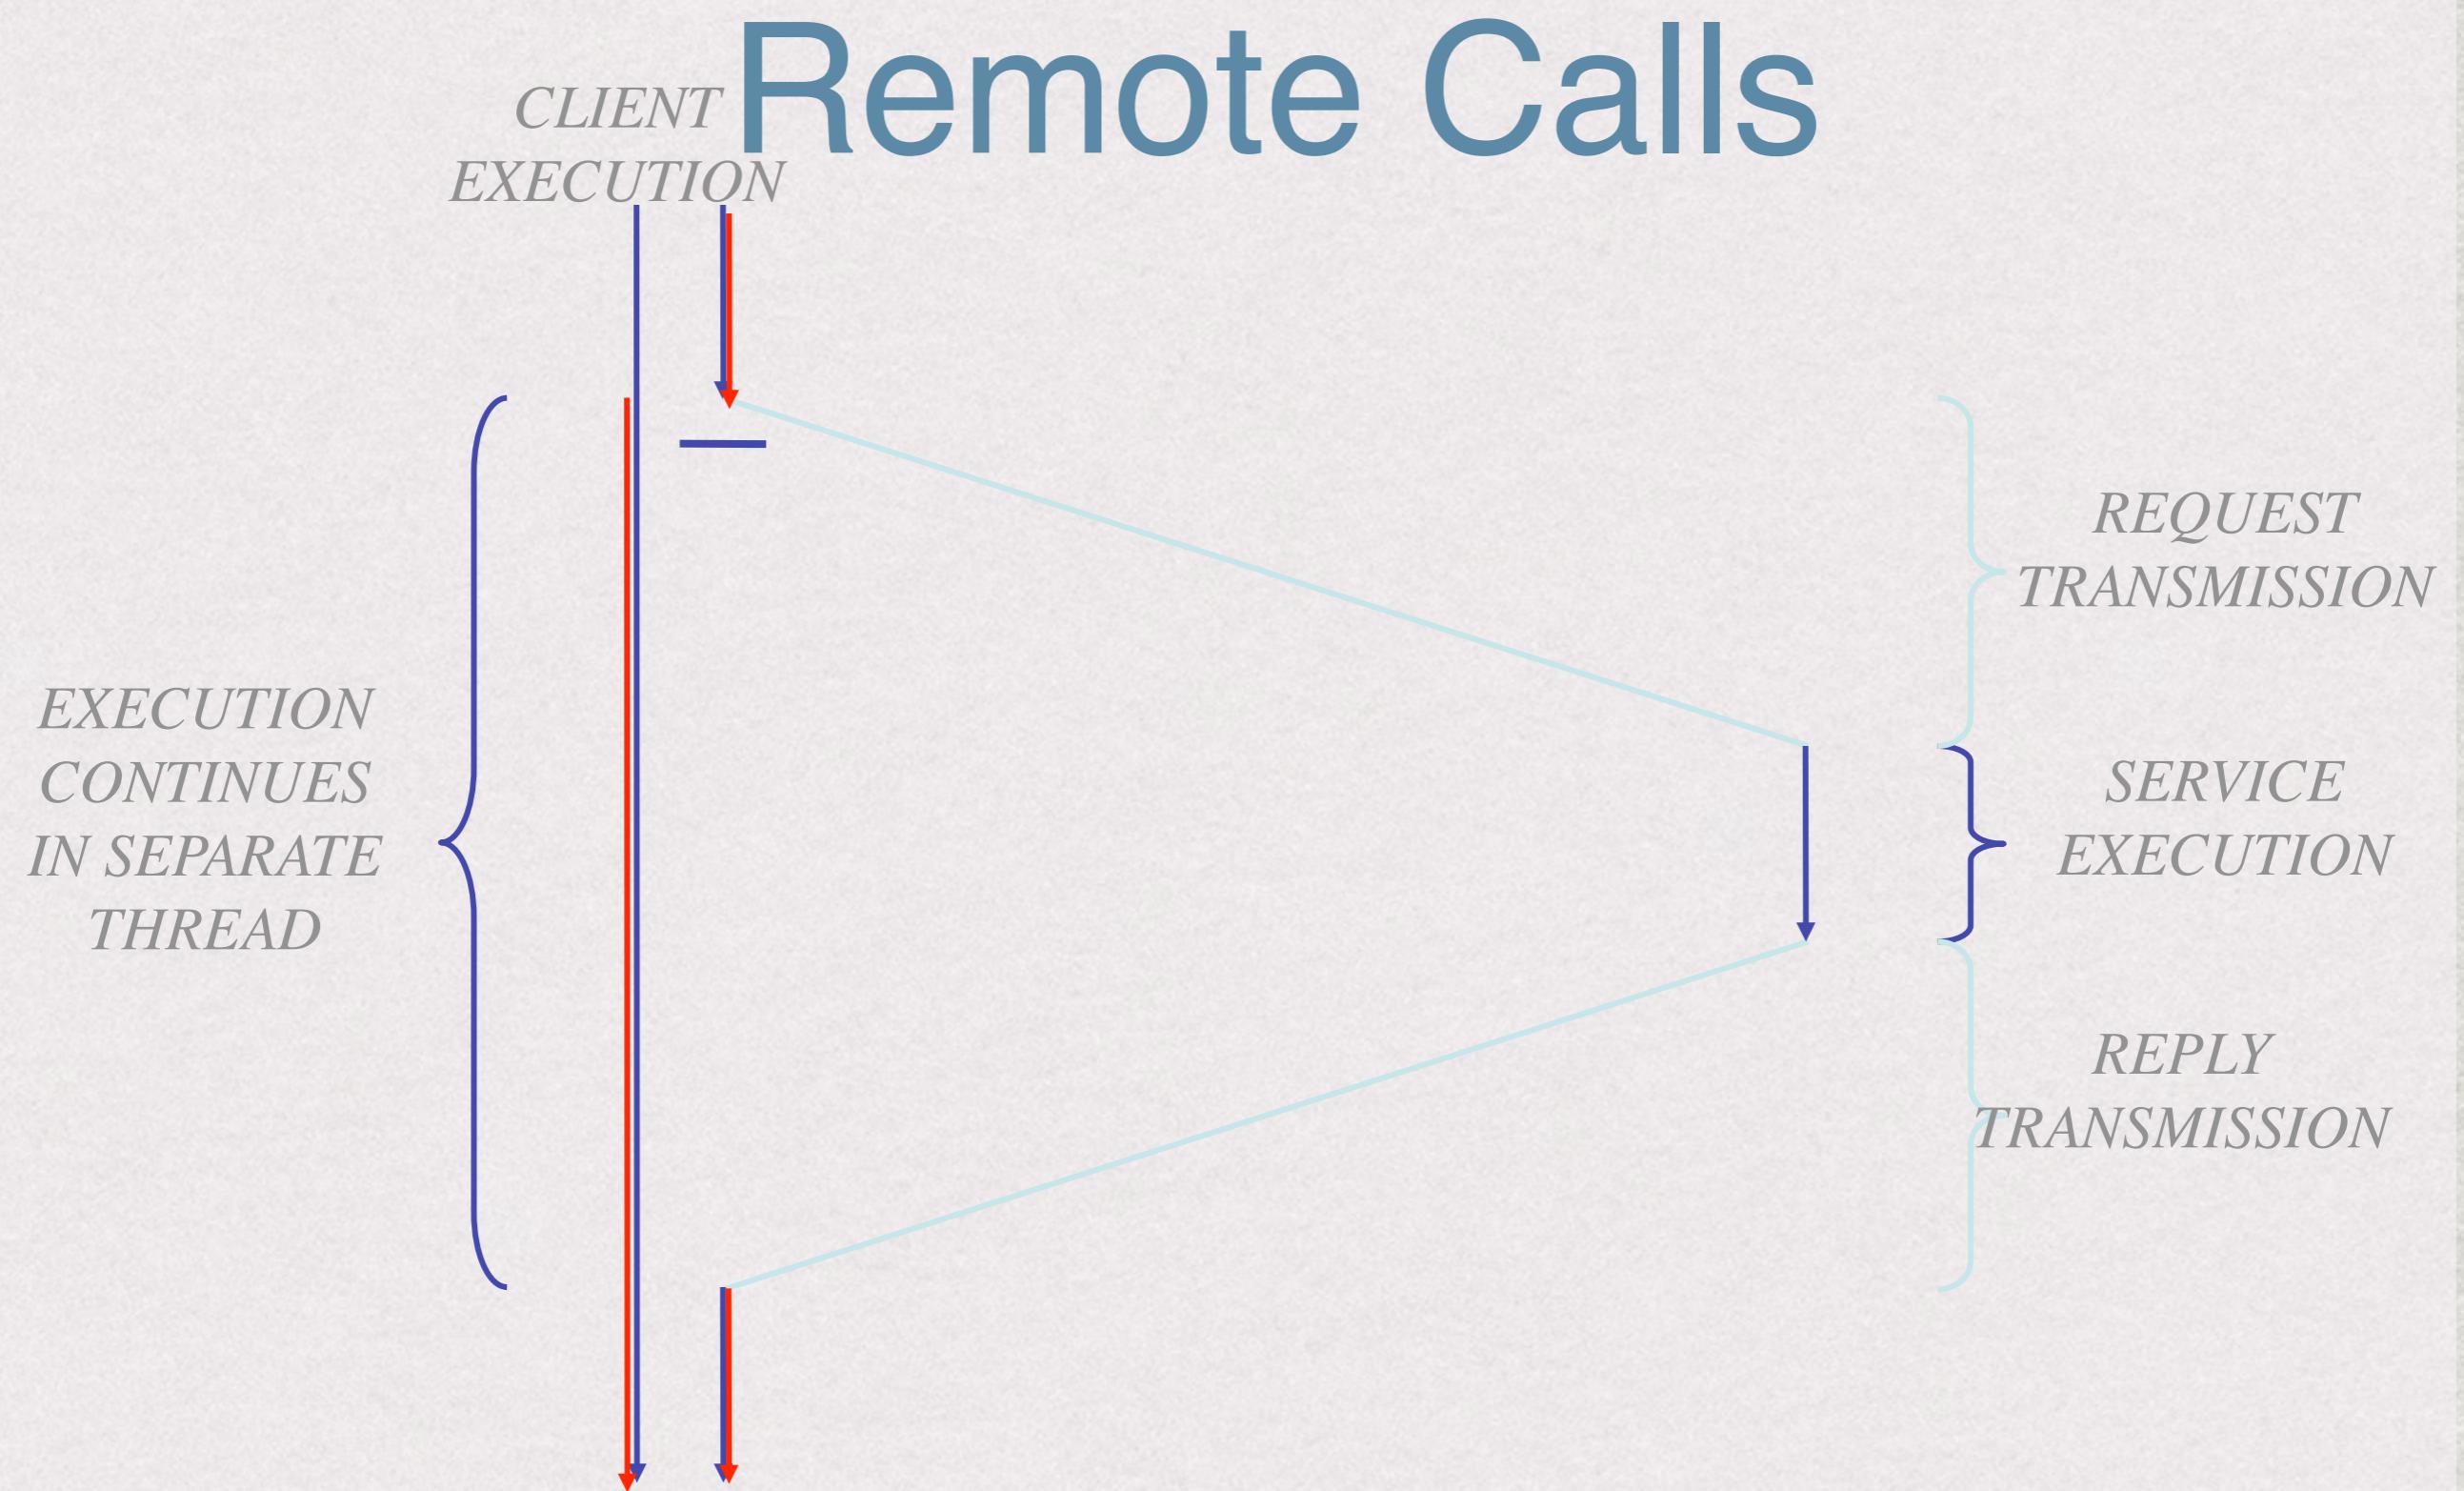
Questions

- ***What can fail between the client starting an RPC and the end of that RPC?***
- ***What happens to the client when each sort of failure occurs?***



Questions

- * **What can fail between the client starting an RPC and the end of that RPC?**
- * **What happens to the client when each sort of failure occurs?**
- * **What happens to the server when each sort of failure occurs?**



Questions

- * What can fail between the client starting an RPC and the end of that RPC?
- * What happens to the client when each sort of failure occurs?
- * What happens to the server when each sort of failure occurs?
- * What happens to the distributed system as a whole in the presence of failures?

Questions

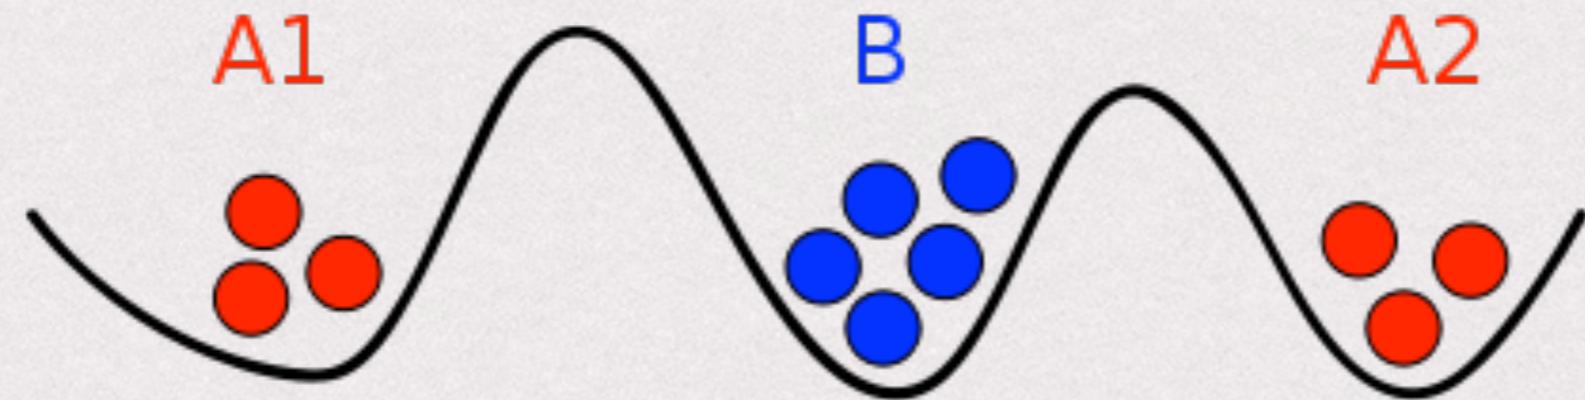
- * What can fail between the client starting an RPC and the end of that RPC?
- * What happens to the client when each sort of failure occurs?
- * What happens to the server when each sort of failure occurs?
- * What happens to the distributed system as a whole in the presence of failures?
 - * Or more positively:
How can we build distributed systems that do useful work in spite of the propensity to failure?

Google failure rates

- * In each cluster's first year, it's typical that 1,000 individual machine failures will occur; thousands of hard drive failures will occur; one power distribution unit will fail, bringing down 500 to 1,000 machines for about 6 hours;
- * 20 racks will fail, each time causing 40 to 80 machines to vanish from the network;
- * 5 racks will “go wonky,” with half their network packets missing in action; 50% chance that the cluster will overheat

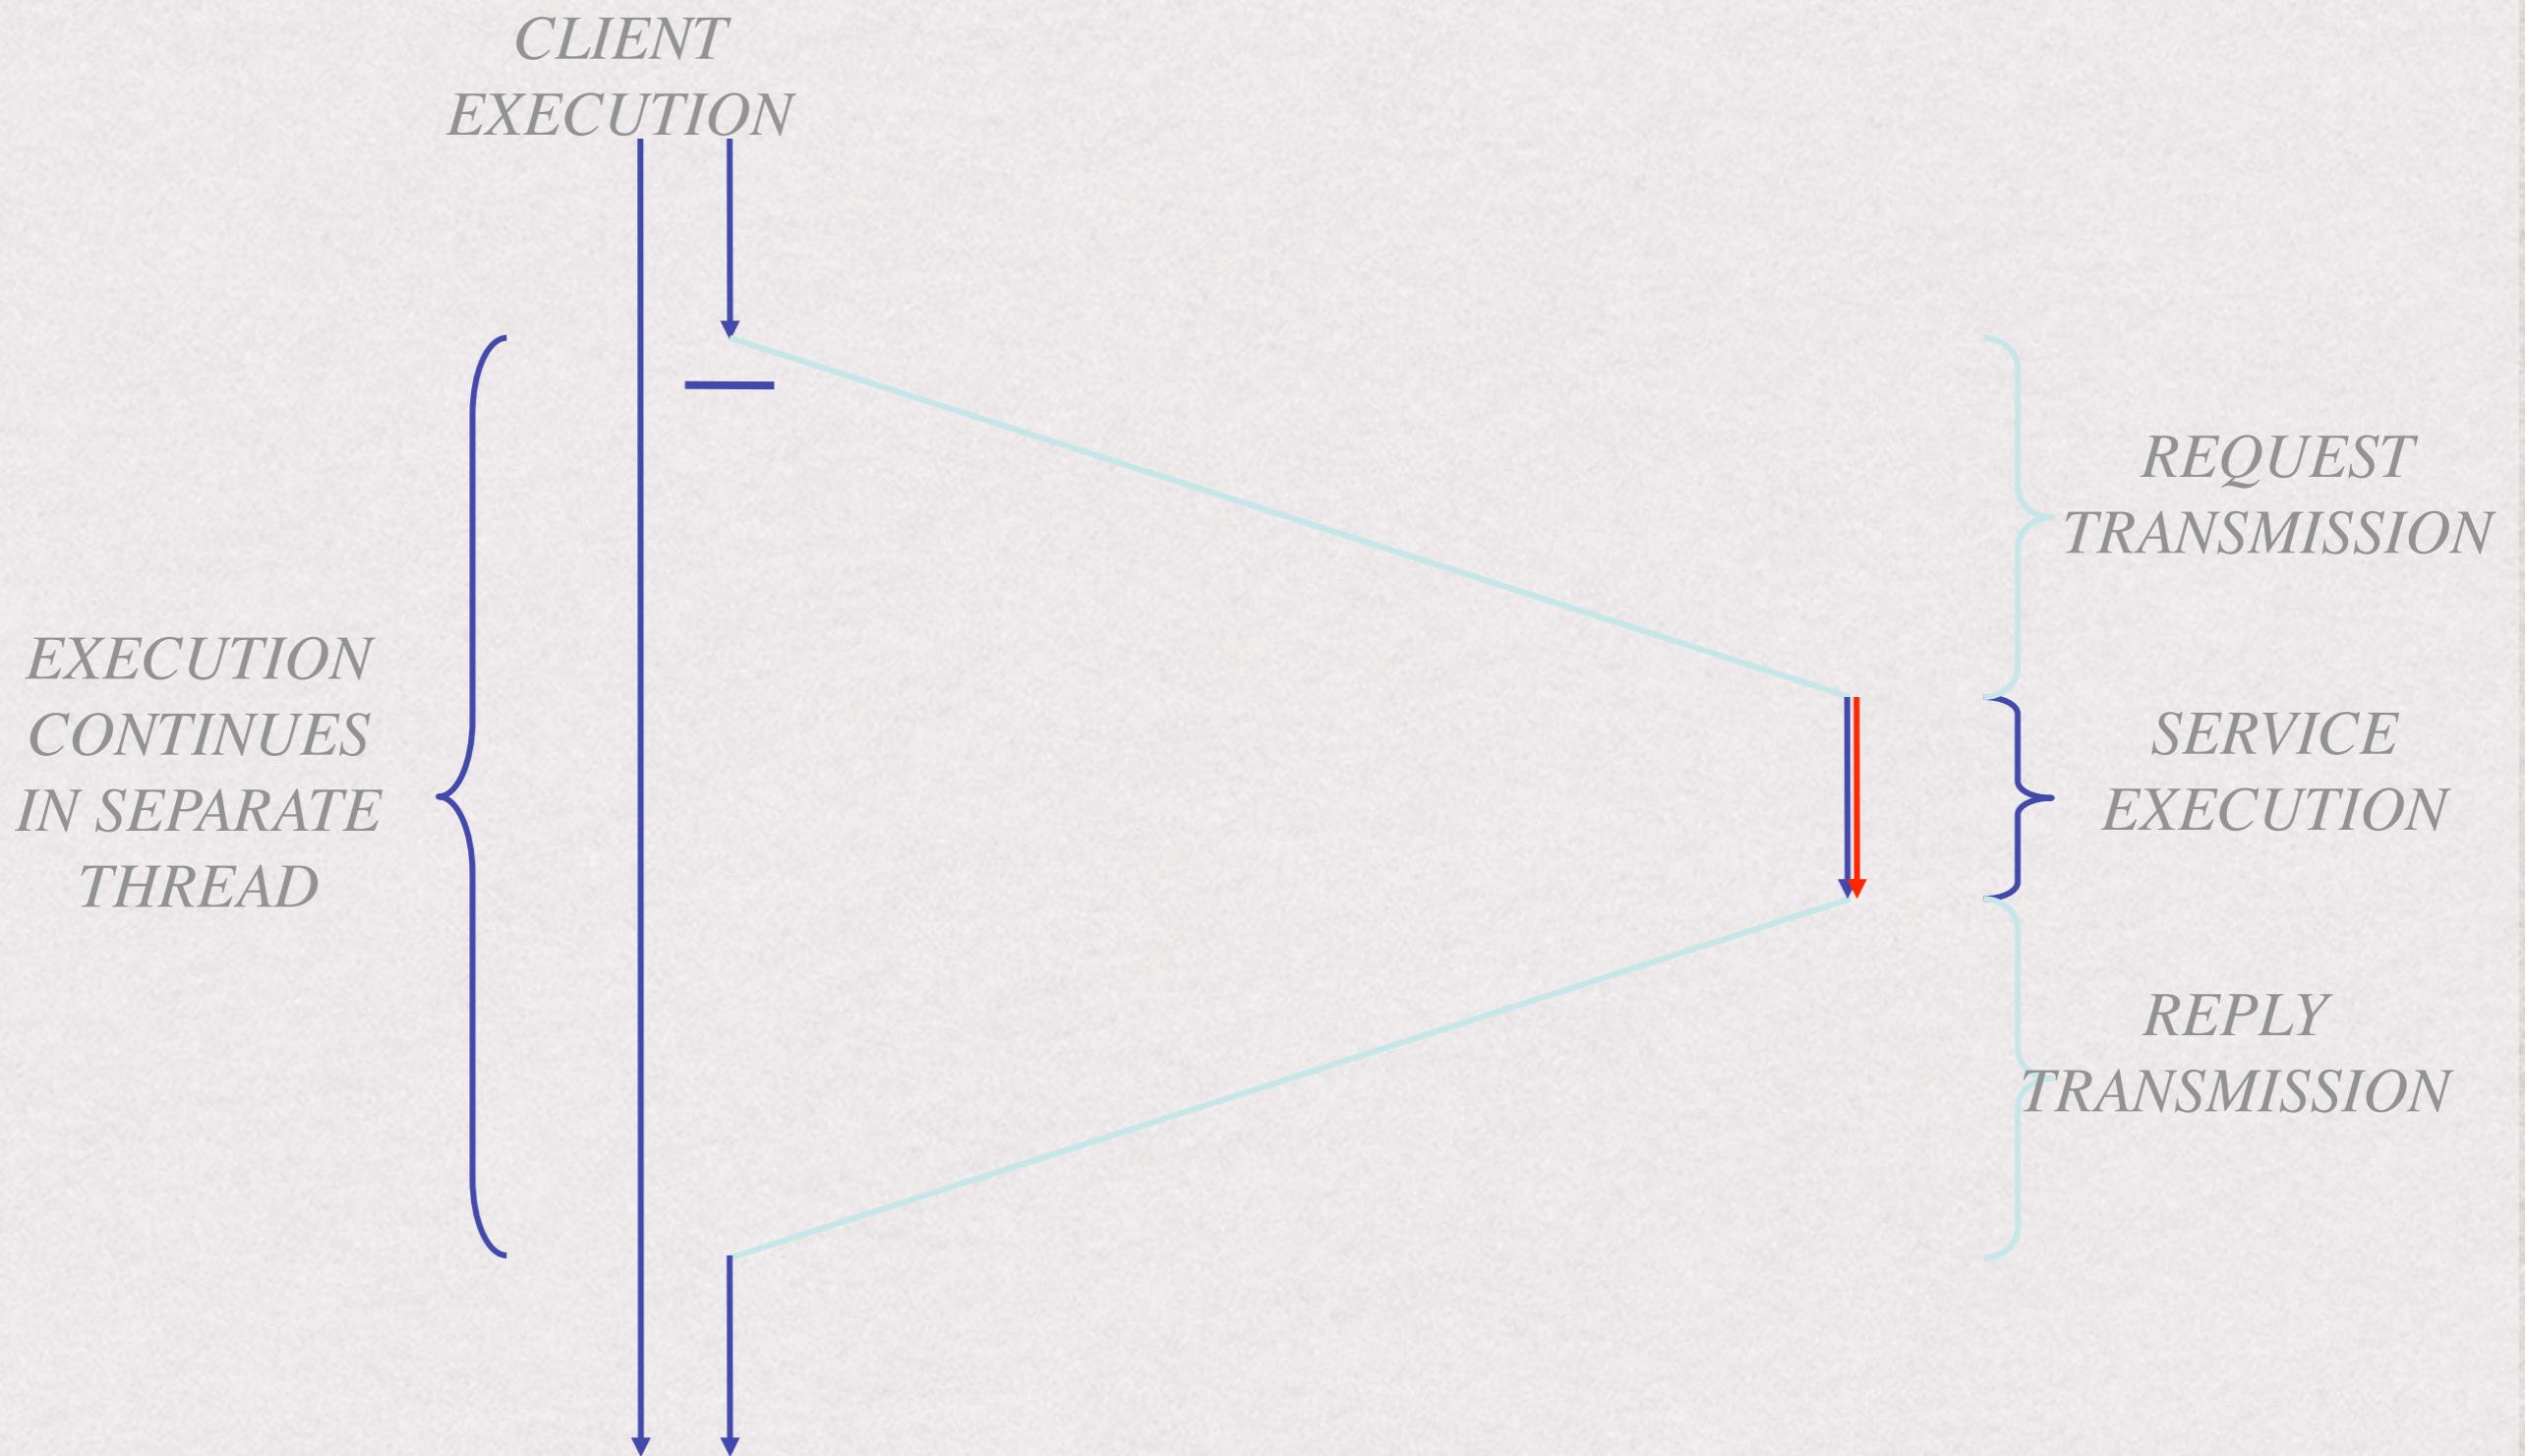
Model of Failure

- * The network may lose both request and reply messages.



- * Servers may crash, possibly whilst RPCs are in progress, and may then subsequently restart and resume receiving requests
- * Clients may crash whilst RPCs are in progress, and may then subsequently restart and resume receiving replies.
- * Messages are delivered in the order sent:
 - * There are no message re-ordering failures

Remote Calls



Server Failure

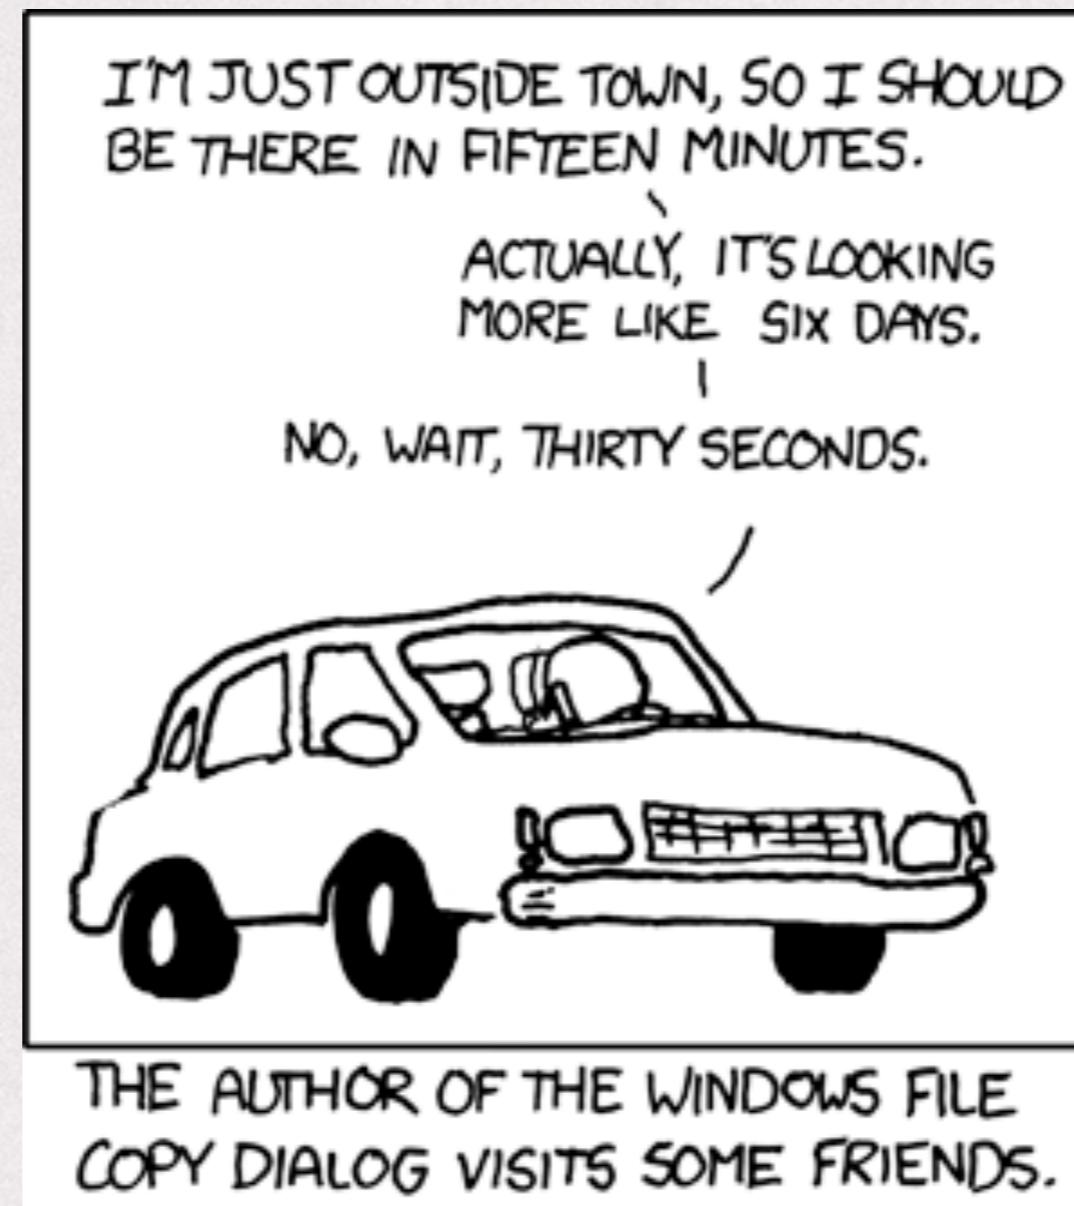
- * Server crashes and the loss of requests and replies by the network creates a problem for the client.
- * The client will use a **FAILURE DETECTOR** to discern when such a problem occurs:
 - * In principle a failure detector will detect when failures occur.
 - * In practice, failure detectors operate by failing to determine that an operation has succeeded within a given time!

Failure Detectors

- * A failure detector returns the following:
 - * It has not been possible to determine that an operation X has succeeded within time T.
- * this is interpreted as:
 - * The operation X has failed.
- * The operation may have failed, or it may be that the time T is too short:
 - * What can one do about this “false failure” result?

Use of Failure Detectors

- * Prior to sending a request message, the client sets a timer to expire at some point in the future:
 - * If the timer expires before a reply to the request arrives, then the call is **DECLARED FAILED**.
 - * However, this declaration might be wrong – it might just be that the timer duration is too short for processing of the particular request.
- * Clearly an infinitely long timer will avoid false failures:
 - * This approach is not useful however, since it will not detect true failures either!



Adaptive Timeouts

- * Instead, RPC systems use an **ADAPTIVE TIMEOUT** approach:
 - * If the initial timer expires, the client sets a new timer and sends the server a **PROBE** message.
 - * The server is configured to acknowledge probes immediately if the request is still being processed.
 - * If the client receives the probe acknowledgement before the new timer expires, it knows the server is still working on the request.
 - * After receiving a probe acknowledgement, the client sets another timer.
 - * Then, at the expiry of this timer, it sets (yet) another timer and sends off another probe.
- * This process continues until either the reply is received or one of the probe acknowledgements fails to arrive:
 - * For efficiency, the inter probe delay may increase up to a point.

Discriminating between Failures

- * The (adaptive) timeout approach does not discriminate between different possible failures:
 - * Request (or probe) loss on the way to the server.
 - * Server crash.
 - * Reply (or probe acknowledgement) loss on the way back to the client.

Discriminating between Failures

- * Client's perspective:
 - * all failures looks essentially the same – an expected message does not arrive in time.
- * It is sometimes possible to find out more detail:
 - * Once the client has received at least one probe acknowledgement from the server, it is possible to deduce that the request made it to the server.
 - * However, one can't conclude from the absence of a probe acknowledgement that the request did not make it to the server – perhaps the probe or probe acknowledgement was lost.

Re-processing of Requests

- * Consider a server operation that adds one to a number:
 - * For the RPC system to re-process such a request in the event of failure detection would be incorrect, since its not possible to tell whether the request was processed (to the extent of adding one) previously.
 - * Re-processing could lead to the number being increased by two!

```
server.increaseCounter();
```

Re-processing of Requests

- * Consider a server operation that adds one to a number:
 - * For the RPC system to re-process such a request in the event of failure detection would be incorrect, since its not possible to tell whether the request was processed (to the extent of adding one) previously.
 - * Re-processing could lead to the number being increased by two!
- * So, re-transmission (and consequent re-processing) of requests runs the risk that a request may be executed more than once:
 - * This could be bad, for example a request to deduct some money from a bank account.
 - * In these situations, the RPC system should provide AT-MOST-ONCE semantics – a given request must never be processed more than once by the server:
 - * The corollary is that sometimes a request will not be executed at all!
- * As a result, it is not always permissible for an (at-most-once) RPC system to re-process a request when a failure is detected.

An Optimisation

- ✳ Server discriminates lost request failures under certain conditions:
 - ✳ Server has an ID that is unique each time it re-starts:
 - ✳ For example, the number of times the server has started.
 - ✳ The client obtains the server's ID each time it binds, and sends that ID with each request to that server.
 - ✳ Each client request is allocated a unique sequence number.
 - ✳ On failure detection, instead of reporting an error, the client resends a copy of the request with its original sequence number.
 - ✳ The server maintains (in memory) a list of requests it has processed.
 - ✳ When the server receives a request, if it has the current server ID:
 - ✳ If in the list of processed requests, send a copy of the reply.
 - ✳ If not, process the request.
 - ✳ If the request has the wrong server ID, the server has recently crashed, and it can't tell whether it has been processed, so discard it and send back an error message.
- ✳ This technique allows the system to overcome some lost request and lost reply failures whilst retaining at-most-once semantics.

An Optimisation

1. CLIENT BINDS TO THE SERVER AND GETS ID

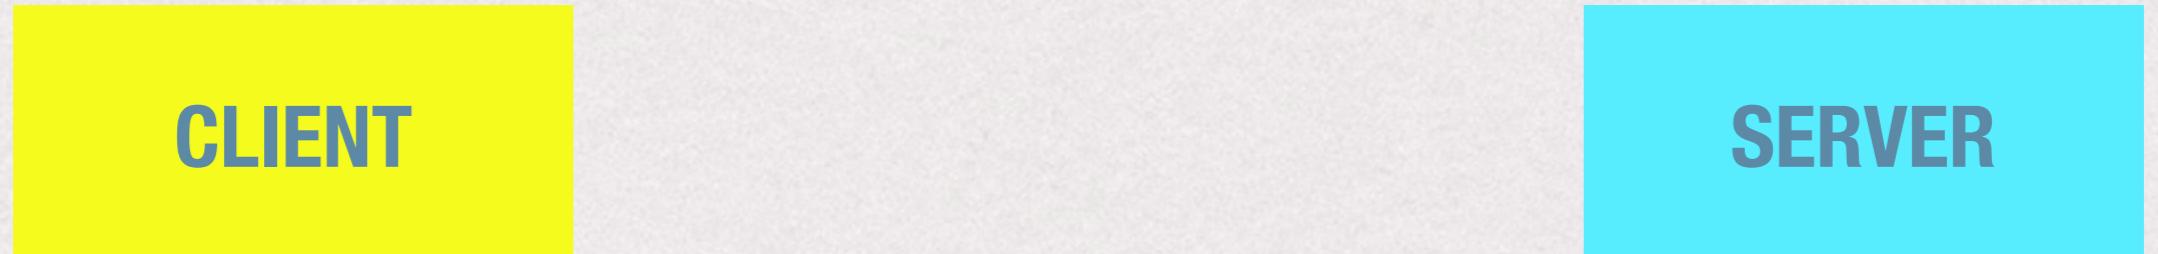


2. WHENEVER THE CLIENT TALKS TO THE SERVER, IT WILL USE SID AND A REQUEST ID



An Optimisation

3. SERVER PROCESSES REQUEST AND SAVES RESULT SOMEWHERE



SERVER ID: 1
RID 1: 42

4. IF CLIENT DOES NOT GET A RESPONSE FROM THE SERVER, IT WILL RESEND



SERVER ID: 1
RID 1: 42

SERVER ID: 1
RID 1: 42

An Optimisation

5. IF THE SERVER ID IS THE SAME, THE SERVER WILL RE-SEND THE REPLY



6. IF SERVER HAS CRASHED IN THE MEANTIME, THE SERVER IDS WILL BE DIFFERENT; ERROR MESSAGE WILL BE SENT



Client Failures

- * When a client crashes, it is wasteful to continue processing any requests from that client.
- * More importantly, the optimisation described previously relies on **unique** sequence numbers, and a client's sequence numbers are reset when the client restarts:
 - * Request messages need to include a unique client ID.
 - * This is derived in the same way as the server ID.
 - * Taken together, these form a session ID:
 - * used instead of the server ID
- * Also, the client needs to differentiate a reply to a request it has sent since it booted from a reply to a request that it sent before it crashed:
 - * Reply messages also need to have session ID and sequence numbers.
 - * Ignore replies from a previous session.

At-least-once Semantics

- * There is an alternative RPC semantics:
 - * With **AT-LEAST-ONCE** semantics, each request is guaranteed to execute to completion.
 - * The corollary is that each request may in fact execute (wholly or partially) more than one time.
- * Implementation of at-least-once semantics is trivial:
 - * Re-send each request until a reply is received!
- * So, why is at-least-once semantics not the answer to all problems?
- * Operations that are suitable for calling with at-least-once semantics are said to be **IDEMPOTENT**

Other Semantics

- * The at-most-once semantics is more generally useful, but:
 - * Even the optimisation to at-most-once doesn't give certainty in all cases – it's still at-most-once, just that the “once” part is more likely.
 - * In the event a failure is detected, the operation may have executed partially and then crashed – the RPC system says nothing about this.
- * Other possible semantics:
 - * **ZERO-OR-ONCE** – each operation either executes completely or not at all – failures due to partial executions are prevented.
 - * **RELIABLE NETWORK** – failures due to loss of request and replies are prevented.
 - * **EXACTLY-ONCE** – each operation executes to completion.
- * How do you think we can go about achieving exactly-once semantics?
 - * Hint, think about the optimisation to at-most-once, and what caused it to break down.